ACSL by Example

Towards a Formally Verified Standard Library

Version 19.1.0
for
Frama-C 19.1 (Potassium)
October 2019

Jens Gerlach
Denis Efremov
Tim Sikatzki
Malte Brodmann

Former Authors

Jochen Burghardt, Andreas Carben
Robert Clausecker, Liangliang Gu
Kerstin Hartig, Timon Lapawczyk
Hans Werner Pohl, Juan Soto
Kim Völlinger

This report was partially funded by the VESSEDIA project.[1]

This document is hosted at

```
https://github.com/fraunhoferfokus/acsl-by-example
```

From there, you can also download the source code of all algorithms discussed here, their contracts, and the employed predicate definitions and lemmas. All examples are developed and proved with the Frama-C/WP [1] plugin.[4] We recommend using the GitHub issue tracker

```
https://github.com/fraunhoferfokus/acsl-by-example/issues
```

to report suggestions or errors. Alternatively, you can email them also to

```
jens.gerlach@fokus.fraunhofer.de
```

---

[4]There is also full support for the Frama-C/AstraVer plugin which is developed at ISP RAS and can be installed with the instruction available on `https://forge.ispras.ru/projects/astraver/wiki`

Changes

This release is intended for Frama-C 19.1 (*Potassium*), issued in September 2019. [2] For changes in previous versions we refer to Appendix B on Page 273.

New in Version 19.1.0 (October 2019)

Aside from the above-mentioned version of Frama-C we are using for this release the Why3 platform [3, v1.2.0] and the provers listed in Table 0.1. Note that all automatic provers are use the Why3 interface. In other words, we do not use anymore the native interface for Alt-Ergo.

Prover	Type	Version	Reference
Alt-Ergo	automatic	2.3.0	[4]
CVC4	automatic	1.6	[5]
CVC3	automatic	2.4.1	[6]
Z3	automatic	4.8.6	[7]
E	automatic	2.3	[8]
Coq	interactive	8.9.1	[9]

Table 0.1.: Information on automatic and interactive theorem provers

Improvements

- Rename arguments of `search` and `find_end` and improve also the description of these algorithms.

- Rename and reorder arguments of `search_n`, make the verification more robust and improve its description.

- Make verification of property `size` of `remove_copy2` more robust.

- Explain role of lemma `RemoveImpliesNotHasValue` in `remove_copy3` and `remove`.

- Simplify definition of `RemoveSize` and `RemovePartition`.

- Make verification of property `reorder` of `partial_sort` more robust.

- Strengthen precondition of `replace_copy`.

- Rename lemma `random_number_modulo` into `RandomNumberModulo`.

- Differentiate between properties `unique` and `solitary` for `unique_copy` examples.

- Simplify the implementation of `is_heap` by calling the new function `is_heap_until`.

- Replace remaining instances of label `Pre` in contracts by `Old`.

- Unify use of `Unchanged` predicate for mutating algorithms.

New examples

- Add the algorithm `clamp` which "clips" a value between a pair of boundary values.

- Add the algorithm `minmax_element` and improve description of other algorithms related to finding minimum and maximum values.

- Add new example `is_heap_until` that generalizes `is_heap`.

- The following examples are not new since they were implicitly used as helper functions for other examples. They are now explicitly listed as examples.

  - `make_pair`

  - `random_number`

  - `heap_parent`

  - `heap_child_max` (formerly known as `heap_maximum_child`

Open issues

The following algorithms and/or lemmas are not completely verified

- `adjacent_difference_inv`

- `partial_sum_inv`

- `pop_heap`

- `ReorderImpliesMatch`

Contents

List of logic definitions

List of Figures

List of Tables

Part I.

Basics

1. Introduction

This report provides various examples for the formal specification, implementation, and deductive verification of C programs using the ANSI/ISO-C Specification Language (ACSL [10]) and the Frama-C/WP plug-in [1] of Frama-C [2] (Framework for Modular Analysis of C programs).

We have chosen our examples from the C++ Standard Library whose initial version is still known as the *Standard Template Library* (STL). The C++ Standard Library contains a broad collection of *generic* algorithms that work not only on C arrays but also on more elaborate container data structures. For the purposes of this document we have selected representative algorithms, and converted their implementation from C++ function templates to C functions that work on arrays of type `int`.

We will continue to extend and refine this report by describing additional STL algorithms and data structures. Thus, step by step, this document will evolve from an ACSL tutorial to a report on a formally specified and deductively verified Standard Library for ANSI/ISO-C. Moreover, as ACSL is extended to a C++ specification language, our work may be extended to a deductively verified C++ Standard Library.

We encourage you to check vigilantly whether our formal specifications capture the essence of the informal description of the STL algorithms. We appreciate your feedback[5] and hope that this document helps foster the adoption of deductive verification techniques.

Acknowledgement

Many members from the Frama-C community provided valuable input and comments during the course of the development of this document. In particular, we wish to thank our project partners Patrick Baudin, Loïc Correnson, Zaynah Dargaye, Florent Kirchner, Virgile Prevosto, and Armand Puccetti from CEA LIST[6] and Pascal Cuoq from TrustInSoft[7].

We also like to express our gratitude to Claude Marché (LRI/INRIA)[8] and Yannick Moy (AdaCore)[9] for their helpful comments and detailed suggestions for improvement. Finally, we would like to thank Aaron Rocha who sent us valuable improvement suggestions and error reports.

---

[5]We suggest GitHub's issue tracker: `https://github.com/fraunhoferfokus/acsl-by-example/issues`
[6]`http://www-list.cea.fr/en`
[7]`http://trust-in-soft.com`
[8]`https://www.lri.fr/index_en.php?lang=EN`
[9]`http://www.adacore.com`

1.1. Frama-C

The Framework for Modular Analyses of C, Frama-C [2], is a suite of software tools dedicated to the analysis of C source code. Its development efforts are conducted and coordinated at two French public institutions: CEA LIST [11], a laboratory of applied research on software-intensive technologies, and INRIA Saclay [12], the French National Institute for Research in Computer Science and Control in collaboration with LRI [13], the Laboratory for Computer Science at Université Paris-Sud.

ACSL (ANSI/ISO-C Specification Language) [10] is a formal language to express behavioral properties of C programs. This language can specify a wide range of functional properties by adding annotations to the code. It allows to create function contracts containing preconditions and postconditions. It is possible to define type and global invariants as well as logic specifications, such as predicates, lemmas, axioms or logic functions. Furthermore, ACSL allows statement annotations such as assertions or loop annotations.

Within Frama-C, the Frama-C/WP plug-in [1] enables deductive verification of C programs that have been annotated with ACSL. The Frama-C/WP plug-in uses Hoare-style weakest precondition computations to formally prove ACSL properties of C code. Verification conditions are generated and submitted to external automatic theorem provers or interactive proof assistants.

The Verification Group at Fraunhofer FOKUS [14] see the great potential for deductive verification using ACSL. However, we recognize that for a novice there are challenges to overcome in order to effectively use the Frama-C/WP plug-in for deductive verification. In order to help users gain confidence, we have written this tutorial that demonstrates how to write annotations for existing C programs. This document provides several examples featuring a variety of annotated functions using ACSL. For an in-depth understanding of ACSL, we strongly recommend users to read the official Frama-C introductory tutorial [15] first. The principles presented in this paper are also documented in the ACSL reference document [16].

1.2. Structure of this document

The functions presented in this document were selected from the C++ Standard Library. The original C++ implementation was stripped from its generic implementation and mapped to C arrays of type `value_type`.

Chapter 2 provides a short introduction into the Hoare Calculus. For a better understanding of Frama-C/WP and the theory behind it, we also recommend Allan Blanchard's ACSL tutorial [17].

We have grouped various standard algorithms in chapters as follows:

- non-mutating algorithms (Chapter 3)

- maximum/minimum algorithms (Chapter 4)

- binary search algorithms (Chapter 5)

- mutating algorithms (Chapter 6)

- numeric algorithms (Chapter 8)

- heap algorithms (Chapter 9)

- sorting algorithms and well-known classical implementations of sorting algorithms (Chapter 10)

The order of these chapters reflects their increasing complexity.

Using the example of a stack, we tackle in Chapter 11 the problem of how a data type and its associated C functions can be specified with ACSL and automatically verified with Frama-C.

Finally, Appendix A lists for each example the results of verification with Frama-C.

1.3. Types, arrays, ranges and valid indices

In order to keep algorithms and specifications as general as possible, we use abstract type names on almost all occasions. We currently defined the following types:

```
typedef int value_type;

typedef unsigned int size_type;

typedef int bool;
```

Programmers who know the types associated with C++ Standard Library containers will not be surprised that `value_type` refers to the type of values in an array whereas `size_type` will be used for the indices of an array.

This approach allows one to modify, say, an algorithm working on an `int` array to work on a `char` array by changing only one line of code, viz. the `typedef` of `value_type`. Moreover, we believe in better readability as it becomes clear whether a variable is used as an index or as a memory for a copy of an array element, just by looking at its type.

The latter reason also applies to the use of `bool`. To denote values of that type, we defined the identifiers `false` and `true` to be `0` and `1`, respectively. While any non-zero value is accepted to denote `true` in ACSL like in C the algorithms shown in this tutorial will always produce `1` for `true`. Due to the above definitions, the ACSL truth-value constant `\false` and `\true` can be used interchangeably with our `false` and `true`, respectively, in ACSL clauses, but not in C code.

1.3.1. Array and ranges

The C Standard describes an array as a "contiguously allocated nonempty set of objects" [18, §6.2.5.20]. If `n` is a constant integer expression with a value greater than zero, then

```
int a[n];
```

describes an array of type `int`. In particular, for each `i` that is greater than or equal to 0 and less than `n`, we can dereference the pointer `a+i`.

Let the following prototype represent a function, whose first argument is the address to a range and whose second argument is the length of this range.

```
void example(value_type* a, size_type n);
```

To be very precise, we have to use the term *range* instead of *array*. This is due to the fact, that functions may be called with empty ranges, i.e., with `n == 0`. Empty arrays, however, are not permitted according to the definition stated above. Nevertheless, we often use the term *array* and *range* interchangeably.

1.3.2. Specification of valid ranges in ACSL

The following ACSL fragment expresses the precondition that the function `example` expects that for each `i`, such that `0 <= i < n`, the pointer `a+i` may be safely dereferenced.

```
/*@
    requires 0 <= n;
    requires \valid(a + (0.. n-1));
*/
void example(value_type* a, size_type n);
```

In this case we refer to each index `i` with `0 <= i < n` as a *valid index* of `a`.

ACSL's built-in predicates `\valid(a + (0.. n))` and `\valid_read(a + (0.. n))` refer to all addresses `a+i` where `0 <= i <= n`. However, the array notation `int a[n]` of the C programming language refers only to the elements `a+i` where `i` satisfies `0 <= i < n`. Users of ACSL must therefore use the range notation `a+(0.. n-1)` in order to express a valid array of length `n`.

2. The Hoare calculus

In 1969, C.A.R. Hoare introduced a calculus for formal reasoning about properties of imperative programs [19], which became known as "Hoare Calculus".

The basic notion is

```
//@ assert P;
Q;
//@ assert R;
```

where `P` and `R` denote logical expressions and `Q` denotes a source-code fragment. Informally, this means

If `P` holds before the execution of `Q`, then `R` will hold after the execution.

Usually, `P` and `R` are called *precondition* and *postcondition* of `Q`, respectively. The syntax for logical expressions is described in [16, §2.2] in full detail. For the purposes of this tutorial, the notions shown in Table 2.1 are sufficient. Note that they closely resemble the logical and relational operators in C.

ACSL syntax	Name	Reading
`!P`	negation	`P` is not true
`P && Q`	conjunction	`P` is true and `Q` is true
<code>P || Q</code>	disjunction	`P` is true or `Q` is true
`P ==> Q`	implication	if `P` is true, then `Q` is true
`P <==> Q`	equivalence	if, and only if, `P` is true, then `Q` is true
`x < y == z`	relation chain	`x` is less than `y` and `y` is equal to `z`
`\forall int x; P(x)`	universal quantifier	`P(x)` is true for every **int** value of `x`
`\exists int x; P(x)`	existential quantifier	`P(x)` is true for some **int** value of `x`

Table 2.1.: Some ACSL formula syntax

Here we show three example source-code fragments and annotations.

```//@ assert x % 2 == 1;```
```++x;```
```//@ assert x % 2 == 0;```
	If `x` has an odd value before execution of the code `++x` then `x` has an even value thereafter.

```//@ assert 0 <= x <= y;```
```++x;```
```//@ assert 0 <= x <= y + 1;```
	If the value of `x` is in the range $\{0, \ldots, y\}$ before execution of the same code, then `x`'s value is in the range $\{0, \ldots, y + 1\}$ after execution.

```c //@ assert true; while (--x != 0) sum += a[x]; //@ assert x == 0; ```	Under any circumstances, the value of x is zero after execution of the loop code.

Any C programmer will confirm that these properties are valid.[10] The examples were chosen to demonstrate also the following issues:

- For a given code fragment, there does not exist one fixed pre- or postcondition. Rather, the choice of formulas depends on the actual property to be verified, which comes from the application context. The first two examples share the same code fragment, but have different pre- and postconditions.

- The postcondition need not be the most restricting possible formula that can be derived. In the second example, it is not an error that we stated only that $0 \;<=\; x$ although we know that even $1 \;<=\; x$.

- In particular, pre- and postconditions need not contain all variables appearing in the code fragment. Neither sum nor a[] is referenced in the formulas of the loop example.

- We can use the predicate **true** to denote the absence of a properly restricting precondition, as we did before the **while** loop.

- It is not possible to express by pre- and postconditions that a given piece of code will always terminate. The loop example only states that *if* the loop terminates, then $x \;==\; 0$ will hold. In fact, if x has a negative value on entry, the loop will run forever. However, if the loop terminates, $x \;==\; 0$ will hold, and that is what the loop example claims.

  Usually, termination issues are dealt with separately from correctness issues. Termination proofs may, however, refer to properties stated (and verified) using the Hoare Calculus.

Hoare provided the rules shown in Listing 2.2 to 2.12 in order to reason about programs. We will comment on them in the following sections.

---

[10]We leave the important issues of overflow aside for a moment.

2.1. The assignment rule

We start with the rule that is probably the least intuitive of all Hoare-Calculus rules, viz. the assignment rule. It is depicted in Listing 2.2, where

$$P\{x \mapsto e\}$$

denotes the result of substituting each occurrence of the variable `x` in the predicate `P` by the expression `e`.

```
//@ assert P {x |--> e};
x = e;
//@ assert P;
```

Listing 2.2: The assignment rule

For example, if `P` is the predicate

```
x > 0 && a[2*x] == 0
```

then $P\{x \mapsto y + 1\}$ is the predicate

```
y+1 > 0 && a[2*(y+1)] == 0
```

Hence, we get Listing 2.3 as an example instance of the assignment rule. Note that parentheses are required in the index expression to get the correct `2*(y+1)` rather than the faulty `2*y+1`.

```
//@ assert y+1 > 0 && a[2*(y+1)] == 0;
x = y+1;
//@ assert x > 0 && a[2*x] == 0;
```

Listing 2.3: An assignment rule example instance

Note that after a substitution several different predicates `P` may result in the same predicate $P\{x \mapsto e\}$. For example, after applying the substitution $P\{x \mapsto y + 1\}$ each of the following four predicates

```
x   > 0 && a[2*x]      == 0
x   > 0 && a[2*(y+1)] == 0
y+1 > 0 && a[2*x]      == 0
y+1 > 0 && a[2*(y+1)] == 0
```

turns into

```
y+1 > 0 && a[2*(y+1)] == 0
```

For this reason, the same precondition and statement may result in several different postconditions (All four above expressions are valid postconditions in Listing 2.3, for example). However, given a postcondition and a statement, there is only one precondition that corresponds.

When first confronted with Hoare's assignment rule, most people are tempted to think of a simpler and more intuitive alternative, shown in Listing 2.4.

```
//@ assert P;
x = e;
//@ assert P && x == e;
```

Listing 2.4: Simpler, but *faulty* assignment rule

Listings 2.5–2.7 show some example instances of this faulty rule.

```
//@ assert y > 0;
x = y+1;
//@ assert y > 0 && x == y+1;
```

Listing 2.5: An example instance of the faulty rule from Listing 2.4

While Listing 2.5 happens to be ok, Listing 2.6 and 2.7 lead to postconditions that are obviously nonsensical formulas.

```
//@ assert true;
x = x+1;
//@ assert x == x+1;
```

Listing 2.6: An example instance of the faulty rule from Listing 2.4

The reason is that in the assignment in Listing 2.6 the left-hand side variable x also appears in the right-hand side expression e, while the assignment in Listing 2.7 just destroys the property from its precondition.

```
//@ assert x < 0;
x = 5;
//@ assert x < 0 && x == 5;
```

Listing 2.7: An example instance of the faulty rule from Listing 2.4

Note that the correct example Listing 2.5 can as well be obtained as an instance of the correct rule from Listing 2.2, since replacing x by y+1 in its postcondition yields y > 0 && y+1 == y+1 as precondition, which is logically equivalent to just y > 0.

2.2. The sequence rule

The sequence rule, shown in Listing 2.8, combines two code fragments `Q` and `S` into a single one `Q ; S`. Note that the postcondition for `Q` must be identical to the precondition of `S`. This just reflects the sequential execution ("first do `Q`, then do `S`") on a formal level. Thanks to this rule, we may "annotate" a program with interspersed formulas, as it is done in Frama-C.

```
//@ assert P;
Q;
//@ assert R;
```
and
```
//@ assert R;
S;
//@ assert T;
```
⤳
```
//@ assert P;
Q ; S;
//@ assert T;
```
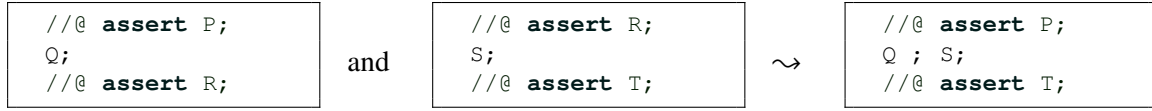
Listing 2.8: The sequence rule

2.3. The implication rule

The implication rule, shown in Listing 2.9, allows us at any time to sharpen a precondition `P` and to weaken a postcondition `R`. More precisely, if we know that `P' ==> P` and `R ==> R'` then the we can replace the left contract in of Listing 2.9 by the right one.

```
//@ assert P;
Q;
//@ assert R;
```
⤳
```
//@ assert P';
Q;
//@ assert R';
```

Listing 2.9: The implication rule

2.4. The choice rule

The choice rule, depicted in Listing 2.10, is needed to verify conditional statements of the form

```
if (C) X;
else   Y;
```

Both the then and else branch must establish the same postcondition, viz. S. The implication rule can be used to weaken differing postconditions S1 of a then-branch and S2 of an else-branch into a unified postcondition S1 || S2, if necessary. In each branch, we may use what we know about the condition C. For example, in the else-branch, we may use that C is false. If the else-branch is missing, it can be considered as consisting of an empty sequence, having the postcondition P && !C.

```
//@ assert P && C;                    //@ assert P && !C;                    //@ assert P;
X;                and                 Y;                ⤳                     if (C) X;
//@ assert S;                         //@ assert S;                          else   Y;
                                                                             //@ assert S;
```

Listing 2.10: The choice rule

Listing 2.11 shows an example application of the choice rule.

```
//@ assert 0 <= i < n;          // given precondition
if (i < n-1) {
  //@ assert 0 <= i < n - 1;    // using that i < n-1 holds in this branch
  //@ assert 1 <= i+1 < n;      // by the implication rule
  i = i+1;
  //@ assert 1 <= i < n;        // by the assignment rule
  //@ assert 0 <= i < n;        // weakened by the implication rule
} else {
  //@ assert 0 <= i == n-1 < n; // using that !(i < n-1) holds in else part
  //@ assert 0 == 0 && 0 < n;   // weakened by the implication rule
  i = 0;
  //@ assert i == 0 && 0 < n;   // by the assignment rule
  //@ assert 0 <= i < n;        // weakened by the implication rule
}
//@ assert 0 <= i < n;          // by the choice rule from both branches
```

Listing 2.11: An example application of the choice rule

The variable i may be used as an index into a ring buffer int a[n]. The shown code fragment just advances the index i appropriately. We verified that i remains a valid index into a[] provided it was valid before. Note the use of the implication rule to establish preconditions for the assignment rule as needed, and to unify the postconditions of the then and else branches, as required by the choice rule.

2.5. The loop rule

The loop rule, shown in Listing 2.12, is used to verify a **while** loop. This requires to find an appropriate formula, P, which is preserved by each execution of the loop body. P is also called a loop invariant.

```
//@ assert P && B;
S;
//@ assert P;
```
$\leadsto$
```
//@ assert P;
while (B) {
  S;
}
//@ assert !B && P;
```

Listing 2.12: The loop rule

To find it requires some intuition in many cases; for this reason, automatic theorem provers usually have problems with this task.

As said above, the loop rule does not guarantee that the loop will always eventually terminate. It merely assures us that, if the loop has terminated, the postcondition holds. To emphasize this, the properties verifiable with the Hoare Calculus are usually called "partial correctness" properties, while properties that include program termination are called "total correctness" properties.

As an example application, let us consider an abstract ring-buffer. Listing 2.13 shows a verification proof for the index i lying always within the valid range [0..n-1] during, and after, the loop. It uses the proof from Listing 2.11 as a sub-part.

```
//@ assert  0 < n;                  // given precondition

int i = 0;
//@ assert  0 <= i < n;             // by the assignment rule

while (!done) {
  //@ assert 0 <= i < n && !done;   // may be assumed by the loop rule

  a[i] = getchar();
  //@ assert 0 <= i < n && !done;   // required property of getchar
  //@ assert 0 <= i < n;            // weakened by the implication rule

  i = (i < n-1) ? i+1 : 0;
  //@ assert 0 <= i < n;            // follows by the choice rule

  process(a, i, &done);
  //@ assert 0 <= i < n;            // required property of process
}
//@ assert 0 <= i < n;              // by the loop rule
```

Listing 2.13: An abstract ring buffer loop

To reuse the proof from Listing 2.11, we had to drop the conjunct `!done`, since we didn't consider it in Listing 2.11. In general, we may *not* infer

```
//@ assert P && S;
Q;
//@ assert R && S;
```
from
```
//@ assert P;
Q;
//@ assert R;
```

since the code fragment `Q` may just destroy the property `S`.

This is obvious for `Q` being the fragment from Listing 2.11, and `S` being e.g. `i != 0`.

Suppose for a moment that `process` had been implemented in a way such that it refuses to set `done` to **true** unless it is **false** at entry. In this case, we would really need that `!done` still holds after execution of Listing 2.11. We would have to do the proof again, looping-through an additional conjunct `!done`.

We have similar problems to carry the property `0 <= i < n && !done` and `0 <= i < n` over the statement `a[i] = getchar()` and `process(a, i, &done)`, respectively. We need to specify that neither `getchar` nor `process` is allowed to alter the value of `i` or `n`. In ACSL, there is a particular language construct `assigns` for that purpose, which is introduced in Section 6.4 on Page 95.

In our example, the loop invariant can be established between any two statements of the loop body. However, this need not be the case in general. The loop rule only requires the invariant holds before the loop and at the end of the loop body. For example, `process` could well change the value of `i`[11] and even `n` intermediately, as long as it re-establishes the property `0 <= i < n` immediately prior to returning.

The loop invariant, `0 <= i < n`, is established by the proof in Listing 2.11 also after termination of the loop. Thus, e.g., a final `a[i] = '\0'` after the loop would be guaranteed not to lead to a bounds violation.

Even if we would need the property `0 <= i < n` to hold only immediately before the assignment `a[i] = getchar()`, for example since `process`'s body didn't use `a` or `i`, we would still have to establish `0 <= i < n` as a loop invariant by the loop rule, since there is no other way to obtain any property inside a loop body. Apart from this formal reason it is obvious that `0 <= i < n` wouldn't hold during the second loop iteration unless we re-established it at the end of the first one, and that is just what the while rule requires.

---

[11] We would have to change the call to `process(a, &i, &done)` and the implementation of `process` appropriately. In this case we couldn't rely on the above-mentioned `assigns` clause for `process`.

2.6. Derived rules

The above rules do not cover all kinds of statements allowed in C. However, missing C-statements can be rewritten into a form that is semantically equivalent and covered by the Hoare rules.

For example, if the expression E doesn't have side-effects, then

```
switch (E) {
    case E1: Q1; break; ...
    case En: Qn; break;
    default: Q0; break;
}
```

is semantically equivalent to

```
if (E == E1) {
    Q1;
} else ... if (E == En) {
    Qn;
} else {
    Q0;
}
```

While the **if-else** form is usually slower in terms of execution speed on a real computer, this doesn't matter for verification purposes, which are separate from execution issues.

Similarly, a loop statement of the form

```
for (P; Q; R) {
  S;
}
```

can be re-expressed as

```
P;
while (Q) {
  S;
  R;
}
```

and so on.

It is then possible to derive a Hoare rule for each kind of statement not previously discussed, by applying the classical rules to the corresponding re-expressed code fragment. However, we do not present these derived rules here.

Although procedures cannot be re-expressed in the above way if they are (directly or mutually) recursive, it is still possible to derive Hoare rules for them. This requires the finding of appropriate "procedure invariants" similar to loop invariants. Non-recursive procedures can, of course, just be inlined to make the classical Hoare rules applicable.

Note that **goto** cannot be rewritten in the above way; in fact, programs containing **goto** statements cannot be verified with the Hoare Calculus. See [20] for a similar calculus that can deal with arbitrary flowcharts, and hence arbitrary jumps. In fact, Hoare's work was based on that calculus. Later calculi inspired from Hoare's work have been designed to re-integrate support for arbitrary jumps. However, in this tutorial, we will not discuss example programs containing a **goto**.

Part II.

Nonmutating and simple search algorithms

3. Non-mutating algorithms

In this chapter, we consider *non-mutating* algorithms of the C++ Standard Library [21, §28.5]. These algorithms neither change their arguments nor any objects outside their scope. This requirement can be formally expressed with the following *assigns clause*:

```
assigns \nothing;
```

Each algorithm in this chapter therefore uses this assigns clause in its specification.

The specifications of these algorithms are not very complex. Nevertheless, we have tried to arrange them so that the earlier examples are simpler than the later ones. Each algorithm works on one-dimensional arrays.

- `find` (Section 3.1 on Page 34) provides *sequential* or *linear search* and returns the smallest index at which a given value occurs in a given range. In Section 3.2, on Page 36, a predicate is introduced in order to simplify the specification. We refer to the simplified version as `find2`.

- `find_first_of` (Section 3.3, on Page 38) provides similar to `find` a *sequential search*. However, unlike `find` it does not search for a particular value, but for an arbitrary member of a set.

- `adjacent_find` (Section 3.4 on Page 40) can be used to find equal neighbors in an array.

- `equal` and `mismatch` (Section 3.5 on Page 42) are useful for comparing two ranges element-by-element and identifying where they differ.

- `search` and `search_n` (Sections 3.6 and 3.7) find a subsequence that is identical to a given sequence when compared element-by-element and returns the position of the first occurrence.

- `count` (Section 3.9, on Page 54) returns the number of occurrences of a given value in a range. Here we will explicitly define a logic function for elements counting and show that the implementation comply with it.

- `count2` (Section 3.10, on Page 58) contains different specification for the `count` function. In this case an inductive predicate defined for elements counting. The section allows one to compare different approaches of writing specifications and demonstrates the ACSL inductive predicates.

3.1. The `find` algorithm

The `find` algorithm in the C++ Standard Library [21, §28.5.5] implements *sequential search* for general sequences. We have modified the generic implementation, which relies heavily on C++ templates, to that of a range of type `value_type`. The signature now reads:

```
size_type find(const value_type* a, size_type n, value_type val);
```

The function `find` returns the least *valid* index `i` of `a` where the condition `a[i] == val` holds. If no such index exists then `find` returns the length `n` of the array.

As an example, we consider in Figure 3.1 an array. The arrows indicate which indices will be returned by `find` for a given value. Note that the index 9 points *one past end* of the array. Values that are not contained in the array are colored in gray.

Figure 3.1.: Some simple examples for `find`

3.1.1. Formal specification of `find`

The formal specification of `find` in ACSL is shown in Listing 3.2.

```
/*@
  requires  \valid_read(a + (0..n-1));

  assigns   \nothing;

  ensures   0 <= \result <= n;

  behavior some:
    assumes  \exists integer i; 0 <= i < n && a[i] == val;
    ensures  0 <= \result < n;
    ensures  a[\result] == val;
    ensures  \forall integer i; 0 <= i < \result ==> a[i] != val;

  behavior none:
    assumes  \forall integer i; 0 <= i < n ==> a[i] != val;
    ensures  \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type
find(const value_type* a, size_type n, value_type val);
```

Listing 3.2: Formal specification of `find`

The `requires`-clause indicates that n is non-negative and that the pointer a points to *n* contiguously allocated objects of type `value_type` (see Section 1.3). The `assigns`-clause indicates that `find` (as a non-mutating algorithm), does not modify any memory location outside its scope (see Page 33).

Generally, we only know that `find` returns a non-negative index that is less or equal the length of the array. However, once we assume more specific situations, we can also make more precise statements about the returned valued. This is the reason why we have subdivided the specification of `find` into two behaviors (named `some` and `none`).

- The behavior `some` applies if the sought-after value is contained in the array. We express this condition by using the `assumes`-clause. The next line expresses that if the assumptions of the behavior are satisfied then `find` will return a valid index. The algorithm also ensures that the returned (valid) index i, `a[i] == val` holds. Therefore we define this property in the second postcondition of behavior `some`. Finally, it is important to express that `find` returns the smallest index i for which `a[i] == val` holds (see last postcondition of behavior `some`).

- The behavior `none` covers the case that the sought-after value is *not* contained in the array (see `assumes`-clause of behavior `none` in Listing 3.2). In this case, `find` must return the length n of the range a.

Note that the formula in the `assumes`-clause of the behavior `some` is the negation of the `assumes`-clause of the behavior `none`. Therefore, we can express that these two behaviors are *complete* and *disjoint*.

3.1.2. Implementation of `find`

Listing 3.3 shows a straightforward implementation of `find`. The only noteworthy elements of this implementation are the *loop annotations*.

```
size_type
find(const value_type* a, size_type n, value_type val)
{
  /*@
    loop invariant 0 <= i <= n;
    loop invariant \forall integer k; 0 <= k < i ==> a[k] != val;
    loop assigns i;
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; i++) {
    if (a[i] == val) {
      return i;
    }
  }

  return n;
}
```

Listing 3.3: Implementation of `find`

The first loop *invariant* is needed to prove that accesses to a only occur with valid indices. The second loop *invariant* is needed for the proof of the postconditions of the behavior `some` (see Listing 3.2). It expresses that for each iteration the sought-after value is not yet found up to that iteration step.

Finally, the loop *variant* `n-i` is needed to generate correct verification conditions for the termination of the loop.

3.2. The `find` algorithm reconsidered

In this section we specify the `find` algorithm in a slightly different way when compared to Section 3.1. Our approach is motivated by a considerable number of closely related formulas. We have in Listings 3.2 and 3.3 the following formulas

```
\exists integer i; 0 <= i < n        &&  a[i] == val;
\forall integer i; 0 <= i < \result ==> a[i] != val;
\forall integer i; 0 <= i < n        ==> a[i] != val;
\forall integer k; 0 <= k < i        ==> a[k] != val;
```

Note that the first formula is the negation of the third one.

3.2.1. The predicate `HasValue`

In order to be more explicit about the commonalities of these formulas we define a predicate, called `HasValue` (see Listing 3.4), which describes the situation that there is a valid index `i` where `a[i]` equals `val`.

```
/*@
  predicate
    HasValue{A}(value_type* a, integer m, integer n, value_type v) =
      \exists integer i; m <= i < n && a[i] == v;

  predicate
    HasValue{A}(value_type* a, integer n, value_type v) =
      HasValue(a, 0, n, v);
*/
```

Listing 3.4: The predicate `HasValue`

Note that we needed to provide a label, viz. `A`, to the predicate, since the evaluation the predicate depends on a memory state, viz. the contents of `a[0..n-1]`. In general, we have to write

```
\exists integer i; 0 <= i < n && \at(a[i],A) == v;
```

in order to express that we refer to the value `a[i]` in the program state `A`. However, ACSL allows to abbreviate `\at(a[i],A)` by `a[i]` if, as in `HasValue`, the label `A` is the only available label.

With this predicate we can encapsulate all uses of the universal and existential quantifiers in both the function contract of `find` and in its loop annotations. The result is shown in Listings 3.5 and 3.6.

3.2.2. Formal specification of `find`

Th revised contract for `find` in Listing 3.5 is more concise than the previous one in Listing 3.2. In particular, it can be seen immediately that the conditions in the assumes clauses of the two behaviors `some` and `none` are mutually exclusive since one is the literal negation of the other. Moreover, the requirement that `find` returns the smallest index can also be expressed using the `HasValue` predicate, as depicted with the last postcondition of behavior `some` shown in Listing 3.5.

```
/*@
  requires  valid:  \valid_read(a + (0..n-1));

  assigns   \nothing;

  ensures   result: 0 <= \result <= n;

  behavior some:
    assumes  HasValue(a, n, val);
    ensures  bound:   0 <= \result < n;
    ensures  result:  a[\result] == val;
    ensures  first:   !HasValue(a, \result, val);

  behavior none:
    assumes  !HasValue(a, n, val);
    ensures  result:  \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type
find2(const value_type* a, size_type n, value_type val);
```

Listing 3.5: Formal specification of `find` using the `HasValue` predicate

We also enriched the specification of `find` by user-defined names (sometimes called *labels*, too, the distinction to program state identifiers being obvious) to refer to the `requires` and `ensures` clauses. We highly recommend this practice in particular for more complex annotations. For example, Frama-C can be instructed to verify only clauses with a given name.

3.2.3. Implementation of **find**

The predicate `HasValue` is also used in the loop annotation inside the implementation of `find`. Note that, as in the case of the specification, we use labels to name individual annotations.

```
size_type
find2(const value_type* a, size_type n, value_type val)
{
  /*@
    loop invariant bound:     0 <= i <= n;
    loop invariant not_found: !HasValue(a, i, val);
    loop assigns i;
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; i++) {
    if (a[i] == val) {
      return i;
    }
  }

  return n;
}
```

Listing 3.6: Implementation of `find` with loop annotations based on `HasValue`

3.3. The `find_first_of` algorithm

The `find_first_of` algorithm [21, §28.5.7] is closely related to `find` (see Sections 3.1 and 3.2).

```
size_type find_first_of(const value_type* a, size_type m,
                        const value_type* b, size_type n);
```

Like `find`, it performs a sequential search. However, while `find` searches for a particular value, the function `find_first_of` returns the least index `i` such that `a[i]` is equal to one of the values `b[0..n-1]`.

Figure 3.7.: A simple example for `find_first_of`

As an example, we consider in Figure 3.7 two arrays. The arrow indicates the smallest index where one of the elements of the three-element array occurs.

3.3.1. The predicate `HasValueOf`

Similar to our approach in Section 3.2, we define a predicate `HasValueOf` that formalizes the fact that there are valid indices `i` and `j` of the respective arrays `a` and `b` such that `a[i] == b[j]` holds. We have chosen to reuse the predicate `HasValue` (Listing 3.4) to define `HasValueOf` (Listing 3.8).

```
/*@
  predicate
    HasValueOf{A}(value_type* a, integer m, value_type* b, integer n) =
      \exists integer i; 0 <= i < m && HasValue{A}(b, n, a[i]);
*/
```

Listing 3.8: The predicate `HasValueOf`

3.3.2. Formal specification of `find_first_of`

The formal specification of `find_first_of` is shown Listing 3.9. The function contract uses the predicates `HasValueOf` and `HasValue` thereby making it very similar the specification `find` (Listing 3.5).

```
/*@
  requires valid: \valid_read(a + (0..m-1));
  requires valid: \valid_read(b + (0..n-1));

  assigns \nothing;

  ensures result: 0 <= \result <= m;

  behavior found:
    assumes HasValueOf(a, m, b, n);
    ensures bound:   0 <= \result < m;
    ensures result: HasValue(b, n, a[\result]);
    ensures first:  !HasValueOf(a, \result, b, n);

  behavior not_found:
    assumes !HasValueOf(a, m, b, n);
    ensures result: \result == m;

  complete behaviors;
  disjoint behaviors;
*/
size_type
find_first_of(const value_type* a, size_type m,
              const value_type* b, size_type n);
```

Listing 3.9: Formal specification of find_first_of

3.3.3. Implementation of `find_first_of`

Our implementation of find_first_of is shown in Listing 3.10.

```
size_type
find_first_of (const value_type* a, size_type m,
               const value_type* b, size_type n)
{
  /*@
    loop invariant bound:     0 <= i <= m;
    loop invariant not_found: !HasValueOf(a, i, b, n);
    loop assigns i;
    loop variant m-i;
  */
  for (size_type i = 0u; i < m; i++) {
    if (find2(b, n, a[i]) < n) {
      return i;
    }
  }

  return m;
}
```

Listing 3.10: Implementation of find_first_of

Note the call of the find function. We opted for an implementation of find_first_of that emphasizes reuse. Besides, leading to a more concise implementation, we also have to write fewer loop annotations.

3.4. The `adjacent_find` algorithm

The `adjacent_find` algorithm of the C++ Standard Library [21, §28.5.8]

```
 size_type adjacent_find(const value_type* a, size_type n);
```

returns the smallest valid index i, such that i+1 is also a valid index and such that

```
    a[i] == a[i+1]
```

holds. The `adjacent_find` algorithm returns n if no such index exists.

The arrow in Figure 3.11 indicates the smallest index where two adjacent elements are equal.
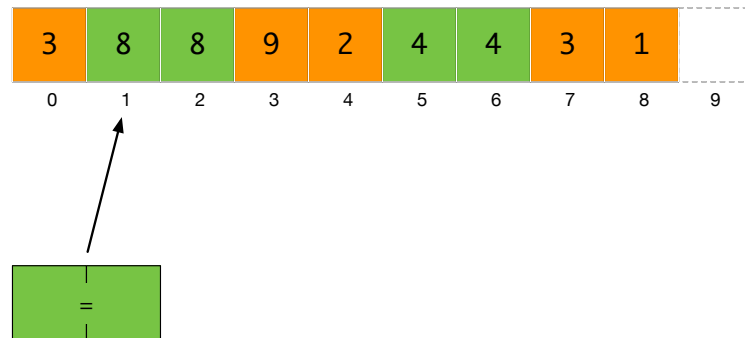
Figure 3.11.: A simple example for `adjacent_find`

3.4.1. The predicate `HasEqualNeighbors`

As in the case of other search algorithms, we first define a predicate `HasEqualNeighbors` (see Listing 3.12) that captures the essence of finding two adjacent indices at which the array holds equal values.

```
/*@
  predicate
    HasEqualNeighbors{L}(value_type* a, integer n) =
      \exists integer i; 0 <= i < n-1 && a[i] == a[i+1];
*/
```

Listing 3.12: The predicate `HasEqualNeighbors`

3.4.2. Formal specification of `adjacent_find`

We use the predicate `HasEqualNeighbors` to define the formal specification of `adjacent_find` (see Listing 3.13).

```
/*@
  requires valid: \valid_read(a + (0..n-1));

  assigns \nothing;

  ensures result:  0 <= \result <= n;

  behavior some:
    assumes  HasEqualNeighbors(a, n);
    ensures  result:    0 <= \result < n-1;
    ensures  adjacent:  a[\result] == a[\result+1];
    ensures  first:     !HasEqualNeighbors(a, \result);

  behavior none:
    assumes  !HasEqualNeighbors(a, n);
    ensures  result:    \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type
adjacent_find(const value_type* a, size_type n);
```

Listing 3.13: Formal specification of `adjacent_find`

3.4.3. Implementation of `adjacent_find`

The implementation of `adjacent_find`, including loop annotations is shown in Listing 3.14. At the beginning we check whether the array contains at least two elements. Otherwise, there is no point in looking for adjacent neighbors . . .

```
size_type
adjacent_find(const value_type* a, size_type n)
{
  if (n > 1u) {
    /*@
      loop invariant bound:  0 <= i < n;
      loop invariant none:   !HasEqualNeighbors(a, i+1);
      loop assigns i;
      loop variant n-i;
    */
    for (size_type i = 0u; i + 1u < n; ++i) {
      if (a[i] == a[i + 1u]) {
        return i;
      }
    }
  }

  return n;
}
```

Listing 3.14: Implementation of `adjacent_find`

Note the use of the predicate `HasEqualNeighbors` in the loop invariant to match the similar postcondition of behavior `some`.

3.5. The `equal` and `mismatch` algorithms

The `equal` [21, §28.5.11] and `mismatch` [21, §28.5.10] algorithms in the C++ Standard Library compare two generic sequences. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signatures read

```
bool        equal(const value_type* a, size_type n, const value_type* b);

size_type mismatch(const value_type* a, size_type n, const value_type* b);
```

The function `equal` returns **true** if and only if `a[i] == b[i]` holds for each `0 <= i < n`. Otherwise, `equal` returns **false**.

The `mismatch` algorithm is slightly more general than the negation of `equal`: it returns the smallest index where the two ranges `a` and `b` differ. If no such index exists, that is, if both ranges are equal, then `mismatch` returns the (common) length `n` of the two ranges.

3.5.1. The `EqualRanges` predicate

The fact that two arrays `a[0]..a[n-1]` and `b[0]..b[n-1]` are equal when compared element by element, is a property we might need again in other specifications, as it describes a very basic property.

The motto *don't repeat yourself* is not just good programming practice.[12] It is also true for concise and easy to understand specifications. We will therefore introduce specification elements that we can apply to the `equal` algorithm as well as to other specifications and implementations with the described property.

We start with introducing in Listing 3.15 several *overloaded* versions of the predicate `EqualRanges`.

```
/*@
  predicate
    EqualRanges{K,L}(value_type* a, integer n, value_type* b) =
      \forall integer i; 0 <= i < n ==> \at(a[i],K) == \at(b[i],L);

  predicate
    EqualRanges{K,L}(value_type* a, integer m, integer n, value_type* b) =
      \forall integer i; m <= i < n ==> \at(a[i],K) == \at(b[i],L);

  predicate
    EqualRanges{K,L}(value_type* a, integer m, integer n,
                     value_type* b, integer p) =
    EqualRanges{K,L}(a+m, n-m, b+p);

  predicate
    EqualRanges{K,L}(value_type* a, integer m, integer n, integer p) =
      EqualRanges{K,L}(a, m, n, a, p);
*/
```

Listing 3.15: Overloaded versions of predicate `EqualRanges`

The letters `K` and `L` in the definition of `EqualRanges` are so-called *labels*[13] that refer to program states in which the ranges `a[..]` and `b[..]` are evaluated. Frama-C defines several standard labels, e.g. `Old` and `Post`, a programmer can use to refer to the pre-state or post-state, respectively, of a function. For more details on labels we refer to the ACSL specification [16, §2.6.9].

---

[12]Compare `http://en.wikipedia.org/wiki/Don't_repeat_yourself`
[13]Labels are used in C to name the target of the *goto* jump statement.

3.5.2. Formal specification of `equal` and `mismatch`

Using predicate EqualRanges we can formulate the specification of equal in Listing 3.16, using the predefined label Here. When used in an ensures clause, the label Here refers to the post-state of a function. Note that the equivalence is needed in the ensures clause. Putting an equality instead is not legal in ACSL, because EqualRanges is a predicate, not a function.

```
/*@
  requires valid:  \valid_read(a + (0..n-1));
  requires valid:  \valid_read(b + (0..n-1));

  assigns \nothing;

  ensures result:  \result <==> EqualRanges{Here,Here}(a, n, b);
*/
bool
equal(const value_type* a, size_type n, const value_type* b);
```

Listing 3.16: Formal specification of equal

```
/*@
  requires valid:  \valid_read(a + (0..n-1));
  requires valid:  \valid_read(b + (0..n-1));

  assigns \nothing;

  ensures  result: 0 <= \result <= n;

  behavior all_equal:
    assumes EqualRanges{Here,Here}(a, n, b);
    ensures result: \result == n;

  behavior some_not_equal:
    assumes !EqualRanges{Here,Here}(a, n, b);
    ensures bound:   0 <= \result < n;
    ensures result:  a[\result] != b[\result];
    ensures first:   EqualRanges{Here,Here}(a, \result, b);

  complete behaviors;
  disjoint behaviors;
*/
size_type
mismatch(const value_type* a, size_type n, const value_type* b);
```

Listing 3.17: Formal specification of mismatch

The formal specification of mismatch in Listing 3.17 is more complex than that of equal because the return value of mismatch provides more information than just reporting whether the two arrays are equal. On the other, the specification is conceptually quite similar to that of find (Listing 3.5). While find returns the smallest index i where a[i] == val holds, mismatch finds the smallest index a[i] != b[i].

Note in particular the use of EqualRanges in the specification of mismatch. As in the specification of find the completeness and disjointness of mismatch's behaviors is quite obvious, because the assumes clauses of all_equal and some_not_equal are negations of each other.

3.5.3. Implementation of `equal` and `mismatch`

Listing 3.18 shows an implementation of the `equal` algorithm by a simple call of `mismatch`.

```
bool
equal(const value_type* a, size_type n, const value_type* b)
{
  return mismatch(a, n, b) == n;
}
```

Listing 3.18: Implementation of `equal` with `mismatch`

Listing 3.19 shows an implementation of `mismatch` that we have enriched with some loop annotations to support the deductive verification.

```
size_type
mismatch(const value_type* a, size_type n, const value_type* b)
{
  /*@
    loop invariant bound:  0 <= i <= n;
    loop invariant equal:  EqualRanges{Here,Here}(a, i, b);
    loop assigns i;
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; i++) {
    if (a[i] != b[i]) {
      return i;
    }
  }

  return n;
}
```

Listing 3.19: Implementation of `mismatch`

We use the predicate `EqualRanges` in order to express that all indices `k` that are less than the current index `i` satisfy the condition `a[k] == b[k]`. This is necessary to prove that `mismatch` indeed returns the smallest index where the two ranges differ.

3.6. The `search` algorithm

The `search` algorithm in the C++ Standard Library [21, §28.5.13] finds a subsequence that is identical to a given sequence when compared element-by-element. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
size_type search(const value_type* a, size_type n,
                 const value_type* b, size_type p)
```

The function `search` returns the first index `s` of the array `a` where the condition `a[s+k] == b[k]` holds for each index `k` with `0 <= k < p` (see Figure 3.20). If no such index exists, then `search` returns the length `n` of the array `a`.

Figure 3.20.: Searching the first occurrence of `b[0..p-1]` in `a[0..n-1]`

3.6.1. The predicate `HasSubRange`

Our specification of `search` starts with introducing the predicate `HasSubRange` in Listing 3.21. This predicate formalizes, using the predicate `EqualRanges` defined in Listing 3.15, that the sequence `a` contains a subsequence which equal the sequence `b`. Of course, in order to contain a subsequence of length `p`, `a` must be at least that large; this is expressed by lemma `HasSubRangeSizes`.

```
/*@
  predicate
    HasSubRange{L}(value_type* a, integer m, integer n, value_type* b, integer p) =
      \exists integer k; (m <= k <= n-p) && EqualRanges{L,L}(a+k, p, b);

  predicate
    HasSubRange{L}(value_type* a, integer n, value_type* b, integer p) =
      HasSubRange{L}(a, 0, n, b, p);

  lemma
    HasSubRangeSizes:
      \forall value_type *a, *b, integer m, n, p;
        HasSubRange(a, m, n, b, p) ==> p <= n-m;
*/
```

Listing 3.21: The predicate `HasSubRange`

3.6.2. Formal specification of `search`

The ACSL specification of `search` is shown in Listing 3.22. Conceptually, the specification of `search` is very similar to that of `find` (Section 3.1). We therefore use again two behaviors to capture the essential aspects of `search`.

- The behavior `has_match` applies if the sequence `a` contains a subsequence identical to `b`. We express this condition with `assumes` using the predicate `HasSubRange`.

  The ensures clause `bound` of behavior `has_match` indicates that the returned index value must be in the range `[0..n-p]`. The clause `result` expresses that `search` returns an index where a copy of `b` can be found in `a`. Clause `first` indicates that the least index with that property is returned, i.e. that `b` can't be found in `a[0..\result+p-2]`.

```
/*@
  requires \valid_read(a + (0..n-1));
  requires \valid_read(b + (0..p-1));

  assigns  \nothing;

  ensures  result: 0 <= \result <= n;

  behavior has_match:
    assumes HasSubRange(a, n, b, p);
    ensures bound:  0 <= \result <= n-p;
    ensures result: EqualRanges{Here,Here}(a+\result, p, b);
    ensures first:  !HasSubRange(a, \result+p-1, b, p);

  behavior no_match:
    assumes !HasSubRange(a, n, b, p);
    ensures result: \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type
search(const value_type* a, size_type n,
       const value_type* b, size_type p);
```

Listing 3.22: Formal specification of `search`

- The behavior `no_match` covers the case that there is no subsequence a that equals b. In this case, `search` must return the length n of the range a. If the ranges a or b are empty then the return value will be 0.

The formula in the assumes clause of the behavior `has_match` is the negation of the assumes clause of the behavior `no_match`. Therefore, we can express that these two behaviors are *complete* and *disjoint*.

3.6.3. Implementation of `search`

Our implementation of `search` is shown in Listing 3.23. It follows the C++ Standard Library implementation in being easy to understand, but needing an order of magnitude of $n*p$ operations. In contrast, the sophisticated algorithm from [22] needs only $n+p$ operations.[14]

The loop invariant `not_found` is needed for the proof of the postconditions of the behavior `has_match` (see Listing 3.22). It expresses that the subsequence b has not been found up to the current iteration step. Neither p == 0 nor n == 0 need to be handled separately, not even for efficiency reasons: in the former case, `equal(a+i, p, b)` will succeed in the first iteration, while in the latter, p > n will apply.

---

[14] The efficiency question has been also discussed by the C++ standardization committee, see http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n3905.html

46

```
size_type
search(const value_type* a, size_type n,
       const value_type* b, size_type p)
{
  if (p <= n) {
    /*@
      loop invariant bound:     i <= n-p+1;
      loop invariant not_found: !HasSubRange(a, p+i-1, b, p);
      loop assigns i;
      loop variant n-i;
    */
    for (size_type i = 0u; i <= n - p; ++i) {
      if (equal(a + i, p, b)) {
        //@ assert has_match: HasSubRange(a, n, b, p);
        return i;
      }
    }
  }

  //@ assert no_match: !HasSubRange(a, n, b, p);
  return n;
}
```

Listing 3.23: Implementation of `search`

3.7. The `search_n` algorithm

The `search_n` algorithm in the C++ Standard Library [21, §28.5.13] finds the first place where a given value starts to occur a given number of times in a given sequence. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
size_type
search_n(const value_type* a, size_type n, size_type p, value_type v)
```

Note the similarity to the signature of `search` (Section 3.6). The only difference is that `v` now is a single value rather than an array.

Figure 3.24.: Searching the first occurrence a given constant sequence in `a[0..n-1]`

The function `search_n` returns the first index `s` of the array `a` where the condition `a[s+k] == v` holds for each index `k` with `0 <= k < p` (see Figure 3.24). If no such index exists, then `search_n` returns the length `n` of the array `a`.

3.7.1. The predicates `ConstantRange` and `HasConstantSubRange`

Our specification of search_n starts with introducing the predicate ConstantRange in Listing 3.25, which expresses that each member a[m..n-1] equals v.

```
/*@
  predicate
    ConstantRange(value_type* a, integer m, integer n, value_type v) =
      \forall integer i; m <= i < n ==> a[i] == v;

  predicate
    ConstantRange(value_type* a, integer m, integer n) =
      ConstantRange(a, m, n, a[m]);

  predicate
    ConstantRange(value_type* a, integer n, value_type v) =
      ConstantRange(a, 0, n, v);
*/
```

Listing 3.25: The predicate ConstantRange

There are two additional overloaded versions of ConstantRange. The first version uses the value a[m] as v. We will use this particular version later on during the specification of unique_copy (Listing 7.3). The second version is just a shortcut when the first index m equals 0.

Based on that, the predicate HasConstantSubRange in Listing 3.26 formalizes that the sequence a of length n contains a subsequence of p times the value v. Similar to HasSubRange, in order to contain p repetitions, a must be at least that large; this is what lemma HasConstantSubRangeSizes says.

```
/*@
  predicate
    HasConstantSubRange{L}(value_type* a, integer m, integer n, value_type v, integer
        p) =
      \exists integer k; m <= k <= n-p && ConstantRange(a, k, k+p, v);

  predicate
    HasConstantSubRange{L}(value_type* a, integer n, value_type v, integer p) =
      HasConstantSubRange(a, 0, n, v, p);

  lemma
    HasConstantSubRangeSizes:
      \forall value_type *a, v, integer n, p;
        HasConstantSubRange(a, n, v, p) ==> p <= n;
*/
```

Listing 3.26: The predicate HasConstantSubRange

3.7.2. Formal specification of `search_n`

The ACSL specification of `search_n` is shown in Listing 3.27. Like for `search`, the specification of `search_n` is very similar to that of `find`.

```
/*@
  requires valid: \valid_read(a + (0..n-1));

  assigns  \nothing;

  ensures  result:  0 <= \result <= n;

  behavior has_match:
    assumes HasConstantSubRange(a, n, v, p);
    assigns \nothing;
    ensures result:  0 <= \result <= n-p;
    ensures match:    ConstantRange(a, \result, \result+p, v);
    ensures first:    !HasConstantSubRange(a, \result+p-1, v, p);

  behavior no_match:
    assumes !HasConstantSubRange(a, n, v, p);
    assigns \nothing;
    ensures result:  \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type
search_n(const value_type* a, size_type n, value_type v, size_type p);
```

Listing 3.27: Formal specification of `search_n`

We again use two behaviors to capture the essential aspects of `search_n`.

- The behavior `has_match` applies if the sequence `a` contains an n-fold repetition of b. We express this condition with `assumes` by using the predicate `HasConstantSubRange`.

  The `result` ensures clause of behavior `has_match` indicates that the return value must be in the range `[0..n-p]`. The `match` ensures clause expresses that the return value of `search_n` actually points to an index where b can be found p or more times in a. The `first` ensures clause expresses that the minimal index with this property is returned.

- The behavior `no_match` covers the case that there is no matching subsequence in sequence `a`. In this case, `search_n` must return the length n of the range a.

The formula in the assumes clause of the behavior `has_match` is the negation of the assumes clause of the behavior `no_match`. Therefore, we can express that these two behaviors are *complete* and *disjoint*.

3.7.3. Implementation of `search_n`

Although the specification of `search_n` strongly resembles that of `search`, their implementations in the C++ Standard Library are significantly different. The former has a time complexity of $O(n)$, whereas the latter employs an easy, but a non-optimal algorithm needing $O(n \cdot p)$ time, cf. Section 3.6.3.

```
size_type
search_n(const value_type* a, size_type n, value_type v, size_type p)
{
  if (0u < p) {
    if (p <= n) {
      size_type start = 0u;

      /*@
        loop invariant match:     ConstantRange(a, start, i, v);
        loop invariant start:     0 < start ==> a[start-1] != v;
        loop invariant bound:     start <= i + 1 <= start + p;
        loop invariant not_found: !HasConstantSubRange(a, i, v, p);
        loop assigns i, start;
        loop variant n - i;
      */
      for (size_type i = 0u; i < n; ++i) {
        if (a[i] != v) {
          start = i + 1u;
          //@ assert not_found: !HasConstantSubRange(a, i+1, v, p);
        }
        else {
          //@ assert match: a[i] == v;
          //@ assert match: ConstantRange(a, start, i+1, v);
          if (p == i + 1u - start) {
            //@ assert bound: start + p == i + 1;
            //@ assert match: ConstantRange(a, start, start+p, v);
            //@ assert match: \exists integer k; 0 <= k <= n-p && ConstantRange(a, k
                , k+p, v);
            //@ assert match: HasConstantSubRange(a, n, v, p);
            return start;
          }
          else {
            //@ assert bound: i + 1 < start + p;
            continue;
          }
        }

        //@ assert not_found: !HasConstantSubRange(a, i+1, v, p);
      }

      //@ assert not_found: !HasConstantSubRange(a, n, v, p);
      return n;
    }
    else {
      //@ assert not_found: n < p;
      //@ assert not_found: !HasConstantSubRange(a, n, v, p);
      return n;
    }
  }
  else {
    //@ assert bound: p == 0;
    //@ assert match: ConstantRange(a, 0, 0, v);
    //@ assert match: HasConstantSubRange(a, n, v, 0);
    return 0u;
  }
}
```

Listing 3.28: Implementation of search_n

Our implementation maintains in the variable `start` the beginning of the most recent consecutive range of values `b`. This property is expressed by three loop invariants:

- `constant` states that `b` is the only value that can occur in `a[start..i-1]`;

- `start` states that it cannot be left-extended, i.e. `a[start-1]` (if defined) is different from `b`;

- `bound` states that the sequence's bounds are actually given left to right.

The loop invariant `not_found` states that we didn't find an `p`-fold repetition of `b` up to now; if we find one, we terminate the loop, returning `start`.

We handle the boundary cases `n < p` and `p == 0` in explicit else branches. We found this easier when trying to ensure a verification by automatic provers.

3.8. The `find_end` algorithm

The `find_end` algorithm in the C++ Standard Library [21, §28.5.6] searches for the last subsequence that is identical to a given sequence when compared element-by-element. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
size_type find_end(const value_type* a, size_type n,
                   const value_type* b, size_type p)
```

The function `find_end` returns the greatest index `s` of the array `a` where the condition `a[s+k] == b[k]` holds for each index `k` with `0 <= k < p` (see Figure 3.29). If no such index exists, then `find_end` returns the length `n` of the array `a`. One has to remark the special case `p == 0`. In this case the last position of the empty string is found (the length `n`) and returned.

Figure 3.29.: Finding the last occurrence `b[0..p-1]` in `a[0..n-1]`

3.8.1. Formal specification of `find_end`

The ACSL specification of `find_end` is shown in Listing 3.30. Conceptually, the specification of the function `find_end` is very similar to that of `find` in Section 3.2. We therefore use again behaviors to capture the essential aspects of `find_end`. It is quite clear that these behaviors are *complete* and *disjoint*.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid_read(b + (0..p-1));

  assigns \nothing;

  ensures result: 0 <= \result <= n;

  behavior has_match:
    assumes HasSubRange(a, n, b, p);
    ensures bound:  0 <= \result <= n-p;
    ensures result: EqualRanges{Here,Here}(a + \result, p, b);
    ensures last:   !HasSubRange(a, \result + 1, n, b, p);

  behavior no_match:
    assumes !HasSubRange(a, n, b, p);
    ensures result: \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type
find_end(const value_type* a, size_type n,
         const value_type* b, size_type p);
```

Listing 3.30: Formal specification of `find_end`

The behavior `has_match` applies if the sequence `a` contains a subsequence identical to `b`. We express this condition with `assumes` using the predicate `HasSubRange`. The `ensures` clause `bound` indicates that the return value must be in the range `0..n-p`. The clause `result` of behavior `has_match` expresses that `find_end` returns an index where `b` can be found in `a`. Finally, the clause `last` indicates that the sequence `a` does not contain `b` beginning at a position larger than `\result`.

The behavior `no_match` covers the case that there is no subsequence of `a` that equals `b`. In this case, `find_end` must return the length `n` of the range `a`.

3.8.2. Implementation of `find_end`

Our implementation of `find_end` is shown in Listing 3.31. Similar to our `search` implementation (Section 3.6), it follows the C++ Standard Library implementation in being easy to understand, but needing an order of magnitude of `n*p` rather than only `n+p` operations.

```
size_type
find_end(const value_type* a, size_type n,
         const value_type* b, size_type p)
{
  size_type r = n;

  if ((0u < p) && (p <= n)) {
    /*@
      loop invariant bound   :  r <= n - p || r == n;
      loop invariant not_found: r == n ==> !HasSubRange(a, p+i-1, b, p);
      loop invariant found:     r < n  ==> EqualRanges{Here,Here}(a+r, p, b);
      loop invariant last:      r < n  ==> !HasSubRange(a, r+1, i+p-1, b, p);
      loop assigns i, r;
      loop variant n - i;
    */
    for (size_type i = 0u; i <= n - p; ++i) {
      if (equal(a + i, p, b)) {
        r = i;
      }
    }
  }

  return r;
}
```

Listing 3.31: Implementation of `find_end`

We maintain in the variable `r` the prospective value to be returned, according to the current knowledge. Initially, it is set to `n`, meaning "no occurrence of `b` found yet". Whenever an occurrence is found, `r` is updated to its starting position.

Invariant `bound` states that `r` either still has the value `n` or has a value up to `n-p`. For the former case, invariant `not_found` indicates that no occurrence of `b` has been found. For the latter case, the loop invariant `found` indicates that an occurrence `b[0..p-1]` at `r` has indeed been found. The invariant `last`, on the other hand states that none was found *after* the index `r`.

3.9. The `count` algorithm

The `count` algorithm in the C++ Standard Library [21, §28.5.9] counts the frequency of occurrences for a particular element in a sequence. For our purposes we have modified the generic implementation to that of arrays of type `value_type`. The signature now reads:

```
size_type count(const value_type* a, size_type n, value_type val);
```

Informally, the function returns the number of occurrences of `val` in the array `a`.

3.9.1. The logic function `Count`

When trying to specify `count` we are faced with the situation that ACSL does not provide a definition of counting a value in an array.[15] We therefore start with an axiomatic definition of *logic function* Count that captures the basic intuitive features of counting on an array section. The expression `CountSection(a, m,n,v)` returns the number of occurrences of `v` in `a[m],...,a[n-1]`.

The specification of `count` will then be fairly short because it employs our *logic function* Count whose (considerably) longer definition is given in Listing 3.32.[16]

```
/*@
  axiomatic CountAxiomatic
  {
    logic integer
    Count{L}(value_type* a, integer m, integer n, value_type v) =
      n <= m ? 0 : Count(a, m, n-1, v) + (a[n-1] == v ? 1 : 0);

    lemma
      CountSectionEmpty{L}:
        \forall value_type *a, v, integer m, n;
          n <= m  ==>  Count(a, m, n, v) == 0;

    lemma
      CountSectionHit{L}:
        \forall value_type *a, v, integer n, m;
          m < n  ==> a[n-1] == v  ==>  Count(a, m, n, v) == Count(a, m, n-1, v) + 1;

    lemma
      CountSectionMiss{L}:
        \forall value_type *a, v, integer n, m;
          m < n  ==> a[n-1] != v  ==>  Count(a, m, n, v) == Count(a, m, n-1, v);

    lemma
      CountSectionRead{K,L}:
        \forall value_type *a, v, integer m, n;
          Unchanged{K,L}(a, m, n)  ==>
            Count{K}(a, m, n, v) == Count{L}(a, m, n, v);
  }
*/
```

Listing 3.32: The logic function Count

---

[15]This statement is not quite true because the ACSL documentation lists `numof` as one of several *higher order logic constructions* [16, §2.6.7]. However, these *extended quantifiers* are mentioned only as experimental features.

[16]This definition of Count is a generalization of the *logic function* `nb_occ` of the ACSL specification [16, p. 55].

- The ACSL keyword `axiomatic` is used to structure the specification and gather the logic function `Count` and related lemmas. Note that the interval bounds `m` and `n` and the return value for `Count` are of type `integer`.

- The logic functions `Count` is defined explicitly with the recursion. It consist of two checks: for the range emptiness and for the value of the "current" element in the array. The recursion goes down on the range length.

- Lemmas `CountSectionEmpty`, `CountSectionHit` and `CountSectionMiss` are not required. They set the implicit definition for `Count`. But as we already has the explicit definition these lemmas simplify replicate the properties of the function under different conditions. Thus these lemmas are proved automatically by solvers. We keep them as an comparative illustration of an implicit definition.

- Lemma `CountSectionEmpty` covers the case of an empty range.

- Lemmas `CountSectionHit` and `CountSectionMiss` reduce counting of a range of length n to a range of length $n - 1$.

- The `Count` depends on the `a[0..n-1]` memory location set.

  Lemma `CountSectionRead` makes this claim explicit by ensuring that `Count` produces the same result if the values `a[0..n-1]` do not change between two program states indicated by the labels `L1` and `L2`. We use predicate `Unchanged` (Listing 6.1 in Section 6.1) to express the premise of Lemma `CountSectionRead`. Lemma `CountSectionRead` is helpful if one has to verify *mutating* algorithms that rely on `Count`, e.g., `remove_copy` in Section 6.14. It is an inductive consequence of the definition and the lemmas `CountSectionEmpty`, `CountSectionHit`, `CountSectionMiss`, but we don't prove it here.

Listing 3.33 shows additional properties of `Count`.

```
/*@
  lemma
    CountSectionOne:
      \forall value_type *a, v, integer m, n;
        m <= n  ==>
          Count(a, m, n+1, v) == Count(a, m, n, v) + Count(a, n, n+1, v);


  lemma
    CountSectionUnion:
      \forall value_type *a, v, integer k, m, n;
        0 <= k <= m <= n  ==>
          Count(a, k, n, v) == Count(a, k, m, v) + Count(a, m, n, v);
*/
```

Listing 3.33: The lemmas `CountSectionOne` and `CountSectionUnion`

3.9.2. Counting on a whole array

We also provide in Listing 3.34 an overloaded version of the logic function `Count` that only takes the starting address and the size of an array. This is just a convenience function for the use in specifications that do not need to consider array sections. Note how the accompanying lemmas in Listing 3.34 correspond to the lemmas in Listing 3.32.

```
/*@
  logic integer
    Count{L}(value_type* a, integer n, value_type v) = Count{L}(a, 0, n, v);

  lemma
    CountEmpty:
      \forall value_type *a, v, integer n;
        n <= 0  ==>  Count(a, n, v) == 0;

  lemma
    CountHit:
      \forall value_type *a, v, integer n;
        0 < n  ==>  a[n-1] == v  ==>  Count(a, n, v) == Count(a, n-1, v) + 1;

  lemma
    CountMiss:
      \forall value_type *a, v, integer n;
        0 < n  ==>  a[n-1] != v  ==>  Count(a, n, v) == Count(a, n-1, v);

  lemma
    CountRead{L1,L2}:
      \forall value_type *a, v, integer n;
        Unchanged{L1,L2}(a, n)  ==>  Count{L1}(a, n, v) == Count{L2}(a, n, v);
*/
```

Listing 3.34: The logic function `Count`

The lemmas for `Count` in Listing 3.35 are just reformulated versions of those in Listing 3.33.

```
/*@
  lemma
    CountOne:
      \forall value_type *a, v, integer n;
        0 <= n ==>
          Count(a, n+1, v) == Count(a, n, v) + Count(a, n, n+1, v);

  lemma
    CountUnion:
      \forall value_type *a, v, integer m, n;
        0 <= m <= n  ==>
        Count(a, n, v) == Count(a, 0, m, v) + Count(a, m, n, v);
*/
```

Listing 3.35: Some lemmas for `Count`

3.9.3. Formal specification of `count`

Listing 3.36 shows how we use the logic function `Count` from Listing 3.34 to specify `count` in ACSL. Note that our specification also states that the result of `count` is non-negative and less than or equal the size of the array.

```
/*@
  requires valid: \valid_read(a + (0..n-1));

  assigns \nothing;

  ensures bound:  0 <= \result <= n;
  ensures count: \result == Count(a, n, val);
*/
size_type
count(const value_type* a, size_type n, value_type val);
```

Listing 3.36: Formal specification of `count`

3.9.4. Implementation of `count`

Listing 3.37 shows a possible implementation of `count`. Note that we refer to the logic function `Count` in one of the loop invariants.

```
size_type
count(const value_type* a, size_type n, value_type val)
{
  size_type counted = 0u;

  /*@
    loop invariant bound: 0 <= i <= n;
    loop invariant bound: 0 <= counted <= i;
    loop invariant count: counted == Count(a, i, val);
    loop assigns i, counted;
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    if (a[i] == val) {
      counted++;
    }
  }

  return counted;
}
```

Listing 3.37: Implementation of `count`

3.10. The `count2` algorithm

In this section, we specify the `count` algorithm in differently compared to Section 3.9. The example demonstrates how inductive predicates could be defined and used in ACSL specifications.

Sections 3.9 and 3.10 share the same `count` implementation.

3.10.1. The inductive predicate `Count`

To specify the counting behavior, we define the inductive predicate `Count`. The definition consists of three cases. The "Nil" case states the predicate holds then the size "n" of the array "a" is not a positive value and the number "sum" of the counted elements is equal to zero. The "Hit" and "Miss" cases define given the `Count` holds for the array "a" of size "n-1" the conditions the `Count` holds for the size "n". The only difference between these two cases is that in the first one the array has the value "v" in position "n-1" and the number "sum" of the counted elements should be increased by one the predicate `Count` to hold for size "n".

One might notice that the cases are very similar to the lemmas CountEmpty, CountHit, CountMiss from Listing 3.32, except we use the additional argument "sum" to save the number of counted elements since this is the predicate.

In the example we intentionally use the scheme "n-1" => "n" instead of "n" => "n+1". In this particular case, it allows solvers to match loop indices with premises without additional hints to prove loop invariants.

It should be noted that this definition is enough to prove the code corresponds to its contract. No additional lemmas required.

```
/*@
  inductive Count2{L}(value_type *a, integer n, value_type v, integer sum) {
    case Nil{L}:
      \forall value_type *a, v, integer n;
        n <= 0 ==>
          Count2{L}(a, n, v, 0);

    case Hit{L}:
      \forall value_type *a, v, integer n, sum;
        0 < n && a[n-1] == v && Count2{L}(a, n-1, v, sum) ==>
          Count2{L}(a, n, v, sum + 1);

    case Miss{L}:
      \forall value_type *a, v, integer n, sum;
        0 < n && a[n-1] != v && Count2{L}(a, n-1, v, sum) ==>
          Count2{L}(a, n, v, sum);
  }
*/
```

Listing 3.38: The inductive predicate `Count`

3.10.2. Additional lemmas for the inductive predicate

These lemmas are not required to prove the `count` function, but we provide them to complete the illustrative example of how inductive predicates could be utilized in the specifications.

Listing 3.39 complements Listing 3.32 and demonstrates how exactly the same lemmas could be rewritten for the inductive predicate.

```
/*@
  axiomatic Count2ImplicitAxiomatic {
    lemma
      Count2Empty{L}:
        \forall value_type *a, v, integer n;
          n <= 0 ==> Count2(a, n, v, 0);

    lemma
      Count2Miss{L}:
        \forall value_type *a, v, integer n, sum;
          0 < n ==> a[n-1] != v ==>
            (Count2(a, n-1, v, sum) ==> Count2(a, n, v, sum));

    lemma
      Count2Hit{L}:
        \forall value_type *a, v, integer n, sum;
          0 < n ==> a[n-1] == v ==>
            (Count2(a, n-1, v, sum) ==> Count2(a, n, v, sum+1));

    lemma
      Count2Read{K,L}:
        \forall value_type *a, v, integer n, sum;
          Unchanged{K,L}(a, n) ==>
            (Count2{K}(a, n, v, sum) <==> Count2{L}(a, n, v, sum));

  }
*/
```

Listing 3.39: Implicit definition of `Count`

The inductive definition is the "complete" definition. This means that a predicate does not hold in every case outside the definition. We state this property explicitly for `Count` with Lemma 3.40. Frama-C does not add such axiom in the context, and thus the Lemma is not proved automatically. The reason for not adding such an axiom that it "could confuse first-order theorem provers".[17] The Coq proof could be as simple as the application of the inversion to the definition.

```
/*@
  lemma Count2Inverse:
    \forall value_type *a, v, integer n, sum;
      Count2(a, n, v, sum) ==>
        (n <= 0 && sum == 0) ||
        (0 < n && a[n-1] != v && Count2(a, n-1, v, sum)) ||
        (0 < n && a[n-1] == v && Count2(a, n-1, v, sum-1));
*/
```

Listing 3.40: Inverse lemma for `Count`

The lower bound for the number of the counted elements is zero. Lemma 3.41 could be proved in Coq in

---

[17]https://stackoverflow.com/a/32457870

one line by induction on the `Count` definition.

```
/*@
  lemma
    Count2NonNegativeSum{L}:
      \forall value_type *a, v, integer n, sum;
        Count2(a, n, v, sum) ==> sum >= 0;
*/
```

Listing 3.41: The result of the `Count` is always greater that zero

The relation between the definitions through the inductive predicate 3.38 and the logic function 3.32 is expressed in Lemma 3.42. The Lemma is proved with Coq.

```
/*@
  lemma
    CountCount2Relation{L}:
      \forall value_type *a, v, integer n;
        Count2(a, n, v, Count(a, n, v));
*/
```

Listing 3.42: The relation between definitions

3.10.3. Specification of `count`

Listing 3.43 shows how we use the inductive predicate `Count` from Listing 3.38 to specify `count` in ACSL.

```
/*@
  requires valid: \valid_read(a + (0..n-1));

  assigns \nothing;

  ensures bound:  0 <= \result <= n;
  ensures count:  Count2(a, n, val, \result);
*/
size_type
count2(const value_type* a, size_type n, value_type val);
```

Listing 3.43: The `count` contract reconsidered

3.10.4. Implementation of `count`

The only difference here with Listing 3.37 is that counted is used as an argument for the `Count` predicate.

```
size_type
count2(const value_type* a, size_type n, value_type val)
{
  size_type counted = 0u;

  /*@
    loop invariant bound: 0 <= i <= n;
    loop invariant bound: 0 <= counted <= i;
    loop invariant count: Count2(a, i, val, counted);
    loop assigns i, counted;
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    if (a[i] == val) {
      counted++;
    }
  }

  return counted;
}
```

Listing 3.44: The `count` implementation reconsidered

4. Maximum and minimum algorithms

In this chapter we discuss the formal specification of algorithms in the C++ Standard Library [21, §28.7.8] that compute the maximum or minimum values of their arguments. As the algorithms in Chapter 3, they also do not modify any memory locations outside their scope. The most important new feature of the algorithms in this chapter is that they compare values using binary operators such as <.

We consider in this chapter the following algorithms.

- `clamp`, which is discussed in Section 4.2, is a very simple algorithms that "clamps" (or "clips") a value between a pair of boundary values.

- `max_element` returns an index to a maximum element in a range. Similar to `find` it also returns the smallest of all possible indices. It is discussed in Section 4.3, on Page 66. In Section 4.4, on Page 68, we will introduce an alternative specification `max_element2` which relies on user-defined predicates.

- `max_seq` (Section 4.5, on Page 70) is very similar to `max_element` and will serve as an example of *modular verification*. It returns the maximum value itself rather than an index to it.

- `min_element` can be used to find the smallest element in an array (Section 4.7).

- `minmax_element` is used to find simultaneously the smallest and largest element in a given range (Section 4.8). This algorithms relies on the auxiliary function `make_pair` (Section 4.6).

First, however, we discuss in Section 4.1 general properties that must be satisfied by the relational operators.

4.1. A note on relational operators

Note that in order to compare values, algorithms in the C++ Standard Library [21, §28.7.8] usually rely solely on the *less than* operator < or special function objects. To be precise, the operator < must be a *partial order*,[18] which means that the following rules must hold.

$$
\begin{array}{lll}
\text{irreflexivity} & \forall x & : \neg(x < x) \\
\text{asymmetry} & \forall x, y & : x < y & \implies \neg(y < x) \\
\text{transitivity} & \forall x, y, z : x < y \wedge y < z & \implies x < z
\end{array}
$$

If you wish to check that the operator < of our `value_type`[19] satisfies these properties you can formulate lemmas in ACSL and verify them with Frama-C (see Listing 4.1).

---

[18]See `http://en.wikipedia.org/wiki/Partially_ordered_set`
[19]See Section 1.3

```
/*@
  lemma
    LessIrreflexivity:
      \forall value_type a; !(a < a);

  lemma
    LessAntisymmetry:
      \forall value_type a, b; (a < b) ==> !(b < a);

  lemma
    LessTransitivity:
      \forall value_type a, b, c; (a < b) && (b < c) ==> (a < c);
*/
```

Listing 4.1: Requirements for a partial order on `value_type`

It is of course possible to specify and implement the algorithms of this chapter by only using operator <. For example, `a <= b` can be written as `a < b || a == b`, or, for our particular ordering on `value_type`, as `!(b < a)`.

Listing 4.2 formulates conditions on the semantics of the derived operator >, <=, and >=.

```
/*@
  lemma
    Greater:
      \forall value_type a, b; (a > b) <==> (b < a);

  lemma
    LessOrEqual:
      \forall value_type a, b; (a <= b) <==> !(b < a);

  lemma
    GreaterOrEqual:
      \forall value_type a, b; (a >= b) <==> !(a < b);
*/
```

Listing 4.2: Semantics of derived comparison operators

We also introduce in this chapter the predicates

- `UpperBound` in Listing 4.7
- `StrictUpperBound` in Listing 4.8
- `LowerBound` in Listing 4.16
- `StrictLowerBound` in Listing 4.17

These overloaded predicates concisely express the comparison of the elements in an array (segment) with a given value. We will heavily rely on these predicates both in this chapter and in Chapter 5.

4.2. The `clamp` algorithm

The `clamp` algorithm in the C++ Standard Library [21, §28.7.9] "clamps" a value between a pair of boundary values. The signature of our version of `clamp` reads:

```
value_type
clamp(value_type v, value_type lower, value_type upper);
```

The function `clamp` returns `v` if the value is greater than `lower` and smaller than `upper`. Otherwise, if `v` is smaller than `lower`, the return value is `lower` instead. Similar to that if `v` is greater than `upper`, `upper` will be the return value of `clamp`.

4.2.1. Formal specification of `clamp`

The ACSL specification of `clamp` is shown in Listing 4.3. Note that we state in the precondition `proper` that `lower` must less or equal to `upper`.

```
/*@
  requires bound: lower < upper;

  assigns \nothing;

  ensures bound: lower <= \result <= upper;

  behavior lower_bound:
    assumes v < lower;
    assigns \nothing;
    ensures result: \result == lower;

  behavior between:
    assumes lower <= v <= upper;
    assigns \nothing;
    ensures result: \result == v;

  behavior upper_bound:
    assumes upper < v;
    assigns \nothing;
    ensures result: \result == upper;

  complete behaviors lower_bound, upper_bound, between;
  disjoint behaviors lower_bound, upper_bound, between;
*/
value_type
clamp(value_type v, value_type lower, value_type upper);
```

Listing 4.3: Formal specification of `clamp`

4.2.2. Implementation of `clamp`

The Listing 4.4 shows an implementation of `clamp`.

```
value_type
clamp(value_type v, value_type lower, value_type upper)
{
  return (v < lower) ? lower : (upper < v) ? upper : v;
}
```

Listing 4.4: Implementation of `clamp`

4.3. The `max_element` algorithm

The `max_element` algorithm in the C++ Standard Library [21, §28.7.8] searches the maximum of a general sequence. The signature of our version of `max_element` reads:

```
size_type max_element(const value_type* a, size_type n);
```

The function finds the largest element in the range `a[0..n-1]`. More precisely, it returns the unique valid index `i` such that:

1. for each index `k` with `0 <= k < n` the condition `a[k] <= a[i]` holds and

2. for each index `k` with `0 <= k < i` the condition `a[k] < a[i]` holds.

The return value of `max_element` is `n` if and only if there is no maximum, which can only occur if `n == 0`.

4.3.1. Formal specification of `max_element`

A formal specification of `max_element` in ACSL is shown in Listing 4.5. Note that we have subdivided the specification of `max_element` into the two behaviors `empty` and `not_empty`. The behavior `empty` contains the specification for the case that the range contains no elements. The behavior `not_empty` applies if the range has a positive length.

The ensures clause `max` of behavior `not_empty` indicates that the returned valid index `k` refers to a maximum value of the array. The postcondition `first` expresses that `k` is indeed the *first* occurrence of a maximum value in the array.

4.3.2. Implementation of `max_element`

Listing 4.6 shows an implementation of `max_element`. In our description, we concentrate on the *loop annotations*.

Loop invariant `max` is needed to prove postcondition `result` of behavior `not_empty` in Listing 4.5. Using loop invariant `upper` we prove postcondition `upper` of behavior `not_empty` in Listing 4.5. Finally, postcondition `first` of this behavior can be proved with loop invariant `first`.

66

```
/*@
  requires valid:   \valid_read(a + (0..n-1));
  assigns  \nothing;
  ensures  result:  0 <= \result <= n;

  behavior empty:
    assumes n == 0;
    ensures result:  \result == 0;

  behavior not_empty:
    assumes 0 < n;
    ensures result:  0 <= \result < n;
    ensures upper:   \forall integer i; 0 <= i < n      ==> a[i] <= a[\result];
    ensures first:   \forall integer i; 0 <= i < \result ==> a[i] <  a[\result];

  complete behaviors;
  disjoint behaviors;
*/
size_type
max_element(const value_type* a, size_type n);
```

Listing 4.5: Formal specification of `max_element`

```
size_type
max_element(const value_type* a, size_type n)
{
  if (0u < n) {
    size_type max = 0u;

    /*@
      loop invariant bound:  0 <= i <= n;
      loop invariant max:    0 <= max <  n;
      loop invariant upper:  \forall integer k; 0 <= k < i   ==> a[k] <= a[max];
      loop invariant first:  \forall integer k; 0 <= k < max ==> a[k] <  a[max];
      loop assigns max, i;
      loop variant n-i;
    */
    for (size_type i = 1u; i < n; i++) {
      if (a[max] < a[i]) {
        max = i;
      }
    }

    return max;
  }

  return n;
}
```

Listing 4.6: Implementation of `max_element`

4.4. The `max_element` algorithm with predicates

In this section we present another specification of the max_element algorithm. The main difference is that we employ several user-defined predicates. First, we define in Listing 4.7 the overloaded predicate UpperBound which basically expresses that a given value is greater or equal than all elements of a given array (section).

```
/*@
  predicate
    UpperBound{L}(value_type* a, integer m, integer n, value_type v) =
      \forall integer i; m <= i < n  ==>  a[i] <= v;

  predicate
    UpperBound{L}(value_type* a, integer n, value_type v) =
      UpperBound{L}(a, 0, n, v);
*/
```

Listing 4.7: Definition of the UpperBound predicate

Closely related to the predicate UpperBound is the overloaded predicate StrictUpperBound from Listing 4.8.

```
/*@
  predicate
    StrictUpperBound{L}(value_type* a, integer m, integer n, value_type v) =
      \forall integer i; m <= i < n  ==>  a[i] < v;

  predicate
    StrictUpperBound{L}(value_type* a, integer n, value_type v) =
      StrictUpperBound{L}(a, 0, n, v);
*/
```

Listing 4.8: Definition of the StrictUpperBound predicate

We then define in Listing 4.9 the predicate MaxElement by stating that the element at a given index max is an *upper bound* of the sequence a[0..n-1], and, by construction, a member of that sequence.

```
/*@
  predicate
    MaxElement{L}(value_type* a, integer n, integer max) =
      0 <= max < n && UpperBound(a, n, a[max]);
*/
```

Listing 4.9: Definition of the MaxElement predicate

4.4.1. Formal specification of `max_element`

The new formal specification of max_element in ACSL is shown in Listing 4.10. Note that we also use the predicate StrictUpperBound from Listing 4.8 in order to express that max_element returns the *first* maximum position in a[0..n-1].

```
/*@
  requires valid:    \valid_read(a + (0..n-1));
  assigns  \nothing;
  ensures  result:   0 <= \result <= n;

  behavior empty:
    assumes n == 0;
    ensures result:   \result == 0;

  behavior not_empty:
    assumes 0 < n;
    ensures result:  0 <= \result < n;
    ensures max:      MaxElement(a, n, \result);
    ensures first:   StrictUpperBound(a, \result, a[\result]);

  complete behaviors;
  disjoint behaviors;
*/
size_type
max_element2(const value_type* a, size_type n);
```

Listing 4.10: Formal specification of max_element

4.4.2. Implementation of `max_element`

Listing 4.11 shows implementation of max_element with only the loop invariants changed.

```
size_type
max_element2(const value_type* a, size_type n)
{
  if (0u < n) {
    size_type max = 0u;

    /*@
      loop invariant bound:   0 <= i <= n;
      loop invariant max:     0 <= max < n;
      loop invariant upper:   UpperBound(a, i, a[max]);
      loop invariant first:   StrictUpperBound(a, max, a[max]);
      loop assigns max, i;
      loop variant n-i;
    */
    for (size_type i = 0u; i < n; i++) {
      if (a[max] < a[i]) {
        max = i;
      }
    }

    return max;
  }

  return n;
}
```

Listing 4.11: Implementation of max_element

4.5. The `max_seq` algorithm

In this section we consider the function max_seq (see Chapter 3, [15]) which is very similar to the function max_element of Section 4.3. The main difference between max_seq and max_element is that max_seq returns the maximum value (not just the index of it). Therefore, it requires a *non-empty* range as an argument.

Of course, max_seq can easily be implemented using max_element (see Listing 4.13). Moreover, using only the formal specification of max_element in Listing 4.10 we are also able to deductively verify the correctness of this implementation. Thus, we have a simple example of *modular verification* in the following sense:

> Any implementation of max_element that is separately proven to implement the contract in Listing 4.10 makes max_seq behave correctly. Once the contracts have been defined, the function max_element could be implemented in parallel, or just after max_seq, without affecting the verification of max_seq.

4.5.1. Formal specification of `max_seq`

A formal specification of max_seq in ACSL is shown in Listing 4.12.

```
/*@
  requires n > 0;
  requires \valid_read(p + (0..n-1));

  assigns  \nothing;

  ensures \forall integer i; 0 <= i <= n-1 ==> \result >= p[i];
  ensures \exists integer e; 0 <= e <= n-1 &&  \result == p[e];
*/
value_type
max_seq(const value_type* p, size_type n);
```

Listing 4.12: Formal specification of max_seq

Using the first requires-clause we express that max_seq needs a *non-empty* range as input. Our post-conditions formalize that max_seq indeed returns the maximum value of the range.

4.5.2. Implementation of `max_seq`

Listing 4.13 shows the trivial implementation of max_seq using max_element. Since max_seq requires a non-empty range the call of max_element returns an index to a maximum value in the range. The fact that max_element returns the smallest index is of no importance in this context.

```
value_type
max_seq(const value_type* p, size_type n)
{
  return p[max_element2(p, n)];
}
```

Listing 4.13: Implementation of max_seq

4.6. The auxiliary function `make_pair`

In order to be able to specify functions that work on pairs of indices we introduce here the type `size_type_pair` (Listing 4.14).

```
struct size_type_pair {
  size_type first;
  size_type second;
};

typedef struct size_type_pair size_type_pair;
```

Listing 4.14: The type `size_type_pair`

We will also use the auxiliary function `make_pair` which turns two indices `first` and `second` into an object of `size_type_pair`. The specification and implementation of `make_pair` is shown in Listing 4.15.

```
/*@
    assigns \nothing;

    ensures \result.first  == first;
    ensures \result.second == second;
*/
static inline
size_type_pair
make_pair(size_type first, size_type second)
{
  size_type_pair pair;

  pair.first  = first;
  pair.second = second;

  return pair;
}
```

Listing 4.15: The function `make_pair`

4.7. The `min_element` algorithm

The `min_element` algorithm in the C++ Standard Library [21, §28.7.8] searches the minimum in a general sequence. The signature of our version of `min_element` reads:

```
size_type min_element(const value_type* a, size_type n);
```

The function `min_element` finds the smallest element in the range `a[0..n-1]`. More precisely, it returns the unique valid index `i` such that `a[i]` is minimal among the values `a[0], ..., a[n-1]`, and `i` is the first position with that property. The return value of `min_element` is `n` if and only if `n == 0`.

First we define in Listing 4.16 the overloaded predicate `LowerBound` that basically expresses that a given value is less or equal than all elements of a given array (section).

```
/*@
  predicate
    LowerBound{L}(value_type* a, integer m, integer n, value_type v) =
      \forall integer i; m <= i < n  ==>  v <= a[i];

  predicate
    LowerBound{L}(value_type* a, integer n, value_type v) =
      LowerBound{L}(a, 0, n, v);
*/
```

Listing 4.16: Definition of the `LowerBound` predicate

Closely related to the predicate `LowerBound` is the overloaded predicate `StrictLowerBound` from Listing 4.17.

```
/*@
  predicate
    StrictLowerBound{L}(value_type* a, integer m, integer n, value_type v) =
      \forall integer i; m <= i < n  ==>  v < a[i];

  predicate
    StrictLowerBound{L}(value_type* a, integer n, value_type v) =
      StrictLowerBound{L}(a, 0, n, v);
*/
```

Listing 4.17: Definition of the `StrictLowerBound` predicate

We then define in Listing 4.18 the predicate `MinElement` by stating that the element at a given index `min` is a *lower bound* of the sequence `a[0..n-1]`, and, by construction, a member of that sequence.

```
/*@
  predicate
    MinElement{L}(value_type* a, integer n, integer min) =
      0 <= min < n && LowerBound(a, n, a[min]);
*/
```

Listing 4.18: Definition of the `MinElement` predicate

4.7.1. Formal specification of `min_element`

The ACSL specification of `min_element` is shown in Listing 4.19. Note that we also use the predicate `StrictLowerBound` from Listing 4.17 in order to express that `min_element` returns the *first* minimum position in `a[0..n-1]`.

```
/*@
  requires valid:   \valid_read(a + (0..n-1));
  assigns  \nothing;
  ensures  result:  0 <= \result <= n;

  behavior empty:
    assumes n == 0;
    ensures result:  \result == 0;

  behavior not_empty:
    assumes 0 < n;
    ensures result: 0 <= \result < n;
    ensures min:    MinElement(a, n, \result);
    ensures first:  StrictLowerBound(a, \result, a[\result]);

  complete behaviors;
  disjoint behaviors;
*/
size_type
min_element(const value_type* a, size_type n);
```

Listing 4.19: Formal specification of min_element

4.7.2. Implementation of **min_element**

Listing 4.20 shows the implementation of min_element with loop invariants where we also employ the predicates LowerBound and StrictLowerBound.

```
size_type
min_element(const value_type* a, size_type n)
{
  if (0u < n) {
    size_type min = 0u;

    /*@
      loop invariant bound:  0 <= i    <= n;
      loop invariant min:    0 <= min <  n;
      loop invariant lower:  LowerBound(a, i, a[min]);
      loop invariant first:  StrictLowerBound(a, min, a[min]);
      loop assigns min, i;
      loop variant n-i;
    */
    for (size_type i = 0u; i < n; i++) {
      if (a[i] < a[min]) {
        min = i;
      }
    }

    return min;
  }

  return n;
}
```

Listing 4.20: Implementation of min_element

4.8. The `minmax_element` algorithm

The `minmax_element` algorithm in the C++ Standard Library [21, §28.7.8] searches the minimum and the maximum in a sequence. The signature of our version of `min_element` reads:

```
size_type_pair minmax_element(const value_type* a, size_type n);
```

Note that `minmax_element` returns a *pair* of indices (see Section 4.6). This pair contains the *first* position where the minimum occurs in the sequence `a[0..n-1]` and the *last* position where maximum occurs.

The properties of the index for the minimum value are the same as the properties for the `min_element` in Section 4.7. In opposition to that, the properties of the index that marks the maximum element, are slightly altered to the properties of `max_element` in Section 4.3. The `max_element` algorithm returns the position of the *first* occurrence of the maximum element if it occurs multiple times in the sequence. The `minmax_element` algorithm returns the position of the last occurrence of the maximum element.

4.8.1. Formal specification of `minmax_element`

The ACSL specification of `minmax_element` is shown in Listing 4.21. Note that we use the predicates `StrictLowerBound` from Listing 4.17 and `StrictUpperBound` from Listing 4.8 to express that the first minimum and last maximum position is returned. We also use the predicates `MinElement` in Listing 4.18 and `MaxElement` in Listing 4.9 analogue to their usage in the specification of their algorithms `min_element` and `max_element`.

```
/*@
  requires valid: \valid_read(a + (0..n-1));

  assigns \nothing;

  ensures result: 0 <= \result.first  <= n;
  ensures result: 0 <= \result.second <= n;

  behavior empty:
    assumes 0 == n;
    ensures result: \result.first == 0;
    ensures result: \result.second == 0;

  behavior not_empty:
    assumes 0 < n;
    ensures result:  0 <= \result.first < n;
    ensures result:  0 <= \result.second < n;

    ensures min:    MinElement(a, n, \result.first);
    ensures first:  StrictLowerBound(a, \result.first, a[\result.first]);
    ensures max:    MaxElement(a, n, \result.second);
    ensures last:   StrictUpperBound(a, \result.second+1, n, a[\result.second]);
*/
size_type_pair
minmax_element(const value_type* a, size_type n);
```

Listing 4.21: Formal specification of `minmax_element`

The specification is similar to the specifications of `min_element` and `max_element`. The only difference lies in the postcondition `last`. Here the postcondition states that after the position of the maximum

element there is no value greater or equal the maximum element. This differs from the specification of max_element, where the first occurrence of the maximum value has to be returned.

4.8.2. Implementation of `minmax_element`

Listing 4.22shows an implementation of minmax_element. It uses the auxiliary function make_pair from Section 4.6 to construct a pair of indices. We will focus on the loop invariant last, because it is the only loop invariant differs from the implementations of min_element and max_element.

```
size_type_pair
minmax_element(const value_type* a, size_type n)
{
  if (0u < n) {
    size_type min = 0u;
    size_type max = 0u;

    /*@
      loop invariant bound: 0 <= i    <= n;
      loop invariant min:   0 <= min  <  n;
      loop invariant max:   0 <= max  <  n;
      loop invariant lower: LowerBound(a, i, a[min]);
      loop invariant upper: UpperBound(a, i, a[max]);
      loop invariant first: StrictLowerBound(a, min, a[min]);
      loop invariant last:  StrictUpperBound(a, max+1, i, a[max]);
      loop assigns min, max, i;

      loop variant n-i;
    */
    for (size_type i = 0u; i < n; i++) {
      if (a[i] >= a[max]) {
        max = i;
      }

      if (a[i] < a[min]) {
        min = i;
      }
    }

    return make_pair(min, max);
  }

  return make_pair(n, n);
}
```

Listing 4.22: Implementation of minmax_element

As already mentioned we had to alter the range for the predicate StrictUpperBound to fit into the property of returning the last maximum position that occurred.

5. Binary search algorithms

In this chapter, we consider the four *binary search* algorithms of the C++ Standard Library [21, §28.7.3], namely

- `lower_bound` in Section 5.1

- `upper_bound` in Section 5.2

- two variants for the implementation of `equal_range` which we refer to as `equal_range` and `equal_range2` in Sections 5.3

- two variants for the formal specification of `binary_search` which we refer to as `binary_search` and `binary_search2` in Section 5.4

All binary search algorithms require that their input array is sorted in ascending order. There are two versions of predicate `Sorted` in Listing 5.1. The first one defines when a section of an array is sorted in ascending order. The second version uses the first one to express that the whole array is sorted.

```
/*@
  predicate
    Sorted{L}(value_type* a, integer m, integer n) =
      \forall integer i, j; m <= i < j < n ==> a[i] <= a[j];

  predicate
    Sorted{L}(value_type* a, integer n) = Sorted{L}(a, 0, n);
*/
```

Listing 5.1: The predicate `Sorted`

As in the case of the of maximum/minimum algorithms from Chapter 4 the binary search algorithms primarily use the less-than operator < (and the derived operators <=, > and >=) to determine whether a particular value is contained in a sorted range. Thus, different to the `find` algorithm in Section 3.1, the equality operator == will play only a supporting part in the specification of binary search.

In order to make the specifications of the binary search algorithms more compact and (arguably) more readable we re-use the following predicates

- `UpperBound` in Listing 4.7

- `StrictUpperBound` in Listing 4.8

- `LowerBound` in Listing 4.16

- `StrictLowerBound` in Listing 4.17

5.1. The `lower_bound` algorithm

The `lower_bound` algorithm is one of the four binary search algorithms of the C++ Standard Library [21, §28.7.3.1]. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
size_type lower_bound(const value_type* a, size_type n, value_type val);
```

As with the other binary search algorithms `lower_bound` requires that its input array is sorted in ascending order. The index `lb`, that `lower_bound` returns satisfies the inequality

$$0 \leq \texttt{lb} \leq n \tag{5.1}$$

and has the following properties for a valid index `k` of the array under consideration

$$
\begin{aligned}
0 \leq k < \texttt{lb} &\quad\Longrightarrow\quad a[k] < \texttt{val} &\tag{5.2}\\
\texttt{lb} \leq k < n &\quad\Longrightarrow\quad \texttt{val} \leq a[k] &\tag{5.3}
\end{aligned}
$$

Conditions (5.2) and (5.3) imply that `val` can only occur in the array section `a[lb..n-1]`. In this sense `lower_bound` returns a *lower bound* for the potential indices.

As an example, we consider in Figure 5.2 a sorted array. The arrows indicate which indices will be returned by `lower_bound` for a given value. Note that the index 9 points *one past end* of the array. Values that are not contained in the array are colored in gray.

Figure 5.2.: Some examples for `lower_bound`

Figure 5.2 also clarifies that care must be taken when interpreting the return value of `lower_bound`. An important difference to the algorithms in Chapter 3 is that a return value of `lower_bound` that is less than *n* does not necessarily implies `a[lb] == val`. We can only be sure that `val <= a[lb]` holds.

5.1.1. Formal specification of `lower_bound`

The ACSL specification of `lower_bound` is shown in Listing 5.3. The preconditions `sorted` expresses that the values in the (valid) array need to be sorted in ascending order. The postconditions reflect the conditions listed above and can be expressed using predicates `LowerBound` from Listing 4.16 and `StrictLowerBound` from Listing 4.17, namely,

- Condition (5.1) becomes postcondition `result`

- Condition (5.2) becomes postcondition `left`

- Condition (5.3) becomes postcondition `right`

```
/*@
  requires valid:  \valid_read(a + (0..n-1));
  requires sorted: Sorted(a, n);

  assigns \nothing;

  ensures result:  0 <= \result <= n;
  ensures left:    StrictUpperBound(a, 0, \result, val);
  ensures right:   LowerBound(a, \result, n, val);
*/
size_type
lower_bound(const value_type* a, size_type n, value_type val);
```

Listing 5.3: Formal specification of `lower_bound`

5.1.2. Implementation of `lower_bound`

Our implementation of `lower_bound` is shown in Listing 5.4. Each iteration step narrows down the range that contains the sought-after result. The loop invariants express that in each iteration step all indices less than the temporary left bound `left` contain values that are less than `val` and all indices not less than the temporary right bound `right` contain values that are greater or equal than `val`. The expression to compute `middle` is slightly more complex than the naïve `(left+right)/2`, but it avoids potential overflows.

```
size_type
lower_bound(const value_type* a, size_type n, value_type val)
{
  size_type left  = 0u;
  size_type right = n;

  /*@
    loop invariant bound:  0 <= left <= right <= n;
    loop invariant left:   StrictUpperBound(a, 0,  left, val);
    loop invariant right:  LowerBound(a, right, n, val);

    loop assigns left, right;
    loop variant right - left;
  */
  while (left < right) {
    const size_type middle = left + (right - left) / 2u;

    if (a[middle] < val) {
      left = middle + 1u;
    }
    else {
      right = middle;
    }
  }

  return left;
}
```

Listing 5.4: Implementation of `lower_bound`

5.2. The `upper_bound` algorithm

The `upper_bound` algorithm of the C++ Standard Library [21, §28.7.3.2] is a variant of binary search and closely related to `lower_bound` of Section 5.1. The signature reads:

```
size_type upper_bound(const value_type* a, size_type n, value_type val)
```

As with the other binary search algorithms, `upper_bound` requires that its input array is sorted in ascending order. The index `ub` returned by `upper_bound` satisfies the inequality

$$0 \leq \text{ub} \leq n \tag{5.4}$$

and is involved in the following implications for a valid index `k` of the array under consideration

$$0 \leq k < \text{ub} \quad \Longrightarrow \quad a[k] \leq \text{val} \tag{5.5}$$
$$\text{ub} \leq k < n \quad \Longrightarrow \quad \text{val} < a[k] \tag{5.6}$$

Conditions (5.5) and (5.6) imply that `val` can only occur in the array section `a[0..ub-1]`. In this sense `upper_bound` returns a *upper bound* for the potential indices where `val` can occur. It also means that the searched-for value `val` can *never* be located at the index `ub`.

Figure 5.5 is a variant of Figure 5.2 for the case of `upper_bound` and the same example array. The arrows indicate which indices will be returned by `upper_bound` for a given value. Note how, compared to Figure 5.2, only the arrows from values that *are present* in the array change their target index.

Figure 5.5.: Some examples for `upper_bound`

5.2.1. Formal specification of `upper_bound`

The ACSL specification of `upper_bound` is shown in Listing 5.6. The specification is quite similar to the specification of `lower_bound` (see Listing 5.3). The precondition `sorted` expresses that the values in the (valid) array need to be sorted in ascending order.

The postconditions reflect the conditions listed above and can be expressed using predicates `UpperBound` from Listing 4.7 and `StrictUpperBound` from Listing 4.8, namely,

- Condition (5.4) becomes postcondition `result`
- Condition (5.5) becomes postcondition `left`
- Condition (5.6) becomes postcondition `right`

```
/*@
  requires valid:  \valid_read(a + (0..n-1));
  requires sorted: Sorted(a, n);

  assigns \nothing;

  ensures result:  0 <= \result <= n;
  ensures left:    UpperBound(a, 0, \result, val);
  ensures right:   StrictLowerBound(a, \result, n, val);
*/
size_type
upper_bound(const value_type* a, size_type n, value_type val);
```

Listing 5.6: Formal specification of upper_bound

5.2.2. Implementation of **upper_bound**

Our implementation of upper_bound is shown in Listing 5.7.

The loop invariants express that for each iteration step all indices less than the temporary left bound left contain values not greater than val and all indices not less than the temporary right bound right contain values greater than val.

```
size_type
upper_bound(const value_type* a, size_type n, value_type val)
{
  size_type left  = 0u;
  size_type right = n;

  /*@
    loop invariant bound:  0 <= left <= right <= n;
    loop invariant left:   UpperBound(a, 0, left, val);
    loop invariant right:  StrictLowerBound(a, right, n, val);

    loop assigns left, right;
    loop variant right - left;
  */
  while (left < right) {
    const size_type middle = left + (right - left) / 2u;

    if (a[middle] <= val) {
      left = middle + 1u;
    }
    else {
      right = middle;
    }
  }

  return right;
}
```

Listing 5.7: Implementation of upper_bound

5.3. The `equal_range` algorithm

The `equal_range` algorithm is one of the four binary search algorithms of the C++ Standard Library [21, §28.7.3.3]. As with the other binary search algorithms `equal_range` requires that its input array is sorted in ascending order. The specification of `equal_range` states that it *combines* the results of the algorithms `lower_bound` (Section 5.1) and `upper_bound` (Section 5.2).

For our purposes we have modified `equal_range` to take an array of type `value_type`. Moreover, instead of a pair of iterators, our version returns a pair of indices. To be more precise, the return type of `equal_range` is the struct `size_type_pair` from Listing 4.14. Thus, the signature of `equal_range` now reads:

```
size_type_pair
equal_range(const value_type* a, size_type n, value_type val);
```

Figure 5.8 combines Figure 5.2 with Figure 5.5 in order visualize the behavior of `equal_range` for select test cases. The two types of arrows → and ⇢ represent the respective fields `first` and `second` of the return value. For values that are not contained in the array, the two arrows point to the same index. More generally, if `equal_range` returns the pair $(\mathrm{lb}, \mathrm{ub})$, then the difference $\mathrm{ub} - \mathrm{lb}$ is equal to the number of occurrences of the argument `val` in the array.

Figure 5.8.: Some examples for `equal_range`

We will provide two implementations of `equal_range` and verify both of them. The first implementation just straightforwardly calls `lower_bound` and `upper_bound` and simply returns their results (see Listing 5.10). The second, more elaborate, implementation follows the original STL code by attempting to minimize duplicate computations (see Listing 5.11).

Let $(\mathrm{lb}, \mathrm{ub})$ be the return value `equal_range`, then the conditions (5.1)–(5.6) can be merged into the inequality

$$0 \leq \mathrm{lb} \leq \mathrm{ub} \leq n \tag{5.7}$$

and the following three implications for a valid index k of the array under consideration

$$0 \leq k < \mathrm{lb} \quad \Longrightarrow \quad a[k] < \mathtt{val} \tag{5.8}$$
$$\mathrm{lb} \leq k < \mathrm{ub} \quad \Longrightarrow \quad a[k] = \mathtt{val} \tag{5.9}$$
$$\mathrm{ub} \leq k < n \quad \Longrightarrow \quad a[k] > \mathtt{val} \tag{5.10}$$

Here are some justifications for these conditions.

- Conditions (5.8) and (5.10) are just the Conditions (5.2) and (5.6), respectively.

- The Inequality (5.7) follows from the Inequalities (5.1) and (5.4) and the following considerations: If ub were less than lb, then according to (5.8) we would have $a[\text{ub}] < \text{val}$. One the other hand, we know from (5.10) that opposite inequality $\text{val} < a[\text{ub}]$ holds. Therefore, we have $\text{lb} \leq \text{ub}$.

- Condition (5.9) follows from the combination of (5.3) and (5.5) and the fact that $\leq$ is a total order on the integers.

5.3.1. Formal specification of `equal_range`

The ACSL specification of `equal_range` is shown in Listing 5.9.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires sorted: Sorted(a, n);

  assigns \nothing;

  ensures result:  0 <= \result.first <= \result.second <= n;
  ensures left:    StrictUpperBound(a, 0, \result.first, val);
  ensures middle:  ConstantRange(a, \result.first, \result.second, val);
  ensures right:   StrictLowerBound(a, \result.second, n, val);
*/
size_type_pair
equal_range(const value_type* a, size_type n, value_type val);
```

Listing 5.9: Formal specification of `equal_range`

The ACSL specification of `equal_range` is shown in Listing 5.9. The precondition `sorted` expresses that the values in the (valid) array need to be sorted in ascending order.

The postconditions reflect the conditions listed above and can be expressed using the well-known predicates `ConstantRange` (Listing 3.25), `StrictUpperBound` (Listing 4.8) and `StrictLowerBound` (Listing 4.17), namely,

- Condition (5.7) becomes postcondition `result`

- Condition (5.8) becomes postcondition `left`

- Condition (5.9) becomes postcondition `middle`

- Condition (5.10) becomes postcondition `right`

5.3.2. First implementation of `equal_range`

Our first implementation of equal_range is shown in Listing 5.10. We just call the two functions lower_bound and upper_bound and return their respective results as a pair.

```
size_type_pair
equal_range(const value_type* a, size_type n, value_type val)
{
  size_type first  = lower_bound(a, n, val);
  size_type second = upper_bound(a, n, val);
  //@ assert aux: second < n  ==>  val < a[second];
  return make_pair(first, second);
}
```

Listing 5.10: First implementation of equal_range

In an earlier version of this document we had proven the similar assertion first <= second with the interactive theorem prover Coq. After reviewing this proof we formulated the new assertion aux that uses a fact from the postcondition of upper_bound (Listing 5.6). The benefit of this reformulation is that both the assertion aux and the postcondition first <= second can now be verified automatically.

5.3.3. Second implementation of `equal_range`

The first implementation of equal_range does more work than needed. STL uses a slightly more complicated implementation of equal_range that performs as much range reduction as possible before calling upper_bound and lower_bound on the reduced ranges. In Listing 5.11 we translated the STL implementation to C code and verified it. It is a drop-in replacement for the first implementation and implements the same formal specification, provided in Listing 5.9.

```
size_type_pair
equal_range2(const value_type* a, size_type n, value_type val)
{
  size_type first  = 0u;
  size_type middle = 0u;
  size_type last   = n;

  /*@
    loop invariant bounds: 0 <= first <= last <= n;
    loop invariant left:  StrictUpperBound(a, 0, first, val);
    loop invariant right: StrictLowerBound(a, last, n, val);
    loop assigns first, last, middle;
    loop variant last - first;
   */
  while (last > first) {
    middle = first + (last - first) / 2u;

    if (a[middle] < val) {
      first = middle + 1u;
    }
    else if (val < a[middle]) {
      last = middle;
    }
    else {
      break;
    }
  }

  if (first < last) {
    //@ assert sorted: Sorted(a, first, middle);
    size_type left = first + lower_bound(a + first, middle - first, val);
    //@ assert constant: LowerBound(a, left, middle, val);
    //@ assert strict: StrictUpperBound(a, first, left, val);
    ++middle;
    //@ assert sorted: Sorted(a, middle, last);
    size_type right = middle + upper_bound(a + middle, last - middle, val);
    //@ assert constant: UpperBound(a, middle, right, val);
    //@ assert strict: StrictLowerBound(a, right, last, val);
    return make_pair(left, right);
  }
  else {
    return make_pair(first, first);
  }
}
```

Listing 5.11: Second implementation of equal_range

Due to the higher complexity of the second implementation, additional assertions had to be added to ensure that Frama-C is able to verify the correctness of the code. All of these are related to pointer arithmetic and shifting base pointers. They fall into three groups and are briefly discussed below. In order to enable the automatic verification of these properties we added the ACSL lemmas in Listing 5.12.

```
/*@
  lemma
    SortedShift{L}:
      \forall value_type *a, integer l, r;
      0 <= l <= r  ==>  Sorted{L}(a, l, r)  ==>  Sorted{L}(a+l, r-l);

  lemma
    LowerBoundShift{L}:
      \forall value_type *a, val, integer b, c, d;
        LowerBound{L}(a+b, c,    d,    val)  ==>
        LowerBound{L}(a,    c+b, d+b, val);

  lemma
    StrictLowerBoundShift{L}:
      \forall value_type *a, val, integer b, c, d;
        StrictLowerBound{L}(a+b, c,    d,    val)  ==>
        StrictLowerBound{L}(a,    c+b, d+b, val);

  lemma
    UpperBoundShift{L}:
    \forall value_type *a, val, integer b, c;
      UpperBound{L}(a+b,  0, c-b, val)  ==>
      UpperBound{L}(a,     b, c,    val);

  lemma
    StrictUpperBoundShift{L}:
      \forall value_type *a, val, integer b, c;
        StrictUpperBound{L}(a+b, 0, c-b, val)  ==>
        StrictUpperBound{L}(a,    b, c,    val);
*/
```

Listing 5.12: Some lemmas to support the verification of `equal_range`

The `sorted` properties

Both `upper_bound` and `lower_bound` require that they operate on sorted data. This is also true for `equal_range`, however, inside our second implementation we need a more specific formulation, namely,

```
Sorted(a + middle, last - middle)
```

A three-argument form of the `Sorted` predicate from Listing 5.1 was added so we can spell out an intermediate step. This enables the provers to verify the preconditions of the call to `lower_bound` automatically. A similar assertion is present before the call to `upper_bound`.

The `strict` and `constant` properties

Part of the post conditions of `equal_range` is that `val` is both a strict upper and a strict lower bound. However, the calls to `upper_bound` and `lower_bound` only give us

```
StrictUpperBound(a + first, 0, left - first, val)

StrictLowerBound(a + middle, right - middle, last - middle, val)
```

which is not enough to reach the desired post conditions automatically. One intermediate step for each of the assertions was sufficient to guide the prover to the desired result.

Conceptually similar to the `strict` properties the `constant` properties guide the prover towards

```
LowerBound(a, left, n, val)

UpperBound(a, 0, right, val)
```

Combining these properties allow the postcondition `middle` to be derived automatically.

5.4. The `binary_search` algorithm

The `binary_search` algorithm is one of the four binary search algorithms of the C++ Standard Library [21, §28.7.3.4]. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
bool binary_search(const value_type* a, size_type n, value_type  val);
```

Again, `binary_search` requires that its input array is sorted in ascending order. It will return **true** if there exists an index `i` in `a` such that `a[i]  == val` holds.[20]

Figure 5.13.: Some examples for `binary_search`

In Figure 5.13 we do not need to use arrows to visualize the effects of `binary_search`. The colors orange and grey of the sought-after values indicate whether the algorithm returns true or false, respectively.

5.4.1. Formal specification of `binary_search`

The ACSL specification of `binary_search` is shown in Listing 5.14.

---

[20]To be more precise: The C++ Standard Library requires that `(a[i]  <= val)&& (val <= a[i])` holds. For our definition
   of `value_type` (see Section 1.3) this means that `val` equals `a[i]`.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires sorted: Sorted(a, n);

  assigns  \nothing;

  ensures  result:  \result <==> \exists integer i; 0 <= i < n && a[i] == val;
*/
bool
binary_search(const value_type* a, size_type n, value_type val);
```

Listing 5.14: Formal specification of `binary_search`

Note that instead of the somewhat lengthy existential quantification in Listing 5.14 we can use our previously introduced predicate `HasValue` (see Listing 3.4) in to achieve the more concise formal specification in Listing 5.15.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires sorted: Sorted(a, n);

  assigns  \nothing;

  ensures  result: \result <==> HasValue(a, n, val);
*/
bool
binary_search2(const value_type* a, size_type n, value_type val);
```

Listing 5.15: Formal specification of `binary_search` using the `HasValue` predicate

It is interesting to compare this specification with that of `find` shown in Listing 3.5. Both `find` and `binary_search` allow to determine whether a value is contained in an array. The fact that the C++ Standard Library requires that `find` has *linear* complexity whereas `binary_search` must have a *logarithmic* complexity can currently not be expressed with ACSL.

5.4.2. Implementation of `binary_search`

Our implementation of `binary_search` is shown in Listing 5.16.

```
bool
binary_search(const value_type* a, size_type n, value_type val)
{
  const size_type i = lower_bound(a, n, val);
  return (i < n) && (a[i] <= val);
}
```

Listing 5.16: Implementation of `binary_search`

The function `binary_search` first calls `lower_bound` from Section 5.1. Remember that if the latter returns an index `0 <= i < n` then we can be sure that `val <= a[i]` holds.

Part III.

Mutating and numeric algorithms

6. Mutating algorithms

Let us now turn our attention to another class of algorithms, viz. *mutating* algorithms of the C++ Standard Library [21, §28.6], i.e., algorithms that change one or more ranges. In Frama-C, you can explicitly specify that, e.g., entries in an array `a` may be modified by a function `f`, by including the following *assigns clause* into the `f`'s specification:

```
assigns a[0..length-1];
```

The expression `length-1` refers to the value of `length` when `f` is entered, see [16, §2.3.2]. Below are the algorithms we will discuss in this chapter. First, however, we introduce in Sections 6.1 and 6.2 the auxiliary predicates `Unchanged` and `MultisetUnchanged`, respectively.

- `fill` in Section 6.3 initializes each element of an array by a given fixed value.

- `swap` in Section 6.4 exchanges two values.

- `swap_ranges` in Section 6.5 exchanges the contents of the arrays of equal length, element by element. We use this example to present "modular verification", as `swap_ranges` reuses the verified properties of `swap`.

- `copy` in Section 6.6 copies a source array to a destination array.

- `copy_backward` in Section 6.7 also copies a source array to a destination array. This version, however, uses another separation condition than `copy`.

- `reverse_copy` and `reverse` in Sections 6.8 and 6.9, respectively, reverse an array. Whereas `reverse_copy` copies the result to a separate destination array, the `reverse` algorithm works in place.

- `rotate_copy` in Section 6.10 rotates a source array by `m` positions and copies the results to a destination array.

- `rotate` in Section 6.11 rotates *in place* a source array by `m` positions.

- `replace_copy` and `replace` in Sections 6.12 and 6.13, respectively, substitute each occurrence of a value by a given new value. Whereas `replace_copy` copies the result to a separate array, the `replace` algorithm works in place.

- `remove_copy` and `remove` in Sections 6.14–6.17 *filter* all occurrences of a given value from an array. Whereas `remove_copy` copies the result to a separate array, the `remove` algorithm works in place. Note that we provide altogether three versions of how to specify `remove_copy`. This shall help the reader to understand that finding appropriate contracts is an iterative process and that it is usually a good idea to *not* strive for a "complete" contract right from the beginning.

  Note that in Chapter 7 we give an even more detailed presentation of a related algorithm.

- `random_shuffle` Section 6.18 re-arranges the elements of an array in a random way.

6.1. The predicate `Unchanged`

Many of the algorithms in this section iterate sequentially over one or several sequences. For the verification of such algorithms it is often important to express that a section of an array, or the complete array, have remained *unchanged*; this cannot always be expressed by an `assigns` clause. In Listing 6.1 we therefore introduce the overloaded predicate Unchanged together with some simple lemmas. The expression Unchanged{K,L}(a,f,l) is true if the range a[f..l-1] in state K is element-wise equal to that range in state L.

```
/*@
  predicate
    Unchanged{K,L}(value_type* a, integer m, integer n) =
      \forall integer i; m <= i < n ==> \at(a[i],K) == \at(a[i],L);

  predicate
    Unchanged{K,L}(value_type* a, integer n) =
      Unchanged{K,L}(a, 0, n);
*/
```

Listing 6.1: The predicate Unchanged

We also provide a few lemmas for Unchanged that we need for the verification of some algorithms.

Lemma UnchangedSection in Listing 6.2 states that if the range a[m..n-1] does not change when going from state K to state L, then a[p..q-1] does not change either, provided the latter is a subrange of the former, i.e. provided $0 \leq m \leq p \leq q \leq n$ holds.

```
/*@
  lemma
    UnchangedSection{K,L}:
      \forall value_type *a, integer m, n, p, q;
        m <= p <= q <= n          ==>
        Unchanged{K,L}(a, m, n)   ==>
        Unchanged{K,L}(a, p, q);
*/
```

Listing 6.2: The lemma UnchangedSection

Lemma UnchangedStep in Listing 6.3 expresses the simple fact that "unchangedness" is an inductive property.

```
/*@
  lemma
    UnchangedStep{K,L}:
      \forall value_type *a, integer n;
        Unchanged{K,L}(a, n)        ==>
        \at(a[n],K) == \at(a[n],L)  ==>
        Unchanged{K,L}(a, n+1);
*/
```

Listing 6.3: The lemma UnchangedStep

Lemma `UnchangedTransitive` in Listing 6.4 expresses the transitivity of `Unchanged` with respect to program states.

```
/*@
  lemma
    UnchangedTransitive{K,L,M}:
      \forall value_type *a, integer n;
        Unchanged{K,L}(a, n)  ==>
        Unchanged{L,M}(a, n)  ==>
        Unchanged{K,M}(a, n);
*/
```

Listing 6.4: The lemma `UnchangedTransitive`

6.2. The predicate `MultisetUnchanged`

Various algorithms in this document *rearrange* or *reorder* the elements of a given range such that the number of each element remains unchanged. In other words, reordering leaves the *multiset*[21] of elements in the range unchanged.

We use the predicate `MultisetUnchanged`, defined in Listing 6.5, to formally describe this property. This predicate, which is given in two overloaded versions, relies on the logic function `Count` that is defined in Listing 3.34.

```
/*@
  predicate
    MultisetUnchanged{L1,L2}(value_type* a, integer first, integer last) =
      \forall value_type v;
        Count{L1}(a, first, last, v) == Count{L2}(a, first, last, v);


  predicate
    MultisetUnchanged{L1,L2}(value_type* a, integer n) =
      MultisetUnchanged{L1,L2}(a, 0, n);
*/
```

Listing 6.5: The predicate `MultisetUnchanged`

---

[21]See `http://en.wikipedia.org/wiki/Multiset`

6.3. The `fill` algorithm

The `fill` algorithm in the C++ Standard Library [21, §28.6.6] initializes general sequences with a particular value. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
void fill(value_type* a, size_type n, value_type val);
```

6.3.1. Formal specification of `fill`

Listing 6.6 shows the formal specification of `fill` in ACSL. We can express the postcondition of `fill` simply by using the overloaded predicate `ConstantRange` from Listing 3.25.

```
/*@
  requires valid: \valid(a + (0..n-1));

  assigns a[0..n-1];

  ensures constant: ConstantRange(a, n, val);
*/
void
fill(value_type* a, size_type n, value_type val);
```

Listing 6.6: Formal specification of `fill`

The `assigns`-clauses formalize that `fill` modifies only the entries of the range `a[0..n-1]`. In general, when more than one *assigns clause* appears in a function's specification, it is permitted to modify any of the referenced memory locations. However, if no *assigns clause* appears at all, the function is free to modify any memory location, see [16, §2.3.2]. To forbid a function to do any modifications outside its scope, a clause `assigns \nothing;` must be used, as we practised in the example specifications in Chapter 3.

6.3.2. Implementation of `fill`

Listing 6.7 shows an implementation of `fill`.

```
void
fill(value_type* a, size_type n,  value_type val)
{
  /*@
    loop invariant bound:    0 <= i <= n;
    loop invariant constant: ConstantRange(a, i, val);
    loop assigns i, a[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    a[i] = val;
  }
}
```

Listing 6.7: Implementation of `fill`

The loop invariant `constant` expresses that for each iteration the array is filled with the value of `val` up to the index `i` of the iteration. Note that we use here again the predicate `ConstantRange` from Listing 3.25.

6.4. The **swap** algorithm

The `swap` algorithm [21, §28.6.3] in the C++ Standard Library exchanges the contents of two variables. Similarly, the `iter_swap` algorithm [21, §28.6.3] exchanges the contents referenced by two pointers. Since C and hence ACSL, does not support an `&` type constructor ("declarator"), we will present an algorithm that processes pointers and refer to it as `swap`.

6.4.1. Formal specification of **swap**

The ACSL specification for the `swap` function is shown in Listing 6.8. The preconditions are formalized by the `requires`-clauses which state that both pointer arguments of the `swap` function must be dereferenceable.

```
/*@
  requires \valid(p);
  requires \valid(q);

  assigns *p;
  assigns *q;

  ensures *p == \old(*q);
  ensures *q == \old(*p);
*/
void
swap(value_type* p, value_type* q);
```

Listing 6.8: Formal specification of `swap`

Upon termination of `swap` the entries must be mutually exchanged. We can express those postconditions by using the `ensures`-clause. The expression `\old(*p)` refers to the pre-state of the function contract, whereas by default, a postcondition refers the values after the functions has been terminated.

6.4.2. Implementation of **swap**

Listing 6.9 shows the usual straight-forward implementation of `swap`. No interspersed ACSL is needed to get it verified by Frama-C.

```
void
swap(value_type* p, value_type* q)
{
  value_type save = *p;
  *p = *q;
  *q = save;
}
```

Listing 6.9: Implementation of `swap`

6.5. The `swap_ranges` algorithm

The `swap_ranges` algorithm in the C++ Standard Library [21, §28.6.3] exchanges the contents of two expressed ranges element-wise. After translating C++ reference types and iterators to C, our version of the original signature reads:

```
void swap_ranges(value_type* a, size_type n, value_type* b);
```

We do not return a value since it would equal n, anyway.

This function refers to the previously discussed algorithm `swap`. Thus, `swap_ranges` serves as another example for "modular verification". The specification of `swap` will be automatically integrated into the proof of `swap_ranges`.

6.5.1. Formal specification of `swap_ranges`

Listing 6.10 shows an ACSL specification for the `swap_ranges` algorithm.

```
/*@
  requires valid:  \valid(a + (0..n-1));
  requires valid:  \valid(b + (0..n-1));
  requires sep:    \separated(a+(0..n-1), b+(0..n-1));

  assigns a[0..n-1];
  assigns b[0..n-1];

  ensures equal:  EqualRanges{Old,Here}(a, n, b);
  ensures equal:  EqualRanges{Old,Here}(b, n, a);
*/
void
swap_ranges(value_type* a, size_type n, value_type* b);
```

Listing 6.10: Formal specification of `swap_ranges`

The `swap_ranges` algorithm works correctly only if a and b do not overlap. Because of that fact we use the clause `sep` to tell Frama-C that a and b must not overlap.

With the `assigns`-clause we postulate that the `swap_ranges` algorithm alters the elements contained in two distinct ranges, modifying the corresponding elements and nothing else.

The postconditions of `swap_ranges` specify that the content of each element in its post-state must equal the pre-state of its counterpart. We can use the predicate `EqualRanges` (see Listing 3.15) together with the label `Old` and `Here` to express the postcondition of `swap_ranges`. In our specification in Listing 6.10, for example, we specify that the array a in the memory state that corresponds to the label `Here` is equal to the array b at the label `Old`. Since we are specifying a postcondition `Here` refers to the post-state of `swap_ranges` whereas `Old` refers to the pre-state.

6.5.2. Implementation of `swap_ranges`

Listing 6.11 shows an implementation of swap_ranges together with the necessary loop annotations.

```
void
swap_ranges(value_type* a, size_type n, value_type* b)
{
  /*@
    loop invariant bound:  0 <= i <= n;
    loop invariant equal:  EqualRanges{Pre,Here}(a, i, b);
    loop invariant equal:  EqualRanges{Pre,Here}(b, i, a);

    loop invariant unchanged:  Unchanged{Pre,Here}(a, i, n);
    loop invariant unchanged:  Unchanged{Pre,Here}(b, i, n);

    loop assigns i, a[0..n-1], b[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    swap(a + i, b + i);
  }
}
```

Listing 6.11: Implementation of swap_ranges

For the postcondition of the specification in Listing 6.10 to hold, our loop invariants must ensure that at each iteration all of the corresponding elements that have already been visited are swapped.

Note that there are two additional loop invariants which claim that all the elements that have not visited yet equal their original values. This a workaround that allows us to prove the postconditions of swap_ranges despite the fact that the loop assigns is coarser than it should be. The predicate Unchanged from Listing 6.1 is used to express this property.

6.6. The `copy` algorithm

The `copy` algorithm in the C++ Standard Library [21, §28.6.1] implements a duplication algorithm for general sequences. For our purposes we have modified the generic implementation to that of a range of type `value_type`. The signature now reads:

```
void copy(const value_type* a, size_type n, value_type* b);
```

Informally, the function copies every element from the source range `a[0..n-1]` to the destination range `b[0..n-1]`, as shown in Figure 6.12.

Figure 6.12.: Effects of `copy`

6.6.1. Formal specification of `copy`

Figure 6.12 might suggest that the ranges `a[0..n-1]` and `b[0..n-1]` must not overlap. However, since the informal specification requires that elements are copied in the order of increasing indices only a weaker condition is necessary. To be more specific, it is required that the pointer `b` does not refer to elements of `a[0..n-1]` as shown in the example in Figure 6.13.

Figure 6.13.: Possible overlap of `copy` ranges

The ACSL specification of `copy` is shown in Listing 6.14. The `copy` algorithm expects that the ranges `a` and `b` are valid for reading and writing, respectively. Note the precondition `sep` that expresses the previously discussed non-overlapping property.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid:       \valid(b + (0..n-1));
  requires sep:     \separated(a + (0..n-1), b);

  assigns b[0..n-1];

  ensures equal:   EqualRanges{Old,Here}(a, n, b);
*/
void
copy(const value_type* a, const size_type n, value_type* b);
```

Listing 6.14: Formal specification of `copy`

Again, we can use the `EqualRanges` predicate from Section 3.5 to express that the array `a` equals `b` after `copy` has been called. Nothing else must be altered. To state this we use the `assigns`-clause.

6.6.2. Implementation of `copy`

Listing 6.15 shows an implementation of the `copy` function.

```
void
copy(const value_type* a, size_type n, value_type* b)
{
  /*@
    loop invariant bound:     0 <= i <= n;
    loop invariant equal:     EqualRanges{Pre,Here}(a, i, b);
    loop invariant unchanged: Unchanged{Pre,Here}(a, i, n);
    loop assigns   i, b[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    b[i] = a[i];
  }
}
```

Listing 6.15: Implementation of `copy`

For the postcondition `equal` to be true, we must ensure that for every index `i`, the value `a[i]` must not yet have been changed before it is copied to `b[i]`. We express this by using the `Unchanged` predicate.[22]

The `assigns` clause ensures that nothing but the range `b[0..n-1]` and the loop variable `i` is modified. Keep in mind, however, that parts of the source range `a[0..n-1]` might change due to its potential overlap with the destination range.

---

[22] Alternatively, this could also be expressed by changing the `loop assigns` clause to `i, b[0..i-1]`; however, Frama-C doesn't yet support `loop assigns` clauses containing the loop variable.

6.7. The `copy_backward` algorithm

The `copy_backward` algorithm in the C++ Standard Library [21, §28.6.1] implements another duplication algorithm for general sequences. For our purposes we have modified the generic implementation to that of a range of type `value_type`. The signature now reads:

```
void copy_backward(const value_type* a, size_type n, value_type* b);
```

The main reason for the existence of `copy_backward` is to allow copying when the start of the destination range `a[0..n-1]` is contained in the source range `b[0..n-1]`. In this case, `copy` can't be employed since its precondition `sep` is violated, as can be see in Listing 6.14.

The informal specification of `copy_backward` states that copying starts at the end of the source range. For this to work, however, the pointer `b+n` must not be contained in the source range. Note that the order of operation (or procedure) calls cannot be specified in ACSL.[23] A similar remark about order of operations tacitly applied to earlier functions as well, e.g. to `copy`, where the C++ order was prescribed by confining the signature to a `ForwardIterator`.

Figure 6.16 gives an example where `copy_backward`, but *not* `copy`, can be applied.

Figure 6.16.: Possible overlap of `copy_backward` ranges

Note that in the original signature the argument `b` refers to one past the end of the destination range. Here, however, it refers to its start. The reason for this change is that in C++ `copy_backward` is defined for *bidirectional iterators* which do not provide random access operations such as adding or subtracting an index. Since our C version works on pointers we do not consider it as necessary to use the one past the end pointer.

6.7.1. Formal specification of `copy_backward`

The ACSL specification of `copy_backward` is shown in Listing 6.17. The `copy_backward` algorithm expects that the ranges `a[0..n-1]` and `b[0..n-1]` are valid for reading and writing, respectively. Precondition `sep` formalizes the constraints on the overlap of the source and destination ranges as discussed at the beginning of this section.

---

[23]The Aoraï specification language and the corresponding Frama-C plugin are provided to specify and verify temporal properties of code; however, they are beyond the scope of this tutorial.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid:      \valid(b + (0..n-1));
  requires sep:    \separated(a + (0..n-1), b + n);

  assigns  b[0..n-1];

  ensures equal: EqualRanges{Old,Here}(a, n, b);
*/
void
copy_backward(const value_type* a, size_type n, value_type* b);
```

Listing 6.17: Formal specification of copy_backward

The function copy_backward assigns the elements from the source range a to the destination range b, modifying the memory of the elements pointed to by b. Again, we can use the EqualRanges predicate from Section 3.5 to express that the array a equals b after copy_backward has been called.

6.7.2. Implementation of copy_backward

Listing 6.18 shows an implementation of the copy_backward function.

```
void
copy_backward(const value_type* a, size_type n, value_type* b)
{
  /*@
    loop invariant bound:     0 <= i <= n;
    loop invariant equal:     EqualRanges{Pre,Here}(a, i, n, b);
    loop invariant unchanged: Unchanged{Pre,Here}(a, i);
    loop assigns i, b[0..n-1];
    loop variant i;
   */
  for (size_type i = n; i > 0u; --i) {
    b[i - 1u] = a[i - 1u];
  }
}
```

Listing 6.18: Implementation of copy_backward

We have loop invariants similar to copy, stating the loop variable's range (bound) and the area that has already been copied in each cycle (equal).

6.8. The `reverse_copy` algorithm

The `reverse_copy` algorithm of the C++ Standard Library [21, §28.6.10] inverts the order of elements in a sequence. `reverse_copy` does not change the input sequence, and copies its result to the output sequence. For our purposes we have modified the generic implementation to that of a range of type `value_type`. The signature now reads:

```
void reverse_copy(const value_type* a, size_type n, value_type* b);
```

Informally, `reverse_copy` copies the elements from the array `a` into array `b` such that the copy is a reverse of the original array.

```
/*@
  predicate
    Reverse{K,L}(value_type* a, integer n, value_type* b) =
      \forall integer i; 0 <= i < n ==> \at(a[i],K) == \at(b[n-1-i], L);

  predicate
    Reverse{K,L}(value_type* a, integer m, integer n,
                 value_type* b, integer p) = Reverse{K,L}(a+m, n-m, b+p);

  predicate
    Reverse{K,L}(value_type* a, integer m, integer n, value_type* b) =
      Reverse{K,L}(a, m, n, b, m);

  predicate
    Reverse{K,L}(value_type* a, integer m, integer n, integer p) =
      Reverse{K,L}(a, m, n, a, p);

  predicate
    Reverse{K,L}(value_type* a, integer m, integer n) =
      Reverse{K,L}(a, m, n, m);

  predicate
    Reverse{K,L}(value_type* a, integer n) = Reverse{K,L}(a, 0, n);
*/
```

Listing 6.19: The predicate `Reverse`

In order to concisely formalize these conditions we define in Listing 6.19 the predicate `Reverse` (see also Figure 6.20). We also define several overloaded variants of `Reverse` that provide default values for some of the parameters. These overloaded versions enable us to write more concise ACSL annotations.

Figure 6.20.: Sketch of predicate `Reverse`

6.8.1. Formal specification of `reverse_copy`

The ACSL specification of reverse_copy is shown in Listing 6.21. We use the second version of predicate Reverse from Listing 6.19 in order to formulate the postcondition of reverse_copy.

```
/*@
  requires valid:  \valid_read(a + (0..n-1));
  requires valid:      \valid(b + (0..n-1));
  requires sep:     \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..(n-1)];

  ensures reverse:     Reverse{Old,Here}(a, n, b);
  ensures unchanged: Unchanged{Old,Here}(a, n);
*/
void
reverse_copy(const value_type* a, size_type n, value_type* b);
```

Listing 6.21: Formal specification of reverse_copy

6.8.2. Implementation of `reverse_copy`

Listing 6.22 shows an implementation of the reverse_copy function. For the postcondition to be true, we must ensure that for every element i, the comparison b[i] == a[n-1-i] holds. This is formalized by the loop invariant reverse where we employ the first version of Reverse from Listing 6.19.

```
void
reverse_copy(const value_type* a, size_type n, value_type* b)
{
  /*@
    loop invariant bound:   0 <= i <= n;
    loop invariant reverse: Reverse{Here,Pre}(b, 0, i, a, n-i);
    loop assigns i, b[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    b[i] = a[n - 1u - i];
  }
}
```

Listing 6.22: Implementation of reverse_copy

6.9. The `reverse` algorithm

The `reverse` algorithm of the C++ Standard Library [21, §28.6.10] inverts the order of elements in a sequence. Note that `reverse` works *in place*, meaning that it modifies its input sequence. Our modified signature reads:

```
void reverse(value_type* a, size_type n);
```

6.9.1. Formal specification of `reverse`

The ACSL specification for the `reverse` function is shown in listing 6.23.

```
/*@
  requires valid: \valid(a + (0..n-1));

  assigns a[0..n-1];

  ensures reverse: Reverse{Old,Here}(a, n);
*/
void
reverse(value_type* a, size_type n);
```

Listing 6.23: Formal specification of `reverse`

6.9.2. Implementation of `reverse`

Listing 6.24 shows an implementation of the `reverse` function. Since the `reverse` algorithm operates *in place* we use the `swap` function from Section 6.4 in order to exchange the elements of the first half of the array with the corresponding elements of the second half. We reuse the predicates `Reverse` (Listing 6.19) and `Unchanged` (Listing 6.1) in order to write concise loop invariants.

```
void
reverse(value_type* a, size_type n)
{
  const size_type half = n / 2u;

  //@ assert half: half <= n - half;
  //@ assert half: 2*half <= n <= 2*half + 1;
  /*@
    loop invariant bound:   0 <= i <= half <= n-i;
    loop invariant left:    Reverse{Pre,Here}(a, 0, i, n-i);
    loop invariant middle: Unchanged{Pre,Here}(a, i, n-i);
    loop invariant right:   Reverse{Pre,Here}(a, n-i, n, 0);
    loop assigns i, a[0..n-1];
    loop variant half - i;
  */
  for (size_type i = 0u; i < half; ++i) {
    swap(&a[i], &a[n - 1u - i]);
  }
}
```

Listing 6.24: Implementation of `reverse`

6.10. The `rotate_copy` algorithm

The `rotate_copy` algorithm of the C++ Standard Library [21, §28.6.11] copies, in a particular way, the elements of one sequence of length *n* into a separate sequence. More precisely,

- the first *m* elements of the first sequence become the last *m* elements of the second sequence, and

- the last *n* − *m* elements of the first sequence become the first *n* − *m* elements of the second sequence.

Figure 6.25 illustrates the effects of `rotate_copy` by highlighting how the initial and final segments of the array a[0..n−1] are mapped to corresponding subranges of the array b[0..n−1].

Figure 6.25.: Effects of `rotate_copy`

For our purposes we have modified the generic implementation to that of a range of type `value_type`. The signature now reads:

```
void rotate_copy(const value_type* a, size_type m, size_type n, value_type* b);
```

6.10.1. Formal specification of `rotate_copy`

The ACSL specification of `rotate_copy` is shown in Listing 6.26. Note that we require explicitly that both ranges do not overlap and that we are only able to *read* from the range a[0..n−1].

```
/*@
  requires bound: 0 <= m <= n;
  requires valid: \valid_read(a + (0..n-1));
  requires valid:      \valid(b + (0..n-1));
  requires sep:    \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..(n-1)];

  ensures left:     EqualRanges{Old,Here}(a, 0, m,   b, n-m);
  ensures right:    EqualRanges{Old,Here}(a, m, n-m, b, 0);
  ensures unchanged:  Unchanged{Old,Here}(a, n);
*/
void
rotate_copy(const value_type* a, size_type m, size_type n, value_type* b);
```

Listing 6.26: Formal specification of `rotate_copy`

6.10.2. Implementation of `rotate_copy`

Listing 6.27 shows an implementation of the `rotate_copy` function. The implementation simply calls the function `copy` twice.

```
void
rotate_copy(const value_type* a, size_type m, size_type n, value_type* b)
{
  copy(a,   m, b + (n - m));
  copy(a + m, n - m, b);
}
```

Listing 6.27: Implementation of `rotate_copy`

6.11. The `rotate` algorithm

The algorithm `rotate` is an *in-place* variant of the algorithm `rotate_copy` of Section 6.10. We have modified the generic specification of `rotate` [21, §28.6.11] such that it refers to a range of objects of `value_type`. The signature now reads:

```
size_type rotate(const value_type* a, size_type m, size_type n);
```

6.11.1. Formal specification of `rotate`

Figure 6.28 shows informally the behavior of `rotate`. The figure is of course very similar to the one for `rotate_copy` (see Figure 6.25). The notable difference is that `rotate` operates *in place* of the array `a[0..n-1]`.

Figure 6.28.: Effects of `rotate`

The ACSL specification of `rotate` is shown in Listing 6.29.

```
/*@
  requires valid: \valid(a + (0..n-1));
  requires bound: m <= n;

  assigns a[0..n-1];

  ensures result: \result == n-m;
  ensures left:    EqualRanges{Old,Here}(a, 0, m, n-m);
  ensures right:   EqualRanges{Old,Here}(a, m, n, 0);
*/
size_type
rotate(value_type* a, size_type m, size_type n);
```

Listing 6.29: Formal specification of `rotate`

6.11.2. Implementation of **rotate**

Listing 6.30 shows an implementation of the `rotate` function together with several ACSL annotations. Actually, there are several ways to implement `rotate`. We have chosen a particularly simple one that is derived from an implementation of `std::rotate` for *bidirectional iterators* [21, §27.2.6] and which essentially consists of several calls to the algorithm `reverse` of Section 6.9.

Note the statement contract of the final call of `reverse` Listing 6.30. Here we use both the labels `Pre` and `Old` which refer to the pre-states of `reverse` and the function `rotate` itself, respectively.

```
size_type
rotate(value_type* a, size_type m, size_type n)
{
  // if one subrange is empty, then nothings needs to be done
  if ((0u < m) && (m < n)) {
    reverse(a,   m);
    reverse(a + m, n - m);
    /*@
      requires left:   Reverse{Pre,Here}(a, 0, m, 0);
      requires right:  Reverse{Pre,Here}(a, m, n, m);

      assigns          a[0..n-1];

      ensures left:    Reverse{Old,Here}(a, 0, m, n-m);
      ensures right:   Reverse{Old,Here}(a, m, n, 0);
    */
    reverse(a, n);
    //@ assert left:   EqualRanges{Pre,Here}(a, 0, m, n-m);
    //@ assert right:  EqualRanges{Pre,Here}(a, m, n, 0);
  }

  return n - m;
}
```

Listing 6.30: Implementation of `rotate`

6.12. The `replace_copy` algorithm

The `replace_copy` algorithm of the C++ Standard Library [21, §28.6.5] substitutes specific elements from general sequences. Here, the general implementation has been altered to process `value_type` ranges. The new signature reads:

```
size_type replace_copy(const value_type* a, size_type n, value_type* b,
                        value_type v, value_type w);
```

The `replace_copy` algorithm copies the elements from the range `a[0..n]` to range `b[0..n]`, substituting every occurrence of `v` by `w`. The return value is the length of the range. As the length of the range is already a parameter of the function this return value does not contain new information.

Figure 6.31.: Effects of `replace`

Figure 6.31 shows the behavior of `replace_copy` at hand of an example where all occurrences of the value 3 in `a[0..n-1]` are replaced with the value 2 in `b[0..n-1]`.

6.12.1. The predicate `Replace`

We start with defining in Listing 6.32 the predicate `Replace` that describes the intended relationship between the input array `a[0..n-1]` and the output array `b[0..n-1]`. Note the introduction of *local bindings* `\let ai = ...` and `\let bi = ...` in the definition of `Replace` (see [16, §2.2]).

```
/*@
  predicate
    Replace{K,L}(value_type* a, integer n, value_type* b,
                 value_type v, value_type w) =
        \forall integer i; 0 <= i < n ==>
          \let ai = \at(a[i],K); \let bi = \at(b[i],L);
              (ai == v ==> bi == w) && (ai != v ==> bi == ai) ;


  predicate
    Replace{K,L}(value_type* a, integer n, value_type v, value_type w) =
      Replace{K,L}(a, n, a, v, w);
*/
```

Listing 6.32: The predicate `Replace`

Listing 6.32 also contains a second, overloaded version of `Replace` which we will use for the specification of the related in-place algorithm `replace` in Section 6.13.

6.12.2. Formal specification of `replace_copy`

Using predicate `Replace` the ACSL specification of `replace_copy` is as simple as in Listing 6.33. We also require that the input range `a[0..n-1]` and output range `b[0..n-1]` do not overlap.

```
/*@
  requires valid:   \valid_read(a + (0..n-1));
  requires valid:      \valid(b + (0..n-1));
  requires sep:      \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..n-1];

  ensures result:    \result == n;
  ensures replace:   Replace{Old,Here}(a, n, b, v, w);
  ensures unchanged: Unchanged{Old,Here}(a, n);
*/
size_type
replace_copy(const value_type* a, size_type n, value_type* b,
             value_type v, value_type w);
```

Listing 6.33: Formal specification of the `replace_copy`

6.12.3. Implementation of `replace_copy`

An implementation (including loop annotations) of `replace_copy` is shown in Listing 6.34. Note how the structure of the loop annotations resembles the specification of Listing 6.33.

```
size_type
replace_copy(const value_type* a, size_type n, value_type* b, value_type v,
             value_type w)
{
  /*@
    loop invariant bounds:   0 <= i <= n;
    loop invariant replace:  Replace{Pre,Here}(a, i, b, v, w);
    loop assigns i, b[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    b[i] = (a[i] == v ? w : a[i]);
  }

  return n;
}
```

Listing 6.34: Implementation of the `replace_copy` algorithm

6.13. The `replace` algorithm

The `replace` algorithm of the C++ Standard Library [21, §28.6.5] substitutes specific values in a general sequence. Here, the general implementation has been altered to process `value_type` ranges. The new signature reads

```
void replace(value_type* a, size_type n, value_type v, value_type w);
```

The `replace` algorithm substitutes all elements from the range `a[0..n-1]` that equal `v` by `w`.

6.13.1. Formal specification of `replace`

Using the second predicate `Replace` from Listing 6.32 the ACSL specification of `replace` can be expressed as in Listing 6.35.

```
/*@
  requires valid: \valid(a + (0..n-1));

  assigns a[0..n-1];

  ensures replace:  Replace{Old,Here}(a, n, v, w);
*/
void
replace(value_type* a, size_type n, value_type v, value_type w);
```

Listing 6.35: Formal specification of the `replace`

6.13.2. Implementation of `replace`

An implementation of `replace` is shown in Listing 6.36. The loop invariant `unchanged` expresses that when entering iteration `i` the elements `a[i..n-1]` have not yet changed.

```
void
replace(value_type* a, size_type n, value_type v, value_type w)
{
  /*@
    loop invariant bounds:    0 <= i <= n;
    loop invariant replace:   Replace{Pre,Here}(a, i, v, w);
    loop invariant unchanged: Unchanged{Pre,Here}(a, i, n);
    loop assigns i, a[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    if (a[i] == v) {
      a[i] = w;
    }
  }
}
```

Listing 6.36: Implementation of the `replace` algorithm

6.14. The `remove_copy` algorithm (basic contract)

The `remove_copy` algorithm of the C++ Standard Library [21, §28.6.8] copies all elements of a sequence other than a given value. Here, the general implementation has been altered to process `value_type` ranges. The new signature reads:

```
size_type
remove_copy(const value_type* a, size_type n, value_type* b, value_type v);
```

The requirements of `remove_copy` are:

Requirements	Description
Remove Copy Size	The output range has to fit in all the elements of the input range, except the ones that equal the value `v` by `remove_copy`.
Remove Copy Separated	The input range and the output range do not overlap
Remove Copy Elements	The `remove_copy` algorithm copies elements that are not equal to `v` from range `a[0..n-1]` to the range `b[0..\ result-1]`.
Remove Copy Stability	The algorithm is stable, that is, the relative order of the elements in `b` is the same as in `a`.
Remove Copy Return	The return value is the length of the resulting range.
Remove Copy Complexity	The algorithm takes *n* comparisons in every case.

Table 6.37.: Properties of `remove_copy`

Figure 6.38 shows an example of how `remove_copy` is supposed to copy elements that differ from `4` from the input range to the output range.

Figure 6.38.: Effects of `remove_copy`

6.14.1. Formal specification of `remove_copy`

Listing 6.39 shows our first attempt to specify `remove_copy`. In postcondition `discard` we use of the negation of the predicate `HasValue` from Listing 3.4 to show that the value `v` does not occur in the range `b[0..\result]`.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep:   \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..n-1];

  ensures bound:     0 <= \result <= n;
  ensures discard:   !HasValue(b, \result, v);
  ensures unchanged: Unchanged{Old, Here}(a, n);
  ensures unchanged: Unchanged{Old, Here}(b, \result, n);
*/
size_type
remove_copy(const value_type *a, size_type n, value_type *b, value_type v);
```

Listing 6.39: Formal specification of `remove_copy`

One shortcoming of this specification is that the postcondition `bound` only makes very general and not very precise statements about the number of copied elements. We will address this problem in Section 6.15. A more serious shortcoming is, however, that we haven't specified what the relationship between the elements of the input range `a[0..n-1]` and the output range `b[0..\result-1]` looks like. This problem will be discussed in Section 6.16.

6.14.2. Implementation of `remove_copy`

An implementation of remove_copy is shown in Listing 6.40.

```
size_type
remove_copy(const value_type *a, size_type n, value_type *b, value_type v)
{
  size_type k = 0u;

  /*@
    loop invariant bound:     0 <= k <= i <= n;
    loop invariant discard:   !HasValue(b, k, v);
    loop invariant unchanged: Unchanged{Pre,Here}(b, k, n);
    loop assigns    k, i, b[0..n-1];
    loop variant    n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    if (a[i] != v) {
      b[k++] = a[i];
    }
  }

  return k;
}
```

Listing 6.40: Implementation of remove_copy

Here we also need to add another loop invariant discard which basically checks if v occurs in b[0..k] for each iteration of the loop.

6.15. The `remove_copy` algorithm (number of copied elements)

This section rests on Section 6.14 and improves on the contract from Listing 6.39 by exactly specifying the number `\result` of elements copied by `remove_copy`.

Actually, it is not difficult to formally describe this number because it is equal to the size n of the input range minus the number of occurrences of the value v in the input range. We therefore define the corresponding logic function `RemoveSize` in Listing 6.41 simply by resorting to the logic function `Count` from Listing 3.32

```
/*@
  logic integer
    RemoveSize(value_type* a, integer n, value_type v) = n - Count(a, n, v);

  lemma
    RemoveSizeEmpty:
      \forall value_type *a, v, integer n;
        n <= 0  ==>  RemoveSize(a, n, v) == n;

  lemma
    RemoveSizeHit:
      \forall value_type *a, v, integer n;
        0 <= n  ==>  a[n] == v  ==>
        RemoveSize(a, n+1, v) == RemoveSize(a, n, v);

  lemma
    RemoveSizeMiss:
     \forall value_type *a, v, integer n;
        0 <= n  ==>  a[n] != v  ==>
        RemoveSize(a, n+1, v) == RemoveSize(a, n, v) + 1;

  lemma
    RemoveSizeRead{L1,L2}:
     \forall value_type *a, v, integer n;
        Unchanged{L1,L2}(a, n)  ==>
        RemoveSize{L1}(a, n, v) == RemoveSize{L2}(a, n, v);
*/
```

Listing 6.41: The predicate `RemoveSize`

6.15.1. Formal specification of `remove_copy`

To extend our formal specification by using `RemoveSize` we add the new postcondition `size`, which states that the returning value of `remove_copy` equals `RemoveSize`. Listing 6.42 shows the extended formal specification with the new postcondition.

While we now can precisely speak of the number of copied elements, it is still not possible to say something about the exact relationship between the elements of range `a[0..n-1]` and range `b[0..n-1]`. We will address this question in Section 6.16.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep:   \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..n-1];

  ensures size:      \result == RemoveSize(a, n, v);
  ensures bound:     0 <= \result <= n;
  ensures discard:   !HasValue(b, \result, v);
  ensures unchanged: Unchanged{Old, Here}(a, n);
  ensures unchanged: Unchanged{Old, Here}(b, \result, n);
*/
size_type
remove_copy2(const value_type* a, size_type n, value_type* b, value_type v);
```

Listing 6.42: Specification with `RemoveSize`

6.15.2. Implementation of `remove_copy`

Listing 6.43 shows the implementation of our extended specification of `remove_copy`. Here we added the loop invariant `size` which corresponds to the postcondition in Listing 6.42. In order to ensure that the loop invariant `size` can be verified we have added the assertions `size` and `unchanged`.

```
size_type
remove_copy2(const value_type* a, size_type n, value_type* b, value_type v)
{
  size_type k = 0u;

  /*@
    loop invariant size:      k == RemoveSize(a, i, v);
    loop invariant bound:     0 <= k <= i <= n;
    loop invariant discard:   !HasValue(b, k, v);
    loop invariant unchanged: Unchanged{Pre,Here}(b, k, n);
    loop assigns   k, i, b[0..n-1];
    loop variant   n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    if (a[i] != v) {
      b[k++] = a[i];
      //@ assert unchanged: Unchanged{Pre,Here}(a, i+1);
      //@ assert size:      k == RemoveSize(a, i+1, v);
    }
  }

  return k;
}
```

Listing 6.43: Implementation with `RemoveSize`

6.16. The `remove_copy` algorithm (final contract)

This section will consider the achievements of the last sections 6.14 and 6.15 while introducing a new logic function, which is able to describe the relationship between the elements of `a[0..n-1]` and the elements of `b[0..RemoveSize(a, n, v)-1]`. Note that we have shown in Section 6.15 that the result of `RemoveSize(a, n, v)` is equal to the value of `\result`. Therefore the ranges `b[0..\result -1]` and `b[0..RemoveSize(a, n, v)-1]` share the same length for a given range `a[0..n-1]`. This property will be used in Section 6.16.4.

6.16.1. A closer look on the properties of `remove_copy`

Figure 6.44 shows a modified version of the Figure 6.38. We left out the indices of values that were not copied into the target array. Furthermore we have added a dashed arrow which points to the index that corresponds to the *one past the end* location of the input and output range.

Figure 6.44.: Partitioning the input of `remove_copy`

These arrows between the indices of the array `b` and array `a` define the following sequence *p* of seven indices. The index of the *one past the end* is underlined. $p = (1, 2, 5, 7, 8, 10, \underline{11})$

More generally, we refer to the sequence p as *partitioning sequence* of `remove_copy` for the array `a[0..n-1]`. For the **length of a partitioning sequence** m we get the following **strictly monotone increasing** sequence:

$$0 \le p_0 < ... < p_m = n \tag{6.1}$$

and the open index intervals

$$(p_i, p_{i+1}) \qquad\qquad \forall i : 0 \le i < m \tag{6.2}$$

mark **consecutive ranges** of the value v in the source array, that is,

$$a[k] = v \qquad\qquad \forall k : p_i < k < p_{i+1} \tag{6.3}$$

Additionally, the half open interval

$$[0, p_0) \tag{6.4}$$

also marks another **consecutive range** of the value v in the source array:

$$a[k] = v \qquad\qquad \forall k : 0 \le k < p_0 \tag{6.5}$$

Another observation is that

$$a[p_i] \ne v \qquad\qquad \forall i : 0 \le i < m \tag{6.6}$$

holds. Finally, we have

$$a[p_i] = b[i] \qquad\qquad \forall i : 0 \le i < m \tag{6.7}$$

which, together with the inequality (6.6) states, that

$$b[i] \ne v \qquad\qquad \forall i : 0 \le i < m$$

This means that the target does not contain the value v.

6.16.2. Formalizing the properties of the partitions of `remove_copy`

The function `RemovePartition`, whose axiomatic definition is given in Listing 6.45, defines the partitioning sequence *p* from Section 6.16.1. Before we begin to relate the axioms from Listing 6.45 to the formulas from Section 6.16.1 we want to remind the reader that logic functions (and predicates) must be total that is they must be defined for all possible argument values.

```
/*@
  axiomatic RemovePartitionAxiomatic
  {
    logic integer
      RemovePartition(value_type* a,  integer n,
                      value_type  v, integer p) reads a[0..n-1];

    axiom RemovePartitionEmpty:
      \forall value_type *a, v, integer n, p;
        n <= 0 ==> RemovePartition(a, n, v, p) == 0;

    axiom RemovePartitionLeft:
      \forall value_type *a, v, integer n, p;
        p < 0 < n ==> RemovePartition(a, n, v, p) == 0;

    axiom RemovePartitionRight:
      \forall value_type *a, v, integer n, p;
        0 < n ==> RemoveSize(a, n, v) <= p ==> RemovePartition(a, n, v, p) == n;

    axiom RemovePartitionMonotone:
      \forall value_type *a, v, integer n, p, q;
        0 <= p < q <= RemoveSize(a, n, v) ==>
          0 <= RemovePartition(a, n, v, p) < RemovePartition(a, n, v, q) < n;

    axiom RemovePartitionSegment:
      \forall value_type *a, v, integer i, n, p;
          0 <= p < RemoveSize(a, n, v) ==>
            RemovePartition(a, n, v, p) < i < RemovePartition(a, n, v, p+1)
            ==> a[i] == v;

    axiom RemovePartitionInitial:
      \forall value_type *a, v, integer i, m, n, p;
        0 <= i < RemovePartition(a, n, v, m) ==> a[i] == v;

    axiom RemovePartitionNotValue:
      \forall value_type *a, v, integer n, p;
        0 <= p < RemoveSize(a, n, v) ==> a[RemovePartition(a, n, v, p)] != v;

    axiom RemovePartitionEqual:
      \forall value_type *a, v, integer n1, n2, p;
        0 <= n1 < n2  ==>  0 <= p  < RemoveSize(a, n1, v)  ==>
        RemovePartition(a, n1, v, p) == RemovePartition(a, n2, v, p);

    axiom RemovePartitionRead{K,L}:
      \forall value_type *a, v, integer n, p;
        Unchanged{K,L}(a, n) ==>
          RemovePartition{K}(a, n, v, p) == RemovePartition{L}(a, n, v, p);
  }
*/
```

Listing 6.45: The logic function `RemovePartition`

Our definition of `RemovePartition` reflects the following properties properties:

- Axioms `RemovePartitionEmpty`, `RemovePartitionLeft` `RemovePartitionRight`, and `RemovePartitionMonotone` describe the monotonicity property (6.1)

- Axiom `RemovePartitionSegment` represents the property of the Equation (6.3)

- Axiom `RemovePartitionInitial` represents the property shown in Equation (6.5)

- Axiom `RemovePartitionNotValue` reflects Equation (6.6)

- Axiom `RemovePartitionEqual` expresses that the value of `RemovePartition(a, n, v, p)` does not rely on the size of the array

- Axiom `RemovePartitionRead`, finally, indicates the independence of `RemovePartition` from different memory states when used with the same parameters

In addition to these some lemmas are used to describe the bounds of `RemovePartition`. These are shown in Listing 6.46.

```
/*@
  lemma RemovePartitionLowerBound:
    \forall value_type *a, v, integer p, n;
      0 < n ==>
      0 <= p < RemoveSize(a, n, v) ==>
      0 <= RemovePartition(a, n, v, p);

  lemma RemovePartitionUpperBound:
    \forall value_type *a, v, integer p, n;
      0 < n ==>
      0 <= p < RemoveSize(a, n, v) ==>
      RemovePartition(a, n, v, p) < n;
*/
```

Listing 6.46: Some lemmas for `RemovePartition`

6.16.3. The predicate `Remove`

In Listing 6.47 we introduce the new predicate `Remove` to improve the readability of our specification.

```
/*@
  predicate
    Remove{A,B}(value_type* a, integer n,
                value_type* b, value_type v) =
      \forall integer k; 0 <= k < RemoveSize{A}(a, n, v) ==>
        \at(b[k],B) == \at(a[RemovePartition(a, n, v, k)],A);

  predicate
    Remove{A,B}(value_type* a, integer n, value_type v) =
      Remove{A,B}(a, n, a, v);
*/
```

Listing 6.47: The predicate `Remove`

The predicate expresses that for every element of the range `b[0..RemoveSize(a, n, v)-1]`

```
b[k] == RemovePartition(a, n, v, k)
```

holds.

In addition, we define `Remove` for two labels that describe memory states. This improvement provides us with the ability to use the predicate `Remove` for the `remove` algorithm in Section 6.17, where the memory state before the execution of `remove` has to be compared to the memory state after its execution. The predicate also contains an overloaded version of `Remove`. This version is used in the specification of the inplace algorithm `remove` later in Section 6.17. We call the predicate with the parameter `a` two times because we want to access `a` at two different memory states as explained above. Therefore we improve the readability of the contract by the parameter for the second range if we want only want to look at one range.

6.16.4. Formal specification of `remove_copy`

The Listing 6.48 shows the extended version of the formal specification of `remove_copy`.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep:   \separated(a + (0..n-1), b);

  assigns b[0..n-1];

  ensures size:      \result == RemoveSize{Old}(a, n, v);
  ensures bound:     0 <= \result <= n;
  ensures remove:    Remove{Old,Here}(a, n, b, v);
  ensures discard:   !HasValue(b, \result, v);
  ensures unchanged: Unchanged{Old, Here}(a, n);
  ensures unchanged: Unchanged{Old, Here}(b, \result,   n);
*/
size_type
remove_copy3(const value_type* a, size_type n, value_type* b, value_type v);
```

Listing 6.48: An extended specification for `remove_copy`

The additional postcondition `remove` makes use of the predicate `Remove` which we have just described. We also use the lemma `RemoveImpliesNotHasValue` shown in Listing 6.49. The lemma basically connects the postconditions `remove` and `discard`. The use of this lemma will become clearer when we talk about the necessary annotations in Section 6.16.5. Furthermore, we have again the postcondition `unchanged` which states that the source array `a[0..n-1]` does not change.

```
/*@
  lemma RemoveImpliesNotHasValue{A,B}:
    \forall value_type *a, *b, v, integer n;
      Remove{A,B}(a, n, b, v) ==> !HasValue{B}(b, RemoveSize{A}(a, n, v), v);
*/
```

Listing 6.49: The lemma `RemoveImpliesNotHasValue`

6.16.5. Implementation of `remove_copy`

The listing 6.50 shows the more detailed implementation of remove_copy. In addition to the loop invariant remove and unchanged which are used by analogy with the postconditions remove and unchanged Note that thanks to lemma RemoveImpliesNotHasValue we can omit the loop invariant discard.

```
size_type
remove_copy3(const value_type* a, size_type n, value_type* b, value_type v)
{
  size_type k = 0u;

  /*@
    loop invariant bound:     0 <= k <= i <= n;
    loop invariant size:      k == RemoveSize{Pre}(a, i, v);
    loop invariant mapping:   i <= RemovePartition{Pre}(a, n, v, k);
    loop invariant remove:    Remove{Pre,Here}(a, i, b, v);
    loop invariant unchanged: Unchanged{Pre,Here}(a, n);
    loop invariant unchanged: Unchanged{Pre,Here}(b, k, n);
    loop assigns    k, i, b[0..n-1];
    loop variant    n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    if (a[i] != v) {
      b[k++] = a[i];
      //@ assert mapping: RemovePartition{Pre}(a, n, v, k) == i;
    }
  }

  return k;
}
```

Listing 6.50: The extended annotations for remove_copy

We also use another loop invariant mapping. The invariant states that the variable i will always be smaller or equal to the result of RemovePartition(a, n, v, k). We also add the assertion mapping to our implementation as stepping stone for the provers to verify the correctness of this loop invariant.

6.17. The `remove` algorithm

The C++ Standard Library also contains a function `remove` [21, 28.6.8] performing the same operation as `remove_copy` as an in-place algorithm. Its signature is very similar to that of `remove_copy`, except that there is no need for an output array.

```
size_type remove(value_type* a, size_type n, value_type v);
```

Figure 6.51 shows how `remove` is supposed to remove all occurrences of the given value 4 from a range.

Figure 6.51.: Effects of `remove`

6.17.1. Formal specification of `remove`

Listing 6.52 shows a formal specification of the `remove` function. Our specification will share many properties with the one of `remove_copy` from Section 6.16.

```
/*@
  requires valid: \valid(a + (0..n-1));

  assigns a[0..n-1];

  ensures size:      \result == RemoveSize{Old}(a, n, v);
  ensures bound:     0 <= \result <= n;
  ensures remove:    Remove{Old, Here}(a, n, v);
  ensures discard:   !HasValue{Here}(a, \result, v);
  ensures unchanged: Unchanged{Old, Here}(a, \result, n);
*/
size_type
remove(value_type* a, size_type n, value_type v);
```

Listing 6.52: Formal specification of `remove`

For the postcondition `remove` we use one of the overloaded versions of the predicate `Remove` from Listing 6.47. We also add the two labels `Old` and `Here` to refer to the memory states before and after the execution of `remove`.

6.17.2. Implementation of `remove`

Listing 6.53 shows our implementation of `remove` together with the additional loop annotations.

```
size_type
remove(value_type* a, size_type n, value_type v)
{
  size_type k = 0u;

  /*@
    loop invariant size:       k == RemoveSize{Pre}(a,i,v);
    loop invariant bound:      0 <= k <= i <= n;
    loop invariant remove:     Remove{Pre, Here}(a, i, v);
    loop invariant mapping:    i <= RemovePartition{Pre}(a, n, v, k);
    loop invariant unchanged:  Unchanged{Pre,Here}(a, k, n);
    loop assigns k, i, a[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i ) {
    if (a[i] != v) {
      a[k++] = a[i];
    }
  }

  return k;
}
```

Listing 6.53: Implementation of `remove`

As loop annotations we add the loop invariants `size`, `bound`, `remove` and `unchanged` which reflect the postconditions from Listing 6.52. In addition we, similar to `remove_copy` in Section 6.16.5, also rely on Lemma `RemoveImpliesNotHasValue` from Listing 6.49, which allows us to omit a dedicated loop invariant to verify the postcondition `discard`.

6.18. The `random_shuffle` algorithm

The `random_shuffle` algorithm in the C++ Standard Library [21, §C.4.12] randomly rearranges the elements of a given range, that is, it randomly picks one of its possible orderings. For our purposes we have modified the generic implementation to that of a range of type `value_type`. The signature now reads:

```
void random_shuffle(value_type* a, size_type n);
```

Figure 6.54 illustrates an example run of `random_shuffle`. In this figure, the values 1, 2, 3, and 4 occur twice, once, once, and three times, respectively, both before and after the `random_shuffle` run. This expresses that the range has been reordered.

Figure 6.54.: Effects of `random_shuffle`

6.18.1. Formal specification of `random_shuffle`

The ACSL specification of `random_shuffle` is shown in Listing 6.55. The `random_shuffle` algorithm expects that the range a is valid for reading and writing. We use the predicate `MultisetUnchanged` defined in Listing 6.5 to express that the contents of `a[0..n-1]` is just permuted, i.e., the number of occurrences of each of its members remains unchanged. The array `random_seed` contains a seed for the random number generator used to randomize the shuffle. By specifying that the function assigns to `random_seed` we capture that the function may return a different permutation every time.

```
extern unsigned short random_seed[3];

/*@
  requires \valid(a + (0..n-1));

  assigns a[0..n-1];
  assigns random_seed[0..2];

  ensures  MultisetUnchanged{Old,Here}(a,n);
*/
void
random_shuffle(value_type* a, size_type n);
```

Listing 6.55: Formal specification of `random_shuffle`

Note that our specification only states that the resulting range is a reordering of the input range; nothing more and nothing less. Ideally, we would also specify that sequence of reorderings obtained by repeated calls of `random_shuffle` is required to be random. However, ACSL does currently not support the specification of temporal properties related to repeated call results.

More generally speaking, it is not trivial to capture the notion of randomness in a mathematically precise way. As a typical example, we refer to a paper [23, p.6–8], which just gives four statistical tests indicating the randomness of the permutations computed with their algorithm. From a theoretical point of view, a sequence of permutations can be called "random" if its Kolmogorov complexity exceeds a certain measure, however, this property is undecidable [24].

6.18.2. Implementation of `random_shuffle`

Listing 6.56 shows an implementation of the `random_shuffle` function. It repeatedly calls the function `swap` from Section 6.4 to *transpose* (randomly) selected elements. For details of the hereby used function `random_number` we refer to Section 6.18.3.

```
void
random_shuffle(value_type* a, size_type n)
{
  if (0u < n) {
    /*@
      loop invariant bounds:    1 <= i <= n;
      loop invariant reorder:   MultisetUnchanged{Pre,Here}(a, 0, i);
      loop invariant unchanged: Unchanged{Pre,Here}(a, i, n);
      loop assigns   i, a[0..n-1], random_seed[0..2];
      loop variant   n - i;
    */
    for (size_type i = 1u; i < n; ++i) {
      const size_type j = random_number(i) + 1u;
      //@ ghost Before:
      swap(&a[j], &a[i]);
      //@ assert reorder: MultisetUnchanged{Before,Here}(a, 0, j);
      //@ assert reorder:           Unchanged{Before,Here}(a, j+1, i);
      //@ assert reorder: MultisetUnchanged{Before,Here}(a, j+1, i);
    }
  }
}
```

Listing 6.56: Implementation of `random_shuffle`

The loop invariants `reorder` and `unchanged` of `random_shuffle` are necessary for the verification of the postcondition `reorder`: in the `i`th loop cycle, the subrange `a[0..i-1]` has been reordered, while the remaining subrange `a[i..n-1]` is yet unchanged. We also formulate two auxiliary assertions `reorder` which use the *ghost label* `Before`, to guide the automatic verification the loop invariant `reorder`.

6.18.3. Verifying a random number generator

Our implementation presupposes a random-number generator named `random_number` which is specified in Listing 6.57. As in the case of `random_shuffle` itself, we do not formulate specific properties of randomness and only require its result to be in the specified range `[0..n-1]`. Again, the `assigns` clause to the array `random_seed` models the dependency on an external state.

```
extern unsigned short random_seed[3];

/*@
  requires 0 < n;
  assigns random_seed[0..2];
  ensures 0 <= \result < n;
*/
size_type
random_number(size_type n);
```

Listing 6.57: Formal specification of `random_number`

Internally, the function `random_number` uses a custom implementation of the POSIX.1 random number generator `lrand48()`.[24] For sake of completeness, we have included our implementation of `lrand48()` (see Listing 6.58).

```
unsigned short random_seed[3] = { 0x243f, 0x6a88, 0x85a3 };

// see IEEE 1003.1-2008, 2016 Edition for specification
/*@
  assigns random_seed[0..2];
  ensures lower: 0 <= \result;
  ensures upper: \result <= 0x7fffffff;
*/
static long
my_lrand48(void)
{
  unsigned long long state = (unsigned long long)random_seed[0] << 32
                           | (unsigned long long)random_seed[1] << 16
                           | (unsigned long long)random_seed[2];
  state = (0x5deece66dull * state + 0xbull) % (1ull << 48);
  //@ assert lower: state < (1ull << 48);
  long result = state / (1ull << 17);
  //@ assert lower: 0 <= result;
  random_seed[0u] = state >> 32 & 0xffff;
  random_seed[1u] = state >> 16 & 0xffff;
  random_seed[2u] = state >>  8 & 0xffff;
  return result;
}

size_type
random_number(size_type n)
{
  return my_lrand48() % n;
}
```

Listing 6.58: Implementation of `random_number`

---

[24]See `http://pubs.opengroup.org/onlinepubs/9699919799/functions/lrand48.html`

Note that we use the custom ACSL lemma `RandomNumberModulo` which we introduce in Listing 6.59 to support the verification of the assertions `lower`.

```
/*@
  lemma
    RandomNumberModulo:
      \forall unsigned long long a;
        (a % (1ull << 48)) < (1ull << 48);
*/
```

Listing 6.59: The lemma `RandomNumberModulo`

The random number generator is a linear congruence generator with a 48 bit state and the iteration procedure

$$x_{n+1} = ax_n + c \mod 2^{48} \tag{6.8}$$

where $a = 25214903917$ and $c = 11$ are relatively prime integers.

As a part of the iteration procedure in Equation (6.8) an unsigned overflow may occur. This does not affect the result as we are only interested in its lowest 48 bits. However, as one of the options we use, `-warn-unsigned-overflow`, causes Frama-C/WP assert the absence of unsigned overflow this algorithm does not verify under the same options used for the other algorithms. As an exception, we have therefore decided to disable `-warn-unsigned-overflow` for this function as the unsigned overflow is both benign and well-defined (cf. [18, §6.2.5, 9]).

7. A closer look at `unique_copy`

In this chapter we take a closer look on testing and formal verification of the algorithm `unique_copy` from the C++ standard library. We start in Section 7.1 with a careful analysis of `unique_copy`'s informal requirements. We use the identified requirements to guide our testing and verification artefacts in the subsequent chapters.

Although testing is rightly considered easier than formal verification, *designing* both good test code and convincing test data is far from trivial. This is particularly true when one has to deal with testing properties that rely on *implicit* relationships between the input and output data of an algorithm. In the case of `unique_copy`, the main problem is to show that the first (and only the first) element of each consecutive range of equal elements is copied. We tackle this problem in Section 7.2 by exploiting the fact that we deal with a *generic* algorithm. This allows us to transport additional data through the algorithm under test and subsequently use this data to establish that the original data are processed according to the requirements.

In Section 7.3 we then proceed to formally verify a non-generic C version of `unique_copy` with the Frama-C [2] verification platform. In particular, we are writing formal function contract in Frama-C's specification language ACSL [10]. Here, we emphasize that there are different levels of how formal one wants to be. Thus, we consider both the formal verification of the *absence of undefined behavior* and the verification of the *functional correctness* of `unique_copy`.

7.1. Informal specification of `unique_copy`

This section deals with the informal specification of `unique_copy` and its behavior. In Section 7.1.1 we present the requirements of `unique_copy` as they are derived from the C++ standard library. In Section 7.1.2 we look at specific examples to provide a better understanding of `unique_copy`. To take it one step further we finally present in Section 7.1.3 a more formal analysis of `unique_copy`'s behavior.

7.1.1. What does the C++ standard says about `unique_copy`?

The `unique_copy` algorithms of the C++ standard library [21, §28.6.9], whose signature is shown in Listing 7.1, is a template function which copies certain values from a sequence given by the right-open interval of iterators `[first, last)` into a sequence that starts at the iterator `result`.

```
template<class InputIterator, class OutputIterator>
OutputIterator
unique_copy(InputIterator first, InputIterator last,
            OutputIterator result);
```

Listing 7.1: Signature of `unique_copy` from the C++ standard library

For the purposes of this report we do not consider the wide possibilities of ranges covered by this signature, rather we assume the two ranges to be arrays[25] `a[0..n-1]` and `b[0..n-1]` of length n. Listing 7.2

---

[25]We employ here and in the following the ACSL notation `a[0..n-1]` to denote an array of n elements.

shows thus the signature and implementation of a simplified yet still generic function `unique_copy`. Later in this document we will consider an even more specific version of `unique_copy` that is implemented in C.

```cpp
template<typename T>
size_type unique_copy(const T* a, size_type n, T* b)
{
  auto result = std::unique_copy(a, a + n, b);
  return result - b;
}
```

Listing 7.2: A simplified version of `unique_copy`

Besides assuming the input and output ranges to be (generic) arrays, this version of `unique_copy` returns the number of copied elements instead of an iterator that indicates the last copied element in the output range.

Table 7.3 shows our interpretation of the requirements for `unique_copy` from the C++ standard[28.6.9] [21] for the signature of Listing 7.2.

Requirement	Description
Unique Copy Size	The output range must be able to store the same number of elements as the input range.
Unique Copy Separation	The input range and the output range do not overlap.
Unique Copy Consecutive	Only the first element from every consecutive group of equal elements of the input range is copied into the output range.
Unique Copy Return	The algorithm returns the number of copied elements.
Unique Copy Complexity	At most $n - 1$ comparisons of adjacent elements are performed.

Table 7.3.: Requirements of `unique_copy`

Requirement **Unique Copy Consecutive** captures the core functionality of the algorithms. The intention of this requirement, which is not explicitly mentioned in **Unique Copy Consecutive**, is that the copied elements do *not* contain adjacent equal elements. One goal of this report is to show how this requirement can be expressed in Frama-C's specification language ACSL.

Complexity requirements, as formulated in terms of the number of comparison operations in requirement **Unique Copy Complexity**, are essential for specifying efficient algorithms. However, since Frama-C does currently not provide sufficient support for specifying this kind of requirements, we will not consider them in the rest of this document.

7.1.2. Some examples for `unique_copy`

In this section we analyze the requirement **Unique Copy Consecutive** by looking at how `unique_copy` behaves on different inputs.

Figure 7.4 shows the result of `unique_copy` when applied to a short array of integers. The arrows indicate from which index in the array a the respective value in the array b originates. The gray portion in the target array indicates that in our example not all elements from a have been copied.

Figure 7.4.: Example of applying `unique_copy`

Requirement **Unique Copy Consecutive** also implies that if `unique_copy` is applied to sequence that contains no adjacent equal elements in the first place, then it behaves like an ordinary copy algorithms (Figure 7.5).

Figure 7.5.: Applying `unique_copy` to a sequence with no adjacent equal elements

Another, somewhat extreme, example is applying `unique_copy` to a sequence where all elements are equal to each other. In this case, the result of `unique_copy` will consists of a single value (Figure 7.6).

Figure 7.6.: Applying `unique_copy` to a sequence where all elements are equal

A typical use case of `unique_copy` is to apply it to a *sorted* sequence. In this case calling `unique_copy` ensures that each value of the input range occurs exactly once in the output range (Figure 7.7). This is of course known to `Unix` programmers who can use `sort FILE | uniq` to remove all duplicate lines from `FILE`.

Figure 7.7.: Applying `unique_copy` to remove all duplicate elements from a sorted sequence

7.1.3. A first analysis of `unique_copy`

Figure 7.8 is a slight modification of Figure 7.4. We show here only the indices of the source array whose values are copied into the target array. In addition, we have added another (dashed) arrow to link the indices that correspond to the *one past the end* locations of the input and output ranges, respectively. We use this additional arrow in order to be able to describe all sub sequences of consecutive equal elements in the source array.

Figure 7.8.: Partitioning the input of `unique_copy`

These arrows between the indices of the array `b` and the array `a` define the following sequence p of eight indices where the index that points one past the end is underlined.

$$p = (0, 1, 3, 4, 5, 8, 9, \underline{11}) \qquad \text{for Figure 7.8}$$

For the other examples the corresponding index sequences are

$$p = (0, 1, 2, 3, 4, 5, 6, \underline{7}) \qquad \text{for Figure 7.5}$$
$$p = (0, \underline{7}) \qquad \text{for Figure 7.6}$$
$$p = (0, 1, 2, 7, \underline{11}) \qquad \text{for Figure 7.7}$$

Don't forget that the last three figures *do not show* the respective *one past the end* arrows.

More generally, we refer to the sequence p as *partitioning sequence* of `unique_copy` for the array `a[0.. n-1]`. This sequence is characterized by the following properties: If $m + 1$ is the **length of a partitioning sequence**, then we observe that they are **strictly monotone increasing**

$$0 = p_0 < \ldots < p_{m+1} = n \tag{7.1}$$

and that the right-open index intervals

$$[p_i, p_{i+1}) \qquad\qquad \forall i : 0 \le i < m$$

mark **consecutive ranges** of equal elements in the source array, that is,

$$a[p_i] = a[k] \qquad\qquad \forall k : p_i \le k < p_{i+1} \tag{7.2}$$

We also have that the consecutive ranges are **maximal** in the following sense

$$a[p_i] \ne a[p_{i+1}] \qquad\qquad \forall i : 0 \le i < m - 1 \tag{7.3}$$

and last but not least for the **result of `unique_copy`** it must hold

$$b[i] = a[p_i] \qquad\qquad \forall i : 0 \le i < m \tag{7.4}$$

7.2. Unit tests for `unique_copy`

In this section we derive both *test data* and *test code* that shall establish that the implementation of `unique_copy` from Listing 7.2 satisfies the requirements listed in Table 7.3. We are mainly concerned with *unit tests* that capture the functionality of `unique_copy`. This is also the reason why we are not discussing the also important issue of *code coverage*.

We are dealing in this section with tests of a *generic* implementation, that is, with an implementation that is parameterized over the type of the elements stored in arrays. Our test code is therefore also as generic as suitable, see Section 7.2.1 for example. Of course, the actual test execution appears with both specific type parameters and concrete test data.

We start our presentation in Section 7.2.2 with a discussion of suitable test data for `unique_copy`. Thereby we also have a look at the test data for `unique_copy` in *libcxx* [25], an open source implementation of the C++ standard library.

Requirement **Unique Copy Consecutive** captures the core of `unique_copy`, and the design of our tests aims particularly at checking this requirement. This is, however, not an easy undertaking because it is not clear in the beginning how to convincingly demonstrate that the *first* (and only the first) element of a consecutive group of equal elements is copied. Our first tests `unique_copy` in Sections 7.2.3 and 7.2.4 (and the corresponding tests in *libcxx*) therefore only show that there are no adjacent equal elements in the output array.

There is, however, a not too complicated method to show the copying of the first element of the consecutive ranges. As we will show in Section 7.2.5 this method relies both on

1. our semi-formal analysis of the `unique_copy` in Section 7.1.3 and

2. the *generic* nature of the implementation of `unique_copy`.

The latter allows us to replace values of type `T` essentially by `std::pair<T,size_t>` where the `second` field of `std::pair` will hold the index of the copied element in the input sequence of `unique_copy`.

7.2.1. Preparation of test code

In order to facilitate the testing of `unique_copy` from Listing 7.2, we provide a wrapper implementation that uses the container `vector` from the C++ standard library.

```cpp
template<typename T>
std::vector<T>
unique_copy(const std::vector<T>& a)
{
  std::vector<T> b(a.size());

  auto size = unique_copy(a.data(), a.size(), b.data());
  b.resize(size);

  return b;
}
```

Listing 7.9: `unique_copy` for `std::vector`

Using this generic auxiliary function from Listing 7.9 has several advantages for our tests.

- The `vector` container conveniently encapsulates both the memory and the number of elements of a C array.

- A `vector` hides many details of dynamic memory allocation from the user. Listing 7.2 shows that it is also easy to *resize* a `vector` object after it was created.

- The C++ standard ensures that different `vector` objects manage their own memory. Thus, using `vector`, it is easy to satisfy **Unique Copy Separation** which states that its two array arguments do not overlap.

- Also note that we initially declare the output `vector` to have the same size as the input `vector`. We are thus making sure that the requirement **Unique Copy Size** is satisfied.

7.2.2. Test data

Table 7.10 shows our initial test data for `unique_copy`. These are exactly the examples from Section 7.1.2 that have been used there to describe the behavior of `unique_copy`.

Input	Output	Reference
(4, 3, 3, 1, 3, 4, 4, 4, 2, 3, 3)	(4, 3, 1, 3, 4, 2, 3)	Figure 7.4
(4, 3, 1, 3, 4, 2, 3)	(4, 3, 1, 3, 4, 2, 3)	Figure 7.5
(4, 4, 4, 4, 4, 4, 4)	(4)	Figure 7.6
(1, 2, 3, 3, 3, 3, 3, 4, 4, 4)	(1, 2, 3, 4)	Figure 7.7

Table 7.10.: Initial test data

Boundary test data are data for extreme inputs of the specific algorithm. It can be argued that the example from Figure 7.6 where the elements of the input array equal *one* value represents boundary test data. Table 7.11 shows the expected behavior of `unique_copy` for input ranges of size 0 and size 1.

Input	Output	Reference
( )	( )	empty input range
(3)	(3)	one-element input range

Table 7.11.: Boundary test data

It is interesting to compare our test data with those from the functional tests for `unique_copy` in an open source implementation of the C++ standard library [25, `unique_copy.pass.cpp`]. We have listed these test data in Table 7.12.

Input	Output
(0, 1, 2, 2, 4)	(0, 1, 2, 4)
(0)	(0)
(0, 1)	(0, 1)
(0, 0)	(0)
(0, 0, 1)	(0, 1)
(0, 0, 1, 0)	(0, 1, 0)
(0, 0, 1, 1)	(0, 1)
(0, 0, 1)	(0, 1)
(0, 1, 1, 1, 2, 2, 2)	(0, 1, 2)

Table 7.12.: Test data for `unique_copy` from an open source test suite

Here the emphasis is on an arguably more systematic presentation of small input ranges of sizes. Interestingly, however, there is no test case for the empty range.

7.2.3. A basic test

When we present *test code* then we present code that shows that the function under test satisfies certain *properties* which in turn are justified by the requirements.

Listing 7.13, for example, contains code that tests that the elements copied by `unique_copy` contain no adjacent equal elements. This is, as we have explained in Section 7.1.2, a simple consequence of **Unique Copy Consecutive** from Table 7.3.

```cpp
template<typename T>
std::vector<T>
unique_copy_basic_test(const std::vector<T>& input)
{
  auto result = unique_copy(input);
  assert(std::adjacent_find(result.begin(), result.end()) == result.end());

  //std::cout << "test " << __func__ << " succeeded " << std::endl;

  return result;
}
```

Listing 7.13: Testing the absence of adjacent equal elements

The key ingredient of the test in Listing 7.13 consists in calling the C++ standard library function `adjacent_find` which searches its input range for (the first) occurrence of two consecutive equal elements. If there is no such occurrence, `adjacent_find` returns the iterator that indicates the end of the range.

Listing 7.14 shows how the test from Listing 7.13 is executed.

```cpp
int main(int argc, char** argv)
{
  assert(argc == 2);
  std::fstream file(argv[1]);
  std::vector<int> v;

  while (true) {
    file >> v;
    if (file) {
      // std::cout << v << std::endl;
      unique_copy_basic_test(v);
      v.clear();
    }
    else {
      break;
    }
  }

  std::cout << "\tsuccessful execution of " << argv[0] << "\n";

  return EXIT_SUCCESS;
}
```

Listing 7.14: Test execution code for the absence of adjacent equal elements

In our setting the test data step from the file in Listing 7.15 which contains the input data from Tables 7.10, 7.11, and 7.12.

```
    (4, 3, 3, 1, 3, 4, 4, 4, 2, 3, 3)
    (4, 3, 1, 3, 4, 2, 3)
    (4, 4, 4, 4, 4, 4, 4)
    (1, 2, 3, 3, 3, 3, 3, 4, 4, 4)

    (3)
    ()

    (0, 1, 2, 2, 4)
    (0)
    (0, 1)
    (0, 0)
    (0, 0, 1)
    (0, 0, 1, 0)
    (0, 0, 1, 1)
    (0, 0, 1)
    (0, 1, 1, 1, 2, 2, 2)
```

Listing 7.15: Test input data for `unique_copy`

7.2.4. Extending the basic test

This basic test can be extended to the slightly more elaborate test in Listing 7.16 that in addition to Listing 7.14 compares whether

1. the number of copied elements equals the size of the expected output range and

2. the copied elements actually equal the expected output range.

```cpp
template<typename T>
void
unique_copy_compare_test(const std::vector<T>& input,
                         const std::vector<T>& expected)
{
  auto result = unique_copy_basic_test<T>(input);
  assert(result == expected);

  //std::cout << "test " << __func__ << " succeeded " << std::endl;
}
```

Listing 7.16: Comparing with expected result

Listing 7.17 shows how the test from Listing 7.16 is executed with the test data read from a file.

```cpp
int main(int argc, char** argv)
{
  assert(argc == 2);
  std::fstream file(argv[1]);

  while (true) {
    std::vector<int> input, expected;
    file >> input;
    file >> expected;
    if (file) {
      // std::cout << input    << "\t" << expected << std::endl;
      unique_copy_compare_test(input, expected);
    }
    else {
      break;
    }
  }

  std::cout << "\tsuccessful execution of " << argv[0] << "\n";

  return EXIT_SUCCESS;
}
```

Listing 7.17: Test execution code for comparing with expected result

In this setting, the test data step from the file in Listing 7.18 which contains the input data and the expected output data from Tables 7.10, 7.11, and 7.12.

```
(4, 3, 3, 1, 3, 4, 4, 4, 2, 3, 3)    (4, 3, 1, 3, 4, 2, 3)
(4, 3, 1, 3, 4, 2, 3)                 (4, 3, 1, 3, 4, 2, 3)
(4, 4, 4, 4, 4, 4, 4)                 (4)
(1, 2, 3, 3, 3, 3, 3, 4, 4, 4         (1, 2, 3, 4)

()                                    ()
(3)                                   (3)

(0, 1, 2, 2, 4)                       (0, 1, 2, 4)
(0)                                   (0)
(0, 1)                                (0, 1)
(0, 0)                                (0)
(0, 0, 1)                             (0, 1)
(0, 0, 1, 0)                          (0, 1, 0)
(0, 0, 1, 1)                          (0, 1)
(0, 0, 1)                             (0, 1)
(0, 1, 1, 1, 2, 2, 2)                 (0, 1, 2)
```

Listing 7.18: Test input and expected output for `unique_copy`

Executing the test from Listing 7.16 with the test data from Section 7.2.2 allows us to establish whether `unique_copy` satisfies the Requirement **Unique Copy Size** and the above mentioned consequence of Requirement **Unique Copy Consecutive**. However, these tests cannot establish that the *first* (and only the first) element of each consecutive group of equal elements is copied. In this sense, our test is as expressive as the tests for `unique_copy` from [25, `unique_copy.pass.cpp`]

7.2.5. Partition testing

In Section 7.1.3 we had, informally, argued that for each sequence $a = (a_0, \ldots, a_{n-1})$ there is a *partitioning sequence* $p = (p_0, \ldots, p_m)$ that satisfies the conditions (7.1), (7.2), and (7.3) and which, moreover, relates the input and output of `unique_copy` according to Equation (7.4). The term *partition testing* refers here to using these properties in the design of our tests.

The problem is that the partitioning sequence does not *explicitly* occur in `unique_copy`. There is, however, a relatively simple and natural way to make the partitioning sequence explicitly.

We explain the basic idea at hand of the input sequence

$$a = (4, 3, 3, 1, 3, 4, 4, 4, 2, 3, 3)$$

from Figure 7.8. We begin with creating a sequence of pairs (a_i, i) consisting of the original value a_i and its index i, in other words, we *zip* the sequence a with the sequence of its indices. In order to facilitate the readability of lists of pairs we also write $\binom{a_i}{i}$ instead of (a_i, i).

For our example, the augmented sequence of pairs a' reads

$$a' = \left(\binom{4}{0}, \binom{3}{1}, \binom{3}{2}, \binom{1}{3}, \binom{3}{4}, \binom{4}{5}, \binom{4}{6}, \binom{4}{7}, \binom{2}{8}, \binom{3}{9}, \binom{3}{10} \right)$$

If we define the equality of two pairs (a_i, i) and (a_j, j) by the equality of its first components, then `unique_copy` *piggybacks* the original index of an element into the result. In other words, `unique_copy` applied to the sequence a' produces the following sequence

$$b' = \left(\binom{4}{0}, \binom{3}{1}, \binom{1}{3}, \binom{3}{4}, \binom{4}{5}, \binom{2}{8}, \binom{3}{9} \right)$$

where the second components indicate the index in the input sequence.

Finally, we extract the two lists of its first and second components from b' (that is we *unzip* b') and add a final element 11 (the number of elements of a) to the sequence of indices. We thus obtain the sequence

$$b = (4, 3, 1, 3, 4, 2, 3)$$

which is just the result of applying `unique_copy` to a, and the partitioning sequence

$$p = (0, 1, 3, 4, 5, 8, 9, 11)$$

Both sequences can, of course be found in Figure 7.8.

In the following subsections, we discuss the details of the implementation of our partitioning test.

Pairs of values and indices

Listing 7.19 shows our definition of a type for *indexed values* which consists of a pair of generic type `T` and `size_t` as index type. We implement this type with the generic class `std::pair` that has two public fields `first` and `second`, respectively.

The equality operation (`operator`==) of this new type has been defined in such a way that only the first component of the underlying pair type is evaluated. Note that we only provide a special implementation for the equality operation but not for copying and assigning a value p of type `Indexed<T>`. This is essential for keeping the values `p.first` and `p.second` unchanged while they are processed by our generic `unique_copy`.

```
template<typename T>
struct Indexed : public std::pair<T, size_t> {

  // inherit constructors of base class
  using std::pair<T, size_t>::pair;


};

template<typename T>
bool operator==(const Indexed<T>& a, const Indexed<T>& b)
{
  return a.first == b.first;
}
```

Listing 7.19: A type for indexed values

Creating the partition sequence

Listing 7.20 shows our implementation of computing the partitioning sequence of a given input sequence.

```
template<typename T>
std::pair<std::vector<T>, std::vector<size_t>>
    unique_copy_create_partition(const std::vector<T>& input)
{
  std::vector<Indexed<T>> zipped(input.size());
  for (size_t i = 0; i < input.size(); i++) {
    zipped[i] = Indexed<T>(input[i], i);
  }

  auto output = unique_copy(zipped);

  std::pair<std::vector<T>, std::vector<size_t>> unzipped;
  unzipped.first.resize(output.size());
  unzipped.second.resize(output.size() + 1);

  for (size_t i = 0; i < output.size(); ++i) {
    unzipped.first[i]  = output[i].first;
    unzipped.second[i] = output[i].second;
  }
  unzipped.second.back() = input.size();

  return unzipped;
}
```

Listing 7.20: Creating the partition underlying `unique_copy`

1. As explained at the beginning of this section we start with creating the sequence of pairs of values and their respective indices.

2. We then pass this sequence of pairs to our generic version of `unique_copy` for `vector` from Listing 7.9. The resulting sequence contains, thanks to our definition of equality of our pair type, the values with their indices in the original sequence.

3. Finally, we *unzip* the vector of pairs into a pair of vectors and add to the index vector the size of the input sequence as final element.

141

Testing the partition properties

Listing 7.21 shows our generic partition test. We have added comments to facilitate the tracing of our test code to the properties (7.1)–(7.4) from Section 7.1.3.

First `unique_copy` is called and the corresponding partition sequence is computed. Note that computing of this partition sequence naturally also leads to a second computation of the result of `unique_copy`. Thus, initially we check (for sanity) that both computations have the same result.

```cpp
template<typename T>
void unique_copy_partition_test(const std::vector<T>& a)
{
  auto b = unique_copy(a);
  auto unzipped = unique_copy_create_partition(a);
  assert(unzipped.first == b);

  const std::vector<size_t>& p = unzipped.second;

  // partition sequence is one element longer than output array
  assert(p.size() == b.size() + 1);

  // monotonicity (first and last element only)
  assert(p.front() == 0);
  assert(p.back()  == a.size());

  for (size_t i = 0; i < b.size(); ++i) {
    // consider i-th segment of the partition
    auto begin = p[i];
    auto end   = p[i + 1];

    // monotonicity
    assert(begin < end);

    // consecutive range of equal elements
    for (size_t k = begin; k < end; ++k) {
      assert(a[begin] == a[k]);
    }

    // maximal consecutive range of equal elements
    if (i + 1 < b.size()) {
      assert(a[begin] != a[end]);
    }

    // result of unique_copy
    assert(b[i] == a[begin]);
  }
}
```

Listing 7.21: Partition testing of `unique_copy`

We then check whether the first element of partitioning sequence equals 0 and whether the last element equals the size of the input sequence. This is, of course, in accordance with the chain of (in)equalities from Relation 7.1. Finally, we check for each partition segment $[p_i, p_{i+1})$ the rest of the monotonicity conditions from Relation (7.1), and then proceed to verify the properties (7.2), (7.3), and (7.4).

Listing 7.22 shows the code for executing partition tests with test data read from a file. As test data we use again the inputs from Listing 7.15.

```cpp
int main(int argc, char** argv)
{
  assert(argc == 2);
  std::fstream file(argv[1]);

  while (true) {
    std::vector<int> v;
    file >> v;
    if (file) {
      // std::cout << v << std::endl;
      unique_copy_partition_test(v);
    }
    else {
      break;
    }
  }

  std::cout << "\tsuccessful execution of " << argv[0] << "\n";

  return EXIT_SUCCESS;
}
```

Listing 7.22: Test execution code for partition tests

7.3. Formal verification of `unique_copy`

In this Section we discuss the formal verification of `unique_copy` with Frama-C/WP. The first issue is that while `std::unique_copy` is implemented in C++ and heavily relies on C++ templates, Frama-C/WP can only deal with C functions. For this reason we present in Section 7.3.1 an implementation of `unique_copy` in C.

We discuss our first and simplest version of an ACSL specification of `unique_copy` in Section 7.3.2. The main idea of a so-called *minimal contract* is that our formal specification is just strong enough to verify the absence of certain undefined behaviors, such as illegal memory accesses and integer overflows. More specifically this means that our minimal contract for `unique_copy` formalizes **Unique Copy Size** and **Unique Copy Separation** but only partially formalizes **Unique Copy Return**. The formalization of the core requirement **Unique Copy Consecutive** is not addressed at all. Still, the verification of a minimal contract is meaningful since the addressed undefined behaviors are often a cause for security vulnerabilities. The verification of the *absence* of these undesirable behaviors can play an important role for ensuring the robustness of software.

In Section 7.3.3, we extend the minimal contract from Section 7.3.2 by a postcondition that states that the output range of `unique_copy` will not contain any adjacent equal elements. In other words, the new contract also partially addresses **Unique Copy Consecutive**.

Finally, in Section 7.3.4 we present a further extension of our contract that also captures the missing aspects of **Unique Copy Return** and **Unique Copy Consecutive** in the specification of `unique_copy`. Here, we build on top of our analysis of the so-called *partitioning sequence* from Section 7.1.3.

In order to differentiate the those different variants of formal verification we number the functions. For example the functions `unique_copy2` and `unique_copy3` share basically the same implementation but have different formal annotations.

7.3.1. Reformulation of the algorithms in C

Our reformulation of `unique_copy` in the C programming language has a signature that can be considered as a specialisation of our simplified generic implementation of Listing 7.2. Instead of the generic type parameter `T`, however, we employ the integer type alias `value_type`. We also replace the standard unsigned integer type `size_t` by the type alias `size_type`. More information can be found in Section 1.3.

The basic idea of our C-implementation of `unique_copy` in Listing 7.23 is to traverse the input array `a[0...n-1]` and copy an element `a[i]` to the output array `b[0..n-1]` whenever it has been detected that it is different from its predecessor `a[i-1]`. Assuming a non-empty array, the implementation starts with copying `a[0]` to `b[0]`. In order to detect whether the current value `a[i]` is different from its predecessor we compare it with the most recently copied value `b[k]`.

```
size_type
unique_copy(const value_type* a, size_type n, value_type* b)
{
  if (n == 0u) {
    return n;
  }
  else {
    size_type k = 0u;
    b[k] = a[0];

    for (size_type i = 1u; i < n; ++i) {
      const value_type val = a[i];

      if (b[k] != val) {
        b[++k] = val;
      }
    }

    return ++k;
  }
}
```

Listing 7.23: Re-implementation of `unique_copy` in C

Note that the test code in Section 7.2 has be designed in such a way that it can be applied also to a function with the signature in Listing 7.23.

7.3.2. A minimal contract for `unique_copy`

When we talk about a *minimal contract* of a function we mean a small contract that covers only basic properties. One might, for example, only be interested that during the execution of a function no runtime errors such as arithmetic overflows or invalid pointer accesses occur. Since many software security problems are caused by undetected runtime errors, minimal contracts can help to achieve a higher degree of quality assurance. This means that our minimal contract for `unique_copy` formalizes the requirements **Unique Copy Size** and **Unique Copy Separation** but only partially formalizes **Unique Copy Return**. The formalization of the core requirement **Unique Copy Consecutive** is not addressed at all.

Formal specification

Listing 7.24 shows the specification of our minimal contract. We have *labeled* the various preconditions and postconditions of our contract by names, e.g., we use the label `sep` in order to refer to our formal specification of **Unique Copy Separation**. Using these user-supplied labels simplifies the documentation of contracts and can also be helpful during the process of formal verification. In the following we often refer to the various formal properties in a contract by their labels.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep:   \separated(a + (0..n-1), b + (0..n-1));

  assigns  b[0..n-1];

  ensures result:    0 <= \result <= n;
  ensures unchanged: Unchanged{Old, Here}(b, \result, n);
  ensures unchanged: Unchanged{Old, Here}(a, n);

*/
size_type
unique_copy2(const value_type* a, size_type n, value_type* b);
```

Listing 7.24: A "minimal" contract for unique_copy

Using the built-in predicates \valid and \valid_read, the preconditions valid state that

1. the elements of the input range a[0..n-1] can be safely accessed for reading,

2. whereas the elements of the output range b[0..n-1] can be safely accessed both for reading and writing.

Note that in accordance with **Unique Copy Size** both ranges have the same size. The informal specification also states in **Unique Copy Separation** that the input and output ranges do not overlap. This precondition is expressed by our property sep which in turn uses the built-in predicate \separated.

The assigns clause of our minimal contract states that unique_copy can only modify the array b [0..n-1]. Note that this requirement on the side effects of unique_copy cannot be found in the requirements in Table 7.3. This can be attributed to the fact that little is known about internal side effects of the generic type parameter T. In our more specific situation with the concrete type value_type we can use the means of ACSL to restrict the side effects allowed by our contract.

The postcondition result describes the numerical range for the return value of unique_copy. In other words, our minimal contract only provides a rather coarse estimation of **Unique Copy Return** for the number of elements copied by unique_copy. Note the use the ACSL keyword \result in this postcondition to refer to the return value of function. Also note that we have employed a *chained inequality* instead of writing

```
    0 <= \result && \result <= n
```

This is a nice, little feature that helps writing compact contracts. There is a second postcondition unchanged that is formulated using the *user defined* ACSL predicate Unchanged that we show in Listing 6.1. This predicate comes in the form of two overloaded versions. The first one is defined for an array section wheres the second one only requires the length of the array. The arguments K and L of the predicate are *labels* that represent *program states*. The predicate Unchanged says the respective elements of the array have the same value in state K and state L.

We use the predicate `Unchanged` in order to make the assigns clause a bit more precise. Using the predefined labels `Old`, which refers to the pre-state of the contract, and `Post`, which refers to the post-state of the contract, the postcondition `unchanged` says that element of the range `b[\result..n-1]` are the same before and after `unique_copy` has been called. All this would easier if we could use just the assigns clause

```
assigns b[0..\result-1];
```

since this would imply our postcondition

```
ensures unchanged: Unchanged{Old, Here}(b, \result, n);
```

We remark here that the ACSL documentation does not forbid the use of `\result` outside of `ensures` clauses [16, p. 30]. While Frama-C/WP does not reject it either, the corresponding proof obligations are, in any case, not verified.

Graphical presentation of the minimal contract

Figure 7.25 is an attempt to graphically represent the minimal contract from Listing 7.24.

Figure 7.25.: Representation of a minimal contract for `unique_copy`

In contrast to Figure 7.4, it is not indicated which elements of the input array have been copied to which elements of the output array. This is because, it has not been specified, whether any element at all has been copied. On the other hand, the minimal contract ensures that the part of the output array, that is not needed to hold the result, is kept unchanged.

Annotating the implementation

The verification of the contract in Listing 7.24 requires that we add appropriate *loop annotations* to the implementation in Listing 7.26. Among these annotations are so-called *loop invariants*, which are formulas that must hold at the beginning of each loop iteration.

```
size_type
unique_copy2(const value_type* a, size_type n, value_type* b)
{
  if (n == 0u) {
    return n;
  }
  else {
    size_type k = 0u;
    b[k] = a[0];

    /*@
      loop invariant bound:      0 <= k < i <= n;
      loop invariant unchanged: Unchanged{Pre, Here}(b, k+1, n);
      loop assigns i, k, b[0..n-1];
      loop variant n-i;
    */
    for (size_type i = 1u; i < n; ++i) {
      const value_type val = a[i];

      if (b[k] != val) {
        b[++k] = val;
      }
    }

    return ++k;
  }
}
```

Listing 7.26: Annotations for the minimal contract

We now have a closer look at the loop annotations.

- The loop invariant `bound` states that the index `k` is always less than `i` that both are limited from below and above by `0` and the array length `n`, respectively.

- Similar to the postcondition `unchanged` in Listing 7.24 the loop invariant `unchanged` states that in each iteration of the loop the values in the range `b[k+1..n]` will be the same as in the pre-state of `unique_copy`. Note the use of the predefined label `Pre` to denote the program state before the function was called.

- We also need a *loop assigns clause* which lists all memory locations that can be changed during the execution of the loop. These memory locations comprise not only the array `b[0..n-1]` but also the local variables `i` and `k`.

- Finally, we have a *loop variant* which must contain a positive value that is decreased in each loop iteration. Loop variants serve to verify the *termination* of loops.

Understanding the verification of minimal contracts

The WP plugin [26] of Frama-C (in short also Frama-C/WP) is activated by using the option `-wp`. However, before Frama-C/WP starts generating and discharging proof obligations, the Frama-C kernel produces a *normalized version* of the source code. Most Frama-C plugins, including Frama-C/WP, use this semantically equivalent presentation to conduct their respective analyses.

```
/*@ requires valid: \valid_read(a + (0 .. n - 1));
    requires valid: \valid(b + (0 .. n - 1));
    requires sep: \separated(a + (0 .. n - 1), b + (0 .. n - 1));
    ensures result: 0 <= \result <= \old(n);
    ensures unchanged: Unchanged{Old, Here}(\old(b), \result, \old(n));
    ensures unchanged: Unchanged{Old, Here}(\old(a), \old(n));
    assigns *(b + (0 .. n - 1));
 */
size_type
unique_copy2(value_type const *a, size_type n, value_type *b)
{
  size_type __retres;

  if (n == 0u) {
    __retres = n;
    goto return_label;
  }
  else {
    size_type k = 0u;
    *(b + k) = *(a + 0);
    {
      size_type i = 1u;

      /*@ loop invariant bound: 0 <= k < i <= n;
          loop invariant unchanged: Unchanged{Pre, Here}(b, k + 1, n);
          loop assigns i, k, *(b + (0 .. n - 1));
          loop variant n - i;
      */
      while (i < n) {
        {
          value_type const val = *(a + i);

          if (*(b + k) != val) {
            k ++;
            *(b + k) = val;
          }
        }
        i ++;
      }
    }
    k ++;
    __retres = k;
    goto return_label;
  }

return_label:
  return __retres;
}
```

Listing 7.27: Normalized presentation of the minimal contract

```
/*@ requires valid: \valid_read(a + (0 .. n - 1));
    requires valid: \valid(b + (0 .. n - 1));
    requires sep: \separated(a + (0 .. n - 1), b + (0 .. n - 1));
    ensures result: 0 <= \result <= \old(n);
    ensures unchanged: Unchanged{Old, Here}(\old(b), \result, \old(n));
    ensures unchanged: Unchanged{Old, Here}(\old(a), \old(n));
    assigns *(b + (0 .. n - 1));
 */
size_type
unique_copy2(value_type const *a, size_type n, value_type *b)
{
  size_type __retres;

  if (n == 0u) {
    __retres = n;
    goto return_label;
  }
  else {
    size_type k = 0u;
    /*@ assert rte: mem_access: \valid(b + k); */
    /*@ assert rte: mem_access: \valid_read(a + 0); */
    *(b + k) = *(a + 0);
    {
      size_type i = 1u;

      /*@ loop invariant bound: 0 <= k < i <= n;
          loop invariant unchanged: Unchanged{Pre, Here}(b, k + 1, n);
          loop assigns i, k, *(b + (0 .. n - 1));
          loop variant n - i;
      */
      while (i < n) {
        {
          /*@ assert rte: mem_access: \valid_read(a + i); */
          value_type const val = *(a + i);

          /*@ assert rte: mem_access: \valid_read(b + k); */
          if (*(b + k) != val) {
            /*@ assert rte: unsigned_overflow: k + 1 <= 4294967295; */
            k ++;
            /*@ assert rte: mem_access: \valid(b + k); */
            *(b + k) = val;
          }
        }
        /*@ assert rte: unsigned_overflow: i + 1 <= 4294967295; */
        i ++;
      }
    }
    /*@ assert rte: unsigned_overflow: k + 1 <= 4294967295; */
    k ++;
    __retres = k;
    goto return_label;
  }

return_label:
  return __retres;
}
```

Listing 7.28: Normalized presentation of the minimal contract with RTE assertions

Listing 7.27 shows the normalized version of the formal specification from Listing 7.24 and the annotated implementation from Listing 7.26.

In order to extract the normalized version, which is also shown in Frama-C GUI, one can use the option `-print`. One of the most visible differences between the original and the normalized form is that *for loops* are represented as *while loops*. For more details on the normalization process we refer to the respective section of the Frama-C manual [27, §5.3].

Using the option `-wp-rte`, the Frama-C/WP plugin allows to generate additional assertions that are placed as guards before potentially dangerous C constructs, such as pointer dereferencing of integer operations that might overflow. Listing 7.28 shows these additional assertions in the normalized version of `unique_copy` when using the options `-wp-rte -warn-`**unsigned**`-overflow -warn-`**unsigned**`-downcast`. For more details how to customize the generation of RTE (runtime error) guards we refer to the respective manuals [26, 28].

Verifying the minimal contract with the additional run time error assertions essentially shows that a large class of undefined behaviors cannot occur *if* the preconditions of the contract are satisfied. Since undefined behaviors often represent security vulnerabilities, even the verification of the minimal contract can, thus, provide significant evidence that the execution of a function such as `unique_copy` cannot cause security weaknesses.

7.3.3. A more elaborate contract for `unique_copy`

After using the minimal contract to prove the absence of undefined behavior we now show that there are no equal neighbors in the output array b[0..\result-1]. This property reflects an important consequence of **Unique Copy Consecutive**. It does, however, neither express the fact that only the first element of every consecutive range within the input array is copied nor does it state that the elements in the output range are at all related to those from the input range.

In order to formalize this new property we use the new predicate `HasEqualNeighbors` which we show in Listing 3.12. The predicate states that there exists an element in the range a[0..n-1] which is equal to its direct successor.

Formal specification

Listing 7.29 is an extension of our minimal contract. We keep all properties but also add the new postcondition `solitary` to our contract.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep:   \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..n-1];

  ensures result:    0 <= \result <= n;
  ensures solitary:  !HasEqualNeighbors (b, \result);
  ensures unchanged: Unchanged{Old, Here}(b, \result, n);
  ensures unchanged: Unchanged{Old, Here}(a, n);
 */
size_type
unique_copy3(const value_type* a, size_type n, value_type* b);
```

Listing 7.29: A more elaborate contract for `unique_copy`

In order to formally express the new postcondition we use the negation of `HasEqualNeighbors` from Listing 3.12.

Graphical presentation of the more elaborate contract

Figure 7.30 is an attempt to graphically represent the more elaborate specification from Listing 7.29. Compared to Figure 7.25 our new figure highlights that neighbouring elements of the output array are not equal. Our figure, however, still not indicates which elements of the input array have been copied to which elements of the output array.

Figure 7.30.: Representation of a more elaborate contract for `unique_copy`

Annotating the implementation

In order to support the verification of the postcondition `unique` from Listing 7.29 we add the loop invariant `solitary` to our implementation 7.31. This loop invariant states that in each loop iteration there are no adjacent equal elements in the range `b[0..k]` of already copied elements. We use the predicate `HasEqualNeighbors` to formally describe this property.

```
size_type
unique_copy3(const value_type* a, size_type n, value_type* b)
{
  if (n == 0u) {
    return n;
  }
  else {
    size_type k = 0u;
    b[k] = a[0];

    /*@
      loop invariant bound:     0 <= k < i <= n;
      loop invariant solitary:  !HasEqualNeighbors(b, k+1);
      loop invariant unchanged: Unchanged{Pre, Here}(b, k+1, n);
      loop assigns i, k, b[0..n-1];
      loop variant n-i;
    */
    for (size_type i = 1; i < n; ++i) {
      const value_type val = a[i];

      if (b[k] != val) {
        b[++k] = val;
      }
    }

    return ++k;
  }
}
```

Listing 7.31: Annotations for a more elaborate contract of `unique_copy`

7.3.4. A complete contract for `unique_copy`

In this section we finally tackle the issue of formalizing the requirements of **Unique Copy Consecutive** in ACSL. The main idea is that we

1. capture in ACSL the properties of the partitioning sequence from Section 7.1.3

2. use these properties to specify and verify `unique_copy`

Formalizing the number of elements copied by `unique_copy`

The logic function `UniqueSize` in Listing 7.32 computes the number of elements that are to be copied by `unique_copy` from an array `a[0..n-1]` In other words, `UniqueSize` represents the number *m* in the inequalities (7.1) of consecutive sub-ranges of equal elements. The Listing 7.32 contains explicit *recursive* definition of `UniqueSize` and the corresponding implicit definition through the lemmas.

```
/*@
  axiomatic UniqueSizeAxiomatic
  {
    logic integer UniqueSize(value_type* a, integer n) =
      n <= 0 ? 0 : (n == 1 ? 1 : UniqueSize(a, n-1) + (a[n-1] == a[n-2] ? 0 : 1));

    lemma UniqueSizeEmpty:
      \forall value_type *a, integer n;
        n <= 0  ==>  UniqueSize(a, n) == 0;

    lemma UniqueSizeOne:
      \forall value_type *a;
        UniqueSize(a, 1) == 1;

    lemma UniqueSizeEqual:
      \forall value_type *a, integer n;
        0 < n ==> a[n-1] == a[n] ==> UniqueSize(a, n+1) == UniqueSize(a, n);

    lemma UniqueSizeDiffer:
      \forall value_type *a, integer n;
        0 < n ==> a[n-1] != a[n] ==>  UniqueSize(a, n+1) == UniqueSize(a, n) + 1;

    lemma UniqueSizeRead{K,L}:
      \forall value_type *a, integer n, i;
        Unchanged{K,L}(a, n) ==> UniqueSize{K}(a, n) == UniqueSize{L}(a, n);
  }
*/
```

Listing 7.32: The logic function `UniqueSize`

Note that the definition of `UniqueSize` also covers arrays of negative size and in general ignore whether the involved pointers can be dereferenced. The issue here is that the function `UniqueSize` must be defined as a *total function* regardless of the fact whether the involved function arguments make sense in C-code. For more details we refer to the rules for logic definitions in the description of ACSL [16, §2.2.2].

The following ACSL lemmas `UniqueSizeBound` (Listing 7.33) formulates some simple bounds on the number of copied elements. As trivial as these inequalities might look like, their not too complicated proofs rely on mathematical induction. Since automatic theorem provers are often not capable of performing induction proofs, we have proven this lemma with the interactive theorem prover Coq.

```
/*@
  lemma UniqueSizeBound:
    \forall value_type *a, integer n;
      0 <= n  ==>  0 <= UniqueSize(a, n) <= n;
*/
```

Listing 7.33: Lemma `UniqueSizeBound`

Formalizing the properties of the partitions of `unique_copy`

The function `UniquePartition`, whose axiomatic definition is given in Listing 7.34, defines the partitioning sequence p from Section 7.1.3. Before we begin to relate the axioms from Listing 7.34 to the formulas from Section 7.1.3 we want to remind the reader that logic functions (and predicates) must be total that is they must be defined for all possible argument values.

- The monotonicity conditions (7.1) are described by the axioms `UniquePartitionEmpty`, `UniquePartitionLeft`, `UniquePartitionRight` and `UniquePartitionMonotone`.

- Equation (7.2) is represented by the axiom `UniquePartitionSegment`.

- Inequality (7.3) is described by axiom `UniquePartitionMaximal`.

- Axiom `UniquePartitionEqual` expresses that the value of `UniquePartition(a, n, i)` does not depend on the size of the array.

- Axiom `UniquePartitionRead`, finally states that `UniquePartition` is independent from the particular programme state in which it is used—as long as the respective array elements are equal in both states.

```
/*@
  axiomatic UniquePartitionAxiomatic
  {
    logic integer
      UniquePartition(value_type* a, integer n, integer i) reads a[0..n-1];

    axiom UniquePartitionEmpty:
      \forall value_type *a, integer n, i;
        n <= 0  ==>  UniquePartition(a, n, i) == 0;

    axiom UniquePartitionLeft:
      \forall value_type *a, integer n, i;
        0 < n  ==>  i <= 0  ==>  UniquePartition(a, n, i) == 0;

    axiom UniquePartitionRight:
      \forall value_type *a, integer n, i;
        0 < n  ==>  UniqueSize(a, n) <= i  ==>  UniquePartition(a, n, i) == n;

    axiom UniquePartitionMonotone:
      \forall value_type *a, integer n, i, j;
        0 <= i < j <= UniqueSize(a, n)  ==>
        UniquePartition(a, n, i) < UniquePartition(a, n, j);

    axiom UniquePartitionSegment:
      \forall value_type *a, integer n, i, k;
        0 <= i < UniqueSize(a, n) ==>
        ConstantRange(a, UniquePartition(a, n, i), UniquePartition(a, n, i+1));

    axiom UniquePartitionMaximal:
      \forall value_type *a, integer n, i;
        0 <= i < UniqueSize(a, n) - 1  ==>
        a[UniquePartition(a, n, i)] != a[UniquePartition(a, n, i+1)];

    axiom UniquePartitionEqual:
      \forall value_type *a, integer n, m, i;
        n < m  ==>  0 <= i < UniqueSize(a, n)  ==>
        UniquePartition(a, n, i) == UniquePartition(a, m, i);

    axiom UniquePartitionRead{K,L}:
      \forall value_type *a, integer n, i;
        Unchanged{K,L}(a, n) ==>
          UniquePartition{K}(a, n, i) == UniquePartition{L}(a, n, i);
  }
*/
```

Listing 7.34: Axiomatic description of the function `UniquePartition`

With the definitions of the logic functions `UniqueSize` and `UniquePartition` we can now formulate the ACSL predicate `Unique` from Listing 7.35. This predicate reflects Equation (7.4) and therefore will serve a prominent role in our complete contract of `unique_copy`.

```
/*@
  predicate
    Unique(value_type* a, integer n, value_type* b) =
      \forall integer k; 0 <= k < UniqueSize(a, n) ==>
        b[k] == a[UniquePartition(a, n, k)];
*/
```

Listing 7.35: The predicate `Unique`

156

Before we turn, however, our attention to the contract of `unique_copy` we show in Listing 7.36 a couple of simple ACSL lemmas that will be helpful in verifying the new contract.

```
/*@
  lemma UniquePartitionZero:
    \forall value_type *a, integer n;
      UniquePartition(a, n, 0) == 0;

  lemma UniquePartitionLowerBound:
    \forall value_type *a, integer n, i;
      0 < n ==> 0 <= i < UniqueSize(a, n) ==> 0 <= UniquePartition(a, n, i);

  lemma UniquePartitionUpperBound:
    \forall value_type *a, integer n, i;
      0 < n ==> 0 <= i < UniqueSize(a, n) ==> UniquePartition(a, n, i) < n;

  lemma UniquePartitionDiffer:
    \forall value_type *a, integer i, k, n;
      UniquePartition(a, n, k-1) < i <= UniquePartition(a, n, k) ==>
      a[i-1] != a[i]  ==>  i == UniquePartition(a, n, k);
*/
```

Listing 7.36: Some lemmas regarding `UniquePartition`

Formal specification

Listing 7.37 shows how we use the predicate `Unique` in the postcondition `unique` in order to formally specify **Unique Copy Consecutive** for `unique_copy`.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep:   \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..n-1];

  ensures bound:     0 <= \result <= n;
  ensures size:      \result == UniqueSize(a, n);
  ensures unique:    Unique(a, n, b);
  ensures solitary:  !HasEqualNeighbors (b, \result);
  ensures unchanged: Unchanged{Old, Here}(a, n);
  ensures unchanged: Unchanged{Old, Here}(b, \result, n);
*/
size_type
unique_copy4(const value_type* a, size_type n, value_type* b);
```

Listing 7.37: A complete contract for `unique_copy`

Note that we use the Lemma `UniqueImpliesNoEqualNeighbors` from Listing 7.38 in order to establish the postcondition `solitary` from the postcondition `unique`. Thanks to this lemma we will be able to omit the dedicated loop annotation `solitary` in Listing 7.40.

```
/*@
  lemma UniqueImpliesNoEqualNeighbors:
    \forall value_type *a, *b, integer n;
       Unique(a, n, b)  ==>  !HasEqualNeighbors(b, UniqueSize(a, n));
*/
```

Listing 7.38: The predicate `UniqueImpliesNoEqualNeighbors`

We mention here already lemma `UniquePreserve` which we show in Listing 7.39.

```
/*@
  lemma UniquePreserve{K,L}:
    \forall value_type *a, *b, integer i, k, n;
       0 <= k <= i < n          ==>
       k == UniqueSize{L}(a, i) ==>
       Unique{K}(a, i, b)       ==>
       Unchanged{K,L}(b, k)     ==>
       Unchanged{K,L}(a, n)     ==>
       Unique{L}(a, i, b);
*/
```

Listing 7.39: The lemma `UniquePreserve`

The lemma indicates that between two memory states `K` and `L`, the predicate `Unique` holds for the later state `L` if it already held in the first state `K` and the value of k equals the value of `UniqueSize(a, i)` in state `L`. Also the values in the ranges `b[0..k]` and `a[0..n]` must still be unchanged in both states.

Annotating the implementation

Listing 7.40 shows that we need considerably more annotation in order to verify the contract from Listing 7.37. As mentioned before, however, we can omit a dedicated loop invariant to verify the postcondition `solitary`.

The annotations in Listing 7.40 come not only in the form of loop invariants or ACSL assertions. We also employ so-called *ghost code* whose purpose we will explain now. As explained in the ACSL documentations [16, §2.12], variables and statements that appear in comments marked as
`/*@ ghost ... */` or `//@ ghost ...` are treated as C variables and statements, however, they are visible only in the specifications. In Listing 7.40 we declare the label `Before` as ghost. We could also have resorted to ACSL *statement contracts* [16, §2.4.4] but opted here for using ghost code.

```
size_type
unique_copy4(const value_type* a, size_type n, value_type* b)
{
  if (n == 0u) {
    return n;
  }
  else {
    size_type k = 0u;
    b[k] = a[0];
    //@ assert mapping:   0 == UniquePartition(a, n, k);

    /*@
      loop invariant bound:     0 <= k < i <= n;
      loop invariant size:      k+1 == UniqueSize(a, i);
      loop invariant copy:      b[k] == a[i-1];
      loop invariant mapping:   UniquePartition(a, n, k) < i;
      loop invariant mapping:   i <= UniquePartition(a, n, k+1);
      loop invariant unique:    Unique(a, i, b);
      loop invariant unchanged: Unchanged{Pre, Here}(a, n);
      loop invariant unchanged: Unchanged{Pre, Here}(b, k+1, n);
      loop assigns i, k, b[0..n-1];
      loop variant n-i;
    */
    for (size_type i = 1u; i < n; ++i) {
      const value_type val = a[i];

      if (b[k] != val) {
        //@ assert distinct: a[i-1] != a[i];
        //@ ghost Before:
        b[++k] = val;
        //@ assert unchanged: Unchanged{Before,Here}(b, k);
        //@ assert unchanged: Unchanged{Before,Here}(a, n);
        //@ assert unchanged: Unchanged{Before,Here}(a, i);
        //@ assert mapping:   i == UniquePartition(a, n, k);
        //@ assert size:      k == UniqueSize(a, i);
        //@ assert unique:    Unique(a, i, b);
      }

      //@ assert mapping:   i <= UniquePartition(a, n, k+1);
      //@ assert unchanged: Unchanged{Pre, Here}(b, k+1, n);
    }

    return ++k;
  }
}
```

Listing 7.40: Annotations for the complete contract of `unique_copy`

8. Numeric algorithms

The algorithms that we considered so far only *compared*, *read* or *copied* values in sequences. In this chapter, we consider so-called *numeric* algorithms of the C++ Standard Library [21, §29.8] that use arithmetic operations on `value_type` to combine the elements of sequences.

```
#define VALUE_TYPE_MAX  INT_MAX
#define VALUE_TYPE_MIN  INT_MIN
```

Listing 8.1: Limits of `value_type`

In order to refer to potential arithmetic overflows we introduce the two constants shown in Listing 8.1 which refer to the numeric limits of `value_type` (see also Section 1.3).

We consider the following algorithms.

- `iota` writes sequentially increasing values into a range (Section 8.1 on Page 162)

- `accumulate` computes the sum of the elements in a range (Section 8.2 on Page 164)

- `inner_product` computes the inner product of two ranges (Section 8.3 on Page 168)

- `partial_sum` computes the sequence of partial sums of a range (Section 8.4 on Page 171)

- `adjacent_difference` computes the differences of adjacent elements in a range (Section 8.5 on Page 175)

The formal specifications of these algorithms raise new questions. In particular, we now have to deal with arithmetic overflows in `value_type`.

8.1. The `iota` algorithm

The `iota` algorithm in the C++ Standard Library [21, §29.8.12] assigns sequentially increasing values to a range, where the initial value is user-defined. Our version of the original signature reads:

```
void iota(value_type* a, size_type n, value_type val);
```

Starting at `val`, the function assigns consecutive integers to the elements of the range `a`. When specifying `iota` we must be careful to deal with possible overflows of the argument `val`.

8.1.1. Formal specification of `iota`

The specification of `iota` relies on the logic function `Iota` that is defined in Listing 8.2.

```
/*@
  predicate
   Iota(value_type* a, integer n, value_type v) =
     \forall integer i; 0 <= i < n ==> a[i] == v+i;
*/
```

Listing 8.2: The logic function `Iota`

The ACSL specification of `iota` is shown in Listing 8.3. It uses the logic function `Iota` in order to express the postcondition `increment`.

```
/*@
  requires valid:  \valid(a + (0..n-1));
  requires limit:  val + n <= VALUE_TYPE_MAX;

  assigns a[0..n-1];

  ensures increment: Iota(a, n, val);
*/
void
iota(value_type* a, size_type n, value_type val);
```

Listing 8.3: Formal specification of `iota`

The specification of `iota` refers to `VALUE_TYPE_MAX` which is the maximum value of the underlying integer type (see Listing 8.1). In order to avoid integer overflows the sum `val+n` must not be greater than the constant `VALUE_TYPE_MAX`.

8.1.2. Implementation of `iota`

Listing 8.4 shows an implementation of the `iota` function.

```
void
iota(value_type* a, size_type n, value_type val)
{
  /*@
    loop invariant bound:    0 <= i <= n;
    loop invariant limit:    val == \at(val, Pre) + i;
    loop invariant increment: Iota(a, i, \at(val, Pre));

    loop assigns i, val, a[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    a[i] = val++;
  }
}
```

Listing 8.4: Implementation of `iota`

The loop invariant `increment` describes that in each iteration of the loop the current value `val` is equal to the sum of the value `val` in state of function entry and the loop index `i`. We have to refer here to `\at(val, Pre)` which is the value on entering `iota`.

8.2. The `accumulate` algorithm

The `accumulate` algorithm in the C++ Standard Library [21, §29.8.2] computes the sum of an given initial value and the elements in a range. Our version of the original signature reads:

```
value_type
accumulate(const value_type* a, size_type n, value_type init);
```

The result of `accumulate` shall equal the value

$$\text{init} + \sum_{i=0}^{n-1} a[i]$$

This implies that `accumulate` will return `init` for an empty range.

8.2.1. The logic function `Accumulate`

As in the case of `count` (see Section 3.9) we specify `accumulate` by first defining a *logic function* `Accumulate` that formally defines the summation of elements in an array.

```
/*@
  axiomatic AccumulateAxiomatic
  {
    logic integer
      Accumulate{L}(value_type* a, integer n, value_type init) =
        n <= 0 ? init : Accumulate(a, n-1, init) + a[n-1];

    lemma
      AccumulateRead{K,L}:
        \forall value_type *a, init, integer n;
          Unchanged{K,L}(a, n) ==>
            Accumulate{K}(a, n, init) == Accumulate{L}(a, n, init);
  }
*/
```

Listing 8.5: The logic function `Accumulate`

With this definition the following equation holds for $n \geq 0$

$$\text{Accumulate}(a, n, \text{init}) = \text{init} + \sum_{i=0}^{n-1} a[i] \tag{8.1}$$

The lemma `AccumulateRead` in Listing 8.5 express that the result of the `Accumulate` function only depends on the values of `a[0..n-1]`.

Listing 8.6 shows an overloaded version of `Accumulate` that uses 0 as default value of `init`. This listing also includes additional properties of observable `Accumulate` behavior, here given as a lemmas. These lemmas are proved automatically based on the definitions from Listing 8.5. We will use this version for the specification of the algorithm `partial_sum` (see Section 8.4).

Thus, for the overloaded version of `Accumulate` we have

$$\text{Accumulate}(a, n) = \sum_{i=0}^{n-1} a[i] \tag{8.2}$$

```
/*@
  logic integer Accumulate{L}(value_type* a, integer n) =
    Accumulate{L}(a, n, (value_type) 0);

  lemma
    AccumulateDefault0{L}:
      \forall value_type* a; Accumulate(a, 0) == 0;

  lemma
    AccumulateDefault1{L}:
      \forall value_type* a; Accumulate(a, 1) == a[0];

  lemma
    AccumulateDefaultNext{L}:
      \forall value_type* a, integer n;
        0 <= n  ==>  Accumulate(a, n+1) == Accumulate(a, n) + a[n];

  lemma
    AccumulateDefaultRead{L1,L2}:
      \forall value_type *a, integer n;
        Unchanged{L1,L2}(a, n)  ==>
        Accumulate{L1}(a, n) == Accumulate{L2}(a, n);
*/
```

Listing 8.6: An overloaded version of `Accumulate`

8.2.2. Preventing numeric overflows for `accumulate`

Before we present our formal specification of `accumulate` we introduce in Listing 8.7 a predicate `AccumulateBounds` that we will subsequently use in order to compactly express requirements that exclude numeric overflows while accumulating value.

```
/*@
  predicate
    AccumulateBounds{L}(value_type* a, integer n, value_type init) =
      \forall integer i; 0 <= i <= n ==>
        VALUE_TYPE_MIN <= Accumulate(a, i, init) <= VALUE_TYPE_MAX;

  predicate
    AccumulateBounds{L}(value_type* a, integer n) =
      AccumulateBounds{L}(a, n, (value_type) 0);
*/
```

Listing 8.7: The overloaded predicate `AccumulateBounds`

Predicate `AccumulateBounds` expresses that for $0 \leq i < n$ the *partial sums*

$$\mathtt{init} + \sum_{k=0}^{i} \mathtt{a}[k] \tag{8.3}$$

do not overflow. If one of them did, one couldn't guarantee that the result of `accumulate` equals the mathematical description of `Accumulate`.

Note that we also provide a second (overloaded) version of `AccumulateBounds` which uses a default value 0 for `init`.

8.2.3. Formal specification of `accumulate`

Using the logic function `Accumulate` and the predicate `AccumulateBounds`, the ACSL specification of `accumulate` is then as simple as shown in Listing 8.8.

```
/*@
  requires valid:  \valid_read(a + (0..n-1));
  requires bounds: AccumulateBounds(a, n, init);

  assigns  \nothing;

  ensures  result: \result == Accumulate(a, n, init);
*/
value_type
accumulate(const value_type* a, size_type n, value_type init);
```

Listing 8.8: Formal specification of `accumulate`

8.2.4. Implementation of `accumulate`

Listing 8.9 shows an implementation of the `accumulate` function with corresponding loop annotations.

```
value_type
accumulate(const value_type* a, size_type n, value_type init)
{
  /*@
    loop invariant index:    0 <= i <= n;
    loop invariant partial:  init == Accumulate(a, i, \at(init,Pre));
    loop assigns i, init;
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    //@ assert rte_help: init + a[i] == Accumulate(a, i+1, \at(init,Pre));
    init = init + a[i];
  }

  return init;
}
```

Listing 8.9: Implementation of `accumulate`

Note that loop invariant `partial` claims that in the *i*-th iteration step `result` equals the accumulated value of Equation (8.3). This depends on the property `bounds` in Listing 8.8 which expresses that there is no numeric overflow when updating the variable `init`.

8.3. The `inner_product` algorithm

The `inner_product` algorithm in the C++ Standard Library [21, §29.8.4] computes the *inner product*[26] of two ranges. Our version of the original signature reads:

```
value_type
inner_product(const value_type* a, const value_type* b,
              size_type n, value_type init);
```

The result of `inner_product` equals the value

$$\texttt{init} + \sum_{i=0}^{n-1} \texttt{a}[i] \cdot \texttt{b}[i]$$

thus, `inner_product` will return `init` for empty ranges.

8.3.1. The logic function `InnerProduct`

As in the case of `accumulate` (see Section 8.2) we specify `inner_product` by first defining a *logic function* `InnerProduct` that formally defines the summation of the element-wise product of two arrays.

```
/*@
  axiomatic InnerProductAxiomatic
  {
    logic integer
      InnerProduct{L}(value_type* a, value_type* b, integer n,
                      value_type init) =
        n <= 0 ? init : InnerProduct(a, b, n-1, init) + (a[n-1] * b[n-1]);

    lemma
      InnerProductRead{K,L}:
        \forall value_type *a, *b, init, integer n;
          Unchanged{K,L}(a, n) ==>
            Unchanged{K,L}(b, n) ==>
              InnerProduct{K}(a, b, n, init) == InnerProduct{L}(a, b, n, init);
  }
*/
```

Listing 8.10: The logic function `InnerProduct`

Lemma `InnerProductRead` express that the result of the `InnerProduct` only depends on the values of `a[0..n-1]` and `b[0..n-1]`.

---

[26]Also referred to as *dot product*, see `http://en.wikipedia.org/wiki/Dot_product`

8.3.2. Preventing numeric overflows for `inner_product`

Before we present our formal specification of `inner_product` we introduce in Listing 8.11 two predicates that we will use subsequently in order to compactly express requirements that exclude numeric overflows while computing the inner product.

```
/*@
  predicate
    ProductBounds(value_type* a, value_type* b, integer n) =
      \forall integer i; 0 <= i < n ==>
        VALUE_TYPE_MIN <= a[i] * b[i] <= VALUE_TYPE_MAX;

  predicate
    InnerProductBounds(value_type* a, value_type* b, integer n,
                       value_type init) =
      \forall integer i; 0 <= i <= n ==>
        VALUE_TYPE_MIN <= InnerProduct(a, b, i, init) <= VALUE_TYPE_MAX;
*/
```

Listing 8.11: The predicates `ProductBounds` and `InnerProductBounds`

Predicate `ProductBounds` expresses that for $0 \le i < n$ the products

$$a[i] \cdot b[i] \tag{8.4}$$

do not overflow. Predicate `InnerProductBounds`, on the other hand, states that for $0 \le i < n$ the *partial sums*

$$\texttt{init} + \sum_{k=0}^{i} a[k] \cdot b[k] \tag{8.5}$$

do not overflow.

Otherwise, one cannot guarantee that the result of `inner_product` equals the mathematical description of `InnerProduct`.

8.3.3. Formal specification of `inner_product`

Using the logic function `InnerProduct`, we specify `inner_product` as shown in Listing 8.12. Note that we needn't require that `a` and `b` are separated.

```
/*@
  requires valid:    \valid_read(a + (0..n-1));
  requires valid:    \valid_read(b + (0..n-1));
  requires bounds:  ProductBounds(a, b, n);
  requires bounds:  InnerProductBounds(a, b, n, init);

  assigns \nothing;

  ensures result:    \result == InnerProduct(a, b, n, init);
  ensures unchanged: Unchanged{Old,Here}(a, n);
  ensures unchanged: Unchanged{Old,Here}(b, n);
*/
value_type
inner_product(const value_type* a, const value_type* b, size_type n,
              value_type init);
```

Listing 8.12: Formal specification of `inner_product`

8.3.4. Implementation of `inner_product`

Listing 8.13 shows an implementation of `inner_product` with corresponding loop annotations.

```
value_type
inner_product(const value_type* a, const value_type* b, size_type n,
              value_type init)
{
  /*@
    loop invariant index: 0 <= i <= n;
    loop invariant inner: init == InnerProduct(a, b, i, \at(init,Pre));
    loop assigns i, init;
    loop variant n-i;
  */
  for (size_type i = 0u; i < n; ++i) {
    /*@
      assert rte_help: init + a[i] * b[i] ==
                       InnerProduct(a, b, i+1, \at(init,Pre));
    */
    init = init + a[i] * b[i];
  }

  return init;
}
```

Listing 8.13: Implementation of `inner_product`

Note that the loop invariant `inner` claims that in the *i*-th iteration step the current value of `init` equals the accumulated value of Equation (8.5). This depends of course on the properties `bounds` in Listing 8.12, which express that there is no arithmetic overflow when computing the updates of the variable `init`.

8.4. The `partial_sum` algorithm

The `partial_sum` algorithm in the C++ Standard Library [21, §29.8.6] computes the sum of a given initial value and the elements in a range. Our version of the original signature reads:

```
size_type
partial_sum(const value_type* a, size_type n, value_type* b);
```

After executing the function `partial_sum` the array `b[0..n-1]` holds the following values

$$b[0] = a[0]$$
$$b[1] = a[0] + a[1]$$
$$\vdots$$
$$b[n-1] = a[0] + a[1] + \ldots + a[n-1]$$

More concisely, for $0 \leq i < n$ holds

$$b[i] = \sum_{k=0}^{i} a[k] \qquad (8.6)$$

8.4.1. The predicate `PartialSum`

Equations (8.6) and (8.2) suggest that we define the ACSL predicate `PartialSum` in Listing 8.14 by using the logic function `Accumulate` from Listing 8.6.

```
/*@
  predicate
    PartialSum{L}(value_type* a, integer n, value_type* b) =
      \forall integer i; 0 <= i < n ==> Accumulate(a, i+1) == b[i];
*/
```

Listing 8.14: The predicate `PartialSum`

8.4.2. Formal specification of `partial_sum`

Using the predicates `PartialSum` and `AccumulateBounds`, we specify `partial_sum` as shown in Listing 8.15.

```
/*@
  requires valid:     \valid_read(a + (0..n-1));
  requires valid:     \valid(b + (0..n-1));
  requires sep:       \separated(a + (0..n-1), b + (0..n-1));
  requires bounds:    AccumulateBounds(a, n+1);

  assigns b[0..n-1];

  ensures result:     \result == n;
  ensures partialsum: PartialSum(a, n, b);
  ensures unchanged:  Unchanged{Old,Here}(a, n);
*/
size_type
partial_sum(const value_type* a, size_type n, value_type* b);
```

Listing 8.15: Formal specification of `partial_sum`

Our specification requires that the arrays `a[0..n-1]` and `b[0..n-1]` are separated, that is, they do not overlap. Note that is a stricter requirement than in the case of the original C++ version of `partial_sum`, which allows that `a` equals `b`, thus allowing the computation of partial sums *in place*.

8.4.3. Implementation of `partial_sum`

Listing 8.16 shows an implementation of `partial_sum` with corresponding loop annotations.

```
size_type
partial_sum(const value_type* a, size_type n, value_type* b)
{
  if (0u < n) {
    b[0u] = a[0u];

    /*@
       loop invariant bound:      1 <= i <= n;
       loop invariant unchanged:  Unchanged{Pre,Here}(a, n);
       loop invariant accumulate: b[i-1] == Accumulate(a, i);
       loop invariant partialsum: PartialSum(a, i, b);
       loop assigns i, b[1..n-1];
       loop variant n - i;
    */
    for (size_type i = 1u; i < n; ++i) {
      //@ ghost Enter:
      b[i] = b[i - 1u] + a[i];
      //@ assert unchanged:  a[i] == \at(a[i],Enter);
      //@ assert unchanged:  Unchanged{Enter,Here}(a, i);
      //@ assert unchanged:  Unchanged{Enter,Here}(b, i);
    }
  }

  return n;
}
```

Listing 8.16: Implementation of `partial_sum`

In order to facilitate the automatic verification of `partial_sum`, we had to add the assertions `unchanged` and provide the lemmas of Listing 8.17.

8.4.4. Additional lemmas

The lemmas shown in Listing 8.17 are needed for the verification of `partial_sum` and the algorithms in Sections 8.6 and 8.7.

```
/*@
  lemma
    PartialSumSection{K}:
      \forall value_type *a, *b, integer m, n;
      0 <= m <= n            ==>
      PartialSum{K}(a, n, b) ==>
      PartialSum{K}(a, m, b);

  lemma
    PartialSumUnchanged{K,L}:
      \forall value_type *a, *b, integer n;
        0 <= n ==>
      PartialSum{K}(a, n, b) ==>
      Unchanged{K, L}(a, n)  ==>
      Unchanged{K, L}(b, n)  ==>
      PartialSum{L}(a, n, b);

  lemma
    PartialSumStep{L}:
      \forall value_type *a, *b, integer n;
        1 <= n                    ==>
      PartialSum(a, n, b)         ==>
      b[n] == Accumulate(a, n+1)  ==>
      PartialSum(a, n+1, b);

  lemma
    PartialSumStep2{K,L}:
      \forall value_type *a, *b, integer n;
        1 <= n                          ==>
      PartialSum{K}(a, n, b)            ==>
      Unchanged{K,L}(a, n+1)            ==>
      Unchanged{K,L}(b, n)              ==>
      \at(b[n] == Accumulate(a, n+1),L) ==>
      PartialSum{L}(a, n+1, b);
*/
```

Listing 8.17: The lemma `PartialSumStep`

8.5. The `adjacent_difference` algorithm

The `adjacent_difference` algorithm in the C++ Standard Library [21, §29.8.11] computes the differences of adjacent elements in a range. Our version of the original signature reads:

```
size_type
adjacent_difference(const value_type* a, size_type n, value_type* b);
```

After executing the function `adjacent_difference` the array `b[0..n-1]` holds the following values

$$b[0] = a[0]$$
$$b[1] = a[1] - a[0]$$
$$\vdots$$
$$b[n-1] = a[n-1] - a[n-2]$$

If we form the partial sums of the sequence `b` we find that

$$a[0] = b[0]$$
$$a[1] = b[0] + b[1]$$
$$\vdots$$
$$a[n-1] = b[0] + b[1] + \ldots + b[n-1]$$

Thus, we have for $0 \le i < n$

$$a[i] = \sum_{k=0}^{i} b[k] \tag{8.7}$$

which means that applying `partial_sum` on the output of `adjacent_difference` produces the original input of `adjacent_difference`.

Conversely, if `a[0..n-1]` and `b[0..n-1]` are the input and output of `partial_sum`, then we have

$$b[0] = a[0]$$
$$b[1] = a[0] + a[1]$$
$$\vdots$$
$$b[n-1] = a[0] + b[1] + \ldots + b[n-1]$$

from which we can conclude

$$a[0] = b[0]$$
$$a[1] = b[1] - b[0]$$
$$\vdots \tag{8.8}$$
$$a[n-1] = b[n-1] - b[n-2]$$

We will verify these claims in Sections 8.6 and 8.7.

8.5.1. The predicate **AdjacentDifference**

We define the predicate AdjacentDifference in Listing 8.19 by first introducing the logic function Difference (Listing 8.18).

```
/*@
  axiomatic DifferenceAxiomatic
  {
    logic integer
      Difference{L}(value_type* a, integer n) =
        n <= 0 ? a[0] : a[n] - a[n-1];

    lemma
      DifferenceRead{K,L}:
        \forall value_type *a, integer n;
          n >= 0 ==> Unchanged{K,L}(a, 1+n) ==>
            Difference{K}(a, n) == Difference{L}(a, n);
  }
*/
```

Listing 8.18: The logic function Difference

```
/*@
  predicate
    AdjacentDifference{L}(value_type* a, integer n, value_type* b) =
      \forall integer i; 0 <= i < n ==> b[i] == Difference(a, i);
*/
```

Listing 8.19: The predicate AdjacentDifference

8.5.2. Formal specification of `adjacent_difference`

We introduce here the predicate AdjacentDifferenceBounds (Listing 8.20) that captures conditions that prevent numeric overflows while computing difference of the form a[i] - a[i-1].

```
/*@
  predicate
    AdjacentDifferenceBounds(value_type* a, integer n) =
      \forall integer i; 1 <= i < n ==>
        VALUE_TYPE_MIN <= Difference(a, i) <= VALUE_TYPE_MAX;
*/
```

Listing 8.20: The predicate AdjacentDifferenceBounds

Using the predicates AdjacentDifference and AdjacentDifferenceBounds we can provide a concise formal specification of adjacent_difference (Listing 8.21). As in the case of the specification of partial_sum we require that the arrays a[0..n-1] and b[0..n-1] are separated.

```
/*@
  requires valid:      \valid_read(a + (0..n-1));
  requires valid:      \valid(b + (0..n-1));
  requires sep:        \separated(a + (0..n-1), b + (0..n-1));
  requires bounds:     AdjacentDifferenceBounds(a, n);

  assigns b[0..n-1];

  ensures result:      \result == n;
  ensures difference:  AdjacentDifference(a, n, b);
  ensures unchanged:   Unchanged{Old,Here}(a, n);
*/
size_type
adjacent_difference(const value_type* a, size_type n, value_type* b);
```

Listing 8.21: Formal specification of adjacent_difference

8.5.3. Implementation of `adjacent_difference`

Listing 8.22 shows an implementation of `adjacent_difference` with corresponding loop annotations.

In order to achieve the verification of the loop invariant `difference` we added

- the assertions `bound` and `difference`

- the lemmas `AdjacentDifferenceStep` and `AdjacentDifferenceSection` from Listing 8.23

- a statement contract with the two postconditions labeled as `step`

```
size_type
adjacent_difference(const value_type* a, size_type n, value_type* b)
{
  if (0u < n) {
    b[0u] = a[0u];

    /*@
        loop invariant index:      1 <= i <= n;
        loop invariant unchanged:  Unchanged{Pre,Here}(a, n);
        loop invariant difference: AdjacentDifference(a, i, b);
        loop assigns i, b[1..n-1];
        loop variant n - i;
    */
    for (size_type i = 1u; i < n; ++i) {
      //@ assert bound: VALUE_TYPE_MIN <= Difference(a, i) <= VALUE_TYPE_MAX;
      /*@
        assigns b[i];
        ensures step: Unchanged{Old,Here}(b, i);
        ensures step: b[i] == Difference(a, i);
      */
      b[i] = a[i] - a[i - 1u];
      //@ assert difference: AdjacentDifference(a, i+1, b);
    }
  }

  return n;
}
```

Listing 8.22: Implementation of `adjacent_difference`

8.5.4. Additional Lemmas

The lemmas shown in Listing 8.23 are also needed for the verification of the algorithm in Section 8.7.

```
/*@
  lemma
    AdjacentDifferenceStep{K,L}:
      \forall value_type *a, *b, integer n;
        AdjacentDifference{K}(a, n, b)        ==>
        Unchanged{K,L}(b, n)                  ==>
        Unchanged{K,L}(a, n+1)                ==>
        \at(b[n],L) == Difference{L}(a, n)    ==>
        AdjacentDifference{L}(a, 1+n, b);

  lemma
    AdjacentDifferenceSection{K}:
      \forall value_type *a, *b, integer m, n;
        0 <= m <= n                           ==>
        AdjacentDifference{K}(a, n, b)  ==>
        AdjacentDifference{K}(a, m, b);
*/
```

Listing 8.23: The lemma `AdjacentDifferenceStep`

8.6. Inverting `partial_sum` with `adjacent_difference`

In Section 8.5 we had informally argued that `partial_sum` and `adjacent_difference` are inverse to each other (see Equations (8.7) and (8.8)). In the current section, we are going to verify the second of these claims with the help of Frama-C, viz. that applying `adjacent_difference` to the output of `partial_sum` produces the original array. In Section 8.7, we will verify the converse first claim.

Listing 8.24 expresses the property from Equation (8.8) as lemma, on the ACSL logic level. This lemma is verified by Frama-C with the help of automatic theorem provers.

```
/*@
  lemma
    PartialSumInv{K,L}:
      \forall value_type *a, *b, integer n;
        0 <= n                         ==>
        PartialSum{K}(a, n, b)         ==>
        Unchanged{K,L}(b, n)           ==>
        AdjacentDifference{L}(b, n, a) ==>
        Unchanged{K,L}(a, n);
*/
```

Listing 8.24: The lemma `PartialSumInv`

Since the lemma does not deal with arithmetic overflows or potential aliasing of data, we give a corresponding auxiliary C function which takes these issues into account.

Function `partial_sum_inv`, shown in Listing 8.25, calls first `partial_sum` and then `adjacent_difference`. The contract of this function formulates preconditions that shall guarantee that during the computation neither arithmetic overflows (property `bound`) nor unintended aliasing of arrays (property `sep`) occur. Under these precondition, Frama-C automatically verifies that the final `adjacent_difference` call just restores the original contents of a used for the initial `partial_sum` call.

```
/*@
  requires valid:     \valid(a + (0..n-1));
  requires valid:     \valid(b + (0..n-1));
  requires sep:       \separated(a + (0..n-1), b + (0..n-1));
  requires bounds:    AccumulateBounds(a, n+1);

  assigns a[0..n-1], b[0..n-1];

  ensures unchanged:  Unchanged{Pre,Here}(a, n);
*/
void
partial_sum_inv(value_type* a, size_type n, value_type* b)
{
  partial_sum(a, n, b);
  adjacent_difference(b, n, a);
}
```

Listing 8.25: `partial_sum` and then `adjacent_difference`

8.7. Inverting `adjacent_difference` with `partial_sum`

In this section, we prove the converse property, viz. that applying `adjacent_difference`, and thereafter `partial_sum`, restores the original data array. Listing 8.26 expresses this property as a lemma on the level of ACSL predicates. It had to be proven interactively with Coq, by induction on `n`.

```
/*@
  lemma
    AdjacentDifferenceInv{K,L}:
      \forall  value_type *a, *b, integer n;
        0 <= n                           ==>
        AdjacentDifference{K}(a, n, b)   ==>
        Unchanged{K,L}(b, n)             ==>
        PartialSum{L}(b, n, a)           ==>
        Unchanged{K,L}(a, n);
*/
```

Listing 8.26: The lemma `AdjacentDifferenceInv`

As in the case discussed in Section 8.6, we give a corresponding C function in order to account for possible arithmetic overflows and potential aliasing of data. Function `adjacent_difference_inv`, shown in Listing 8.27, calls first `adjacent_difference` and then `partial_sum`. The contract of this function formulates preconditions that shall guarantee that during the computation neither arithmetic overflows (property `bound`) nor unintended aliasing of arrays (property `sep`) occur.

```
/*@
  requires valid:     \valid(a + (0..n-1));
  requires valid:     \valid(b + (0..n-1));
  requires sep:       \separated(a + (0..n-1), b + (0..n-1));
  requires bounds:    AdjacentDifferenceBounds(a, n+1);

  assigns a[0..n-1], b[0..n-1];

  ensures unchanged:  Unchanged{Old,Here}(a, n);
*/
void
adjacent_difference_inv(value_type* a, size_type n, value_type* b)
{
  adjacent_difference(a, n, b);
  partial_sum(b, n, a);
}
```

Listing 8.27: `adjacent_difference` and then `partial_sum`

In order to improve the automatic verification rate of the function `adjacent_difference_inv` we also use lemma `UnchangedTransitive` from Listing 6.4. Both lemmas itself and (with its additional help) the contract of `adjacent_difference_inv` are proven by Frama-C without further manual intervention. This finishes the formal proof of our inversity claims from Section 8.5.

Part IV.

Sorting algorithms

9. Heap Algorithms

The heap algorithms of the C++ Standard Library [21, 28.7.7] were already part of *ACSL by Example* from 2010–2012. In this chapter we re-introduce them and discuss—based on the bachelor thesis of one of the authors—the verification efforts in some detail [29].

The C++ standard[27] introduces the concept of a *heap* as follows:

1. A *heap* is a particular organization of elements in a range between two random access iterators `[a,b)`. Its two key properties are:

   a) There is no element greater than `*a` in the range and

   b) `*a` may be removed by `pop_heap()`, or a new element added by `push_heap()`, in $O(\log(N))$ time.

2. These properties make heaps useful as priority queues.

3. `make_heap()` converts a range into a heap and `sort_heap()` turns a heap into a sorted sequence.

Figure 9.1 gives an overview on the five heap algorithms by means of an example. Algorithms, which in a typical implementation are in a caller-callee relation, have the same color.

Figure 9.1.: Overview on heap algorithms

---

[27]See `http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2011/n3242.pdf`

Roughly speaking, the algorithms from Figure 9.1 have the following behavior.

- `is_heap_until` and `is_heap` from Sections 9.4 and 9.5 allow to test at run time whether a given array is arranged as a heap

- `push_heap` from Section 9.6 *adds* an element to a given heap in such a way that resulting array is again a heap

- `pop_heap`, which is currently not included in this document, *removes* an element from a given heap in such a way that the resulting array is again a heap

- `make_heap` from Section 9.8 turns a given array into a heap.

- `sort_heap` from Section 9.9 transforms a given heap into a sorted range.

In Section 9.1 we present in more detail how heaps are defined. The ACSL logic functions and predicate that formalize the basic heap properties of heaps are introduced in Section 9.2.

```
#define SIZE_TYPE_MAX  UINT_MAX
```

Listing 9.2: Upper limits of `size_type`

In order to admit maximally large heaps, we had to catch border cases in ACSL as well as in C, cf. e.g. Listing 9.12 and 9.13. To this end, we introduced the constant from Listing 9.2. to refer to the upper bound of `size_type`. We don't need a corresponding constant `SIZE_TYPE_MIN` for the lower bound, since it is trivial.

9.1. Basic heap concepts

The description of heaps at the beginning of this chapter is of course fairly vague. It outlines only the most important properties of various operations but does not clearly state what specific and verifiable properties a range must satisfy such that it may be called a heap.

A more detailed description can be found in the Apache C++ Standard Library User's Guide:[28]

> A heap is a binary tree in which every node is larger than the values associated with either child. A heap and a binary tree, for that matter, can be very efficiently stored in a vector, by placing the children of node i at positions $2i + 1$ and $2i + 2$.

We have, in other words, the following basic relations between indices of a heap:

$$\text{left child for index } i \qquad \text{child}_l : i \mapsto 2i + 1 \qquad (9.1)$$
$$\text{right child for index } i \qquad \text{child}_r : i \mapsto 2i + 2 \qquad (9.2)$$

and

$$\text{parent index for index } i \qquad \text{parent} : i \mapsto \frac{i-1}{2} \qquad (9.3)$$

These function are related through the following two equations that hold for all integers i. Note that in ACSL integer division rounds towards zero (cf. [16, §2.2.4]).

---

[28] See `http://stdcxx.apache.org/doc/stdlibug/14-7.html`

$$\text{parent}(\text{child}_{\text{l}}(i)) = i \tag{9.4}$$
$$\text{parent}(\text{child}_{\text{r}}(i)) = i \tag{9.5}$$

In order to given an example for the usefulness of heaps we consider the following multiset of integers X.

$$X = \{2, 3, 3, 3, 6, 7, 8, 8, 9, 11, 13, 14\} \tag{9.6}$$

Figure 9.3 shows how the multiset from Equation (9.6) can, according to the parent-child relations of a heap, be represented as a tree.

Figure 9.3.: Tree representation of the multiset X

The numbers outside the nodes in Figure 9.3 are the indices at which the respective node value is stored in the underlying array of a heap (cf. Figure 9.4).

Figure 9.4.: Underlying array of a heap

It is important to understand that there can be various representations of a multiset as a heap. Figure 9.5, for example, arranges the elements of the multiset X as a heap in a different tree.

187

Figure 9.5.: An alternative representation of the multiset *X*

Figure 9.6 then shows the underlying array that corresponds to the tree in Figure 9.5.

Figure 9.6.: Underlying array of the alternative representation

9.2. Presentation of heap concepts in ACSL

Listing 9.7 shows three logic functions `HeapLeft`, `HeapRight` and `HeapParent` that correspond to the definitions (9.1), (9.2) and (9.3), respectively. Listing 9.7 also contains a number of ACSL lemma that state among other things that

- the `HeapParent` function satisfies the equations (9.4) and (9.5) and
- the function `HeapParent` is the *left inverse* to the `HeapLeft` and `HeapRight` functions.[29]

---

[29]See Section *Left and right inverses* at `http://en.wikipedia.org/wiki/Inverse_function`

```
/*@
  logic integer HeapLeft(integer i) = 2*i + 1;

  logic integer HeapRight(integer i) = 2*i + 2;

  logic integer HeapParent(integer i) = (i-1) / 2;

  lemma
    HeapParentOfLeft:
      \forall integer p; 0 <= p ==> HeapParent(HeapLeft(p)) == p;

  lemma
    HeapParentOfRight:
      \forall integer p; 0 <= p ==> HeapParent(HeapRight(p)) == p;

  lemma
    HeapParentChild:
      \forall integer c, p;
        0 < c ==> HeapParent(c) == p ==>
        (c == HeapLeft(p) || c == HeapRight(p));

  lemma
    HeapChilds:
      \forall integer a, b;
        0 < a ==> 0 < b ==>
        HeapParent(a) == HeapParent(b) ==>
        (a == b || a+1 == b || a == b+1);

  lemma
    HeapParentBounds:
      \forall integer c; 0 < c ==> 0 <= HeapParent(c) < c;

  lemma
    HeapChildBounds:
      \forall integer p;
        0 <= p ==> p < HeapLeft(p) < HeapRight(p);
*/
```

Listing 9.7: Some logic functions for describing heaps

On top of these basic definitions we introduce in Listing 9.8 the predicate `IsHeap`.

```
/*@
  predicate
    IsHeap{L}(value_type* a, integer n) =
      \forall integer i; 0 < i < n  ==> a[i] <= a[HeapParent(i)];
*/
```

Listing 9.8: The predicate `IsHeap`

The root of a heap, that is the element at index 0, is always the largest element of the heap. Lemma `HeapMaximum` in Listing 9.9 expresses this property using the `MaxElement` predicate from Listing 4.9.

```
/*@
  lemma
    HeapMaximum{L} :
      \forall value_type* a, integer n;
        1 <= n ==> IsHeap(a, n) ==> MaxElement(a, n, 0);
*/
```

Listing 9.9: The lemma `HeapMaximum`

We use the following fact about division in C in the proof of lemma `HeapMaximum`.

```
/*@
  lemma
    C_Division_2:
      \forall integer a; 0 <= a  ==>  0 <= a/2 <= a;
*/
```

Listing 9.10: The lemma `C_Division_2`

9.3. The auxiliary functions `heap_parent` and `heap_child_max`

This section features the two auxiliary heap functions `heap_parent` and `heap_child_max` that are used in the verification of the heap algorithms. The following Listing 9.11 contains the C counterpart of a logic heap function in Listing 9.7. Note that here we have to take into account the limitations of C types.

```
/*@
   requires bound:  0 < child;
   assigns          \nothing;
   ensures  parent: \result == HeapParent(child);
 */
static inline size_type
heap_parent(size_type child)
{
  return (child - 1u) / 2u;
}
```

Listing 9.11: The auxiliary heap function `heap_parent`

Our implementation relies on the auxiliary function `heap_child_max` from Listing 9.13 (Section 9.3). As discussed there this helper function computes the child of an heap node with the the largest index.

Our second auxiliary function is `heap_child_max`. It computes the child of an heap node with the the largest index. This function will be used in the implementation of `pop_heap` 9.7. On the logic level this notion is expressed in ACSL by the predicate `HeapChildMax` from Listing 9.12.

```
/*@
  predicate
    HeapChildMax{L}(value_type* a, integer n,
                    integer p,     integer c) =
      0 <= p < n-1                                       &&
      (p < (SIZE_TYPE_MAX-1)/2  ==>  p == HeapParent(c))      &&
      (HeapLeft(p)  < n-1       ==>  a[HeapLeft(p)]  <= a[c]) &&
      (HeapRight(p) < n-1       ==>  a[HeapRight(p)] <= a[c]);
*/
```

Listing 9.12: The predicate `HeapChildMax`

Listing 9.13 shows the contract and the implementation of the function `heap_child_max`. This function returns a child index `c` of `parent` with the property that for a potential other child `d` of `parent` with `d < n − 1` the condition `a[d] <= a[c]` holds. Note that it explicitly handles the case that the child index computation would overflow; in which case it returns `n`.

```
/*@
   requires bound: 2 <= n;
   requires bound: 0 <= parent < n − 1;
   requires valid: \valid(a + (0..n-1));
   requires heap:  IsHeap(a, n);

   assigns         \nothing;

   ensures heap:   IsHeap(a, n);
   ensures max:    HeapChildMax(a, n, parent, \result);
   ensures less:   parent < \result;
   ensures less:   \result < n − 1  ==>  parent == HeapParent(\result);
*/
static inline size_type
heap_child_max(const value_type* a, size_type n, size_type parent)
{
  if (parent < (SIZE_TYPE_MAX − 1u) / 2u) {
    const size_type right = 2u * parent + 2u;
    const size_type left  = right − 1u;

    if (right < n − 1u) {
      // case of two children: select child with maximum value
      return a[left] >= a[right] ? left : right;
    }
    else {
      // at most one child that comes before n−1 can exist
      return left;
    }
  }
  else {
    return n;
  }
}
```

Listing 9.13: The auxiliary function `heap_child_max`

9.4. The `is_heap_until` algorithm

The `is_heap_until` algorithm of the C++ Standard Library [21, §28.7.7.5] works on generic sequences. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
size_type is_heap_until(const value_type* a, int n);
```

The algorithm `is_heap_until` returns the largest range of an array, beginning at the first position, where it still satisfies the heap properties we have semi-formally described in the beginning of this chapter. In particular, `is_heap_until` will return the size of the array, called with the array argument from Figure 9.4.

9.4.1. Formal specification of `is_heap_until`

The ACSL specification of `is_heap_until` is shown in Listing 9.14. The function returns an index `i` if and only if its array `a[0..i-1]` satisfies the predicate `IsHeap` introduced in Section 9.2. In addition the postcondition `last` states, that for all indices greater than or equal to `i` the predicate `IsHeap` is not satisfied.

```
/*@
   requires valid: \valid_read(a + (0..n-1));

   assigns \nothing;

   ensures bound: 0 <= \result <= n;
   ensures heap:  IsHeap(a, \result);
   ensures last:  \forall integer i; \result < i <= n ==> !IsHeap(a, i);
*/
size_type
is_heap_until(const value_type* a, size_type n);
```

Listing 9.14: Function Contract of `is_heap_until`

9.4.2. Implementation of `is_heap_until`

Listing 9.15 shows one way to implement the function `is_heap_until`. The algorithms starts at the index 1, which is the smallest index, where a child node of the heap might reside. The algorithms checks for each (child) index whether the value at the corresponding parent index is greater than or equal to the value at the child index. If the value at a parent index is smaller than the value at a (child) index, `is_heap_until` returns the (child) index. Otherwise, if the algorithm iterates through the whole array, the size of the array is returned.

```
size_type
is_heap_until(const value_type* a, size_type n)
{
  size_type parent = 0u;

  /*@
    loop invariant bound:    0 <= parent < child <= n+1;
    loop invariant parent:   parent == HeapParent(child);
    loop invariant heap:     IsHeap(a, child);
    loop invariant not_heap: a[parent] < a[child] ==> \forall integer i; child < i <=
        n ==> !IsHeap(a, i);

    loop assigns child, parent;
    loop variant n - child;
  */
  for (size_type child = 1u; child < n; ++child) {
    if (a[parent] < a[child]) {
      return child;
    }

    if ((child % 2u) == 0u) {
      ++parent;
    }
  }

  return n;
}
```

Listing 9.15: Implementation of `is_heap_until`

9.5. The `is_heap` algorithm

The `is_heap` algorithm of the C++ Standard Library [21, §28.7.7.5] works on generic sequences. For our purposes we have modified the generic implementation to that of an array of type `value_type`. The signature now reads:

```
bool is_heap(const value_type* a, int n);
```

The algorithm `is_heap` checks whether a given array satisfies the heap properties we have semi-formally described in the beginning of this chapter. In particular, `is_heap` will return **true** called with the array argument from Figure 9.4.

9.5.1. Formal specification of `is_heap`

The ACSL specification of `is_heap` is shown in Listing 9.16. The function returns **true** if and only if its arguments satisfy the predicate `IsHeap` introduced in Section 9.2.

```
/*@
   requires valid: \valid_read(a +(0..n-1));

   assigns \nothing;

   ensures heap:  \result <==> IsHeap(a, n);
*/
bool
is_heap(const value_type* a, size_type n);
```

Listing 9.16: Function Contract of `is_heap`

9.5.2. Implementation of `is_heap`

Listing 9.17 shows one way to implement the function `is_heap` by utilizing the function `is_heap_until` from Section 9.4.

The function `is_heap` returns **true** if its whole given range satisfies the heap properties. So basically, `is_heap` returns **true** if the function `is_heap_until` returns the size of the input array.

```
bool
is_heap(const value_type* a, size_type n)
{
  return is_heap_until(a, n) == n;
}
```

Listing 9.17: Implementation of `is_heap`

9.6. The `push_heap` algorithm

Whereas in the C++ Standard Library [21, §28.7.7.1] push_heap works on a range of random access iterators, our version operates on an array of value_type. We therefore use the following signature for push_heap

```
void push_heap(value_type* a, size_type n);
```

The push_heap algorithm expects that n is greater or equal than 1. It also assumes that the array a[0..n-2] forms a heap. The algorithms then *rearranges* the array a[0..n-1] such that the resulting array is a heap. In this sense the algorithm *pushes* an element on a heap.

9.6.1. Formal Specification of `push_heap`

Listing 9.18 shows our ACSL specification of push_heap. Note that the post condition reorder states that push_heap is not allowed to change the number of occurrences of an array element. Without this post condition, an implementation that assigns 0 to each array element would satisfy the post condition heap—surely not what the user of the algorithm has in mind.

```
1  /*@
2     requires nonempty: 0 < n;
3     requires valid:    \valid(a + (0..n-1));
4     requires heap:     IsHeap(a, n-1);
5
6     assigns a[0..n-1];
7
8     ensures heap:      IsHeap(a, n);
9     ensures reorder:   MultisetUnchanged{Old,Here}(a, n);
10 */
11 void
12 push_heap(value_type* a, size_type n);
```

Listing 9.18: Formal specification of push_heap

Pushing an element on a heap usually *rearranges* several elements of the array (cf. Figures 9.19 and 9.20). We therefore must be able express that push_heap only *reorders* the elements of the array. We re-use the predicate MultisetUnchanged, defined in Listing 6.5, to formally describe this property.

9.6.2. Implementation of `push_heap`

The following two figures illustrate how push_heap affects an array, which is shown as a tree with blue and grey nodes, representing heap and non-heap nodes, respectively. Figure 9.19 shows the heap from Figure 9.3 together with the additional element 12 that is to be on the heap. To be quite clear about it: the new element 12 is the last element of the array and not yet part of the heap.

Figure 9.19.: Heap before the call of `push_heap`

Figure 9.20 shows the array after the call of `push_heap`. We can see that now all nodes are colored in blue, i.e., they are part of the heap. The dashed nodes changed their contents during the function call. The pushed element `12` is now at its correct position in the heap.

Figure 9.20.: Heap after the call of `push_heap`

Challenges during the verification

In order to properly describe different stages of `push_heap` and to accommodate the sheer size of our implementation we split the source code into three separate parts, to which we refer as

- *prologue* (see Section 9.6.2)
- *main act* (see Section 9.6.2)
- *epilogue* (see Section 9.6.2)

We will illustrate the changes to the array after each stage by figures of the array in tree form, based on the `push_heap` example from Figure 9.19.

Verifying `push_heap` is a non-trivial undertaking, and we will proceed, roughly speaking, as follows:

> We can establish the `heap` property of Listing 9.18 already in the prologue. However, the `reorder` property only holds at the function boundaries but is violated while `push_heap` manipulates the array. To be more precise: We loose the `reorder` property in the prologue and formally capture and maintain a slightly more general property in the main act. From this we will recover the `reorder` property in the epilogue.

Prologue

Our prologue initializes some important variables, checks whether the initial heap is nonempty, *and* also tries to move the new element upwards within the heap. In other implementations, the latter step is usually performed as part of `push_heap`'s main loop. In order to better understand our implementation decision we can look at Figure 9.21 which shows exemplarily its effects.

Figure 9.21.: Heap after the prologue of `push_heap`

If we compare this tree to the tree in Figure 9.19 we notice that the element in the node with the dashed outlining changed its value. The number of occurrences of 9 increased by one during the prologue, while the number of occurrences of 12 decreased by one. The number of occurrences for all other elements is maintained. We store the element 12 in the variable `val` so we can write it back into the array later. The increased number of occurrences of 9 and the decreased number of elements 12 means that at this stage the postcondition `reorder` is violated. On the other hand, the modified array is now a heap. We express this by coloring all elements blue (cf. Figure 9.19).

More generally speaking, the following properties hold after the prologue:

1. The modified array is now a heap.

2. The "parent value" `a[parent]` now occurs one time more often.

3. The value `a[n-1]`, on the other hand, now occurs one time fewer.

4. No other value changed its number of occurrences.

Listing 9.22 shows the implementation of the prologue. It starts with the listing of an auxiliary function `heap_parent` that computes the parent index of its argument. The prologue also deals with the trivial cases that

- the array contains only one element or
- if `a[n-1]` is less or equal than its parent element

At the end of the prologue we have added four assertions that formally express the properties we have just enumerated above. As we will see (in Listing 9.28), these properties will occur as loop invariants in the main act. Adding these assertions makes the purpose of the prologue more explicit and thus supports the long-term maintenance of the annotated code. The auxiliary predicates that are used in the assertions are:

- predicate `MultisetAdd` (Listing 9.23)
- predicate `MultisetMinus` (Listing 9.24)
- various overloaded versions of predicate `MultisetRetainRest` (Listings 9.25 and 9.26)

```
void
push_heap(value_type* a, size_type n)
{
  // start of prologue
  if (1u < n) { // otherwise nothings needs to be done
    const value_type v    = a[n - 1u];
    size_type  hole = heap_parent(n - 1u);

    if (a[hole] < v) {
      a[n - 1u] = a[hole];

      //@ assert heap:   IsHeap(a, n);
      //@ assert add:    MultisetAdd{Pre,Here}(a, n, a[hole]);
      //@ assert minus:  MultisetMinus{Pre,Here}(a, n, v);
      //@ assert retain: MultisetRetainRest{Pre,Here}(a, n, v, a[hole]);
      // end of prologue
```

Listing 9.22: Prologue of `push_heap` implementation

These auxiliary predicates are discussed in the following subsections.

The predicates `MultisetAdd` and `MultisetMinus`

The predicate `MultisetAdd` in Listing 9.23 expresses that the number of occurrences of a specific element in an array has increased by one between between two program points `K` and `L`.

```
/*@
  predicate
    MultisetAdd{K,L}(value_type* a, integer n, value_type val) =
      Count{L}(a, n, val) == Count{K}(a, n, val) + 1;
*/
```

Listing 9.23: The predicate `MultisetAdd`

The predicate `MultisetMinus` in Listing 9.24, on the other hand, expresses that the number of occurrences of a specific element in an array has decreased by one between two program points `K` and `L`.

```
/*@
  predicate
    MultisetMinus{K,L}(value_type* a, integer n, value_type val) =
      Count{L}(a, n, val) == Count{K}(a, n, val) - 1;
*/
```

Listing 9.24: The predicate `MultisetMinus`

Note that we could have defined `MultisetMinus` also by calling `MultisetAdd` with the labels reversed.

```
  predicate
    MultisetMinus{K,L}(value_type* a, integer n, value_type val) =
      MultisetAdd{L,K}(a, n, val);
```

It is a often only a matter of taste how to decide which of several ways to define a predicate is more appropriate. However, one also has to take into account which definition can be handled more easily by Frama-C/WP and its associated theorem provers.

The predicate `MultisetRetainRest`

In order to achieve a concise specification we introduce the overloaded predicate `MultisetRetainRest` (see Listing 9.25). The expression `MultisetRetainRest{K,L}(a, m, b, n, v)` is true if the range `a[0..n-1]` at time `K` contains the same elements as `b[0..m-1]` at time `L`, except possibly for occurrences of `v`; the elements' order may differ in `a` and `b`. There is also a more general version of `MultisetRetainRest` in Listing 9.25 that is defined over array segments.

```
/*@
  predicate
    MultisetRetainRest{K,L}(value_type* a, integer m1, integer m2,
                            value_type* b, integer n1, integer n2,
                            value_type  v) =

      \forall value_type x;
        x != v ==> Count{K}(a, m1, m2, x) == Count{L}(b, n1, n2, x);


  predicate
    MultisetRetainRest{K,L}(value_type* a, integer m,
                            value_type* b, integer n,
                            value_type  v) =
    MultisetRetainRest{K,L}(a, 0, m, b, 0, n, v);
*/
```

Listing 9.25: The predicate `MultisetRetainRest`

For `push_heap` another overloaded version of `MultisetRetainRest` (Listing 9.26) is particularly useful. The new version holds if the number of occurrences for all elements, except the two given ones, remains unchanged between two program points.

```
/*@
  predicate
    MultisetRetainRest{K,L}(value_type* a, integer n,
                            value_type  v, value_type w) =
      \forall value_type x;
        x != v  ==>  x != w  ==>  MultisetRetain{K,L}(a, n, x);
*/
```

Listing 9.26: An overloaded version of predicate `MultisetRetainRest`

The definition of this new version of the `MultisetRetainRest` predicate uses the simpler predicate `MultisetRetain` shown in Listing 9.27. This predicate holds if the number of occurrences of a given value in an array does not change between two program points. We will later also directly employ `MultisetRetain` to succinctly formulate additional assertions in `push_heap`.

```
/*@
  predicate
    MultisetRetain{K,L}(value_type* a, integer n, value_type v) =
      Count{K}(a, n, v) == Count{L}(a, n, v);
*/
```

Listing 9.27: The predicate `MultisetRetain`

Main act

The goal of the main act is to locate the array index to which the new element can be assigned. Listing 9.28 shows its implementation.

```
      // start of main act
      if (0u < hole) {
        size_type parent = heap_parent(hole);

        /*@
          loop invariant bound:  0 <= hole < n-1;
          loop invariant heap:   IsHeap(a, n);
          loop invariant heap:   parent == HeapParent(hole);
          loop invariant less:   a[hole] < v;
          loop invariant add:    MultisetAdd{Pre,Here}(a, n, a[hole]);
          loop invariant minus:  MultisetMinus{Pre,Here}(a, n, v);
          loop invariant retain: MultisetRetainRest{Pre,Here}(a, n, v, a[hole]);
          loop assigns           hole, parent, a[0..n-1];
          loop variant           hole;
        */
        while ((0u < hole) && (a[parent] < v)) {
          //@ ghost Loop: // LoopEntry not yet supported!
          //@ ghost const value_type old_a = a[hole];
          //@ assert reorder:  old_a == \at(a[hole],Loop);
          if (a[hole] < a[parent]) {
            a[hole] = a[parent];
            //@ assert less:    old_a  < v;
            //@ assert less:    a[hole] < v;
            //@ assert retain:  MultisetUnchanged{Loop,Here}(a, 0, hole);
            //@ assert retain:  MultisetUnchanged{Loop,Here}(a, hole + 1, n);
            //@ assert minus:   MultisetMinus{Loop,Here}(a, n, old_a);
            //@ assert add:     MultisetAdd{Loop,Here}(a, n, a[hole]);
            //@ assert retain:  MultisetRetain{Loop,Here}(a, n, v);
            //@ assert retain:  MultisetRetain{Pre,Here}(a, n, old_a);
            //@ assert retain:  MultisetRetainRest{Pre,Here}(a, n, v, a[hole]);
          }

          hole = parent;

          if (0u < hole) {
            parent = heap_parent(hole);
          }
        }
      }

      // end of main act
```

Listing 9.28: Main act of `push_heap` implementation

The loop invariant `heap` expresses that the predicate `IsHeap` is true for the array throughout the main act. Instead of an invariant `reorder` that reflects the postcondition with the same name, we now consider the invariants `add`, `minus`, and `retain`.

It is important to understand the use of the variable `hole` in these loop invariants. Before each loop iteration, `hole` stores the index of the node whose value was assigned to one of its children in the previous iteration or in the prologue (for the first loop run). Therefore, the value `a[hole]` appears in the loop invariants `add` and `retain`.

Verifying the various loop invariants and assertions has been far from being straightforward, and required additional assertions and the *ghost* label `Loop`. The following remarks highlight some of the issues.

1. The `heap` property implies that `a[hole] <= a[parent]` always holds. Thus, the assignment `a[hole] = a[parent]` might be redundant. We have not check whether guarding this assignment with the condition `a[hole] < a[parent]` is more efficient. Important for us is the following: the guard allows us to put additional assertions where they really matter and where they can more easily be verified.

2. In order to guide the automatic provers, we have also provided the lemmas `MultisetAddDistinct` (Listing 9.29), `MultisetMinusDistinct` (Listing 9.30), and `MultisetPushHeapRetain` (Listing 9.31). These lemmas formalize conditions under which the respective predicates `MultisetAdd`, `MultisetMinus`, and `MultisetRetain` apply.

```
/*@
  lemma
    MultisetAddDistinct{K,L}:
      \forall value_type *a, v, integer i, n;
        0 <= i < n                        ==>
        \at(a[i],K) != v                  ==>
        \at(a[i],L) == v                  ==>
        MultisetUnchanged{K,L}(a, 0, i)   ==>
        MultisetUnchanged{K,L}(a, i+1, n) ==>
        MultisetAdd{K,L}(a, n, v);
*/
```

Listing 9.29: The lemma `MultisetAddDistinct`

```
/*@
  lemma
    MultisetMinusDistinct{K,L}:
      \forall value_type *a, v, integer i, n;
        0 <= i < n                        ==>
        \at(a[i],K) == v                  ==>
        \at(a[i],L) != v                  ==>
        MultisetUnchanged{K,L}(a, 0, i)   ==>
        MultisetUnchanged{K,L}(a, i+1, n) ==>
        MultisetMinus{K,L}(a, n, v);
*/
```

Listing 9.30: The lemma `MultisetMinusDistinct`

```
/*@
 lemma
  MultisetPushHeapRetain{K,L,M}:
    \forall value_type *a, ap, ah, v, integer h, p, n;
      0 <= p < h < n-1                      ==>
      ah < ap < v                           ==>
      \at(a[h],L)  ==  ah                   ==>
      \at(a[p],L)  ==  ap                   ==>
      \at(a[h],M)  ==  ap                   ==>
      MultisetMinus{K,L}(a, n, v)           ==>
      MultisetAdd{K,L}(a, n, ah)            ==>
      MultisetRetainRest{K,L}(a, n, v, ah)  ==>
      MultisetUnchanged{L,M}(a, 0, h)       ==>
      MultisetUnchanged{L,M}(a, h+1, n)     ==>
      MultisetRetainRest{K,M}(a, n, v, ap);
*/
```

Listing 9.31: The lemma `MultisetPushHeapRetain`

Figure 9.32 shows the array after the main act. The contents of the dashed nodes have been overwritten with the values of their parents until `hole` reached a node to which `val` can be assigned, whilst maintaining the `heap` property.

Figure 9.32.: Heap after the main act of `push_heap`

In our example, the loop performs just one assignment, viz. `a[5]=a[2]`, and then stops with `hole` being 2. At this point, the new element 12 can be assigned to the node with the index 2 and the `heap` property stays intact. This assignment takes place in the epilogue.

Epilogue

The last part of the implementation is the epilogue, shown in Listing 9.33. It consists of exactly one assignment which re-establishes the `reorder` property while maintaining the `heap` property already established.

```
      // start of epilogue
      //@ ghost Epi:
      a[hole] = v;
      //@ assert value:      \at(a[hole],Epi)  != v;
      //@ assert value:      \at(a[hole],Here) == v;
      //@ assert unchanged:  MultisetUnchanged{Epi,Here}(a, 0, hole);
      //@ assert unchanged:  MultisetUnchanged{Epi,Here}(a, hole+1, n);
      //@ assert add:        MultisetAdd{Epi,Here}(a, n, v);
      //@ assert minus:      MultisetMinus{Epi,Here}(a, n, \at(a[hole],Epi));
      //@ assert reorder:    MultisetUnchanged{Pre,Here}(a, n);
    }
  }

  // end of epilogue
}
```

Listing 9.33: Epilogue of `push_heap` implementation

The ghost label `Epi`, together with a couple of assertions and the lemma `MultisetPushHeapClosure` are necessary to help the automatic theorem provers to prove the final assertion `reorder`.

```
/*@
 lemma
  MultisetPushHeapClosure{K,L,M}:
    \forall value_type *a, u, v, integer i, n;
      0 <= i < n-1                         ==>
      u != v                               ==>
      \at(a[i],M)  == v                    ==>
      MultisetAdd{K,L}(a, n, u)            ==>
      MultisetMinus{K,L}(a, n, v)          ==>
      MultisetRetainRest{K,L}(a, n, v, u)  ==>
      MultisetUnchanged{L,M}(a, 0, i)      ==>
      MultisetUnchanged{L,M}(a, i+1, n)    ==>
      MultisetAdd{L,M}(a, n, v)            ==>
      MultisetMinus{L,M}(a, n, u)          ==>
      MultisetUnchanged{K,M}(a, n);
*/
```

Listing 9.34: The lemma `MultisetPushHeapClosure`

Concerning the `reorder` property, the main act finished with an increased count of nodes with the value `11` and a decreased count of nodes with the value `12` (cf. Figure 9.32). The heap in Figure 9.35, on the other hand, shows the tree after the epilogue has assigned the value `12` to the node with the index `2`, which contained the value `11`. Hence the `reorder` property is re-established and the function can return. Moreover, since the `heap` property has been inferred, all nodes are now colored blue.

Figure 9.35.: Heap after the epilogue of `push_heap`

9.7. The `pop_heap` algorithm

Whereas in the C++ Standard Library [21, §28.7.7.2] `pop_heap` works on a range of random access iterators, our version operates on an array of `value_type`. We therefore use the following signature for `pop_heap`

```
void pop_heap(value_type* a, size_type n);
```

The `pop_heap` algorithm expects that n is greater or equal than 1 and that the array `a[0..n-1]` forms a heap. The algorithms then *rearranges* the array `a[0..n-1]` such that the resulting array satisfies the following properties.

- `a[n-1] = \old(a[0])`, that is, the largest element[30] of the original heap is transferred to the end of the array

- the subarray `a[0..n-2]` is a heap

In this sense the algorithm *pops* the largest element from a heap.

9.7.1. Formal Specification of `pop_heap`

The ACSL function contract of `pop_heap` is given in Listing 9.36.

```
/*@
   requires bounds: 0 < n;
   requires valid:  \valid(a + (0..n-1));
   requires heap:   IsHeap(a, n);

   assigns a[0..n-1];

   ensures heap:    IsHeap(a, n-1);
   ensures result:  a[n-1] == \old(a[0]);
   ensures max:     MaxElement(a, n, n-1);
   ensures reorder: MultisetUnchanged{Old, Here}(a, n);
*/
void
pop_heap(value_type* a, size_type n);
```

Listing 9.36: Formal specification of `pop_heap`

9.7.2. Implementation of `pop_heap`

Listing 9.37 shows our implementation of `pop_heap` together with ACSL annotations. Note that in this version the postcondition `reorder` of `pop_heap`, which states that the algorithm only *rearranges* the elements of the array, is not verified by Frama-C/WP. Verifying this postcondition would require more elaborate loop invariants which we will supply in a later version of this document.

---

[30]See Lemma `HeapMaximum` in Listing 9.9.

```
void
pop_heap(value_type* a, size_type n)
{
  if (1u < n) { // otherwise nothings needs to be done
    const value_type v = a[0u];

    //@ assert max:  MaxElement(a, n, 0);
    if (a[n - 1u] < v) { // otherwise nothings needs to be done
      //@ assert bounds: 2 <= n;
      size_type hole = 0u;
      size_type child = heap_child_max(a, n, hole);

      //@ assert heap: child < n - 1  ==>  hole == HeapParent(child);
      /*@
          loop invariant bounds: 0 <= hole < n-1;
          loop invariant bounds: hole < child;
          loop invariant heap:   IsHeap(a, n);
          loop invariant heap:   a[n-1] < a[HeapParent(hole)];
          loop invariant heap:   child < n - 1  ==>  hole == HeapParent(child);
          loop invariant child:  HeapChildMax(a, n, hole, child);
          loop invariant max:    UpperBound(a, 0, n, v);
          loop assigns           hole, child, a[0..n-2];
          loop variant           n - hole;
       */
      while ((child < n - 1u) && (a[n - 1u] < a[child])) {
        a[hole] = a[child];
        hole    = child;
        //@ assert heap: IsHeap(a, n);
        child = heap_child_max(a, n, hole);
      }

      //@ assert child: child < n-1 ==> a[n-1] >= a[child];
      //@ assert child: HeapChildMax(a, n, hole, child);
      //@ assert heap: IsHeap(a, n);
      //@ assert heap: a[n-1] < a[HeapParent(hole)];
      a[hole]   = a[n - 1u];
      //@ assert heap: IsHeap(a, n-1);
      a[n - 1u] = v;
      //@ assert heap: IsHeap(a, n-1);
    }
  }
}
```

Listing 9.37: Implementation of `pop_heap`

Our implementation relies on the auxiliary function `heap_child_max` from Listing 9.13 (Section 9.3). As discussed there this helper function computes the child of an heap node with the the largest index. On the logic level this notion is expressed in ACSL by the predicate `HeapChildMax` from Listing 9.12.

9.8. The `make_heap` algorithm

Whereas in the C++ Standard Library [21, §28.7.7.3] make_heap works on a pair of generic random access iterators, our version operators on a range of value_type. Thus the signature of make_heap reads

```
void make_heap(value_type* a, size_type n);
```

The function make_heap rearranges the elements of the given array a[0..n-1] such that they form a heap.

As an examples we look at the array in Figure 9.38. The elements of this array do not form a heap, as indicated by the grey colouring. Executing the make_heap algorithm on this array rearranges its elements so that they form a heap as shown in Figure 9.4.

8	8	7	3	14	11	3	6	2	3	13	9

Figure 9.38.: Array before the call of make_heap

9.8.1. Formal Specification of `make_heap`

Listing 9.39 shows the ACSL specification of make_heap.

```
/*@
   requires valid:  \valid(a + (0..n-1));

   assigns a[0..n-1];

   ensures heap:    IsHeap(a, n);
   ensures reorder: MultisetUnchanged{Old,Here}(a, n);
*/
void
make_heap(value_type* a, size_type n);
```

Listing 9.39: The Specification of make_heap

Like with push_heap the formal specification of make_heap must ensure that the resulting array is a heap of size n and contains the same multiset of elements as in the pre-state of the function. These properties are expressed by the heap and reorder postconditions respectively. The reorder postcondition uses the predicate MultisetUnchanged (see Listing 6.5) to ensure make_heap only rearranges the array elements.

9.8.2. Implementation of `make_heap`

The implementation of `make_heap`, shown in Listing 9.40, is straightforward. From low to high the array's elements are pushed to the growing heap. We used `i < n` as loop condition, rather than the more tempting `i <= n`, in order to admit also `n == SIZE_TYPE_MAX`; as a consequence, we had to call `push_heap` with `i+1`. The iteration starts at `i+1 == 2`, because an array with length one is a heap already.

```
void
make_heap(value_type* a, size_type n)
{
  if (0u < n) {
    /*@
        loop invariant bounds:    1 <= i <= n;
        loop invariant heap:      IsHeap(a, i);
        loop invariant reorder:   MultisetUnchanged{Pre,Here}(a, i+1);
        loop invariant unchanged: Unchanged{Pre,Here}(a, i+1, n);
        loop assigns   i, a[0..n-1];
        loop    variant n - i;
    */
    for (size_type i = 1u; i < n; ++i) {
      push_heap(a, i + 1u);
    }
  }

  //@ assert  heap: IsHeap(a, n);
}
```

Listing 9.40: The Implementation of `make_heap`

Since the loop statement consists just of a call to `push_heap` we obtain the both loop invariants `heap` and `reorder` by simply lifting them from the contract of `push_heap` (see Section 9.6.1).

The postcondition of `push_heap` only specifies the multiset of elements from index 0 to `i`. We therefore also have to specify that the elements from index `i+1` to `n-1` are only reordered. This property can be derived from the `unchanged` property of `push_heap`.

9.9. The `sort_heap` algorithm

Whereas in the C++ Standard Library [21, §28.7.7.4] `sort_heap` works on a range of random access iterators, our version operates on an array of `value_type`. We therefore use the following signature for `sort_heap`

```
void sort_heap(value_type* a, size_type n);
```

The function `sort_heap` rearranges the elements of a given heap `a[0..n-1]` into an array that is sorted in increasing order. Thus, applying `sort_heap` to the heap in Figure 9.4 produces the sorted array in Figure 9.41.

2	3	3	3	6	7	8	8	9	11	13	14

Figure 9.41.: Array after the call of `sort_heap`

9.9.1. Formal Specification of `sort_heap`

Listing 9.42 shows our ACSL specification of `sort_heap`. The formal specification of `sort_heap` must ensure that the resulting array is sorted. Furthermore the multiset contained by the array must be the same as in the pre-state of the function. The postconditions `sorted` and `reorder` express these properties, respectively. The specification effort is relatively simple because we can reuse the previously defined predicates `MultisetUnchanged` (Listing 6.5) and `Sorted` (Listing 5.1).

```
1   /*@
2     requires valid:  \valid(a + (0..n-1));
3     requires heap:   IsHeap(a, n);
4
5     assigns a[0..n-1];
6
7     ensures sorted:   Sorted(a, n);
8     ensures reorder:  MultisetUnchanged{Old, Here}(a, n);
9   */
10  void
11  sort_heap(value_type* a, size_type n);
```

Listing 9.42: Formal specification of `sort_heap`

9.9.2. Implementation of `sort_heap`

The implementation of `sort_heap` (Listing 9.43) is relatively simple because it relies on the `pop_heap` algorithm performing essential work.

Our implementation of `sort_heap` repeatedly calls `pop_heap` to extract the maximum of the shrinking heap and adding it to the sorted part of the array.

```
void
sort_heap(value_type* a, size_type n)
{
  /*@
     loop invariant bound:    0 <= i <= n;
     loop invariant heap:     IsHeap(a, i);
     loop invariant sorted:   Sorted(a, i, n);
     loop invariant lower:    LowerBound(a, i, n, a[0]);
     loop invariant reorder:  MultisetUnchanged{Pre,Here}(a, 0, n);
     loop assigns i, a[0..n-1];
     loop variant i;
  */
  for (size_type i = n; i > 1u; --i) {
    /*@
       requires heap:     IsHeap(a, i);
       assigns  a[0..i-1];
       ensures  heap:     IsHeap(a, i-1);
       ensures  max:      a[i-1] == \old(a[0]);
       ensures  max:      MaxElement(a, i, i-1);
       ensures  reorder:  MultisetUnchanged{Old,Here}(a, 0, i);
       ensures  reorder:  Unchanged{Old,Here}(a, i, n);
    */
    pop_heap(a, i);
    //@ assert lower:  LowerBound(a, i, n, a[i-1]);
  }
}
```

Listing 9.43: The Implementation of `sort_heap`

The loop invariants of `sort_heap` describe the content of the array in two parts. The first `i` elements form a heap and are described by the `heap` invariant. The last `n-i` elements are already sorted.

Supporting lemmas

In order to facilitate the automatic verification of the property `sorted` we rely among others on the properties `lower` and `max` and the lemma `SortedUpperBound` from Listing 9.44.

```
/*@
  lemma
    SortedUpperBound{L}:
      \forall value_type *a, integer n;
        UpperBound(a, n, a[n])   ==>
        Sorted(a, n)             ==>
        Sorted(a, n+1);
*/
```

Listing 9.44: The lemma `SortedUpperBound`

To verify the property `reorder` we formulate in Listing 9.45 several lemmas that express that the properties

- `MultisetUnchanged{K,L}(a, 0, i)` and

- `Unchanged{Old,Here}(a, i, n)`

imply the desired loop invariant `MultisetUnchanged{K,L}(a, 0, n)`.

```
/*@
  lemma
    UnchangedImpliesMultisetUnchanged{L1,L2}:
      \forall value_type *a, integer k, n;
        Unchanged{L1,L2}(a, k, n) ==>
        MultisetUnchanged{L1,L2}(a, k, n);


  lemma
    MultisetUnchangedUnion{L1,L2}:
      \forall value_type *a, integer i, k, n;
        0 <= i <= k <= n                    ==>
        MultisetUnchanged{L1,L2}(a, i, k)   ==>
        MultisetUnchanged{L1,L2}(a, k, n)   ==>
        MultisetUnchanged{L1,L2}(a, i, n);


  lemma
    MultisetUnchangedTransitive{L1,L2,L3}:
      \forall value_type *a, integer n;
        MultisetUnchanged{L1,L2}(a, n)   ==>
        MultisetUnchanged{L2,L3}(a, n)   ==>
        MultisetUnchanged{L1,L3}(a, n);
*/
```

Listing 9.45: Some lemmas for `MultisetUnchanged`

10. Sorting Algorithms

In this chapter, we present algorithms of the C++ Standard Library [21, §28.7.1] that are related to the task of sorting a linear array.

- Section 10.1 shows an algorithm to check if a given array is already sorted in ascending order.

- The algorithm in Section 10.2 partitions a given array, and sorts only the resulting lower part.

Many issues in computer science can be exemplified in the field of sorting algorithms; see e.g. [30] for a famous textbook. Therefore we arrange some of the most common classic sorting algorithms.

Following [31], we have also used (C rephrasings of) functions from the C++ Standard Library as far as possible to implement the different algorithmic approaches.

- `heap_sort` in Section 10.3 describes the quite efficient *heap sort*, which relies on the algorithms presented in Chapter 9.[31]

- `selection_sort` in Section 10.4 presents the classic *selection sort* algorithm.[32]

- `insertion_sort` in Section 10.5 the also well-known *insertion sort* algorithm.[33]

- `merge` in Section 10.6 the *merge* algorithm from *merge sort*.[34]

These algorithms essentially share the following contract; it is their implementations that differ fundamentally.

```
/*@
   requires valid: \valid(a + (0..n-1));

   assigns  a[0..n-1];

   ensures sorted:  Sorted(a, n);
   ensures reorder: MultisetUnchanged{Old, Here}(a, n);
*/
void xxx_sort(value_type* a, size_type n);
```

While `heap_sort` achieves a run-time complexity upper bound of $O(n \cdot \log(n))$ due to the efficiency of the heap data structure, both `selection_sort` and `insertion_sort` need $O(n^2)$ in the average case, and also in the worst case.

Note that the `sort` algorithm from the C++ Standard Library is not handled here because it typically relies on *introspection sort* which is sophisticated mix of various classic algorithms.[35] In future releases we plan to handle the more algorithms related sorting.

---

[31]See `https://en.wikipedia.org/wiki/Heapsort`

[32]See `https://en.wikipedia.org/wiki/Selection_sort`

[33]See `https://en.wikipedia.org/wiki/Insertion_sort`

[34]See `https://en.wikipedia.org/wiki/Merge_sort`

[35]See `https://en.wikipedia.org/wiki/Introsort`

10.1. The `is_sorted` algorithm

Our version of the `is_sorted` algorithm compared to the C++ Standard Library [21, §28.7.1.5] has the signature

```
bool is_sorted(const value_type* a, size_type n);
```

It returns **true** if the given array is in ascending order, and **false** else.

10.1.1. Formal Specification of `is_sorted`

Listing 10.1 shows the ACSL specification of `is_sorted`. In the contract, we use the predicate `Sorted`, (see Listing 5.1) which states that any array element is always less or equal to any other element right of it. We'll use an easier-to-handle predicate in the implementation, see below.

```
/*@
  requires valid: \valid_read(a + (0..n-1));

  assigns  \nothing;

  ensures  \result <==> Sorted(a, n);
*/
bool
is_sorted(const value_type* a, size_type n);
```

Listing 10.1: The Specification of `is_sorted`

10.1.2. Implementation of `is_sorted`

The implementation of `is_sorted` is shown in Listing 10.3. As usual, it doesn't compare every array element to all that are right to it, but only to the immediately adjacent one, which is of course more efficient. For this reason, we use the predicate `WeaklySorted`, shown in Listing 10.2, in the loop invariant of the implementation.

Users inexperienced in formal verification often have a blind spot at the difference between `Sorted` and `WeaklySorted`. Both versions are logically equivalent, and proving `Sorted` ⟹ `WeaklySorted` is even trivial. However, proving the converse direction is not, and requires an induction on the array size n, employing the transitivity of <= in the induction step. Humans are trained to perform such inductions unnoticed, but none of the automated provers supported by Frama-C is able to perform induction at all.

```
/*@
  predicate
    WeaklySorted{L}(value_type* a, integer m, integer n) =
      \forall integer i; m <= i < n-1 ==> a[i] <= a[i+1];

  predicate
    WeaklySorted{L}(value_type* a, integer n) = WeaklySorted{L}(a, 0, n);
*/
```

Listing 10.2: The predicate `WeaklySorted`

Since our implementation uses `WeaklySorted` in its loop invariant, and follows the same principle in its code, its verification is straight-forward — except for the final reasoning that `WeaklySorted(a,n)` implies `Sorted(a,n)`. We have an own lemma for that step, shown in Listing 10.4, which needs to be proven manually with Coq. The converse lemma (Listing 10.4) is proven automatically, but isn't actually needed to verify our `is_sorted` implementation.

Alternatively, we could have dragged the predicate `Sorted` along the loop, which happens to cause no particular problems in this case.

```
bool
is_sorted(const value_type* a, size_type n)
{
  if (0u < n) {
    /*@
        loop invariant sorted: WeaklySorted(a, i+1);
        loop assigns i;
        loop variant n - i;
    */
    for (size_type i = 0u; i < n - 1u; ++i) {
      if (a[i] > a[i + 1u]) {
        return false;
      }
    }
  }

  return true;
}
```

Listing 10.3: The implementation of `is_sorted`

```
/*@
  lemma
    SortedImpliesWeaklySorted{L}:
      \forall value_type* a, integer m, n;
        0 <= m <= n     ==>
        Sorted(a, m, n) ==>
        WeaklySorted(a, m, n);

  lemma
    WeaklySortedImpliesSorted{L}:
      \forall value_type* a, integer m, n;
        0 <= m <= n          ==>
        WeaklySorted(a, m, n) ==>
        Sorted(a, m, n);
*/
```

Listing 10.4: The lemmas `SortedImpliesWeaklySorted` and `WeaklySortedImpliesSorted`

10.2. The `partial_sort` algorithm

Our version of the `partial_sort` algorithm compared to the C++ Standard Library [21, §28.7.1.3] has the signature

```
void partial_sort(value_type* a, size_type m, size_type n);
```

The algorithm *reorders* the given array `a` in such a way that it represents a *partition*: each member of the left part `a[0..m-1]` is less or equal to each member of the right part `a[m..n-1]`. Moreover, the algorithm *sorts* the left part in increasing order. The order of elements in the right part, however, is *unspecified*. Figure 10.5 uses a bar chart to depict a typical result of a call `partial_sort(a, m, n)`. In the post-state, the left and the right part is colored in green and orange, respectively.
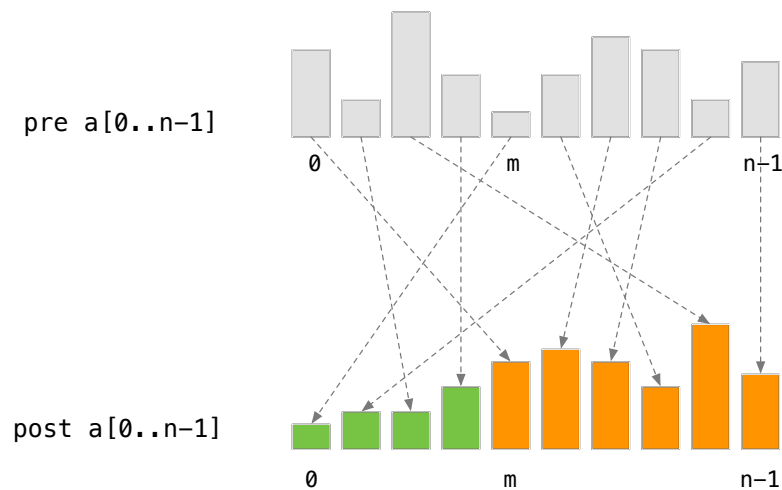
Figure 10.5.: Effects of `partial_sort`
.

10.2.1. Formal specification of `partial_sort`

We start this section by introducing in Listing 10.6 the new predicate `Partition` which formalizes the partitioning property.

```
/*@
  predicate
    Partition{L}(value_type* a, integer m, integer n) =
      0 <= m <= n ==>
      \forall integer i, k; 0 <= i < m <= k < n ==> a[i] <= a[k];
*/
```

Listing 10.6: The predicate `Partition`

The formal specification of the `partial_sort` function is shown in Listing 10.7. It uses the just introduced predicate `Partition` and reuses the previously defined predicates `Sorted` (shown in Listing 5.1) and `MultisetUnchanged` (Listing 6.5).

```
/*@
  requires valid: \valid(a + (0..n-1));
  requires split: 0 <= m <= n;

  assigns  a[0..n-1];

  ensures sorted:     Sorted(a, m);
  ensures partition:  Partition(a, m, n);
  ensures reorder:    MultisetUnchanged{Old,Here}(a, n);
*/
void
partial_sort(value_type* a, size_type m, size_type n);
```

Listing 10.7: The Specification of `partial_sort`

10.2.2. Implementation of `partial_sort`

Our implementation is shown in Listing 10.9 and 10.10. It initially calls `make_heap` (Section 9.8) to rearrange the left part `a[0..m-1]` into a heap. After that, it scans the right part, from left to right, for elements that are too small; each such element is exchanged for the left part's maximum, by applying `pop_heap` (9.7) and `push_heap` (9.6) appropriately. When the scan is done, the smallest elements are collected in the left part. We finally convert it from a heap into an ascending range, by `sort_heap` (9.9).

Figure 10.8.: An iteration of `partial_sort`

In the scan loop, we maintain as invariants

- that the left part is a heap (invariant `heap`);

- that its maximal element, `a[0]`, is a "separating element" between the left part `a[0..m-1]` and the right part `a[m..i-1]`, i.e., an upper bound of the left (invariant `upper`) and a lower bound of the right part (invariant `lower`), respectively;

- that `a[i..m-1]` is yet unchanged (invariant `unchanged`); and

- that only permutation operations have been applied to `a[0..i-1]` (invariant `reorder`).

217

```
void
partial_sort(value_type* a, size_type m, size_type n)
{
  if (m > 0u) {
    make_heap(a, m);

    //@ assert reorder: Unchanged{Pre,Here}(a, m, n);
    /*@
      loop invariant bound:     m <= i <= n;
      loop invariant heap:      IsHeap(a, m);
      loop invariant upper:     UpperBound(a, 0, m, a[0]);
      loop invariant lower:     LowerBound(a, m, i, a[0]);
      loop invariant reorder:   MultisetUnchanged{Pre,Here}(a, i);
      loop invariant unchanged: Unchanged{Pre,Here}(a, i, n);
      loop assigns              i, a[0..n-1];
      loop variant              n-i;
    */
    for (size_type i = m; i < n; ++i)
      if (a[i] < a[0u]) {
        /*@
          assigns              a[0..m-1];
          ensures heap:        IsHeap(a, m-1);
          ensures max:         a[m-1] == \old(a[0]);
          ensures max:         MaxElement(a, m, m-1);
          ensures reorder:     MultisetUnchanged{Old, Here}(a, m);
          ensures unchanged:   Unchanged{Old, Here}(a, m, i);
          ensures unchanged:   Unchanged{Old, Here}(a, m, n);
        */
        pop_heap(a, m);
        //@ assert lower:     a[0] <= a[m-1];
        //@ assert lower:     a[i] <  a[m-1];
        //@ assert lower:     LowerBound(a, m, i, a[m-1]);
        //@ assert upper:     UpperBound(a, 0, m-1, a[0]);
        //@ assert upper:     UpperBound(a, 0, m,   a[m-1]);
        //@ assert partition: Partition(a, m, i);
        //@ assert reorder:   MultisetUnchanged{Pre,Here}(a, i);
```

Listing 10.9: The Implementation of `partial_sort` (1)

In order to preserve the loop invariants after `i` is incremented, nothing has to be done if `a[0]` happens to be also a lower bound for `a[i]`. Otherwise, let us follow the algorithm through the `then` part code, depicting the intermediate states in Figure 10.8. The elements considered so far are shown colored similar to Figure 10.5; in particular the heap part is shown in green. The overlaid transparent red shape indicates the ranges to which `Partition` applies, in each state. The figure assumes the initial contents of `a[0]` and `a[i]` to be 9 and 5, for sake of generality, let us call them p and q, respectively.

After `pop_heap` and `swap`, we have p at `a[i]`, and q at `a[m-1]`. At that point we know

1. $q < p \leq$ `a[k]` for each $m \leq k < i$, since p was a lower bound for `a[m..i-1]`;

2. $q < p =$ `a[i]`;

3. `a[j]` $\leq p \leq$ `a[k]` for each $0 \leq j < m - 1$ and each $m \leq k < i$, since this held on loop entry, and we didn't more than reordering inside the parts; and

4. `a[j]` $\leq p =$ `a[i]` since p was the heap maximum on loop entry.

Altogether, we have `a[j]` $\leq p \leq$ `a[k]` for each $0 \leq j < m$ and each $m \leq k < i + 1$. That is, `Partition` `(a,m,i+1)` holds, although we cannot name a separating element of `a` here.

```
        //@ ghost BeforeSwap:
        swap(a + m - 1u, a + i);
        //@ assert swapped:    SwappedInside{BeforeSwap,Here}(a, m-1, i, n);
        //@ assert reorder:    MultisetUnchanged{BeforeSwap,Here}(a, i+1);
        //@ assert reorder:    MultisetUnchanged{Pre,Here}(a, i+1);
        //@ assert unchanged:  Unchanged{Pre,Here}(a, i+1, n);
        //@ assert lower:      a[m-1] < a[i];
        //@ assert lower:      \forall integer k; 0 <= k < m ==> LowerBound(a, m, i
            +1, a[k]);
        //@ assert upper:      UpperBound(a, 0, m-1, a[0]);

        /*@
          assigns           a[0..m-1];
          ensures heap:      IsHeap(a, m);
          ensures reorder:   MultisetUnchanged{Old,Here}(a, m);
          ensures unchanged: Unchanged{Old,Here}(a, m, i+1);
          ensures unchanged: Unchanged{Old,Here}(a, i+1, n);
        */
        push_heap(a, m);
        //@ assert upper:      UpperBound(a, 0, m,   a[0]);
        //@ assert lower:      LowerBound(a, m, i+1, a[0]);
    }

    //@ assert partition: Partition(a, m, n);
    /*@
      assigns           a[0..m-1];
      ensures sorted:    Sorted(a, m);
      ensures reorder:   MultisetUnchanged{Old,Here}(a, m);
      ensures reorder:   MultisetUnchanged{Old,Here}(a, m, n);
    */
    sort_heap(a, m);
    //@ assert reorder:   MultisetUnchanged{Pre,Here}(a, n);
    //@ assert partition: Partition(a, m, n);
  }
}
```

Listing 10.10: The Implementation of `partial_sort` (2)

After calling `push_heap`, which just performs some more reorderings of the left part,[36] this property is preserved. Moreover, we now know again that `a[0]` has become an upper bound of the left part, and hence a separating element between `a[0..m-1]` and `a[m..i]`; that is, the loop invariants `upper` and `lower` have been re-established. These two invariants together are eventually used to prove the property `partition` of the contract.

Compared to its size, the algorithm makes a lot of procedure calls; in this respect it is closer to real-life software than most other algorithms of this tutorial. Therefore, we use it to illustrate a methodical point: For almost every procedure call, we give the callee's contract, tailored to its actual parameters, as a statement contract of the call. For example, everything we know from the `pop_heap` contract, instantiated to the particular situation, is documented in the first statement contract (Listing 10.9). In contrast, we use `assert` clauses to indicate intermediate reasoning to obtain subsequently needed properties.

Our implementation has a worst-case time complexity of $O((n + m) \cdot \log m)$. On the other hand, an implementation that ignores `m` and just sorts `a[0..n-1]` also satisfies the contract in Listing 10.7, and may have $O(n \cdot \log n)$ complexity. Some arithmetic shows that `partial_sort` performs better than plain sort if, and only if, $\log m < \frac{n}{m} \cdot \log\left(\frac{n}{m}\right)$, that is, if n is sufficiently larger than m.

---

[36]We can't and we needn't tell which position q is moved to; the former is indicated in Figure 10.5 by the vague grey triangle.

Lemmas used in partition proofs

The following lemmas are use in proofs of properties and annotations related to the loop invariants `upper` and `lower`. Lemma `ReorderPreservesUpperBound` (Listing 10.11) informally says that a lower bound `v` of a range `a[0..n-1]` keeps its property even after the range is reordered.

```
/*@
  lemma
    ReorderPreservesUpperBound{K,L}:
      \forall value_type* a, integer n, value_type v;
        0 <= n                      ==>
        MultisetUnchanged{K,L}(a, n) ==>
        UpperBound{K}(a, n, v)       ==>
        UpperBound{L}(a, n, v);
*/
```

Listing 10.11: The lemma `ReorderPreservesUpperBound`

Dually, lemma `ReorderPreservesLowerBound` (Listing 10.11) says that reordering a range doesn't affect any of its upper bounds.

```
/*@
  lemma
    ReorderPreservesLowerBound{K,L}:
      \forall value_type* a, integer n, value_type v;
        0 <= n                      ==>
        MultisetUnchanged{K,L}(a, n) ==>
        LowerBound{K}(a, n, v)       ==>
        LowerBound{L}(a, n, v);
*/
```

Listing 10.12: The lemma `ReorderPreservesLowerBound`

Lemma `PartialReorderPreservesLowerBounds` (Listing 10.13) describes a more particular situation: if each element in `a[0..m-1]` is known to be a lower bound of `a[m..n-1]`, and the former range is reordered while the latter is kept untouched, then `a[0]` will still be a lower bound of `a[m..n-1]`. We employ this lemma to infer that, after `push_heap` was called, the new heap maximum `a[0]`, is a lower bound of `a[m..i]`,

```
/*@
  lemma
    PartialReorderPreservesLowerBounds{K,L}:
      \forall value_type* a, integer m, n;
        0 < m <= n                              ==>
        (\forall integer k; 0 <= k < m
          ==> LowerBound{K}(a, m, n, \at(a[k],K))) ==>
        MultisetUnchanged{K,L}(a, 0, m)          ==>
        Unchanged{K,L}(a, m, n)                  ==>
        LowerBound{L}(a, m, n, \at(a[0],L));
*/
```

Listing 10.13: The lemma `PartialReorderPreservesLowerBounds`

Lemma `ReorderImpliesMatch` states that a value `a[i]` taken from a range `a[0..n-1]` after some reordering must have been in that range already before reordering. It is used to prove the lemmas above.

```
/*@
  lemma
    ReorderImpliesMatch{K,L}:
      \forall value_type *a, integer i, n;
        0 <= i < n                     ==>
        MultisetUnchanged{K,L}(a, n)   ==>
        HasValue{K}(a, n, \at(a[i],L));
*/
```

Listing 10.14: The lemma `ReorderImpliesMatch`

Lemmas handling some effects of `swap`

The next group of lemmas is related to proofs dealing with the effect of the `swap` call. We used the predicate `SwappedInside`, shown in Listing 10.15 rather than the literal postcondition of `swap`, since this leads to better performance of the provers.

```
/*@
  predicate
    SwappedInside{K,L}(value_type* a, integer i, integer k, integer n) =
      0 <= i < k < n            &&
      \at(a[i],K) == \at(a[k],L)  &&
      \at(a[k],K) == \at(a[i],L)  &&
      Unchanged{K,L}(a, 0,   i)   &&
      Unchanged{K,L}(a, i+1, k)   &&
      Unchanged{K,L}(a, k+1, n);
*/
```

Listing 10.15: The predicate `SwappedInside`

Our next lemma, `SwappedInsideMultisetUnchanged`, which is shown in Listing 10.16, employs the predicate `SwappedInside` from Listing 10.15. It states that swapping the elements `a[i]` and `a[k]` is a particular kind of reordering on the range `a[i..k]`.

```
/*@
  lemma
    SwappedInsideMultisetUnchanged{K,L}:
      \forall value_type* a, integer i, k, n;
        SwappedInside{K,L}(a, i, k, n)    ==>
        MultisetUnchanged{K,L}(a, i, k+1);
*/
```

Listing 10.16: The lemma `SwappedInsideMultisetUnchanged`

This lemma is extended to lemma `SwappedInsidePreservesMultisetUnchanged` (Listing 10.17), which additionally considers a left context `a[0..k-1]` and a right context `a[k..n-1]`; if the left and right context is reordered and kept untouched, respectively, and `a[i]` and `a[k]` are swapped as before, then this whole action is a reordering on the range `a[0..k]`. These two lemmas are needed to prove that the loop invariant `reorder` is preserved.

```
/*@
  lemma
    SwappedInsidePreservesMultisetUnchanged{K,L,M}:
      \forall value_type* a, integer i, k, n;
        MultisetUnchanged{K,L}(a, k)       ==>
        Unchanged{K,L}(a, k, n)            ==>
        SwappedInside{L,M}(a, i, k, n)     ==>
        MultisetUnchanged{K,M}(a, k+1);
*/
```

Listing 10.17: The lemma `SwappedInsidePreservesMultisetUnchanged`

10.3. The `heap_sort` algorithm

The `heap_sort` algorithm has the signature

```
void heap_sort(value_type* a, size_type n);
```

It relies upon the heap data structure discussed in Chapter 9 to efficiently bring the given array into an ascending order.

10.3.1. Formal Specification of `heap_sort`

Listing 10.18 shows the ACSL specification of `heap_sort`.

```
/*@
   requires valid:  \valid(a + (0..n-1));

   assigns a[0..n-1];

   ensures sorted:   Sorted(a, n);
   ensures reorder:  MultisetUnchanged{Old, Here}(a, n);
*/
void
heap_sort(value_type* a, size_type n);
```

Listing 10.18: The Specification of `heap_sort`

10.3.2. Implementation of `heap_sort`

The implementation of `heap_sort`, shown in Listing 10.19, is straightforward. Given the unsorted array `a`, we use `make_heap` to arrange it into a heap; after that, we use `sort_heap` to draw from this heap the elements in ascending order.

```
void
heap_sort(value_type* a, size_type n)
{
  make_heap(a, n);
  sort_heap(a, n);
}
```

Listing 10.19: The Implementation of `heap_sort`

10.4. The `selection_sort` algorithm

Our version of the `selection_sort` algorithm has the signature

```
void selection_sort(value_type* a, size_type n);
```

The `selection_sort` algorithm constructs the sorted array, left to right, by selecting in each step the minimum element of the remaining (yet unsorted) part and *swaps* it the first element of the unsorted part. This implies that each member of the sorted initial array segment is less or equal than each member of the unsorted part.
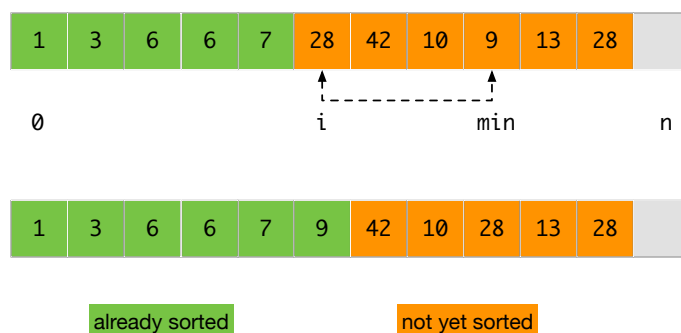
Figure 10.20.: An iteration of `selection_sort`

Figure 10.20 shows a typical situation in an example run. The algorithm will swap the `28` at position `i` with the `9` at position `min` to extend the sorted area one field to the right.

10.4.1. Formal Specification of `selection_sort`

Listing 10.21 shows the ACSL specification of `selection_sort`.

```
/*@
   requires valid:  \valid(a + (0..n-1));

   assigns a[0..n-1];

   ensures sorted:   Sorted(a, n);
   ensures reorder:  MultisetUnchanged{Old, Here}(a, n);
*/
void
selection_sort(value_type* a, size_type n);
```

Listing 10.21: The Specification of `selection_sort`

10.4.2. Implementation of `selection_sort`

The implementation of `selection_sort` is shown in Listing 10.22. We use `min_element` from Section 4.7 to find the minimum element of the unsorted area.

The loop invariants `sorted` and `lower` establish the sortedness of the initial array segment `a[0..i-1]` and, respectively, state that `a[i-1]` is a lower bound of the remaining array segment `a[i..n-1]`. Since the `min_element` call uses an address offset, we had to employ again the *shift lemmas* from Listing 5.12.

```
void
selection_sort(value_type* a, size_type n)
{
  /*@
    loop invariant bound:   0 <= i <= n;
    loop invariant sorted:  Sorted(a, i);
    loop invariant sorted:  0 < i  ==> LowerBound(a, i, n, a[i-1]);
    loop invariant reorder: MultisetUnchanged{Pre,Here}(a, n);
    loop assigns   i, a[0..n-1];
    loop variant   n - i;
  */
  for (size_type i = 0u; i < n; ++i) {
    const size_type sel = i + min_element(a + i, n - i);
    //@ assert reorder: i <= sel < n;
    /*@
       assigns a[sel], a[i];

       ensures reorder: a[i]   == \old(a[sel]);
       ensures reorder: a[sel] == \old(a[i]);
       ensures reorder: Unchanged{Old,Here}(a, 0,     i);
       ensures reorder: Unchanged{Old,Here}(a, i+1,   sel);
       ensures reorder: Unchanged{Old,Here}(a, sel+1, n);
    */
    swap(a + sel, a + i);
    //@ assert reorder: MultisetUnchanged{Pre,Here}(a, n);
  }
}
```

Listing 10.22: The Implementation of `selection_sort`

The loop invariant `reorder`, on the other hand, states that the multiset of values in the array `a` are only *rearranged* during the algorithm. While this is intuitively most obvious (as the call to the `swap` routine, from Section 6.4, is the only code that modifies `a`), it took considerable effort to prove it formally; including a statement contract that captures that captures the effects of calling `swap`. assertions in the loop body.

The main reason for introducing the statement contract is that it *transforms* the postcondition of the call to `swap` from Listing 6.8 into the hypotheses for the new lemma `SwapImpliesMultisetUnchanged` in Listing 10.23. This lemma, which relies on the lemmas from Listing 9.45, captures the fact that *swapping two elements of an array* is a *reordering*.

```
/*@
  lemma
    SwapImpliesMultisetUnchanged{K,L}:
      \forall value_type *a, integer i, k, n;
        0 <= i <= k < n                  ==>
        \at(a[i],K) == \at(a[k],L)       ==>
        \at(a[k],K) == \at(a[i],L)       ==>
        MultisetUnchanged{K,L}(a, 0,   i) ==>
        MultisetUnchanged{K,L}(a, i+1, k) ==>
        MultisetUnchanged{K,L}(a, k+1, n) ==>
        MultisetUnchanged{K,L}(a, n);
*/
```

Listing 10.23: The lemma `SwapImpliesMultisetUnchanged`

10.5. The `insertion_sort` algorithm

Like `selection_sort`, the algorithm `insertion_sort` traverses the given array `a[0..n-1]` left to right, maintaining a left-adjusted, constantly increasing range `a[0..i-1]` that is already sorted. Unlike `selection_sort`, however, `insertion_sort` adds `a[i]` to the sorted range in the `i`th step (see Figure 10.24). It determines the (rightmost) appropriate position to insert `a[i]` by a call to `upper_bound` (Section 5.2), and then uses `rotate` (Section 6.11) to actually perform the insertion.
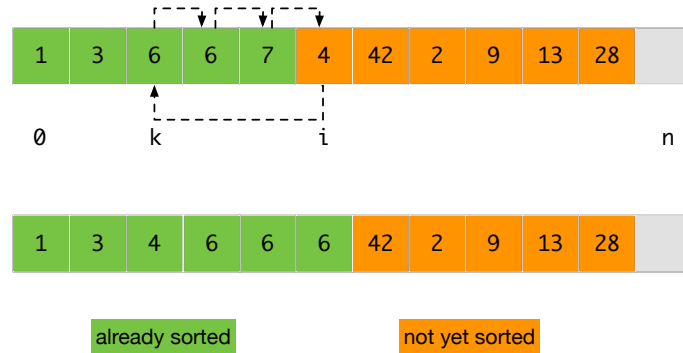
Figure 10.24.: An iteration of `insertion_sort`

10.5.1. Formal Specification of `insertion_sort`

Listing 10.25 shows our (generic sorting) contract for `insertion_sort`.

```
/*@
    requires valid:  \valid(a + (0..n-1));

    assigns a[0..n-1];

    ensures sorted:   Sorted(a, n);
    ensures reorder:  MultisetUnchanged{Old, Here}(a, n);
*/
void
insertion_sort(value_type* a, size_type n);
```

Listing 10.25: The Specification of `insertion_sort`

10.5.2. Implementation of `insertion_sort`

The implementation of `insertion_sort` is shown in Listing 10.26. We used an ACSL statement contract to specify those aspects of the `rotate` contract that are needed here. Properties related to the result of `insertion_sort` being in ascending order are labelled `sorted`. Properties related to the rearrangement of elements are labelled `reorder` and, whenever their order isn't changed, `unchanged`.

```
void
insertion_sort(value_type* a, size_type n)
{
  /*@
    loop invariant bound:     0 <= i <= n;
    loop invariant sorted:    Sorted(a, i);
    loop invariant reorder:   MultisetUnchanged{Pre,Here}(a, 0, i);
    loop invariant unchanged: Unchanged{Pre,Here}(a, i, n);
    loop assigns   i, a[0..n-1];
    loop variant   n - i;
  */
  for (size_type i = 0u; i < n; ++i) {
    const size_type k = upper_bound(a, i, a[i]);
    //@ assert bound:   0 <= k <= i;
    /*@
       requires sorted:  UpperBound(a, k, a[i]);
       requires sorted:  StrictLowerBound(a, k, i, a[i]);
       requires sorted:  Sorted(a, k, i);

       assigns a[k..i];

       ensures unchanged: Unchanged{Old,Here}(a, 0, k);
       ensures unchanged: Unchanged{Old,Here}(a, i+1, n);
       ensures reorder:   MultisetUnchanged{Old,Here}(a, 0, k);
       ensures reorder:   EqualRanges{Old,Here}(a, k, i, k+1);
       ensures reorder:   EqualRanges{Old,Here}(a, i, i+1, k);
       ensures sorted:    Sorted(a, 0, k);
       ensures sorted:    UpperBound(a, k, a[k]);
    */
    rotate(a + k, i - k, i - k + 1u);
    //@ assert sorted:  Sorted(a, k+1, i+1);
    //@ assert sorted:  StrictLowerBound(a, k+1, i+1, a[k]);
    //@ assert sorted:  Sorted(a, i+1);
    //@ assert reorder: MultisetUnchanged{Pre,Here}(a, 0, i+1);
  }
}
```

Listing 10.26: The Implementation of `insertion_sort`

When we originally implemented and verified `rotate`, we hadn't yet in mind to use that function inside of `insertion_sort`. Consequently, the properties needed for the latter aren't directly provided by the former. One approach to solve this problem is to add the new properties to `rotate`'s contract (Listing 6.29) and repeat its verification proof.[37]

However, if `rotate` is assumed to be part of a pre-verified library, this approach isn't feasible, since `rotate`'s implementation may not be available for re-verification. Therefore, we used another approach, viz. to prove that `rotate`'s original specification *implies* all the properties we need in `insertion_sort`.

This is another use of the Hoare calculus' implication rule (Section 2.3). We used several lemmas, shown below, to make the necessary implications explicit, and to help the provers to establish them. Some of them needed manual proofs by induction.

---

[37]ACSL allows to declare a function several times with different contracts; they are merged into a single one. Alternatively, non-disjoint `behaviors`, with empty `assumes` clauses, allow contract merging and provide finer control over the set of hypotheses generated from e.g. an `assert`.

Supporting lemmas related to property `sorted`

Lemma `EqualRangesPreservesSorted` (Listing 10.27) assumes an ordered range `a[m..n-1]` and claims that every (elementwise) equal range range `a[m+p..n+p-1]` is ordered, too. It is needed to establish that the `rotate` call preserves the order of those elements that are shifted upwards (cf. Figure 10.24).

```
/*@
    lemma
      EqualRangesPreservesSorted{K,L}:
        \forall value_type* a, integer m, n, p;
          Sorted{K}(a, m, n)                ==>
          EqualRanges{K,L}(a, m, n, m+p)  ==>
          Sorted{L}(a, m+p, n+p);
*/
```

Listing 10.27: The lemma `EqualRangesPreservesSorted`

Lemma `RotatePreservesStrictLowerBound` (Listing 10.28) is used to prove that the range `a[k..i-1]` having `a[i]` as strict lower bound before our `rotate` call ensures that it has `a[k]` as such a bound after the call. Note that this lemma reflects the situation at our particular `rotate` call site.

```
/*@
    lemma
      RotatePreservesStrictLowerBound{K,L}:
        \forall value_type* a, integer m, n;
          StrictLowerBound{K}(a, m, n, \at(a[n],K)) ==>
          EqualRanges{K,L}(a, m, n, m+1)            ==>
          EqualRanges{K,L}(a, n, n+1, m)            ==>
          StrictLowerBound{L}(a, m+1, n+1, \at(a[m],L));
*/
```

Listing 10.28: The lemma `RotatePreservesStrictLowerBound`

Supporting lemmas related to property `reorder`

Lemma `RotateImpliesMultisetUnchanged` (Listing 10.29) establishes that `rotate` by one produces just a reordering of the range it is applied to. Again, this lemma has been special-tailored to our `rotate` call; one could also think of a more general version that allows for arbitrary rotation amounts.

```
/*@
   lemma
     RotateImpliesMultisetUnchanged{K,L}:
       \forall value_type* a, integer m, n;
         0 <= m <= n                       ==>
         EqualRanges{K,L}(a, m, n, m+1)    ==>
         EqualRanges{K,L}(a, n, n+1, m)    ==>
         MultisetUnchanged{K,L}(a, m, n+1);
*/
```

Listing 10.29: The lemma `RotateImpliesMultisetUnchanged`

Finally, lemma `EqualRangesPreservesCount` (Listing 10.30) says that two elementwise equal ranges `a[m..n-1]` and `a[p..p+n-m-1]` will result in the same occurrence count, for each value `v`. It is useful in the proof of the previous lemma, `RotateImpliesMultisetUnchanged`, since the predicate `MultisetUnchanged` (Listing 6.5) is defined via the `Count` function (Listing 3.32).

```
/*@
   lemma
     EqualRangesPreservesCount{K,L}:
       \forall value_type *a, v, integer m, n, p;
         0 <= m <= n                    ==>
         EqualRanges{K,L}(a, m, n, p)   ==>
         Count{K}(a, m, n, v) == Count{L}(a, p, p + (n-m), v);
*/
```

Listing 10.30: The lemma `EqualRangesPreservesCount`

10.6. The `merge` algorithm

Our version of the `merge` algorithm from the C++ standard library [21, 28.7.5] has the following signature.

```
void
merge(const value_type* a, size_type n,
      const value_type* b, size_type m,
      value_type* result);
```

The `merge` algorithm is a part of the *merge sort* algorithm. It operates on the second step to merge two sorted sub-arrays into a new one. The algorithm merges two sorted arrays `a[0..n-1]` and `b[0..m-1]`, respectively. The merged values are stored in the output array that starts at `result` which must be able to hold *m + n* values of both input arrays.

10.6.1. Formal specification of `merge`

Listing 10.31 shows the ACSL specification of `merge`. The specification requires the sorted input arrays of the proper size and the output array of enough size to contain all the input elements. The input arrays should not overlap with the output array.

In the current edition of this guide, we prove only the sortedness of the resulting array. Future editions will contain additional postconditions stating that the result array consists of reordered input elements and the stability of the algorithm, i.e., the same elements of the input arrays preserve their order in the output array.

```
/*@
  requires bound:  n + m <= SIZE_TYPE_MAX;
  requires valid:  \valid_read(a + (0..n-1));
  requires valid:  \valid_read(b + (0..m-1));
  requires valid:  \valid(result + (0..n+m-1));
  requires sep:    \separated(a + (0..n-1), result + (0..n+m-1));
  requires sep:    \separated(b + (0..m-1), result + (0..n+m-1));
  requires sorted: WeaklySorted(a, n);
  requires sorted: WeaklySorted(b, m);

  assigns result[0 .. n+m-1];

  ensures sorted: Sorted(result, n + m);
*/
void
merge(const value_type* a, size_type n,
      const value_type* b, size_type m,
      value_type* result);
```

Listing 10.31: The specification of `merge`

10.6.2. Lemmas for `WeaklySorted`

Listing 10.32 contains additional lemmas required to prove the sortedness of the output array in the merge algorithm.

```
/*@
  lemma WeaklySortedAddElement{L}:
    \forall value_type *a, integer m;
      1 < m && WeaklySorted(a, m-1) && a[m-2] <= a[m-1] ==> WeaklySorted(a, m);

  lemma WeaklySortedShift{L}:
    \forall value_type *a, integer n, m;
      WeaklySorted(a + n, 0, m) <==> WeaklySorted(a, n, m + n);

  lemma EqualRangesWeaklySorted{L}:
    \forall value_type *a, *b, integer n, m;
      WeaklySorted(a, n, m) && EqualRanges{L,L}(a, n, m, b) ==>
       WeaklySorted(b, n, m);

  lemma WeaklySortedJoin{L}:
    \forall value_type *a, integer n, m;
      0 < n < m            &&
      WeaklySorted(a, n)    &&
      WeaklySorted(a, n, m) &&
      a[n-1] <= a[n] ==>
        WeaklySorted(a, m);
*/
```

Listing 10.32: Lemmas for `WeaklySorted` predicate

- Lemma `WeaklySortedAddElement` defines the way a weakly sorted array can be constructed.

- Lemma `WeaklySortedShift` states the equality with respect to the `WeaklySorted` property of the elements of the array `a[0..n]` and the elements of the array section `a[n..m+n]`.

- Lemma `EqualRangesWeaklySorted` states that if an array is weakly sorted, then another array, whose elements are in a one-to-one correspondence with the original array, is also weakly sorted.

- Lemma `WeaklySortedJoin` defines the conditions that two consequent weakly sorted ranges can be viewed as merged weakly sorted range.

Lemma `WeaklySortedShift` requires a manual proof with Coq. The other lemmas can be proved automatically. Note that we also rely on the lemma from Listing 10.4 to verify the postcondition.

10.6.3. Implementation of `merge`

The merge algorithm operates by traversing both input arrays. On each iteration it copies smaller of both elements to the result array, thus constructing the sorted array. If the algorithm reaches the end of one of the input arrays, it just copies the rest elements of the other array to the result array.

The implementation of merge, shown in Listing 10.33, is straightforward. The listing contains a number of assertions to trigger an application of lemmas by the provers. The **while** loop traverses the input arrays and constructs, in accordance with Lemma WeaklySortedAddElement, the resulting weakly sorted array. After the loop, the algorithm copies the remaining elements to the resulting array.

- Lemma EqualRangesWeaklySorted is used to show that the copied elements from one of the input arrays preserve the WeaklySorted property.

- Lemma WeaklySortedJoin is used to extend the WeaklySorted property of the two sub-ranges of the resulting array over the whole range. Lemma WeaklySortedShift is used in-between for array offset arithmetic.

- Finally, Lemma WeaklySortedImpliesSorted is used to prove the output array is sorted.

```
void
merge(const value_type* a, size_type n,
      const value_type* b, size_type m,
      value_type* result)
{
  size_type i = 0;
  size_type j = 0;
  size_type x = 0;

  if (n == 0 && m == 0) {
    return;
  }

  /*@ loop invariant 0 <= i <= n;
      loop invariant 0 <= j <= m;
      loop invariant x == i + j;
      loop invariant 0 <= x <= n + m - 1;
      loop invariant order:  \forall integer k; 0 <= k < x && i < n ==>
                                      result[k] <= a[i];
      loop invariant order:  \forall integer k; 0 <= k < x && j < m ==>
                                      result[k] <= b[j];
      loop invariant sorted: WeaklySorted(result, x);
      loop assigns i, j, x, result[0 .. n+m-1];
      loop variant (n + m) - (i + j);
   */
  while (i < n && j < m) {
    if (a[i] < b[j]) {
      result[x++] = a[i++];
    }
    else {
      result[x++] = b[j++];
    }
  }

  //@ assert i == n ^^ j == m;
  //@ assert i < n ^^ j < m;
  //@ assert WeaklySorted(result, 0, x);

  if (i < n) {
    //@ assert 0 < x ==> result[x-1] <= a[i];
    //@ assert WeaklySorted(a + i, 0, n - i);
    copy(a + i, n - i, result + x);
    //@ assert result[x] == a[i];
    /*@ assert WeaklySorted(a + i, 0, n - i) &&
            EqualRanges{Here,Here}(a + i, 0, n - i, result + x) ==>
                WeaklySorted(result + x, 0, n - i);
     */
    //@ assert n - i + x == n + m;
  }
  else {
    //@ assert 0 < x ==> result[x-1] <= b[j];
    //@ assert WeaklySorted(b + j, 0, m - j);
    copy(b + j, m - j, result + x);
    //@ assert result[x] == b[j];
    /*@ assert WeaklySorted(b + j, 0, m - j) &&
            EqualRanges{Here,Here}(b + j, 0, m - j, result + x) ==>
                WeaklySorted(result + x, 0, m - j);
     */
    //@ assert m - j + x == n + m;
  }

  //@ assert WeaklySorted(result, x, n + m);
  //@ assert x > 0 ==> result[x-1] <= result[x];
  //@ assert WeaklySorted(result, 0, n + m);
}
```

Listing 10.33: The Implementation of merge

Part V.

Verification of data structures

11. The `Stack` data type

So far we have used the ACSL specification language for the task of specifying and verifying one single C function at a time. However, in practice we are also faced with the task to implement a family of functions, usually around some sophisticated data structure, which have to obey certain rules of interdependence. In this kind of task, we are not only interested in the properties of a single function but also in properties describing how several function play together.

The C++ Standard Library provides a generic container adaptor `stack` [21, §26.6.6] whose signature and behavior we try to follow as far as our C implementation it allows. For a more detailed discussion of our approach to the formal verification of `Stack` we refer to Kim Völlinger's thesis [32].

A *stack* is a data type that can hold objects and has the property that, if an object *a* is *pushed* on a stack *before* object *b*, then *a* can only be removed (*popped*) after *b*. A stack is, in other words, a *first-in, last-out* data type (see Figure 11.1). The *top* function of a stack returns the last element that has been pushed on a stack.

Figure 11.1.: Push and pop on a stack

We consider only stacks that have a finite *capacity*, that is, that can only hold a maximum number *c* of elements that is constant throughout their lifetime. This restriction allows us to define a stack without relying on dynamic memory allocation. When a stack is *created* or *initialized*, it contains no elements, i.e., its *size* is 0. The function *push* and *pop* increases and decreases the size of a stack by at most one, respectively.

11.1. Methodology overview

Figure 11.2 gives an overview of our methodology to specify and verify abstract data types (verification of one axiom shown only).
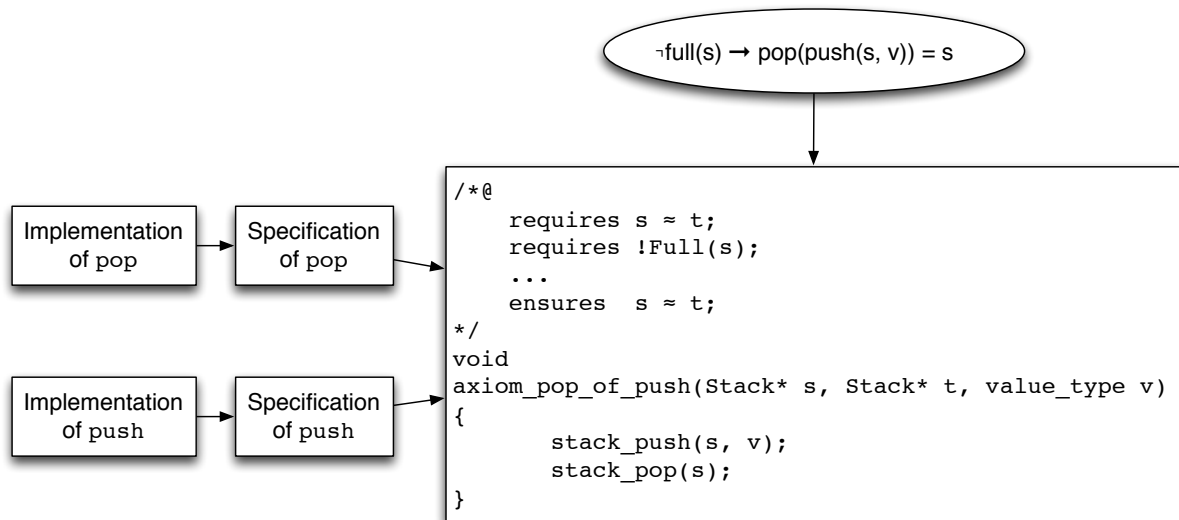
Figure 11.2.: Methodology Overview

What we will basically do is:

1. specify axioms about how the stack functions should interact with each other (Section 11.2),

2. define a basic implementation of C data structures (only one in our example, viz. **struct** Stack; see Section 11.3) and some invariants the instances of them have to obey (Section 11.4),

3. provide for each stack function an ACSL contract and a C implementation (Section 11.7),

4. verify each function against its contract (Section 11.7),

5. transform the axioms into ACSL-annotated C code (Section 11.8), and

6. verify that code, using access function contracts and data-type invariants as necessary (Section 11.8).

Section 11.5 provides an ACSL-predicate deciding whether two instances of a **struct** Stack are considered to be equal (indication by "≈" in Figure 11.2), while Section 11.6 gives a corresponding C implementation. The issue of an appropriate definition of equality of data instances is familiar to any C programmer who had to replace a faulty comparison **if**(s1 == s2) by the correct **if**(strcmp(s1,s2)== 0) to compare two strings **char** *s1,*s2 for equality.

11.2. Stack axioms

To specify the interplay of the stack access functions, we use a set of axioms[38], all but one of them having the form of a conditional equation.

Let V denote an arbitrary type. We denote by S_c the type of stacks with capacity $c > 0$ of elements of type V. The aforementioned functions then have the following signatures.

$$\text{init} : S_c \to S_c,$$
$$\text{push} : S_c \times V \to S_c,$$
$$\text{pop} : S_c \to S_c,$$
$$\text{top} : S_c \to V,$$
$$\text{size} : S_c \to \mathbb{N}.$$

With $\mathbb{B}$ denoting the *boolean* type we will also define two auxiliary functions

$$\text{empty} : S_c \to \mathbb{B},$$
$$\text{full} : S_c \to \mathbb{B}.$$

To qualify as a stack these functions must satisfy the following rules which are also referred to as *stack axioms*.

11.2.1. Stack initialization

After a stack has been initialized its size is 0.

$$\text{size}(\text{init}(s)) = 0. \tag{11.1}$$

The auxiliary functions empty and full are defined as follows

$$\text{empty}(s), \quad \text{iff} \quad \text{size}(s) = 0, \tag{11.2}$$
$$\text{full}(s), \quad \text{iff} \quad \text{size}(s) = c. \tag{11.3}$$

We expect that for every stack s the following condition holds

$$0 \leq \text{size}(s) \leq c. \tag{11.4}$$

11.2.2. Adding an element to a stack

To push an element v on a stack the stack must not be full. If an element has been pushed on an eligible stack, its size increases by 1

$$\text{size}(\text{push}(s, v)) = \text{size}(s) + 1, \qquad \text{if} \quad \neg\text{full}(s). \tag{11.5}$$

Moreover, the element pushed on a stack is the top element of the resulting stack

$$\text{top}(\text{push}(s, v)) = v, \qquad \text{if} \quad \neg\text{full}(s). \tag{11.6}$$

---

[38]There is an analogy in geometry: Euclid (e.g. [33]) invented the use of axioms there, but still kept definitions of *point*, *line*, *plane*, etc. Hilbert [34] recognized that the latter are not only unformalizable, but also unnecessary, and dropped them, keeping only the formal descriptions of relations between them.

11.2.3. Removing an element from a stack

An element can only be removed from a non-empty stack. If an element has been removed from an eligible stack the stack size decreases by 1

$$\text{size}(\text{pop}(s)) = \text{size}(s) - 1, \qquad\qquad \text{if} \quad \neg\text{empty}(s). \qquad (11.7)$$

If an element is pushed on a stack and immediately afterwards an element is removed from the resulting stack then the final stack is equal to the original stack

$$\text{pop}(\text{push}(s, v)) = s, \qquad\qquad \text{if} \quad \neg\text{full}(s). \qquad (11.8)$$

Conversely, if an element is removed from a non-empty stack and if afterwards the top element of the original stack is pushed on the new stack then the resulting stack is equal to the original stack.

$$\text{push}(\text{pop}(s), \text{top}(s)) = s, \qquad\qquad \text{if} \quad \neg\text{empty}(s). \qquad (11.9)$$

11.2.4. A note on exception handling

We don't impose a requirement on `push(s, v)` if s is a full stack, nor on `pop(s)` or `top(s)` if s is an empty stack. Specifying the behavior in such *exceptional* situations is a problem by its own; a variety of approaches is discussed in the literature. We won't elaborate further on this issue, but only give an example to warn about "innocent-looking" exception specifications that may lead to undesired results.

If we'd introduce an additional error value `err` in the element type *V* and require `top(s)= err` if s is empty, we'd be faced with the problem of specifying the behavior of `push(s, err)`. At first glance, it would seem a good idea to have `err` just been ignored by `push`, i.e. to require

$$\text{push}(s, \text{err}) = s. \qquad (11.10)$$

However, we then could derive for any non-full and non-empty stack s, that

$$
\begin{aligned}
\text{size}(s) &= \text{size}(\text{pop}(\text{push}(s, \text{err}))) && \text{by 11.8} \\
&= \text{size}(\text{pop}(s)) && \text{as assumed in 11.10} \\
&= \text{size}(s) - 1 && \text{by 11.7}
\end{aligned}
$$

i.e. no such stacks could exist, or all **int** values would be equal.

11.3. The structure `Stack` and its associated functions

We now introduce one possible C implementation of the above axioms. It is centred around the C structure `Stack` shown in Listing 11.3.

```
struct Stack
{
  value_type* obj;

  size_type   capacity;

  size_type   size;
};

typedef struct Stack Stack;
```

Listing 11.3: Definition of type `Stack`

This struct holds an array `obj` of positive length called `capacity`. The capacity of a stack is the maximum number of elements this stack can hold. The field `size` indicates the number elements that are currently in the stack. See also Figure 11.4 which attempts to interpret this definition according to Figure 11.1.

Figure 11.4.: Interpreting the data structure `Stack`

Based on the stack functions from Section 11.2, we declare in Listing 11.5 the following functions as part of our `Stack` data type.

```
void          stack_init(Stack* s, value_type* a, size_type n);

bool          stack_equal(const Stack* s, const Stack* t);

size_type     stack_size(const Stack* s);

bool          stack_empty(const Stack* s);

bool          stack_full(const Stack* s);

value_type    stack_top(const Stack* s);

void          stack_push(Stack* s, value_type v);

void          stack_pop(Stack* s);
```

Listing 11.5: Declaration of functions of type `Stack`

Most of these functions directly correspond to methods of the C++ `std::stack` template class [21, §26.6.6.1]. The function `stack_equal` corresponds to the comparison operator ==, whereas one use of `stack_init` is to bring a stack into a well-defined initial state. The function `stack_full` has no counterpart in `std::stack`. This reflects the fact that we avoid dynamic memory allocation, while `std::stack` does not.

11.4. Stack invariants

Not every possible instance of type `Stack` is considered a valid one, e.g., with our definition of `Stack` in Listing 11.3, `Stack s = {{0,0,0,0},4,5}` is not. Below, we will define an ACSL-predicate `Invariant` that discriminates valid and invalid instances.

Before, we introduce in Listing 11.6 the auxiliary logic function `Capacity` and `Size` which we can use in specifications to refer to the fields `capacity` and `size` of `Stack`, respectively. This listing also contains the logic function `Top` which defines the array element with index `size-1` as the top place of a stack. The reader can consider this as an attempt to hide implementation details from the specification. We intentionally use here integer as a return value of these logic functions. Inaccurate use of logic functions with bounded types in axioms with arithmetic operations may lead to inconsistencies.

```
/*@
  logic integer Capacity{L}(Stack* s) = s->capacity;

  logic integer Size{L}(Stack* s) = s->size;

  logic value_type* Storage{L}(Stack* s) = s->obj;

  logic integer Top{L}(Stack* s) = s->obj[s->size-1];
*/
```

Listing 11.6: The logic functions `Capacity`, `Size` and `Top`

We also introduce in Listing 11.7 two predicates that express the concepts of empty and full stacks by referring to a stack's size and capacity (see Equations (11.2) and (11.3)).

```
/*@
  predicate
    Empty{L}(Stack* s) =  Size(s) == 0;

  predicate
    Full{L}(Stack* s)  =  Size(s) == Capacity(s);
*/
```

Listing 11.7: Predicates for empty an full stacks

There are some obvious invariants that must be fulfilled by every valid object of type `Stack`:

- The stack capacity shall be strictly greater than zero (an empty stack is ok but a stack that cannot hold anything is not useful).

- The pointer `obj` shall refer to an array of length `capacity`.

- The number of elements `size` of a stack the must be non-negative and not greater than its capacity.

These invariants are formalized in the predicate `Invariant` of Listing 11.8.

```
/*@
  predicate
    Invariant{L}(Stack* s) =
      0 < Capacity(s) &&
      0 <= Size(s) <= Capacity(s) &&
      \valid(Storage(s) + (0..Capacity(s)-1)) &&
      \separated(s, Storage(s) + (0..Capacity(s)-1));
*/
```

Listing 11.8: The predicate `Invariant`

Note how the use of the previously defined logic functions and predicates allows us to define the stack invariant without directly referring to the fields of `Stack`.

11.5. Equality of stacks

Defining equality of instances of non-trivial data types, in particular in object-oriented languages, is not an easy task. The book *Programming in Scala* [35, Chapter 28] devotes to this topic a whole chapter of more than twenty pages. In the following two sections we give a few hints how ACSL and Frama-C can help to correctly define equality for a simple data type.

We consider two stacks as equal if they have the same size and if they contain the same objects. To be more precise, let s and t two pointers of type Stack, then we define the predicate Equal as in Listing 11.9.

```
/*@
   predicate
     Equal{S,T}(Stack* s, Stack* t) =
       Size{S}(s) == Size{T}(t) &&
       EqualRanges{S,T}(Storage{S}(s), Size{S}(s), Storage{T}(t));
*/
```

Listing 11.9: Equality of stacks

Our use of labels in Listing 11.9 makes the specification somewhat hard to read (in particular in the last line where we reuse the predicate EqualRanges from Page 42). However, this definition of Equal will allow us later to compare the same stack object at different points of a program. The logical expression Equal{A,B}(s,t) reads informally as: The stack object *s at program point A equals the stack object *t at program point B.

The reader might wonder why we exclude the capacity of a stack into the definition of stack equality. This approach can be motivated with the behavior of the method capacity of the class std::vector<T>. There, equal instances of type std::vector<T> may very well have different capacities.[39]

If equal stacks can have different capacities then, according to our definition of the predicate Full in Listing 11.7, we can have to equal stacks where one is full and the other one is not.

A finer, but very important point in our specification of equality of stacks is that the elements of the arrays s->obj and t->obj are compared only up to s->size and *not* up to s->capacity. Thus the two stacks s and t in Figure 11.10 are considered equal although there is are obvious differences in their internal arrays.

Figure 11.10.: Example of two equal stacks

[39]See http://www.cplusplus.com/reference/vector/vector/capacity

If we define an equality relation (=) of objects for a data type such as `Stack`, we have to make sure that the following rules hold.

$$\text{reflexivity} \qquad \forall s \in S : s = s, \qquad (11.11a)$$

$$\text{symmetry} \qquad \forall s, t \in S : s = t \implies t = s, \qquad (11.11b)$$

$$\text{transitivity} \qquad \forall s, t, u \in S : s = t \wedge t = u \implies s = u. \qquad (11.11c)$$

Any relation that satisfies the conditions (11.11) is referred to as an *equivalence relation*. The mathematical set of all instances that are considered equal to some given instance `s` is called the equivalence class of `s` with respect to that relation.

Listing 11.11 shows a formalization of these three rules for the relation `Equal`; it can be automatically verified that they are a consequence of the definition of `Equal` in Listing 11.9.

```
/*@
  lemma
    StackEqualReflexive{S} :
      \forall Stack* s; Equal{S,S}(s, s);

  lemma
    StackEqualSymmetric{S,T} :
      \forall Stack *s, *t;
        Equal{S,T}(s, t) ==> Equal{T,S}(t, s);

  lemma
    StackEqualTransitive{S,T,U}:
      \forall Stack *s, *t, *u;
        Equal{S,T}(s, t) && Equal{T,U}(t, u) ==> Equal{S,U}(s, u);
*/
```

Listing 11.11: Equality of stacks is an equivalence relation

The two stacks in Figure 11.10 show that an equivalence class of `Equal` can contain more than one element.[40] The stacks `s` and `t` in Figure 11.10 are also referred to as two *representatives* of the same equivalence class. In such a situation, the question arises whether a function that is defined on a set with an equivalence relation can be defined in such a way that its definition is *independent of the chosen representatives*.[41] We ask, in other words, whether the function is *well-defined* on the set of all equivalence classes of the relation `Equal`.[42] The question of well-definition will play an important role when verifying the functions of the `Stack` (see Section 11.7).

---

[40] This is a common situation in mathematics. For example, the equivalence class of the rational number $\frac{1}{2}$ contains infinitely many elements, viz. $\frac{1}{2}, \frac{2}{4}, \frac{7}{14}, \ldots$.

[41] This is why mathematicians know that $\frac{1}{2} + \frac{3}{5}$ equals $\frac{7}{14} + \frac{3}{5}$.

[42] See http://en.wikipedia.org/wiki/Well-definition.

11.6. Runtime equality of stacks

The function `stack_equal` is the C equivalent for the `Equal` predicate. The specification of the function `stack_equal` is shown in Listing 11.12. Note that this specifications explicitly refers to valid stacks.

```
/*@
    requires valid: \valid(s) && Invariant(s);
    requires valid: \valid(t) && Invariant(t);

    assigns \nothing;

    ensures equal:     \result == 1  <==>  Equal{Here,Here}(s, t);
    ensures not_equal: \result == 0  <==>  !Equal{Here,Here}(s, t);
*/
bool
stack_equal(const Stack* s, const Stack* t);
```

Listing 11.12: Specification of `stack_equal`

The implementation of `stack_equal` in Listing 11.13 compares two stacks according to the same rules of predicate `Equal`.

```
bool
stack_equal(const Stack* s, const Stack* t)
{
  return (s->size == t->size) && equal(s->obj, s->size, t->obj);
}
```

Listing 11.13: Implementation of `stack_equal`

11.7. Verification of stack functions

In this section we verify the functions `stack_init` (Section 11.7.1), `stack_size` (Section 11.7.2), `stack_empty` (Section 11.7.3), `stack_full` (Section 11.7.4), `stack_top` (Section 11.7.5), and `stack_push` (Section 11.7.6) `stack_pop` (Section 11.7.7), of the data type `Stack`. To be more precise, we provide for each of function `stack_foo`:

- an ACSL specification of `stack_foo`

- a C implementation of `stack_foo`

- a C function `stack_foo_wd`[43] accompanied by a an ACSL contract that expresses that the implementation of `stack_foo` is well-defined. Figure 11.14 shows our methodology for the verification of well-definition in the `pop` example, ($\approx$) again indicating the user-defined `Stack` equality.

```
                              /*@
                                requires s ≈ t;
                                requires !Empty(s);
                                  ...
                                ensures  s ≈ t;
  ┌─────────────────┐        */
  │  Specification  │───────▶ void stack_pop_wd(Stack *s, Stack *t)
  │    of pop       │         {
  └─────────────────┘           stack_pop(s);
                                stack_pop(t);
                              }
```

Figure 11.14.: Methodology for the verification of well-definition

Note that the specifications of the various functions will explicitly refer to the *internal state* of `Stack`. In Section 11.8 we will show that the *interplay* of these functions satisfy the stack axioms from Section 11.2.

---

[43]The suffix `_wd` stands for *well definition*

11.7.1. The function `stack_init`

Listing 11.15 shows the ACSL specification of `stack_init`. Note that our specification of the post-conditions contains a redundancy because a stack is empty if and only if its size is zero.

```
/*@
  requires valid:    \valid(s);
  requires capacity: 0 < capacity;
  requires storage:  \valid(storage + (0..capacity-1));
  requires sep:      \separated(s, storage + (0..capacity-1));

  assigns s->obj, s->capacity, s->size;

  ensures valid:     \valid(s);
  ensures invariant: Invariant(s);
  ensures capacity:  Capacity(s) == capacity;
  ensures empty:     Empty(s);
  ensures storage:   Storage(s) == storage;
*/
void
stack_init(Stack* s, value_type* storage, size_type capacity);
```

Listing 11.15: Specification of `stack_init`

Listing 11.15 shows the implementation of `stack_init`. It simply initializes `obj` and `capacity` with the respective value of the array and sets the field `size` to zero.

```
void
stack_init(Stack* s, value_type* storage, size_type capacity)
{
  s->obj      = storage;
  s->capacity = capacity;
  s->size     = 0u;
}
```

Listing 11.16: Implementation of `stack_init`

11.7.2. The function `stack_size`

The function `stack_size` is the runtime version of the logic function `Size` from Listing 11.6 on Page 242. The specification of `stack_size` in Listing 11.17 simply states that `stack_size` produces the same result as `Size`.

```
/*@
    requires valid: \valid(s) && Invariant(s);

    assigns \nothing;

    ensures size: \result == Size(s);
*/
size_type
stack_size(const Stack* s);
```

Listing 11.17: Specification of `stack_size`

As in the definition of the logic function `Size` the implementation of `stack_size` in Figure 11.18 simply returns the field `size`.

```
size_type
stack_size(const Stack* s)
{
  return s->size;
}
```

Listing 11.18: Implementation of `stack_size`

Listing 11.19 shows our check whether `stack_size` is well-defined. Since `stack_size` neither modifies the state of its `Stack` argument nor that of any global variable we only check whether it produces the same result for equal stacks. Note that we simply may use operator == to compare integers since we didn't introduce a nontrivial equivalence relation on that data type.

```
/*@
  requires valid: \valid(s) && Invariant(s);
  requires valid: \valid(t) && Invariant(t);
  requires equal: Equal{Here,Here}(s, t);

  assigns \nothing;

  ensures equal:  \result;
*/
bool
stack_size_wd(const Stack* s, const Stack* t)
{
  return stack_size(s) == stack_size(t);
}
```

Listing 11.19: Well-definition of `stack_size`

11.7.3. The function `stack_empty`

The function `stack_empty` is the runtime version of the predicate `Empty` from Listing 11.7 on Page 243.

```
/*@
    requires valid: \valid(s) && Invariant(s);

    assigns \nothing;

    ensures empty:     \result == 1  <==>  Empty(s);
    ensures not_empty: \result == 0  <==> !Empty(s);
*/
bool
stack_empty(const Stack* s);
```

Listing 11.20: Specification of `stack_empty`

As in the definition of the predicate `Empty` the implementation of `stack_empty` in Figure 11.21 simply checks whether the size of the stack is zero.

```
bool
stack_empty(const Stack* s)
{
  return stack_size(s) == 0u;
}
```

Listing 11.21: Implementation of `stack_empty`

Listing 11.22 shows our check whether `stack_empty` is well-defined.

```
/*@
  requires valid: \valid(s) && Invariant(s);
  requires valid: \valid(t) && Invariant(t);
  requires equal: Equal{Here,Here}(s, t);

  assigns \nothing;

  ensures equal: \result;
*/
bool
stack_empty_wd(const Stack* s, const Stack* t)
{
  return stack_empty(s) == stack_empty(t);
}
```

Listing 11.22: Well-definition of `stack_empty`

11.7.4. The function `stack_full`

The function stack_full is the runtime version of the predicate Full from Listing 11.7 on Page 243.

```
/*@
    requires valid: \valid(s) && Invariant(s);

    assigns \nothing;

    ensures full:     \result == 1  <==>  Full(s);
    ensures not_full: \result == 0  <==>  !Full(s);
*/
bool
stack_full(const Stack* s);
```

Listing 11.23: Specification of stack_full

As in the definition of the predicate Full the implementation of stack_full in Figure 11.24 simply checks whether the size of the stack equals its capacity.

```
bool
stack_full(const Stack* s)
{
  return stack_size(s) == s->capacity;
}
```

Listing 11.24: Implementation of stack_full

Note that with our definition of stack equality (Section 11.5) there can be equal stack with different capacities. As a consequence, there can are equal stacks where one is full while the other is not. In other words, Full is not well-defined.

11.7.5. The function `stack_top`

The function `stack_top` is the runtime version of the logic function `Top` from Listing 11.6 on Page 242. The specification of `stack_top` in Listing 11.25 simply states that for non-empty stacks `stack_top` produces the same result as `Top` which in turn just returns the element `obj[size-1]` of `Stack`.

```
/*@
    requires valid: \valid(s) && Invariant(s);

    assigns \nothing;

    ensures top: !Empty(s) ==> \result == Top(s);
*/
value_type
stack_top(const Stack* s);
```

Listing 11.25: Specification of `stack_top`

For a non-empty stack the implementation of `stack_top` in Figure 11.26 simply returns the element `obj[size-1]`. Note that our implementation of `stack_top` does not crash when it is applied to an empty stack. In this case we return the first element of the internal, non-empty array `obj`. This is consistent with our specification of `stack_top` which only refers to non-empty stacks.

```
value_type
stack_top(const Stack* s)
{
  if (!stack_empty(s)) {
    return s->obj[s->size - 1u];
  }
  else {
    return s->obj[0u];
  }
}
```

Listing 11.26: Implementation of `stack_top`

Listing 11.27 shows our check whether `stack_top` well-defined for non-empty stacks.

```
/*@
  requires valid: \valid(s) && Invariant(s) && !Empty(s);
  requires valid: \valid(t) && Invariant(t) && !Empty(t);
  requires equal: Equal{Here,Here}(s, t);

  assigns \nothing;

  ensures equal:  \result;
*/
bool
stack_top_wd(const Stack* s, const Stack* t)
{
  return stack_top(s) == stack_top(t);
}
```

Listing 11.27: Well-definition of `stack_top`

Since our axioms in Section 11.2 did not impose any behavior on the behavior of `stack_top` for empty stacks, we prove the well-definition of `stack_top` only for nonempty stacks.

11.7.6. The function `stack_push`

Listing 11.28 shows the ACSL specification of the function stack_push. In accordance with Axiom (11.5), stack_push is supposed to increase the number of elements of a non-full stack by one. The specification also demands that the value that is pushed on a non-full stack becomes the top element of the resulting stack (see Axiom (11.6)).

```
/*@
  requires valid: \valid(s) && Invariant(s);

  assigns s->size;
  assigns s->obj[s->size];

  behavior not_full:
    assumes !Full(s);

    assigns s->size;
    assigns s->obj[s->size];

    ensures valid:     \valid(s) && Invariant(s);
    ensures size:      Size(s) == Size{Old}(s) + 1;
    ensures top:       Top(s) == v;
    ensures not_empty: !Empty(s);
    ensures unchanged: Unchanged{Old,Here}(Storage(s), Size{Old}(s));
    ensures storage:   Storage(s) == Storage{Old}(s);
    ensures capacity:  Capacity(s) == Capacity{Old}(s);

  behavior full:
    assumes Full(s);

    assigns \nothing;

    ensures valid:     \valid(s) && Invariant(s);
    ensures full:      Full(s);
    ensures unchanged: Unchanged{Old,Here}(Storage(s), Size(s));
    ensures size:      Size(s) == Size{Old}(s);
    ensures storage:   Storage(s) == Storage{Old}(s);
    ensures capacity:  Capacity(s) == Capacity{Old}(s);

  complete behaviors;
  disjoint behaviors;
*/
void
stack_push(Stack* s, value_type v);
```

Listing 11.28: Specification of stack_push

The implementation of stack_push is shown in Listing 11.29. It checks whether its argument is a non-full stack in which case it increases the field size by one but only after it has assigned the function argument to the element obj[size].

```
void
stack_push(Stack* s, value_type v)
{
  if (!stack_full(s)) {
    //@ assert not_full: s->size < s->capacity;
    s->obj[s->size++] = v;
  }
}
```

Listing 11.29: Implementation of stack_push

The function stack_push does not return a value but rather modifies its argument. For the well-definition of stack_push we therefore check whether it turns equal stacks into equal stacks. However, equality of the stack arguments is not sufficient for a proof that stack_push is well-defined. We must also ensure that there is no *aliasing* between the two stacks. Otherwise modifying one stack could modify the other stack in unexpected ways. In order to express that there is no aliasing between two stacks, we define in Listing 11.30 the predicate Separated.

```
/*@
  predicate
    Separated(Stack* s, Stack* t) =
      \separated(s, s->obj + (0..s->capacity-1),
                 t, t->obj + (0..t->capacity-1));
*/
```

Listing 11.30: The predicate Separated

Listing 11.31 shows our formalization of the well-definition for `stack_push`.

```
/*@
  requires valid:     \valid(s) && Invariant(s);
  requires valid:     \valid(t) && Invariant(t);
  requires equal:     Equal{Here,Here}(s, t);
  requires not_full:  !Full(s) && !Full(t);
  requires sep:       Separated(s, t);

  assigns s->size, s->obj[s->size];
  assigns t->size, t->obj[t->size];

  ensures  valid:     Invariant(s) && Invariant(t);
  ensures  equal:     Equal{Here,Here}(s, t);
*/
void
stack_push_wd(Stack* s, Stack* t, value_type v)
{
  stack_push(s, v);
  stack_push(t, v);
  //@ assert top:   Top(s) == v;
  //@ assert top:   Top(t) == v;
  //@ assert equal: EqualRanges{Here,Here}(Storage(s), Size{Pre}(s), Storage(t));
}
```

Listing 11.31: Well-definition of `stack_push`

In order to achieve an automatic verification of the well-definition of `stack_push` we added in List-ing 11.31 the assertions `top` and `equal` and introduced the lemma `StackPushEqual` from Listing 11.32.

```
/*@
  lemma
    StackPushEqual{K,L}:
      \forall Stack *s, *t;
        Equal{K,K}(s,t)                ==>
        Size{L}(s) == Size{K}(s) + 1 ==>
        Size{L}(s) == Size{L}(t)       ==>
        Top{L}(s)  == Top{L}(t)        ==>
        EqualRanges{L,L}(Storage{L}(s), Size{K}(s), Storage{L}(t)) ==>
        Equal{L,L}(s,t);
*/
```

Listing 11.32: The lemma `StackPushEqual`

11.7.7. The function `stack_pop`

Listing 11.33 shows the ACSL specification of the function `stack_pop`. In accordance with Axiom (11.7) `stack_pop` is supposed to reduce the number of elements in a non-empty stack by one. In addition to the requirements imposed by the axioms, our specification demands that `stack_pop` changes no memory location if it is applied to an empty stack.

```
/*@
  requires valid: \valid(s) && Invariant(s);

  assigns s->size;

  ensures  valid: \valid(s) && Invariant(s);

  behavior not_empty:
    assumes !Empty(s);

    assigns s->size;

    ensures size:      Size(s) == Size{Old}(s) - 1;
    ensures full:      !Full(s);
    ensures unchanged: Unchanged{Old,Here}(Storage(s), Size(s));
    ensures storage:   Storage(s) == Storage{Old}(s);
    ensures capacity:  Capacity(s) == Capacity{Old}(s);

  behavior empty:
    assumes Empty(s);

    assigns \nothing;

    ensures empty:     Empty(s);
    ensures unchanged: Unchanged{Old,Here}(Storage(s), Size(s));
    ensures size:      Size(s) == Size{Old}(s);
    ensures storage:   Storage(s) == Storage{Old}(s);
    ensures capacity:  Capacity(s) == Capacity{Old}(s);

  complete behaviors;
  disjoint behaviors;
*/
void
stack_pop(Stack* s);
```

Listing 11.33: Specification of `stack_pop`

The implementation of `stack_pop` is shown in Listing 11.34. It checks whether its argument is a non-empty stack in which case it decreases the field `size` by one.

```
void
stack_pop(Stack* s)
{
  if (!stack_empty(s)) {
    --s->size;
  }
}
```

Listing 11.34: Implementation of `stack_pop`

Listing 11.35 shows our check whether `stack_pop` is well-defined. As in the case of `stack_push` we use the predicate `Separated` (Listing 11.30) in order to express that there is no aliasing between the two stack arguments.

```
/*@
  requires valid:     \valid(s) && Invariant(s);
  requires valid:     \valid(t) && Invariant(t);
  requires equal:     Equal{Here,Here}(s, t);
  requires sep:       Separated(s, t);

  assigns s->size;
  assigns t->size;

  ensures valid: Invariant(s);
  ensures valid: Invariant(t);
  ensures equal: Equal{Here,Here}(s, t);
*/
void
stack_pop_wd(Stack* s, Stack* t)
{
  stack_pop(s);
  stack_pop(t);
}
```

Listing 11.35: Well-definition of `stack_pop`

11.8. Verification of stack axioms

In this section we show that the stack functions defined in Section 11.7 satisfy the stack Axioms of Section 11.2.

The annotated code has been obtained from the axioms in a fully systematical way. In order to transform a condition equation $p \rightarrow s = t$:

- Generate a clause `requires p`.

- Generate a clause `requires x1 == ... == xn` for each variable `x` with n occurrences in s and t.

- Change the i-th occurrence of `x` to `xi` in s and t.

- Translate both terms s and t to reversed polish notation.

- Generate a clause `ensures y1 == y2`, where `y1` and `y2` denote the value corresponding to the translated s and t, respectively.

This makes it easy to implement a tool that does the translation automatically, but yields a slightly longer contract in our example.

11.8.1. Resetting a stack

Our formulation in ACSL/C of the Axiom in Equation (11.1) on Page 239 is shown in Listing 11.36.

```
/*@
  requires valid:     \valid(s);
  requires size:      0 < n;
  requires valid:     \valid(a + (0..n-1));
  requires sep:       \separated(s, a + (0..n-1));

  assigns s->obj, s->capacity, s->size;

  ensures valid: Invariant(s);
  ensures size:  \result == 0;
*/
size_type
axiom_size_of_init(Stack* s, value_type* a, size_type n)
{
  stack_init(s, a, n);
  return stack_size(s);
}
```

Listing 11.36: Specification of Axiom (11.1)

11.8.2. Adding an element to a stack

Axioms (11.5) and (11.6) describe the behavior of a stack when an element is added.

```
/*@
  requires valid:    \valid(s) && Invariant(s);
  requires not_full: !Full(s);

  assigns s->size;
  assigns s->obj[s->size];

  ensures valid: Invariant(s);
  ensures size:  \result == Size{Old}(s) + 1;
*/
size_type
axiom_size_of_push(Stack* s, value_type v)
{
  stack_push(s, v);
  return stack_size(s);
}
```

Listing 11.37: Specification of Axiom (11.5)

Except for the `assigns` clauses, the ACSL-specification refers only to encapsulating logic functions and predicates defined in Section 11.4. If ACSL would provide a means to define encapsulating logic functions returning also sets of memory locations, the expressions in `assigns` clauses would not need to refer to the details of our `Stack` implementation.[44] As an alternative, `assigns` clauses could be omitted, as long as the proofs are only used to convince a human reader.

```
/*@
  requires valid:    \valid(s) && Invariant(s);
  requires not_full: !Full(s);

  assigns s->size;
  assigns s->obj[s->size];

  ensures top: \result == v;
*/
value_type
axiom_top_of_push(Stack* s, value_type v)
{
  stack_push(s, v);
  return stack_top(s);
}
```

Listing 11.38: Specification of Axiom (11.6)

---

[44]In [16, §2.3.4], a powerful sublanguage to build memory location set expressions is defined. We will explore its capabilities in a later version.

11.8.3. Removing an element from a stack

This section shows the Listings for Axioms 11.7, 11.8 and 11.9 which describe the behavior of a stack when an element is removed.

```
/*@
  requires  valid: \valid(s) && Invariant(s);
  requires  empty: !Empty(s);

  assigns   s->size;

  ensures   size: \result == Size{Old}(s) - 1;
*/
size_type
axiom_size_of_pop(Stack* s)
{
  stack_pop(s);
  return stack_size(s);
}
```

Listing 11.39: Specification of Axiom (11.7)

```
/*@
  requires  valid:    \valid(s) && Invariant(s);
  requires  not_full: !Full(s);

  assigns   s->size;
  assigns   s->obj[s->size];

  ensures   equal:    Equal{Old,Here}(s, s);
*/
void
axiom_pop_of_push(Stack* s, value_type v)
{
  stack_push(s, v);
  stack_pop(s);
}
```

Listing 11.40: Specification of Axiom (11.8)

```
/*@
  requires  valid:     \valid(s) && Invariant(s);
  requires  not_empty: !Empty(s);

  assigns   s->size;
  assigns   s->obj[s->size-1];

  ensures   equal: Equal{Old,Here}(s, s);
*/
void
axiom_push_of_pop_top(Stack* s)
{
  const value_type val = stack_top(s);
  stack_pop(s);
  stack_push(s, val);
}
```

Listing 11.41: Specification of Axiom (11.9)

Part VI.

Appendices

A. Results of formal verification with Frama-C

In this chapter we introduce the formal verification tools used in this tutorial. We will afterwards present to what extent the examples from Chapters 3–11 could be deductively verified.

Within Frama-C, the Frama-C/WP plug-in [1] enables deductive verification of C programs that have been annotated with the ANSI/ISO-C Specification Language (ACSL) [10]. The Frama-C/WP plug-in uses weakest precondition computations to generate proof obligations. To formally prove the ACSL properties, these proof obligations can be submitted to external automatic theorem provers or interactive proof assistants. For the precise settings for Frama-C/WP we employed in this release we refer to Section A.1. In Sections A.2 and A.3 we show detailed verification results for different scenarios how the provers are called.

A.1. Verification settings

This section gives all settings that depend on the software release of Frama-C, Why3, or one if its employed provers.

Here are the most important options of Frama-C that we used in for almost all functions.[45]

```
-pp-annot
-no-unicode
-wp
-wp-rte
-wp-model Typed+ref
-warn-unsigned-overflow
-warn-unsigned-downcast
-wp-timeout 10
-wp-steps 1000
-wp-coq-timeout 10
```

---

[45]For the `my_lrand48()` function in `random_shuffle`, the option `-warn-`**`unsigned`**`-overflow` is disabled as explained in Section 6.18.3.

A.2. Verification results (sequential)

In the *sequential verification scenario* each proof obligation is processed by a set of automatic and interactive theorem provers that are arranged as a *pipe*.[46] This means that each prover passes on to the next prover only those proof obligations that it could not verify. This *verification pipeline* is shown in Figure A.1.

Figure A.1.: Verification pipeline of automatic and interactive theorem provers

For each algorithm we list in the following tables the number of generated verification conditions (VC), the percentage of proven verification conditions, and the number of VC proven by each prover. The value zero is indicated by an empty cell. The tables show that all verification conditions could be verified. Please note that the number of proven verification conditions do *not* reflect on the quality/strength of the individual provers. The reason for that is that we "pipe" each verification condition sequentially through a list of provers (see Figure A.1).

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP	CQ
find	§3.1	19/19	(100%)	10	9	·	·	·	·	·
find2	§3.2	19/19	(100%)	10	9	·	·	·	·	·
find_first_of	§3.3	25/25	(100%)	16	9	·	·	·	·	·
adjacent_find	§3.4	22/22	(100%)	10	12	·	·	·	·	·
mismatch	§3.5	20/20	(100%)	10	10	·	·	·	·	·
equal	§3.5	7/7	(100%)	6	1	·	·	·	·	·
search	§3.6	30/30	(100%)	18	12	·	·	·	·	·
search_n	§3.7	91/91	(100%)	62	29	·	·	·	·	·
find_end	§3.8	28/28	(100%)	15	13	·	·	·	·	·
count	§3.9	28/28	(100%)	7	16	·	·	·	·	5
count2	§3.10	31/31	(100%)	7	19	·	·	·	·	5

Table A.2.: Results for non-mutating algorithms

---

[46]Sequential processing is achieved by passing the option -wp-par 1 to Frama-C/WP.

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP	CQ
properties of operator <	§4.1	6/ 6	(100%)	4	2	·	·	·	·	·
clamp	§4.2	22/22	(100%)	18	4	·	·	·	·	·
max_element	§4.3	24/24	(100%)	13	11	·	·	·	·	·
max_element2	§4.4	24/24	(100%)	12	12	·	·	·	·	·
max_seq	§4.5	8/ 8	(100%)	5	3	·	·	·	·	·
min_element	§4.7	24/24	(100%)	12	12	·	·	·	·	·
make_pair	§4.6	4/ 4	(100%)	4	·	·	·	·	·	·
minmax_element	§4.8	44/44	(100%)	27	17	·	·	·	·	·

Table A.3.: Results for maximum and minimum algorithms

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP	CQ
lower_bound	§5.1	19/19	(100%)	5	14	·	·	·	·	·
upper_bound	§5.2	19/19	(100%)	7	12	·	·	·	·	·
equal_range	§5.3	20/20	(100%)	17	3	·	·	·	·	·
equal_range2	§5.3	64/64	(100%)	24	35	1	·	·	·	4
binary_search	§5.4	10/10	(100%)	8	2	·	·	·	·	·
binary_search2	§5.4	10/10	(100%)	8	2	·	·	·	·	·

Table A.4.: Results for binary search algorithms

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP	CQ
fill	§6.3	12/12	(100%)	4	8	·	·	·	·	·
swap	§6.4	7/ 7	(100%)	7	·	·	·	·	·	·
swap_ranges	§6.5	22/22	(100%)	5	17	·	·	·	·	·
copy	§6.6	15/15	(100%)	4	11	·	·	·	·	·
copy_backward	§6.7	17/17	(100%)	7	10	·	·	·	·	·
reverse_copy	§6.8	17/17	(100%)	4	13	·	·	·	·	·
reverse	§6.9	24/24	(100%)	5	19	·	·	·	·	·
rotate_copy	§6.10	17/17	(100%)	5	12	·	·	·	·	·
rotate	§6.11	24/24	(100%)	10	14	·	·	·	·	·
replace_copy	§6.12	19/19	(100%)	7	12	·	·	·	·	·
replace	§6.13	15/15	(100%)	4	11	·	·	·	·	·
remove_copy	§6.14	21/21	(100%)	9	12	·	·	·	·	·
remove_copy2	§6.15	39/39	(100%)	9	28	1	·	·	·	1
remove_copy3	§6.16	44/44	(100%)	10	31	2	·	·	·	1
remove	§6.17	39/39	(100%)	9	26	3	·	·	·	1
random_shuffle	§6.18	60/60	(100%)	24	27	2	·	·	5	2
random_number	§6.18.3	23/23	(100%)	13	5	·	·	2	3	·

Table A.5.: Results for mutating algorithms

Algorithm		Verification Conditions		Individual Provers						
				QD	AE	C4	C3	Z3	EP	CQ
`unique_copy2`	§7.3.2	24/24	(100%)	8	16	·	·	·	·	·
`unique_copy3`	§7.3.3	27/27	(100%)	8	19	·	·	·	·	·
`unique_copy4`	§7.3.4	62/62	(100%)	10	43	3	2	·	1	3

Table A.6.: Results for variants of `unique_copy`

Algorithm		Verification Conditions		Individual Provers						
				QD	AE	C4	C3	Z3	EP	CQ
`iota`	§8.1	16/16	(100%)	7	9	·	·	·	·	·
`accumulate`	§8.2	19/19	(100%)	6	12	·	·	·	·	1
`inner_product`	§8.3	20/20	(100%)	6	13	·	·	·	·	1
`partial_sum`	§8.4	39/39	(100%)	9	25	4	·	·	·	1
`adjacent_difference`	§8.5	33/33	(100%)	11	20	2	·	·	·	·
`partial_sum_inv`	§8.6	19/20	(95%)	8	9	1	·	·	·	1
`adjacent_difference_inv`	§8.7	27/28	(96%)	7	15	3	·	·	·	2

Table A.7.: Results for numeric algorithms

Algorithm		Verification Conditions		Individual Provers						
				QD	AE	C4	C3	Z3	EP	CQ
`heap_parent`	§9.3	9/ 9	(100%)	2	7	·	·	·	·	·
`heap_child_max`	§9.3	22/ 22	(100%)	8	14	·	·	·	·	·
`is_heap_until`	§9.4	27/ 27	(100%)	6	21	·	·	·	·	·
`is_heap`	§9.5	12/ 12	(100%)	5	7	·	·	·	·	·
`push_heap`	§9.6	100/100	(100%)	33	52	13	·	·	·	2
`pop_heap`	§9.7	96/ 97	(98%)	49	41	4	·	·	·	2
`make_heap`	§9.8	38/ 38	(100%)	15	21	1	·	·	·	1
`sort_heap`	§9.9	54/ 54	(100%)	16	35	1	·	·	·	2

Table A.8.: Results for heap algorithms

Algorithm		Verification Conditions		Individual Provers						
				QD	AE	C4	C3	Z3	EP	CQ
`is_sorted`	§10.1	16/ 16	(100%)	7	8	·	·	·	·	1
`partial_sort`	§10.2	117/117	(100%)	39	59	7	·	·	·	12
`heap_sort`	§10.3	23/ 23	(100%)	8	14	·	·	·	·	1
`selection_sort`	§10.4	58/ 58	(100%)	17	33	2	·	·	·	6
`insertion_sort`	§10.5	67/ 67	(100%)	18	41	1	2	·	·	5
`merge`	§10.6	268/268	(100%)	185	71	10	·	·	·	2

Table A.9.: Results for algorithms related to sorting

Algorithm		Verification Conditions		Individual Provers						
				QD	AE	C4	C3	Z3	EP	CQ
stack_equal	§11.6	18/18	(100%)	7	11	·	·	·	·	·
stack_init	§11.7.1	14/14	(100%)	4	10	·	·	·	·	·
stack_size	§11.7.2	6/ 6	(100%)	1	5	·	·	·	·	·
stack_empty	§11.7.3	10/10	(100%)	5	5	·	·	·	·	·
stack_full	§11.7.4	11/11	(100%)	5	6	·	·	·	·	·
stack_top	§11.7.5	16/16	(100%)	6	10	·	·	·	·	·
stack_push	§11.7.6	44/44	(100%)	28	16	·	·	·	·	·
stack_pop	§11.7.7	32/32	(100%)	20	12	·	·	·	·	·

Table A.10.: Results for Stack functions

Algorithm		Verification Conditions		Individual Provers						
				QD	AE	C4	C3	Z3	EP	CQ
stack_size_wd	§11.7.2	12/12	(100%)	8	4	·	·	·	·	·
stack_empty_wd	§11.7.3	12/12	(100%)	8	4	·	·	·	·	·
stack_top_wd	§11.7.5	12/12	(100%)	8	4	·	·	·	·	·
stack_push_wd	§11.7.6	15/15	(100%)	3	12	·	·	·	·	·
stack_pop_wd	§11.7.7	12/12	(100%)	6	6	·	·	·	·	·

Table A.11.: Results for the well-definition of the Stack functions

Algorithm		Verification Conditions		Individual Provers						
				QD	AE	C4	C3	Z3	EP	CQ
axiom_size_of_init	§11.8.1	15/15	(100%)	11	4	·	·	·	·	·
axiom_size_of_push	§11.8.2	12/12	(100%)	9	3	·	·	·	·	·
axiom_top_of_push	§11.8.2	11/11	(100%)	8	3	·	·	·	·	·
axiom_size_of_pop	§11.8.3	11/11	(100%)	8	3	·	·	·	·	·
axiom_pop_of_push	§11.8.3	10/10	(100%)	6	4	·	·	·	·	·
axiom_push_of_pop_top	§11.8.3	15/15	(100%)	9	6	·	·	·	·	·

Table A.12.: Results for Stack axioms

A.3. Verification results (parallel)

In the *parallel verification scenario* each proof obligation is first passed to Frama-C/WP's built-in simplifier Qed. If Qed cannot discharge a proof obligation it is submitted in parallel to *all* the other automatic provers from Table 0.1.[47] Figure A.13 depicts this arrangement of provers. This arrangement of automatic theorem provers makes it a little bit easier to quantify their strength.

Figure A.13.: Parallel execution of automatic theorem provers

Note that in this scenario we used Frama-C/WP only for the generation and simplification of the proof obligations. For the parallel execution we developed our own (shell) scripts that pass the proof obligations directly through Why3 to the individual provers.

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP
						Individual Provers			
find	§3.1	19/19	(100%)	10	9	9	9	9	7
find2	§3.2	19/19	(100%)	10	9	9	9	9	5
find_first_of	§3.3	25/25	(100%)	16	9	9	9	9	3
adjacent_find	§3.4	22/22	(100%)	10	12	12	12	12	1
mismatch	§3.5	20/20	(100%)	10	10	10	10	10	5
equal	§3.5	7/ 7	(100%)	6	1	1	1	1	1
search	§3.6	29/30	(96%)	18	11	11	11	11	5
search_n	§3.7	90/91	(98%)	62	28	23	22	28	14
find_end	§3.8	27/28	(96%)	15	12	12	12	12	3
count	§3.9	15/28	(53%)	7	8	8	8	8	3
count2	§3.10	15/31	(48%)	7	8	7	7	7	3

Table A.14.: Results for non-mutating algorithms

---

[47]We did not include the interactive theorem prover Coq in this setting.

Algorithm		Verification Conditions		Individual Provers					
				QD	AE	C4	C3	Z3	EP
properties of operator <	§4.1	4/ 6	(66%)	4	·	·	·	·	·
clamp	§4.2	22/22	(100%)	18	4	4	4	4	3
max_element	§4.3	24/24	(100%)	13	11	11	11	11	6
max_element2	§4.4	24/24	(100%)	12	12	12	12	12	5
max_seq	§4.5	8/ 8	(100%)	5	3	3	3	3	1
min_element	§4.7	24/24	(100%)	12	12	12	12	12	5
make_pair	§4.6	4/ 4	(100%)	4	·	·	·	·	·
minmax_element	§4.8	44/44	(100%)	27	17	17	17	17	5

Table A.15.: Results for maximum and minimum algorithms

Algorithm		Verification Conditions		Individual Provers					
				QD	AE	C4	C3	Z3	EP
lower_bound	§5.1	19/19	(100%)	5	14	14	12	14	4
upper_bound	§5.2	19/19	(100%)	7	12	12	10	12	2
equal_range	§5.3	20/20	(100%)	17	3	3	3	3	·
equal_range2	§5.3	59/64	(92%)	24	34	35	32	35	5
binary_search	§5.4	10/10	(100%)	8	2	2	2	2	·
binary_search2	§5.4	10/10	(100%)	8	2	2	2	2	·

Table A.16.: Results for binary search algorithms

Algorithm		Verification Conditions		Individual Provers					
				QD	AE	C4	C3	Z3	EP
fill	§6.3	12/12	(100%)	4	8	8	8	8	3
swap	§6.4	7/ 7	(100%)	7	·	·	·	·	·
swap_ranges	§6.5	22/22	(100%)	5	17	17	17	17	8
copy	§6.6	15/15	(100%)	4	11	11	11	11	4
copy_backward	§6.7	17/17	(100%)	7	10	10	10	10	3
reverse_copy	§6.8	17/17	(100%)	4	13	12	13	13	3
reverse	§6.9	24/24	(100%)	5	19	17	18	19	3
rotate_copy	§6.10	17/17	(100%)	5	12	12	11	12	·
rotate	§6.11	24/24	(100%)	10	14	13	14	14	·
replace_copy	§6.12	19/19	(100%)	7	12	12	12	12	3
replace	§6.13	15/15	(100%)	4	11	11	11	11	5
remove_copy	§6.14	21/21	(100%)	9	12	12	12	12	2
remove_copy2	§6.15	26/39	(66%)	9	16	17	16	17	5
remove_copy3	§6.16	30/44	(68%)	10	19	20	15	20	6
remove	§6.17	26/39	(66%)	9	15	17	14	17	5
random_shuffle	§6.18	47/60	(78%)	24	16	17	17	17	11
random_number	§6.18.3	22/23	(95%)	13	4	4	4	6	8

Table A.17.: Results for mutating algorithms

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP
`unique_copy2`	§7.3.2	24/24	(100%)	8	16	16	16	16	3
`unique_copy3`	§7.3.3	27/27	(100%)	8	19	19	19	19	3
`unique_copy4`	§7.3.4	49/62	(79%)	10	37	36	34	36	12

Table A.18.: Results for variants of `unique_copy`

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP
`iota`	§8.1	16/16	(100%)	7	9	9	9	9	3
`accumulate`	§8.2	14/19	(73%)	6	8	8	8	8	3
`inner_product`	§8.3	19/20	(95%)	6	13	13	13	13	5
`partial_sum`	§8.4	30/39	(76%)	9	19	21	19	17	4
`adjacent_difference`	§8.5	30/33	(90%)	11	18	19	18	19	2
`partial_sum_inv`	§8.6	12/20	(60%)	8	4	4	4	4	3
`adjacent_difference_inv`	§8.7	12/28	(42%)	7	5	5	5	5	4

Table A.19.: Results for numeric algorithms

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP
`heap_parent`	§9.3	3/ 9	(33%)	2	1	1	1	1	1
`heap_child_max`	§9.3	16/ 22	(72%)	8	8	8	8	8	3
`is_heap_until`	§9.4	21/ 27	(77%)	6	15	13	15	15	5
`is_heap`	§9.5	6/ 12	(50%)	5	1	1	1	1	1
`push_heap`	§9.6	75/100	(75%)	33	34	40	33	36	8
`pop_heap`	§9.7	80/ 97	(82%)	49	27	31	28	29	7
`make_heap`	§9.8	24/ 38	(63%)	15	8	9	8	9	3
`sort_heap`	§9.9	32/ 54	(59%)	16	16	16	15	16	7

Table A.20.: Results for heap algorithms

Algorithm		Verification Conditions		QD	AE	C4	C3	Z3	EP
`is_sorted`	§10.1	14/ 16	(87%)	7	7	7	7	7	·
`partial_sort`	§10.2	82/117	(70%)	39	39	39	40	28	6
`heap_sort`	§10.3	9/ 23	(39%)	8	1	1	1	1	·
`selection_sort`	§10.4	37/ 58	(63%)	17	19	17	19	17	5
`insertion_sort`	§10.5	48/ 67	(71%)	18	28	26	28	29	8
`merge`	§10.6	262/268	(97%)	185	67	77	61	74	28

Table A.21.: Results for algorithms related to sorting

Algorithm		Verification Conditions		Individual Provers					
				QD	AE	C4	C3	Z3	EP
stack_equal	§11.6	15/18	(83%)	7	8	8	8	8	2
stack_init	§11.7.1	11/14	(78%)	4	7	7	7	7	2
stack_size	§11.7.2	3/ 6	(50%)	1	2	2	2	2	1
stack_empty	§11.7.3	7/10	(70%)	5	2	2	2	2	2
stack_full	§11.7.4	8/11	(72%)	5	3	3	3	3	2
stack_top	§11.7.5	13/16	(81%)	6	7	7	7	7	2
stack_push	§11.7.6	41/44	(93%)	28	13	13	13	13	2
stack_pop	§11.7.7	29/32	(90%)	20	9	9	9	9	3

Table A.22.: Results for `Stack` functions

Algorithm		Verification Conditions		Individual Provers					
				QD	AE	C4	C3	Z3	EP
stack_size_wd	§11.7.2	9/12	(75%)	8	1	1	1	1	1
stack_empty_wd	§11.7.3	9/12	(75%)	8	1	1	1	1	·
stack_top_wd	§11.7.5	9/12	(75%)	8	1	1	1	1	·
stack_push_wd	§11.7.6	11/15	(73%)	3	8	8	7	3	·
stack_pop_wd	§11.7.7	9/12	(75%)	6	3	3	3	3	·

Table A.23.: Results for the well-definition of the `Stack` functions

Algorithm		Verification Conditions		Individual Provers					
				QD	AE	C4	C3	Z3	EP
axiom_size_of_init	§11.8.1	12/15	(80%)	11	1	1	1	1	1
axiom_size_of_push	§11.8.2	9/12	(75%)	9	·	·	·	·	·
axiom_top_of_push	§11.8.2	8/11	(72%)	8	·	·	·	·	·
axiom_size_of_pop	§11.8.3	8/11	(72%)	8	·	·	·	·	·
axiom_pop_of_push	§11.8.3	7/10	(70%)	6	1	1	1	1	·
axiom_push_of_pop_top	§11.8.3	12/15	(80%)	9	3	3	3	1	2

Table A.24.: Results for `Stack` axioms

B. Changes in previous releases

This chapter describes the changes in previous versions of this document. For the most recent changes see Page 3.

The version numbers of this document are related to the versioning of Frama-C [2]. The versions of Frama-C are named consecutively after the elements of the periodic table. Therefore, our version numbering (X.Y.Z) are constructed as follows:

X the major number of our tutorial is the atomic number[48] of the chemical element after which Frama-C is named.

Y the Frama-C subrelease number

Z the subrelease number of this tutorial

B.1. New in Version 19.0.0 (Potassium, June 2019)

- Structure of document

  - The document is now structured into several parts.

  - The chapter on classic sorting algorithms has been merged into the chapter on sorting.

  - The various variants of `unique_copy` are now grouped into a separate chapter.

- Fix various inconsistencies

  - Change the return types of the logic functions `Accumulate`, `Difference`, `Capacity`, `Size`, `Top` from bounded one (e.g., `value_type`, `size_type`) to integer. A combination of bounded type for a logic function with an arithmetic operations in the logical definitions may lead to inconsistency. This fixes the inconsistencies in the `accumulate`, `stack` and `stack_wd` examples.

  - Fix an inconsistency in `DifferenceRead` axiom: restriction on the array size added to premises.

- Various improvements

  - An important change is the rewriting of the implicit, *axiomatic* definitions of `Accumulate`, `Count`, `Difference`, `InnerProduct` and `UniqueSize` logic functions to explicit, *recursive* ones. Accordingly, all axioms in the respective examples have been rewritten as lemmas.

  - Generalize `CountSectionMonotonic`, `UnchangedSection` lemmas: remove restriction on lower bound for the range.

  - Fix typo in postcondition of `find`.

  - Rewrite specifications of `remove_copy` and `remove` examples.

---

[48]See `http://en.wikipedia.org/wiki/Atomic_number`

- – Rename predicate `RemoveCount` to `RemoveSize`.

  – Gather all versions of `MultisetRetainRest` in section on `push_heap`.

  – Add another figure to highlight simple contract for `unique_copy`.

  – Adapt Coq proofs to the fact that the `Z` scope is not available by default.

- New examples

  – Add `count2` example with an inductive predicate instead of a logic function in `count`.

  – Add `merge` example.

- Infrastructure

  – Travis-CI configuration for the GitHub repository added as an illustrative example of how the verification results could be reproduced.

  – Add support for Frama-C/AstraVer plugin.

B.2. New in Version 18.0.0 (Argon, December 2018)

- Replace the links to the (now abandoned) original site of *Standard Template Library* (STL) by references to the C++ standard.

- Add new algorithm `unique_copy` (two versions).

- Add another assertion `half` for `reverse`.

- Add two overloaded versions of predicate `ConstantRange` and use them for the algorithms `fill` and `unique_copy`, respectively.

B.3. New in Version 17.1.0 (Chlorine, July 2018)

The exact version number of Frama-C originally was Chlorine-20180502. This version number was changed in October 2018 to 17.1

- Slightly change the definition of predicate `HasEqualNeighbors` and its use in the specification of `adjacent_find`.

- Remove the algorithm `remove` and the more elaborate version of `remove_copy`. We are currently working on new specifications of these algorithms.

- Adapt some Coq proofs related to the logic function `Count` in order to reflect changes in output of Frama-C/WP.

- Remove table on ACSL lemmas that had to be proved by Coq.

B.4. New in Version 16.1.1 (Sulfur, March 2018)

- fix several errors reported by Aaron Rocha, including,

  – fix an error in figure for `upper_bound` algorithms

- fix merging of contracts in second version of `binary_search`

- improve and justify the `retain` annotations of in the implementation of `remove`

- Alt-Ergo is now directly called in the parallel setting (instead of going through Why3) to be compatible with the sequential setting

- add a third assertion `reorder` in the `random_shuffle` body to keep verification rate at 100% after prover upgrade

B.5. New in Version 16.1.0 (Sulfur, December 2017)

- special thanks to Aaron Rocha who provided various improvements for Chapters 3, 4, and 5

- improve some mutating algorithms

  - add more assertions to `reverse` to reduce reliance on CVC3

  - improve structure and ACSL annotations of `remove_copy` and `remove`

    * add overloaded version of predicate `MultisetRetainRest`

    * add lemma `HasValueImpliesPositiveCount`

    * add lemma `PositiveCountImpliesHasValue`

    * remove lemma `HasValueShiftInversion`

    * remove lemma `HasValueCountInversion`

  - add custom lemma `random_number_modulo` for `random_shuffle`

- add new Chapter 10 with more algorithms related to sorting

  - add algorithm `is_sorted` including predicate `WeaklySorted`

    * add lemma `SortedImpliesWeaklySorted`

    * add lemma `WeaklySortedImpliesSorted`

  - add algorithm `partial_sort` including predicate `Partition`

    * add lemma `ReorderImpliesMatch`

    * add lemma `ReorderPreservesUpperBound`

    * add lemma `ReorderPreservesLowerBound`

    * add lemma `PartialReorderPreservesLowerBounds`

    * add lemma `SwappedInside`

    * add lemma `SwappedInsideMultisetUnchanged`

    * add lemma `SwappedInsidePreservesMultisetUnchanged`

- improve various lemmas

  - rename lemma `SortedUp` to `SortedUpperBound`

  - generalize lemma `UnchangedSection`

  - refactor lemma `HeapBounds` into `C_Division_2`

B.6. New in Version 15.1.2 (Phosphorus, October 2017)

- fix several typos reported by `seniorlackey@github` (thanks a lot!)
- add a new chapter on classic sorting algorithms which comprises
  - `selection_sort` including lemma `SwapImpliesMultisetUnchanged`
  - `insertion_sort` including lemmas
    * `EqualRangesPreservesSorted`
    * `RotatePreservesStrictLowerBound`
    * `RotateImpliesMultisetUnchanged`
    * `EqualRangesPreservesCount`
  - `heap_sort`
- heap algorithms
  - remove length requirements in `pop_heap`, `sort_heap`, `make_heap`, and `heap_sort`
    * introduce `SIZE_TYPE_MAX` to catch border cases in ACSL and C
  - improve description of `pop_heap`
    * add predicate `HeapChildMax`
    * provide the auxiliary function `heap_child_max`
    * the postcondition `reorder` is still not verified
  - improve description of `push_heap`
  - other, minor improvements
    * add auxiliary function `heap_parent`
    * add predicate `SortedDown` and lemma `SortedDownIsHeap`
    * add lemmas `HeapParentChild` and `HeapChilds`
    * add lemmas `HeapParentBounds` and `HeapChildBounds`

B.7. New in Version 15.1.1 (Phosphorus, September 2017)

- add ensures clause to default behavior of the following algorithms
  - `find`, `find_first_of`, `adjacent_find`, `mismatch`, `search`, `search_n`, `find_end`
  - `max_element`, `min_element`
- rewrite axiomatic definitions to ensure disjoint guards which is better suited for E-ACSL
  - concerns the axiomatic definitions of `Count`, `Accumulate`, `InnerProduct` and `Difference`
  - some Coq proofs related to `Count` had to be adapted as well

- shorten names of some auxiliary algorithms

  - `adjacent_difference_inverse` ↦ `adjacent_difference_inv`

  - `partial_sum_inverse` ↦ `partial_sum_inv`

- heap algorithms

  - fix a typo in Figure 9.4

  - fix a typo in Figure 9.41

  - explain that there can be multiple representations of an array as a heap

  - add a version of `pop_heap` that is, however, not completely verified

B.8. New in Version 15.1.0 (Phosphorus, June 2017)

- The verification results are now part of the appendix.

- Fix an error in the specification of the well-definition of `stack_size`.

- This release of Frama-C/WP could not discharge some of our assertions of `push_heap`. We therefore have completely rewritten the annotations and also tweaked the implementation of `push_heap`. We also added some new predicates and lemmas to maintain a concise specification that can easily be verified by automatic provers.

  - add predicate `MultisetAdd` and lemma `MultisetAddDistinct`

  - add predicate `MultisetMinus` and lemma `MultisetMinusDistinct`

  - add predicate `MultisetRetain` and lemma `MultisetPushHeapRetain`

  - provide an additional version of predicate `MultisetRetainRest`

  - and lemma `MultisetPushHeapClosure`

B.9. New in Version 14.1.1 (Silicon, April 2017)

- changes in verification infrastructure

  - add verification results for the case where each proof obligation is submitted to all automatic theorem provers (see Section A.3)

- changes in algorithms

  - simplify loop invariants of `search_n` and improve description

  - rename predicate `CountOneHit` to `CountHit`

  - rename predicate `CountOneMiss` to `CountMiss`

  - rewrite predicates `EqualRanges` and `Reverse` in order to simplify the task for automatic theorem provers

  - remove lemmas on `Reverse` that were necessary for `rotate` but are not needed anymore

  - rename predicate `Valid(Stack*)` to `Invariant(Stack*)` and remove `\valid` from `Invariant(Stack*)`

- add a simple random number generator to `random_shuffle` and verify it

- fix an inconsistency in the axioms for `Count` (thanks to Denis Efremov for reporting this issue)

  - add more guards to axioms `CountSectionHit` and `CountSectionMiss`

  - add corresponding guards to lemmas

    * `CountSectionOne`, `CountHit`, `CountMiss` and `CountOne`

    * `RemoveCountHit` and `RemoveCountMiss`

  - add lemma `UnchangedShift` and add more assertions to `remove` in order to simplify the task for automatic theorem provers

B.10. New in Version 14.1.0 (Silicon, January 2017)

- use label `Old` instead of `Pre` in function contracts

- add algorithm `rotate`

- rewrite definition of predicates `EqualRanges` and `Reverse` and provide more overloaded versions

- add figures for algorithms `rotate` and `replace_copy`

- update figure for predicate `Reverse`

- update Coq proofs and add a table with more information on the ACSL lemmas that had to be verified with Coq

B.11. New in Version 13.1.1 (Aluminium, November 2016)

- improve layout of tables of verification results

- use two additional automatic theorem provers (CVC3 and E)

- non-mutating algorithms

  - add algorithm `find_end`

  - add definition of predicate `HasSubRange` on subranges

  - add definition of predicate `EqualRanges` on subranges

  - rename lemma `HasSubRange_fit_size` to `HasSubRangeSize`

  - rename lemma `HasConstantSubRange_fit_size` to `HasSubRangeSize`

  - rename logic function `CountSection` to `Count` (using overloading in ACSL)

  - add lemma `HasValueCountInversion`

  - add lemma `HasValueShiftInversion`

  - add lemma `CountShift`

- mutating algorithms

  - add algorithm `copy_backward`

- relax precondition on separation of `copy`, `replace_copy` and `remove_copy`

- provide a more sophisticated implementation of `remove`

- re-introduce a second version of `remove_copy` that also specifies the *stability* of the algorithm

- add algorithm `random_shuffle`

B.12. New in Version 13.1.0 (Aluminium, August 2016)

The most notable changes of this version are the re-introduction of heap algorithms in Chapter 9. This new description of heap algorithms is based to a large extend on the bachelor thesis of one of the authors [29].

- provide names ("labels") for more ACSL annotations

- non-mutating algorithms

  - reorder and improve description in chapter on non-mutating algorithms

  - add more figures to describe algorithms

  - add non-mutating algorithm `search_n`

  - rewrite logic function `Count` with new logic function `CountSection`

  - move lemmas `CountBounds` and `CountMonotonic` to separate files

  - use `integer` instead of `size_type` in `HasSubRange`

  - change index computation in `HasEqualNeighbors`

- maximum and minimum algorithms

  - isolate predicate `ConstantRange` from predicates on lower and upper bounds

  - fix typo in precondition of first version of `max_element`

- binary search algorithms

  - add version `Sorted` for subranges

  - add second (more efficient) version of `equal_range`

    * add lemmas `SortedShift`, `LowerBoundShift`, `StrictLowerBoundShift`, `UpperBoundShift` and `StrictUpperBoundShift` to support the automatic verification of this version of `equal_range`

  - add figures to binary search algorithms and improve description

- mutating algorithms

  - greatly reduce the number of assertions needed to verify the first version `remove_copy`

  - temporarily remove the second version of `remove_copy` which also specified the *stability* of the algorithm

  - add `remove`, an in-place variant of `remove_copy`

  - rename predicate `RetainAllButOne` to `MultisetRetainRest`

- re-introduce chapter on heap algorithms

- includes the heap algorithms `is_heap`, `push_heap`, `make_heap` and `sort_heap`

- for `pop_heap` only a function contract is provided in this version

- add lemma `SortedUp` to support verification of `sort_heap`

- add several lemmas to combine the predicates `Unchanged` and `MultisetUnchanged`

B.13. New in Version 12.1.0 (Magnesium, February 2016)

A main goal of this release is to reduce the number of proof obligations that cannot be verified automatically and therefore must be tackled by an interactive theorem prover such as Coq. To this end, we analyzed the proof obligations (often using Coq) and devised additional assertions or ACSL lemmas to guide the automatic provers. Often we succeeded in enabling automatic provers to discharge the concerned obligations. Specifically, whereas the previous version 11.1.1 of *ACSL by Example* listed *nine* proof obligations that could only be discharged with Coq, the document at hand (version 12.1.0) only counts *five* such obligations. Moreover, all these remaining proof obligations are associated to ACSL lemmas, which are usually easier to tackle with Coq than proof obligations directly related to the C code. The reason for this is that ACSL lemmas usually have a much smaller set of hypotheses.

Adding assertions and lemmas also helps to alleviate a problem in Frama-C/WP Magnesium and Sodium where prover processes are not properly terminated.[49] Left-over "zombie processes" lead to a deterioration of machine performance which sometimes results in unpredictable verification results.

- mutating algorithms

  - simplify annotations of `replace_copy` and add new algorithm `replace`

    * add predicate `Replace` to write more compact post conditions and loops invariants

  - add several lemmas for predicate `Unchanged` and use predicate `Unchanged` in postconditions of mutating and numeric algorithms

  - simplify annotations of `reverse`

    * rename `Reversed` to `Reverse` (again) and provide another overloaded version

    * add figure to support description of the `Reverse` predicate

  - changes regarding `remove_copy`

    * rename `PreserveCount` to `RetainAllButOne`

    * rename `StableRemove` to `RemoveMapping`

    * add statement contracts for both versions of `remove_copy` such that only ACSL lemmas require Coq proofs

- numeric algorithms

  - define limits `VALUE_TYPE_MIN` and `VALUE_TYPE_MAX`

  - simplify specification of `iota` by using new logic function `Iota`

  - simplify implementation of `accumulate`

    * add overloaded predicates `AccumulateBounds`

---

[49]See https://bts.frama-c.com/view.php?id=2154

        * add lemmas `AccumulateDefault0`, `AccumulateDefault1`, `AccumulateDefaultNext`, and `AccumulateDefaultRead`

   – simplify implementation of `inner_product`

        * add predicates `ProductBounds` and `InnerProductBounds`

   – enable automatic verification of `partial_sum`

        * add lemmas `PartialSumSection`, `PartialSumUnchanged`, `PartialSumStep`, and `PartialSumStep2` to automatically discharge loop invariants

   – enable automatic verification of `adjacent_difference`

        * add logic function `Difference` and predicate `AdjacentDifference`

        * add predicate `AdjacentDifferenceBounds`

        * add lemmas `AdjacentDifferenceStep` and `AdjacentDifferenceSection` to automatically discharge proof obligation

   – add two auxiliary functions `partial_sum_inverse` and `adjacent_difference_inverse` in order to verify that `partial_sum` and `adjacent_difference` are inverse to each other

        * add lemmas `PartialSumInverse` and `AdjacentDifferenceInverse` to support the automatic verification of the auxiliary functions

- stack functions

   – add lemma `StackPushEqual` to enable the automatic verification of the well-definition of `stack_push`

B.14. New in Version 11.1.1 (Sodium, June 2015)

- add Chapter on numeric algorithms

   – move `iota` algorithm to numeric algorithms (Section 8.1)

   – add `accumulate` algorithm (Section 8.2)

   – add `inner_product` algorithm (Section 8.3)

   – add `partial_sum` algorithm (Section 8.4)

   – add `adjacent_difference` algorithm (Section 8.5)

B.15. New in Version 11.1.0 (Sodium, March 2015)

- Use built-in predicates `\valid` and `\valid_read` instead of `valid_range`.

- Simplify loop invariants of `find_first_of`.

- Replace two loop invariants of `remove_copy` by ACSL lemmas.

- Rename several predicates

   – `IsEqual` $\mapsto$ `EqualRanges`.

- IsMaximum $\mapsto$ MaxElement.

- IsMinimum $\mapsto$ MinElement.

- Reverse $\mapsto$ Reversed.

- IsSorted $\mapsto$ Sorted.

- Several changes for `Stack`:

  - Rename `Stack` functions from *foo*_stack to stack_*foo*.

  - Equality of stacks now ignores the `capacity` field. This is similar to how equality for objects of type `std::vector<T>` is defined. As a consequence `stack_full` is not well-defined any more. Other stack functions are not effected.

  - Remove all assertions from stack functions (including in axioms).

  - Describe predicate `Separated` in text.

B.16. New in Version 10.1.1 (Neon, January 2015)

- use option `-wp-split` to create simpler (but more) proof obligations

- simplify definition of predicate `Count`

- add new predicates for lower and upper bounds of ranges and use it in

  - `max_element`

  - `min_element`

  - `lower_bound`

  - `upper_bound`

  - `equal_range`

  - `fill`

- use a new auxiliary assertion in `equal_range` to enable the complete *automatic* verification of this algorithm

- add predicate `Unchanged` and use it to simplify the specification of several algorithms

  - `swap_ranges`

  - `reverse`

  - `remove_copy`

  - `stack_push` and `stack_push_wd`

  - `stack_pop` and `stack_pop_wd`

- add predicate `Reverse` and use it for more concise specifications of

  - `reverse_copy`

  - `reverse`

- several changes in the two versions of `remove_copy`

- – use predicate `HasValue` instead of logic function `Count`

- – add predicate `PreserveCount`

- – reformulate logic function `RemoveCount`

- – add predicate `StableRemove`

- – add predicate `RemoveCountMonotonic`

- – add predicate `RemoveCountJump`

- use overloading in ACSL to create shorter logic names for `Stack`

- remove unnecessary labels in several `Stack` functions

B.17. New in Version 10.1.0 (Neon, September 2014)

- remove additional labels in the `assumes` clauses of some stack function that were necessary due to an error in Oxygen

- provide a second version of `remove_copy` in order to explain the specification of the *stability* of the algorithms

- coarsen loop assigns of mutating algorithms

- temporarily remove the `unique_copy` algorithm

B.18. New in Version 9.3.1 (Fluorine, not published)

- specify bounds of the return value of `count` and fix reads clause of `Count` predicate

- use an auxiliary function `make_pair` in the implementation of `equal_range`

- provide more precise loop assigns clauses for the mutating algorithms

  - – simplify implementation of `fill`

  - – removed the `ensures \valid(p)` clause in specification of `swap`

  - – simplify implementation of `swap_ranges`

  - – simplify implementation of `copy`

  - – fix implementation of `reverse_copy` after discovering an undefined behavior

  - – new implementation of `reverse` that uses a simple **for**-loop

  - – simplify implementation of `replace_copy`

  - – refactor specification and simplify implementation of `remove_copy`

- remove work-around with `Pre`-label in `assumes` clauses of `stack_push` and `stack_pop`

B.19. New in Version 9.3.0 (Fluorine, December 2013)

- adjustments for *Fluorine* release of Frama-C

- swap now ensures that its pointer arguments are valid after the function has been called

- change definition of size_type to **unsigned int**

- change implementation of the iota algorithm . The content of the field a is calculated by increasing the value val instead of sum val+i.

- change implementation of fill.

- The specification/implementation of Stack has been revised by Kim Völlinger [32] and now has a much better verification rate.

B.20. New in Version 8.1.0 (Oxygen, not published)

- simplified specification and loop annotations of replace_copy

- add binary search variant equal_range

- greatly simplified specification of remove_copy by using the logic function Count

- remove chapter on heap operations

B.21. New in Version 7.1.1 (Nitrogen, August 2012)

- improvements with respect to several suggestions and comments of Yannick Moy, e.g., specification refinements of remove_copy, reverse_copy and iota

- restricted verification of algorithms to Frama-C/WP with Alt-Ergo

- replaced deprecated \valid_range by \valid

- fixed inconsistencies in the description of the Stack data type

- binary search algorithms can now be proven without additional axioms for integer division

- changed axioms into lemmas to document that provability is expected, even if not currently granted

- adopted new Fraunhofer logo and contact email

B.22. New in Version 7.1.0 (Nitrogen, December 2011)

- changed to Frama-C Nitrogen

- changed to Why 2.30

- discussed both plug-ins Frama-C/WP and Jessie

- removed swap_values algorithm

B.23. New in Version 6.1.0 (Carbon, not published)

- changed definition of Stack

- renamed reset_stack to init_stack

B.24. New in Version 5.1.1 (Boron, February 2011)

- prepared algorithms for checking by the new Frama-C/WP plug-in of Frama-C

- changed to Alt-Ergo Version 0.92, Z3 Version 2.11 and Why 2.27

- added List of user-defined predicates and logic functions

- added remarks on the relation of logical values in C and ACSL

- rewrote section on `equal` and `mismatch`

- used a simpler logical function to count elements in an array

- added `search` algorithm

- added chapter to unite the maximum/minimum algorithms

- added chapter for the new `lower_bound`, `upper_bound` and `binary_search` algorithms

- added `swap_values` algorithm

- used `IsEqual` predicate for `swap_ranges` and `copy`

- added `reverse_copy` and `reverse` algorithms

- added `rotate_copy` algorithm

- added `unique_copy` algorithm

- added chapter on specification of the data type `Stack`

B.25. New in Version 5.1.0 (Boron, May 2010)

- adaption to Frama-C Boron and Why 2.26 releases

- changed from the `-jessie-no-regions` command-line option to using the pragma `SeparationPolicy(value)`

B.26. New in Version 4.2.2 (Beryllium, May 2010)

- changed to latest version of CVC3 2.2

- added additional remarks to our implementation of `find_first_of`

- changed `size_type` (**int**) to `integer` in all specifications

- removed casts in `fill` and `iota`

- renamed `is_valid_range` as `IsValidRange`

- renamed `has_value` as `HasValue`

- renamed predicate `all_equal` as `IsEqual`

- extended timeout to 30 sec.

B.27. New in Version 4.2.1 (Beryllium, April 2010)

- added alternative specification of `remove_copy` algorithm that uses `ghost` variables

- added Chapter on heap operations

- added `mismatch` algorithm

- moved algorithms `adjacent_find` and `min_element` from the appendix to chapter on non-mutating algorithms

- added typedefs `size_type` and `value_type` and used them in all algorithms

- renamed `is_valid_int_range` as `is_valid_range`

B.28. New in Version 4.2.0 (Beryllium, January 2010)

- complete rewrite of pre-release

- adaption to Frama-C Beryllium 2 release

Bibliography

[1] WP Plug-in. `http://frama-c.com/wp.html`.

[2] Frama-C Software Analyzers. `http://frama-c.com`, 2018.

[3] Why – Where Programs Meet Provers. `http://why3.lri.fr`, 2018.

[4] Sylvain Conchon, Evelyne Contejean, and Johannes Kanig. The Alt-Ergo SMT Solver. `http://alt-ergo.lri.fr`, 2018.

[5] Clark Barrett and Cesare Tinelli. Homepage of CVC4. `http://cvc4.cs.stanford.edu/web/`, 2018.

[6] Clark Barrett and Cesare Tinelli. Homepage of CVC3. `http://www.cs.nyu.edu/acsys/cvc3/`, 2010.

[7] Microsoft Research. The Z3 Theorem Prover. `https://github.com/Z3Prover/z3`, 2018.

[8] Stephan Schulz. The E Theorem Prover. `https://wwwlehre.dhbw-stuttgart.de/~sschulz/E/E.html`, 2018.

[9] The Coq Consortium. The Coq Proof Assistant. `https://coq.inria.fr`, 2018.

[10] ANSI/ISO C Specification Language. `http://frama-c.com/acsl.html`, 2018.

[11] CEA LIST, Laboratory of Applied Research on Software-Intensive Technologies. `http://www-list.cea.fr/gb/index_gb.htm`.

[12] INRIA-Saclay, French National Institute for Research in Computer Science and Control . `http://www.inria.fr/saclay/`.

[13] LRI, Laboratory for Computer Science at Université Paris-Sud. `http://www.lri.fr/`.

[14] Fraunhofer-Institut für Offene Kommunikationssysteme (FOKUS). `http://www.fokus.fraunhofer.de`.

[15] Virgile Prevosto. ACSL Mini-Tutorial. `http://frama-c.com/download/acsl-tutorial.pdf`.

[16] Patrick Baudin, Pascal Cuoq, Jean-Christophe Filliâtre, Claude Marché, Benjamin Monate, Yannick Moy, and Virgile Prevosto. ACSL 1.13 Implementation in Argon 18.0. `https://frama-c.com/download/acsl-implementation-18.0-Argon.pdf`, 2018.

[17] Allan Blanchard. Introduction to C Program Proof using Frama-C and its wp plugin. `http://allan-blanchard.fr/publis/frama-c-wp-tutorial-en.pdf`, December 2017.

[18] Programming languages – C, Committee Draft. `http://www.open-std.org/JTC1/SC22/WG14/www/docs/n1362.pdf`, 2009.

[19] C.A.R. Hoare. An axiomatic basis for computer programming. *Communications of the ACM*, 12:576–583, 1969.

[20] Robert W. Floyd. Assigning meanings to programs. In J. T. Schwartz, editor, *Proc. Symposium*

on Applied Mathematics, volume 19 of *Mathematical Aspects of Computer Science*, pages 19–32, Providence, RI, 1967. American Mathematical Society.

[21] Richard Smith. Working Draft, Standard for Programming Language C++. `http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/n4659.pdf`, 2017. publicly available draft of C++ 17 standard.

[22] Donald E. Knuth, James H. Morris, and Vaughan R. Pratt. Fast pattern matching in strings. *SIAM J Comput*, 6(2):323–350, Jun 1977.

[23] Lincoln E. Moses and Robert V. Oakford. *Tables of Randon Permutations*. Stanford University Press, 1963.

[24] Ming Li and Paul Vitányi. *An Introduction to Kolmogorov Complexity and Its Applications*. Graduate texts in computer science. Springer, New York, 1997.

[25] Repository of "libc++" C++ Standard Library. `https://llvm.org/svn/llvm-project/libcxx/trunk`, 2018. Revision 345375: /libcxx/trunk.

[26] Patrick Baudin, François Bobot, Loïc Correnson, and Zaynah Dargaye. WP Plug-in Manual — Frama-C Argon 18.0. `http://frama-c.com/download/frama-c-wp-manual.pdf`, 2018.

[27] Loïc Correnson, Pascal Cuoq, Florent Kirchner, André Maroneze, Virgile Prevosto, Armand Puccetti, Julien Signoles, and Boris Yakobowski. Frama-C User Manual, Release Argon 18.0. https://frama-c.com/download/user-manual-18.0-Argon.pdf.

[28] Philippe Herrmann and Julien Signoles. Frama-C's annotation generator plug-in for Frama-C Argon 18.0. https://frama-c.com/download/rte-manual-18.0-Argon.pdf.

[29] Timon Lapawczyk. Formale Verifikation von Heap-Algorithmen mit Frama-C. bachelor thesis, Humboldt-Universität zu Berlin, July 2016.

[30] D.E. Knuth. *Sorting and Searching*, volume 3 of *The Art of Computer Programming*. Addison-Wesley, 1973.

[31] Sellibitze. How to Implement Classic Sorting Algorithms in Modern C++. `https://stackoverflow.com/questions/24650626/how-to-implement-classic-sorting-algorithms-in-modern-c`, Aug 2014.

[32] Kim Völlinger. Einsatz des Beweisassistenten Coq zur deduktiven Programmverifikation. Diplomarbeit, Humboldt-Universität zu Berlin, August 2013. `https://www2.informatik.hu-berlin.de/top/_media/www/mitarbeiter/diplomarbeit-kim-voellinger.pdf`.

[33] Richard Fitzpatrick J.L. Heiberg. *Euclid's Elements of Geometry*. `http://farside.ph.utexas.edu/euclid.html`, Austin/TX, 2008.

[34] David Hilbert. *Grundlagen der Geometrie*. B.G.Teubner, Stuttgart, 1968.

[35] Martin Odersky, Lex Spoon, and Bill Venners. *Programming in Scala*. Artima, 2008.