

ACSL By Example

Towards a Verified C Standard Library

Version 12.1.0
for
Frama-C (Magnesium)
February 2016

Jochen Burghardt
Jens Gerlach
Timon Lapawczyk


Former Authors

Andreas Carben
Liangliang Gu
Kerstin Hartig
Hans Pohl
Juan Soto
Kim Völlinger



The research leading to these results has received funding from the STANCE project¹ within European Union's Seventh Framework Programme [FP7/2007-2013] under grant agreement number 317753.²

This body of work was completed within the DEVICE-SOFT project, which was supported by the Programme Inter Carnot Fraunhofer from BMBF (Grant 01SF0804) and ANR.³



Except where otherwise noted, this work is licensed under
<http://creativecommons.org/licenses/by-nc-sa/3.0/>

¹ See <http://www.stance-project.eu>

²project duration: 2012–2016

³project duration: 2009–2012

Changes in Version 12.1.0 (February 2016)

This release is intended for Frama-C Magnesium, issued in January 2016.⁴ For changes in previous versions we refer to the Appendix A on Page 149.

A main goal of this release is to reduce the number of proof obligations that cannot be verified automatically and therefore must be tackled by an interactive theorem prover such as Coq. To this end, we analyzed the proof obligations (often using Coq) and devised additional assertions or ACSL lemmas to guide the automatic provers. Often we succeeded in enabling automatic provers to discharge the concerned obligations. Specifically, whereas the previous version 11.1.1 of *ACSL by Example* listed *nine* proof obligations that could only be discharged with Coq, the document at hand (version 12.1.0) only counts *five* such obligations. Moreover, all these remaining proof obligations are associated to ACSL lemmas, which are usually easier to tackle with Coq than proof obligations directly related to the C code. The reason for this is that ACSL lemmas usually have a much smaller set of hypotheses.

Adding assertions and lemmas also helps to alleviate a problem in WP Magnesium and Sodium where prover processes are not properly terminated.⁵ Left-over “zombie processes” lead to a deterioration of machine performance which sometimes results in unpredictable verification results.

List of changes

- mutating algorithms
 - simplify annotations of `replace_copy` and add new algorithm `replace`
 - * add predicate `Replace` to write more compact post conditions and loops invariants
 - add several lemmas for predicate `Unchanged` and use predicate `Unchanged` in postconditions of mutating and numeric algorithms
 - simplify annotations of `reverse`
 - * rename `Reversed` to `Reverse` (again) and provide another overloaded version
 - * add figure to support description of the `Reverse` predicate
 - changes regarding `remove_copy`
 - * rename `PreserveCount` to `RetainAllButOne`
 - * rename `StableRemove` to `RemoveMapping`
 - * add statement contracts for both versions of `remove_copy` such that only ACSL lemmas require Coq proofs
- numeric algorithms
 - define limits `VALUE_TYPE_MIN` and `VALUE_TYPE_MAX`

⁴ See <https://frama-c.com/download/frama-c-Magnesium-20151002.tar.gz>

⁵ See <https://bts.frama-c.com/view.php?id=2154>

- simplify specification of `iota` by using new logic function `Iota`
- simplify implementation of `accumulate`
 - * add overloaded predicates `AccumulateBounds`
 - * add lemmas `AccumulateDefault0`, `AccumulateDefault1`, `AccumulateDefaultNext`, and `AccumulateDefaultRead`
- simplify implementation of `inner_product`
 - * add predicates `ProductBounds` and `InnerProductBounds`
- enable automatic verification of `partial_sum`
 - * add lemmas `PartialSumSection`, `PartialSumUnchanged`, `PartialSumStep`, and `PartialSumStep2` to automatically discharge loop invariants
- enable automatic verification of `adjacent_difference`
 - * add logic function `Difference` and predicate `AdjacentDifference`
 - * add predicate `AdjacentDifferenceBounds`
 - * add lemmas `AdjacentDifferenceStep` and `AdjacentDifferenceSection` to automatically discharge proof obligation
- add two auxiliary functions `partial_sum_inverse` and `adjacent_difference_inverse` in order to verify that `partial_sum` and `adjacent_difference` are inverse to each other
 - * add lemmas `PartialSumInverse` and `AdjacentDifferenceInverse` to support the automatic verification of the auxiliary functions
- stack functions
 - add lemma `StackPushEqual` to enable the automatic verification of the well-definition of `stack_push`

Remarks on verification settings

This section gives all settings that depend on the software release of Frama-C, Why3, or one of its employed provers. For our experiments we used the WP plugin-in of Magnesium release of Frama-C and version 0.86.2 of the Why3 platform.⁶

Each verification conditions was submitted to the provers CVC4, Z3, Alt-Ergo, and Coq (in that order). The versions of the provers are listed in the following table.

Prover	Version	Reference
CVC4	1.4	https://cvc4.cs.nyu.edu/web
Z3	4.4	https://github.com/Z3Prover/z3
Alt-Ergo	0.99.1	http://alt-ergo.lri.fr
Coq	8.4.6	https://coq.inria.fr

Here are the most important options of Frama-C that we used in for almost all functions.

```
-pp-annot
-no-unicode
-wp
-wp-rte
-wp-model Typed+ref
-wp-timeout 10
-wp-steps 1000
-wp-coq-timeout 10
```

The only exceptions are the two versions of `remove_copy` in Sections 6.11 and 6.12, where we additionally used the option `-wp-split` and increased the time-out to 30 and 40 seconds, respectively.

⁶ See <http://why3.lri.fr>

Foreword

This report provides various examples for the formal specification, implementation, and deductive verification of C programs using the ANSI/ISO-C Specification Language (ACSL [1]) and the WP plug-in [2] of Frama-C [3] (Framework for Modular Analysis of C programs).

We have chosen our examples from the C++ standard library whose initial version is still known as the *Standard Template Library* (STL).⁷ The STL contains a broad collection of *generic* algorithms that work not only on C arrays but also on more elaborate container data structures. For the purposes of this document we have selected representative algorithms, and converted their implementation from C++ function templates to C functions that work on arrays of type `int`.⁸

We will continue to extend and refine this report by describing additional STL algorithms and data structures. Thus, step by step, this document will evolve from an ACSL tutorial to a report on a formally specified and deductively verified standard library for ANSI/ISO-C. Moreover, as ACSL is extended to a C++ specification language, our work may be extended to a deductively verified C++ Standard Library.

You may email suggestions, errors or greetings(!) to

`jens.gerlach@fokus.fraunhofer.de`

In particular, we encourage you to check vigilantly whether our formal specifications capture the essence of the informal description of the STL algorithms.

We appreciate your feedback and hope that this document helps foster the adoption of deductive verification techniques.

Acknowledgement

Many members from the Frama-C community provided valuable input and comments during the course of the development of this document. In particular, we wish to thank our project partners Patrick Baudin, Loïc Correnson, Zaynah Dargaye, Florent Kirchner, Virgile Prevosto, and Armand Puccetti from CEA LIST⁹ and Pascal Cuoq from TrustInSoft¹⁰.

We also like to express our gratitude to Claude Marché (LRI/INRIA)¹¹ and Yannick Moy (AdaCore)¹² for their helpful comments and detailed suggestions for improvement.

⁷ See <http://www.sgi.com/tech/stl/>

⁸ We are not directly using `int` in the source code but rather `value_type` which is a `typedef` for `int`.

⁹ <http://www-list.cea.fr/en>

¹⁰ <http://trust-in-soft.com>

¹¹ https://www.lri.fr/index_en.php?lang=EN

¹² <http://www.adacore.com>

Contents

Changes for Version 12.1.0	3
Foreword	6
1. Introduction	13
1.1. Structure of this document	13
1.2. Types, arrays, ranges and valid indices	14
2. The Hoare calculus	17
2.1. The assignment rule	19
2.2. The sequence rule	21
2.3. The implication rule	21
2.4. The choice rule	22
2.5. The loop rule	23
2.6. Derived rules	25
3. Non-mutating algorithms	27
3.1. The equal algorithm	28
3.2. The mismatch algorithm	32
3.3. The find algorithm	34
3.4. The find algorithm reconsidered	36
3.5. The find_first_of algorithm	38
3.6. The adjacent_find algorithm	40
3.7. The search algorithm	42
3.8. The count algorithm	45
4. Maximum and minimum algorithms	49
4.1. A note on relational operators	50
4.2. The max_element algorithm	52
4.3. The max_element algorithm with predicates	54
4.4. The max_seq algorithm	56
4.5. The min_element algorithm	58
5. Binary search algorithms	61
5.1. The lower_bound algorithm	62
5.2. The upper_bound algorithm	64
5.3. The equal_range algorithm	66
5.4. The binary_search algorithm	68
6. Mutating algorithms	71
6.1. The predicate Unchanged	72
6.2. The swap algorithm	72
6.3. The fill algorithm	74
6.4. The swap_ranges algorithm	76

6.5. The <code>copy</code> algorithm	78
6.6. The <code>reverse_copy</code> algorithm	80
6.7. The <code>reverse</code> algorithm	82
6.8. The <code>rotate_copy</code> algorithm	84
6.9. The <code>replace_copy</code> algorithm	86
6.10. The <code>replace</code> algorithm	88
6.11. The <code>remove_copy</code> algorithm	90
6.12. Capturing the stability of <code>remove_copy</code>	94
7. Numeric algorithms	99
7.1. The <code>iota</code> algorithm	100
7.2. The <code>accumulate</code> algorithm	102
7.3. The <code>inner_product</code> algorithm	106
7.4. The <code>partial_sum</code> algorithm	109
7.5. The <code>adjacent_difference</code> algorithm	113
7.6. Inverting <code>partial_sum</code> with <code>adjacent_difference</code>	117
7.7. Inverting <code>adjacent_difference</code> with <code>partial_sum</code>	118
8. The <code>Stack</code> data type	121
8.1. Methodology overview	122
8.2. Stack axioms	123
8.3. The structure <code>Stack</code> and its associated functions	125
8.4. Stack invariants	126
8.5. Equality of stacks	128
8.6. Runtime equality of stacks	130
8.7. Verification of stack functions	131
8.8. Verification of stack axioms	141
9. Results of formal verification	145
A. Changes in previous releases	149
A.1. New in Version 11.1.1 (June 2015)	149
A.2. New in Version 11.1.0 (March 2015)	149
A.3. New in Version 10.1.1 (January 2015)	150
A.4. New in Version 10.1.0 (September 2014)	151
A.5. New in Version 9.3.1 (not published)	151
A.6. New in Version 9.3.0 (December 2013)	151
A.7. New in Version 8.1.0 (not published)	152
A.8. New in Version 7.1.1 (August 2012)	152
A.9. New in Version 7.1.0 (December 2011)	152
A.10. New in Version 6.1.0 (not published)	152
A.11. New in Version 5.1.1 (February 2011)	152
A.12. New in Version 5.1.0 (May 2010)	153
A.13. New in Version 4.2.2 (May 2010)	153
A.14. New in Version 4.2.1 (April 2010)	153
A.15. New in Version 4.2.0 (January 2010)	154
Bibliography	155

List of Logic Specifications

3.2. The predicate <code>EqualRanges</code>	29
3.11. The predicate <code>HasValue</code>	36
3.14. The predicate <code>HasValueOf</code>	38
3.17. The predicate <code>HasEqualNeighbors</code>	40
3.21. The predicate <code>HasSubRange</code>	42
3.24. The logic function <code>Count</code>	45
3.25. Some lemmas for <code>Count</code>	46
4.1. Requirements for a partial order on <code>value_type</code>	50
4.2. Semantics of derived comparison operators	50
4.3. Predicates for comparing array elements with a given value	51
4.6. Definition of the <code>MaxElement</code> predicate	54
4.11. Definition of the <code>MinElement</code> predicate	58
5.1. The predicate <code>Sorted</code>	61
6.1. Predicate <code>Unchanged</code>	72
6.10. The predicate <code>Reverse</code>	80
6.19. The predicate <code>Replace</code>	86
6.25. The predicate <code>RetainAllButOne</code>	90
6.28. A lemma for <code>RetainAllButOne</code>	93
6.30. The logic function <code>RemoveCount</code>	94
6.31. Additional lemmas for <code>RemoveCount</code>	95
6.32. The predicate <code>RemoveMapping</code>	95
7.5. The logic function <code>Accumulate</code>	102
7.6. An overloaded version of <code>Accumulate</code>	103
7.7. The overloaded predicate <code>AccumulateBounds</code>	104
7.10. The logic function <code>InnerProduct</code>	106
7.11. The predicates <code>ProductBounds</code> and <code>InnerProductBounds</code>	107
7.14. The predicate <code>PartialSum</code>	109
7.17. The lemma <code>PartialSumStep</code>	112
7.18. The logic function <code>Difference</code>	114
7.19. The predicate <code>AdjacentDifference</code>	114
7.20. The predicate <code>AdjacentDifferenceBounds</code>	114
7.23. The lemma <code>AdjacentDifferenceStep</code>	116
7.24. The lemma <code>PartialSumInverse</code>	117
7.26. The lemma <code>AdjacentDifferenceInverse</code>	118
7.28. The lemma <code>UnchangedTransitive</code>	119
8.6. The logical functions <code>Capacity</code> , <code>Size</code> and <code>Top</code>	126
8.7. Predicates for empty an full stacks	127
8.8. The predicate <code>Valid</code>	127
8.9. Equality of stacks	128
8.11. Equality of stacks is an equivalence relation	129

8.30. The predicate <code>Separated</code>	138
8.32. The lemma <code>StackPushEqual</code>	138

List of Figures

3.20. Matching $b[0..n-1]$ in $a[0..m-1]$	42
6.11. Sketch of predicate <code>Reverse</code>	81
6.16. Effects of <code>rotate_copy</code>	84
6.24. Effects of <code>remove_copy</code>	90
6.29. Stability of <code>remove_copy</code> with respect to indices	94
8.1. Push and pop on a stack	121
8.2. Methodology Overview	122
8.4. Interpreting the data structure <code>Stack</code>	125
8.10. Example of two equal stacks	128
8.14. Methodology for the verification of well-definition	131

List of Tables

2.1. Some ACSL formula syntax	17
9.1. ACSL lemmas that were proved with Coq	145
9.2. Results for non-mutating algorithms	146
9.3. Results for maximum and minimum algorithms	146
9.4. Results for binary search algorithms	146
9.5. Results for mutating algorithms	147
9.6. Results for numeric algorithms	147
9.7. Results for <code>Stack</code> functions	148
9.8. Results for the well-definition of the <code>Stack</code> functions	148
9.9. Results for <code>Stack</code> axioms	148

1. Introduction

The Framework for Modular Analyses of C, **Frama-C** [3], is a suite of software tools dedicated to the analysis of C source code. Its development efforts are conducted and coordinated at two French public institutions: CEA LIST [4], a laboratory of applied research on software-intensive technologies, and INRIA Saclay[5], the French National Institute for Research in Computer Science and Control in collaboration with LRI [6], the Laboratory for Computer Science at Université Paris-Sud.

ACSL (ANSI/ISO-C Specification Language) [1] is a formal language to express behavioral properties of C programs. This language can specify a wide range of functional properties by adding annotations to the code. It allows to create function contracts containing preconditions and postconditions. It is possible to define type and global invariants as well as logic specifications, such as predicates, lemmas, axioms or logic functions. Furthermore, ACSL allows statement annotations such as assertions or loop annotations.

Within Frama-C, the WP plug-in [2] enables deductive verification of C programs that have been annotated with ACSL. The WP plug-in uses Hoare-style weakest precondition computations to formally prove ACSL properties of C code. Verification conditions are generated and submitted to external automatic theorem provers or interactive proof assistants.

The Verification Group at Fraunhofer FOKUS[7] see the great potential for deductive verification using ACSL. However, we recognize that for a novice there are challenges to overcome in order to effectively use the WP plug-in for deductive verification. In order to help users gain confidence, we have written this tutorial that demonstrates how to write annotations for existing C programs. This document provides several examples featuring a variety of annotated functions using ACSL. For an in-depth understanding of ACSL, we strongly recommend users to read the official Frama-C introductory tutorial [8] first. The principles presented in this paper are also documented in the ACSL reference document [9].

1.1. Structure of this document

The functions presented in this document were selected from the C++ standard library. The original C++ implementation was stripped from its generic implementation and mapped to C arrays of type `value_type`.

Chapter 2 provides a short introduction into the Hoare Calculus.

We have grouped various standard algorithms in chapters as follows:

- non-mutating algorithms (Chapter 3)
- maximum/minimum algorithms (Chapter 4)
- binary search algorithms (Chapter 5)
- mutating algorithms (Chapter 6)
- numeric algorithms (Chapter 7)

The order of these chapters reflects their increasing complexity.

Using the example of a stack, we tackle in Chapter 8 the problem of how a data type and its associated C functions can be specified with ACSL and automatically verified with Frama-C.

1.2. Types, arrays, ranges and valid indices

This section describe several general conventions and basic definitions we use throughout this document.

1.2.1. Types

In order to keep algorithms and specifications as general as possible, we use abstract type names on almost all occasions. We currently defined the following types:

```
typedef int value_type;

typedef unsigned int size_type;

typedef int bool;
```

Programmers who know the types associated with C++ standard library containers will not be surprised that `value_type` refers to the type of values in an array whereas `size_type` will be used for the indices of an array.

This approach allows one to modify e.g. an algorithm working on an `int` array to work on a `char` array by changing only one line of code, viz. the `typedef` of `value_type`. Moreover, we believe in better readability as it becomes clear whether a variable is used as an index or as a memory for a copy of an array element, just by looking at its type.

The latter reason also applies to the use of `bool`. To denote values of that type, we `#defined` the identifiers `false` and `true` to be 0 and 1, respectively. While any non-zero value is accepted to denote `true` in ACSL like in C the algorithms shown in this tutorial will always produce 1 for `true`. Due to the above definitions, the ACSL truth-value constant `\false` and `\true` can be used interchangeably with our `false` and `true`, respectively, in ACSL clauses, but not in C code.

1.2.2. Array and ranges

The C Standard describes an array as a “contiguously allocated nonempty set of objects” [10, §6.2.5.20]. If `n` is a constant integer expression with a value greater than zero, then

```
int a[n];
```

describes an array of type `int`. In particular, for each `i` that is greater than or equal to 0 and less than `n`, we can dereference the pointer `a+i`.

Let the following prototype represent a function, whose first argument is the address to a range and whose second argument is the length of this range.

```
void example(value_type* a, size_type n);
```

To be very precise, we have to use the term *range* instead of *array*. This is due to the fact, that functions may be called with empty ranges, i.e., with `n == 0`. Empty arrays, however, are not permitted according to the definition stated above. Nevertheless, we often use the term *array* and *range* interchangeably.

1.2.3. Specification of valid ranges in ACSL

The following ACSL fragment expresses the precondition that the function `example` expects that for each i , such that $0 \leq i < n$, the pointer $a+i$ may be safely dereferenced.

```
/*@  
    requires 0 <= n;  
    requires \valid(a+(0.. n-1));  
*/  
void example(value_type* a, size_type n);
```

In this case we refer to each index i with $0 \leq i < n$ as a *valid index* of a .

ACSL's built-in predicates `\valid(a + (0.. n))` and `\valid_read(a + (0.. n))` refer to all addresses $a+i$ where $0 \leq i \leq n$. However, the array notation `int a[n]` of the C programming language refers only to the elements $a+i$ where i satisfies $0 \leq i < n$. Users of ACSL must therefore use the range notation `a+(0.. n-1)` in order to express a valid array of length n .

2. The Hoare calculus

In 1969, C.A.R. Hoare introduced a calculus for formal reasoning about properties of imperative programs [11], which became known as “Hoare Calculus”.

The basic notion is

```

//@ assert P;
Q;
//@ assert R;
```

where P and R denote logical expressions and Q denotes a source-code fragment. Informally, this means

If P holds before the execution of Q , then R will hold after the execution.

Usually, P and R are called *precondition* and *postcondition* of Q , respectively. The syntax for logical expressions is described in [9, Section 2.2] in full detail. For the purposes of this tutorial, the notions shown in Table 2.1 are sufficient. Note that they closely resemble the logical and relational operators in C.

ACSL syntax	Name	Reading
$\neg P$	negation	P is not true
$P \ \&\& \ Q$	conjunction	P is true and Q is true
$P \ \ Q$	disjunction	P is true or Q is true
$P \ ==> \ Q$	implication	if P is true, then Q is true
$P \ <==> \ Q$	equivalence	if, and only if, P is true, then Q is true
$x < y == z$	relation chain	x is less than y and y is equal to z
<code>\forall</code> int x ; $P(x)$	universal quantifier	$P(x)$ is true for every int value of x
<code>\exists</code> int x ; $P(x)$	existential quantifier	$P(x)$ is true for some int value of x

Table 2.1.: Some ACSL formula syntax

Here we show three example source-code fragments and annotations.

<pre>//@ assert x % 2 == 1; ++x; //@ assert x % 2 == 0;</pre>	<p>If x has an odd value before execution of the code <code>++x</code> then x has an even value thereafter.</p>
---	---

<pre>//@ assert 0 <= x <= y; ++x; //@ assert 0 <= x <= y + 1;</pre>	<p>If the value of x is in the range $\{0, \dots, y\}$ before execution of the same code, then x's value is in the range $\{0, \dots, y + 1\}$ after execution.</p>
---	---

<pre>//@ assert true; while (--x != 0) sum += a[x]; //@ assert x == 0;</pre>	<p>Under any circumstances, the value of x is zero after execution of the loop code.</p>
--	---

Any C programmer will confirm that these properties are valid.¹³ The examples were chosen to demonstrate also the following issues:

- For a given code fragment, there does not exist one fixed pre- or postcondition. Rather, the choice of formulas depends on the actual property to be verified, which comes from the application context. The first two examples share the same code fragment, but have different pre- and postconditions.
- The postcondition need not be the most restricting possible formula that can be derived. In the second example, it is not an error that we stated only that $0 \leq x$ although we know that even $1 \leq x$.
- In particular, pre- and postconditions need not contain all variables appearing in the code fragment. Neither `sum` nor `a[]` is referenced in the formulas of the loop example.
- We can use the predicate `true` to denote the absence of a properly restricting precondition, as we did before the `while` loop.
- It is not possible to express by pre- and postconditions that a given piece of code will always terminate. The loop example only states that *if* the loop terminates, then $x == 0$ will hold. In fact, if x has a negative value on entry, the loop will run forever. However, if the loop terminates, $x == 0$ will hold, and that is what the loop example claims.

Usually, termination issues are dealt with separately from correctness issues. Termination proofs may, however, refer to properties stated (and verified) using the Hoare Calculus.

Hoare provided the rules shown in Listing 2.2 to 2.12 in order to reason about programs. We will comment on them in the following sections.

¹³ We leave the important issues of overflow aside for a moment.

2.1. The assignment rule

We start with the rule that is probably the least intuitive of all Hoare-Calculus rules, viz. the assignment rule. It is depicted in Listing 2.2, where

$$P \{x \mapsto e\}$$

denotes the result of substituting each occurrence of the variable x in the predicate P by the expression e .

```
//@ assert P {x |--> e};  
x = e;  
//@ assert P;
```

Listing 2.2: The assignment rule

For example, if P is the predicate

$$x > 0 \ \&\& \ a[2*x] == 0$$

then $P \{x \mapsto y + 1\}$ is the predicate

$$y+1 > 0 \ \&\& \ a[2*(y+1)] == 0$$

Hence, we get Listing 2.3 as an example instance of the assignment rule. Note that parentheses are required in the index expression to get the correct $2 * (y+1)$ rather than the faulty $2 * y + 1$.

```
//@ assert y+1 > 0 && a[2*(y+1)] == 0;  
x = y+1;  
//@ assert x > 0 && a[2*x] == 0;
```

Listing 2.3: An assignment rule example instance

Note that after a substitution several different predicates P may result in the same predicate $P \{x \mapsto e\}$. For example, after applying the substitution $P \{x \mapsto y + 1\}$ each of the following four predicates

$$\begin{aligned} x > 0 \ \&\& \ a[2*x] &== 0 \\ x > 0 \ \&\& \ a[2*(y+1)] &== 0 \\ y+1 > 0 \ \&\& \ a[2*x] &== 0 \\ y+1 > 0 \ \&\& \ a[2*(y+1)] &== 0 \end{aligned}$$

turns into

$$y+1 > 0 \ \&\& \ a[2*(y+1)] == 0$$

For this reason, the same precondition and statement may result in several different postconditions (All four above expressions are valid postconditions in Listing 2.3, for example). However, given a postcondition and a statement, there is only one precondition that corresponds.

When first confronted with Hoare's assignment rule, most people are tempted to think of a simpler and more intuitive alternative, shown in Listing 2.4.

```
//@ assert P;  
x = e;  
//@ assert P && x == e;
```

Listing 2.4: Simpler, but *faulty* assignment rule

Listings 2.5–2.7 show some example instances of this faulty rule.

```
//@ assert y > 0;  
x = y+1;  
//@ assert y > 0 && x == y+1;
```

Listing 2.5: An example instance of the faulty rule from Listing 2.4

While Listing 2.5 happens to be ok, Listing 2.6 and 2.7 lead to postconditions that are obviously nonsensical formulas.

```
//@ assert true;  
x = x+1;  
//@ assert x == x+1;
```

Listing 2.6: An example instance of the faulty rule from Listing 2.4

The reason is that in the assignment in Listing 2.6 the left-hand side variable x also appears in the right-hand side expression e , while the assignment in Listing 2.7 just destroys the property from its precondition.

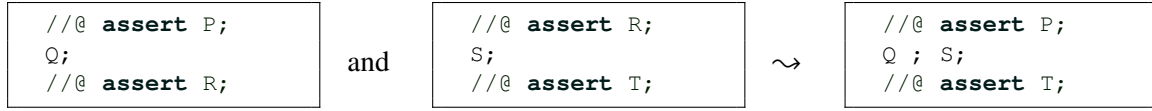
```
//@ assert x < 0;  
x = 5;  
//@ assert x < 0 && x == 5;
```

Listing 2.7: An example instance of the faulty rule from Listing 2.4

Note that the correct example Listing 2.5 can as well be obtained as an instance of the correct rule from Listing 2.2, since replacing x by $y+1$ in its postcondition yields $y > 0 \ \&\& \ y+1 == y+1$ as precondition, which is logically equivalent to just $y > 0$.

2.2. The sequence rule

The sequence rule, shown in Listing 2.8, combines two code fragments Q and S into a single one $Q ; S$. Note that the postcondition for Q must be identical to the precondition of S . This just reflects the sequential execution (“first do Q , then do S ”) on a formal level. Thanks to this rule, we may “annotate” a program with interspersed formulas, as it is done in Frama-C.



Listing 2.8: The sequence rule

2.3. The implication rule

The implication rule, shown in Listing 2.9, allows us at any time to sharpen a precondition P and to weaken a postcondition R . More precisely, if we know that $P' \implies P$ and $R \implies R'$ then we can replace the left contract in of Listing 2.9 by the right one.




Listing 2.9: The implication rule

2.4. The choice rule

The choice rule, depicted in Listing 2.10, is needed to verify conditional statements of the form

```
if (C) X;
else Y;
```

Both the then and else branch must establish the same postcondition, viz. S . The implication rule can be used to weaken differing postconditions $S1$ of a then-branch and $S2$ of an else-branch into a unified postcondition $S1 \mid\mid S2$, if necessary. In each branch, we may use what we know about the condition C , e.g. in the else-branch, that it is false. If the else-branch is missing, it can be considered as consisting of an empty sequence, having the postcondition $P \ \&\& \ !C$.



Listing 2.10: The choice rule

Listing 2.11 shows an example application of the choice rule.


```
//@ assert 0 <= i < n;           // given precondition
if (i < n-1) {
  //@ assert 0 <= i < n - 1;     // using that i < n-1 holds in this branch
  //@ assert 1 <= i+1 < n;       // by the implication rule
  i = i+1;
  //@ assert 1 <= i < n;         // by the assignment rule
  //@ assert 0 <= i < n;         // weakened by the implication rule
} else {
  //@ assert 0 <= i == n-1 < n; // using that !(i<n-1) holds in else part
  //@ assert 0 == 0 && 0 < n;    // weakened by the implication rule
  i = 0;
  //@ assert i == 0 && 0 < n;    // by the assignment rule
  //@ assert 0 <= i < n;        // weakened by the implication rule
}
//@ assert 0 <= i < n;           // by the choice rule from both branches
```

Listing 2.11: An example application of the choice rule

The variable `i` may be used as an index into a ring buffer `int a[n]`. The shown code fragment just advances the index `i` appropriately. We verified that `i` remains a valid index into `a[]` provided it was valid before. Note the use of the implication rule to establish preconditions for the assignment rule as needed, and to unify the postconditions of the then and else branches, as required by the choice rule.

2.5. The loop rule

The loop rule, shown in Listing 2.12, is used to verify a **while** loop. This requires to find an appropriate formula, P , which is preserved by each execution of the loop body. P is also called a loop invariant.



Listing 2.12: The loop rule

To find it requires some intuition in many cases; for this reason, automatic theorem provers usually have problems with this task.

As said above, the loop rule does not guarantee that the loop will always eventually terminate. It merely assures us that, if the loop has terminated, the postcondition holds. To emphasize this, the properties verifiable with the Hoare Calculus are usually called “partial correctness” properties, while properties that include program termination are called “total correctness” properties.

As an example application, let us consider an abstract ring-buffer. Listing 2.13 shows a verification proof for the index i lying always within the valid range $[0..n-1]$ during, and after, the loop. It uses the proof from Listing 2.11 as a sub-part.

```
//@ assert 0 < n; // given precondition  
  
int i = 0;  
//@ assert 0 <= i < n; // by the assignment rule  
  
while (!done) {  
    //@ assert 0 <= i < n && !done; // may be assumed by the loop rule  
  
    a[i] = getchar();  
    //@ assert 0 <= i < n && !done; // required property of getchar  
    //@ assert 0 <= i < n; // weakened by the implication rule  
  
    i = (i < n-1) ? i+1 : 0;  
    //@ assert 0 <= i < n; // follows by the choice rule  
  
    process(a, i, &done);  
    //@ assert 0 <= i < n; // required property of process  
}  
//@ assert 0 <= i < n; // by the loop rule
```

Listing 2.13: An abstract ring buffer loop

To reuse the proof from Listing 2.11, we had to drop the conjunct `!done`, since we didn't consider it in Listing 2.11. In general, we may *not* infer

<pre>//@ assert P && S; Q; //@ assert R && S;</pre>	from	<pre>//@ assert P; Q; //@ assert R;</pre>
---	------	---

since the code fragment `Q` may just destroy the property `S`.

This is obvious for `Q` being the fragment from Listing 2.11, and `S` being e.g. `i != 0`.

Suppose for a moment that `process` had been implemented in a way such that it refuses to set `done` to `true` unless it is `false` at entry. In this case, we would really need that `!done` still holds after execution of Listing 2.11. We would have to do the proof again, looping-through an additional conjunct `!done`.

We have similar problems to carry the property `0 <= i < n && !done` and `0 <= i < n` over the statement `a[i] = getchar()` and `process(a, i, &done)`, respectively. We need to specify that neither `getchar` nor `process` is allowed to alter the value of `i` or `n`. In ACSL, there is a particular language construct `assigns` for that purpose, which is introduced in Section 6.2 on Page 72.

In our example, the loop invariant can be established between any two statements of the loop body. However, this need not be the case in general. The loop rule only requires the invariant holds before the loop and at the end of the loop body. For example, `process` could well change the value of `i`¹⁴ and even `n` intermediately, as long as it re-establishes the property `0 <= i < n` immediately prior to returning.

The loop invariant, `0 <= i < n`, is established by the proof in Listing 2.11 also after termination of the loop. Thus, e.g., a final `a[i] = '\0'` after the loop would be guaranteed not to lead to a bounds violation.

Even if we would need the property `0 <= i < n` to hold only immediately before the assignment `a[i] = getchar()`, since, e.g., `process`'s body didn't use `a` or `i`, we would still have to establish `0 <= i < n` as a loop invariant by the loop rule, since there is no other way to obtain any property inside a loop body. Apart from this formal reason it is obvious that `0 <= i < n` wouldn't hold during the second loop iteration unless we re-established it at the end of the first one, and that is just what the while rule requires.

¹⁴We would have to change the call to `process(a, &i, &done)` and the implementation of `process` appropriately. In this case we couldn't rely on the above-mentioned `assigns` clause for `process`.

2.6. Derived rules

The above rules do not cover all kinds of statements allowed in C. However, missing C-statements can be rewritten into a form that is semantically equivalent and covered by the Hoare rules.

For example, if the expression E doesn't have side-effects, then

```
switch (E) {  
    case E1: Q1; break; ...  
    case En: Qn; break;  
    default: Q0; break;  
}
```

is semantically equivalent to

```
if (E == E1) {  
    Q1;  
} else ... if (E == En) {  
    Qn;  
} else {  
    Q0;  
}
```

While the **if-else** form is usually slower in terms of execution speed on a real computer, this doesn't matter for verification purposes, which are separate from execution issues.

Similarly, a loop statement of the form

```
for (P; Q; R) {  
    S;  
}
```

can be re-expressed as

```
P;  
while (Q) {  
    S;  
    R;  
}
```

and so on.

It is then possible to derive a Hoare rule for each kind of statement not previously discussed, by applying the classical rules to the corresponding re-expressed code fragment. However, we do not present these derived rules here.

Although procedures cannot be re-expressed in the above way if they are (directly or mutually) recursive, it is still possible to derive Hoare rules for them. This requires the finding of appropriate “procedure invariants” similar to loop invariants. Non-recursive procedures can, of course, just be inlined to make the classical Hoare rules applicable.

Note that **goto** cannot be rewritten in the above way; in fact, programs containing **goto** statements cannot be verified with the Hoare Calculus. See [12] for a similar calculus that can deal with arbitrary flowcharts, and hence arbitrary jumps. In fact, Hoare's work was based on that calculus. Later calculi inspired from Hoare's work have been designed to re-integrate support for arbitrary jumps. However, in this tutorial, we will not discuss example programs containing a **goto**.

3. Non-mutating algorithms

In this chapter, we consider *non-mutating* algorithms, i.e., algorithms that neither change their arguments nor any objects outside their scope. This requirement can be formally expressed with the following *assigns clause*:

```
assigns \nothing;
```

Each algorithm in this chapter therefore uses this assigns clause in its specification.

The specifications of these algorithms are not very complex. Nevertheless, we have tried to arrange them so that the earlier examples are simpler than the later ones. All algorithms work on one-dimensional arrays.

- `equal` (Section 3.1 on Page 28) compares two ranges element-by-element. Here, we will present two versions to specify such a function.
- `mismatch` (Section 3.2 on Page 32) returns the smallest index where two ranges differ. An implementation of `equal` using `mismatch` is also presented.
- `find` (Section 3.3 on Page 34) provides *sequential* or *linear search* and returns the smallest index at which a given value occurs in a range. In Section 3.4, on Page 36, a predicate is introduced in order to simplify the specification.
- `find_first_of` (Section 3.5, on Page 38) provides similar to `find` a *sequential search*, but unlike `find` it does not search for a particular value, but for the least index of a given range which occurs in another range.
- `adjacent_find` (Section 3.6 on Page 40) can be used to find equal neighbors in an array.
- `search` (Section 3.7, on Page 42) finds a subsequence that is identical to a given sequence when compared element-by-element and returns the position of the first occurrence.
- `count` (Section 3.8, on Page 45) returns the number of occurrences of a given value in a range. Here we will employ some user-defined axioms to formally specify `count`.

3.1. The `equal` algorithm

The `equal` algorithm in the C++ Standard Library compares two generic sequences. For our purposes we have modified the generic implementation¹⁵ to that of an array of type `value_type`. The signature now reads:

```
bool equal(const value_type* a, size_type n, const value_type* b);
```

The function returns `true` if `a[i] == b[i]` holds for each $0 \leq i < n$. Otherwise, `equal` returns `false`.

3.1.1. Formal specification of `equal`

The ACSL specification of `equal` is shown in Listing 3.1. We discuss the specification now line by line.

```
/*@
requires \valid_read(a + (0..n-1));
requires \valid_read(b + (0..n-1));

assigns \nothing;

behavior all_equal:
  assumes \forall integer i; 0 <= i < n ==> a[i] == b[i];
  ensures \result;

behavior some_not_equal:
  assumes \exists integer i; 0 <= i < n && a[i] != b[i];
  ensures !\result;

complete behaviors;
disjoint behaviors;
*/
bool equal(const value_type* a, size_type n, const value_type* b);
```

Listing 3.1: Formal specification of `equal`

The first part of our specification are the preconditions, which must be satisfied before the algorithm is executed. Those requirements can be specified with the `requires`-clause in ACSL. In case of the `equal` algorithm it is needed that `n` is non-negative (not specified) and that the pointers `a` and `b` point to `n` contiguously allocated objects of type `value_type` (see also Section 1.2).

In the second part of our specification we make a statement about objects and arguments that the function is allowed to change. Since `equal` is a non-mutating algorithm and does not modify any memory location outside its scope we just define `assigns \nothing` (see Page 27).

Finally, we define the postconditions, which must hold after the `equal` algorithm is finished. Corresponding to the informal description from the STL documentation, we have two behaviors:

The behavior `all_equal` applies if an element-wise comparison of the two ranges yields that they are all equal (this is formalized in the first `assumes` clause. In this case the function `equal` is expected to return `true`; we express this by “`ensures \result`”. The behavior `some_not_equal` applies if there is at least one valid index `i` where the elements `a[i]` and `b[i]` differ (second `assumes` clause). In this case the function `equal` is expected to return `false`, expressed as “`ensures !\result`”.

The negation of the formula

¹⁵ See <http://www.sgi.com/tech/stl/equal.html>

```
\forall i: \text{integer}; 0 \leq i < n \implies a[i] == b[i];
```

in behavior `all_equal` is just the formula

```
\exists i: \text{integer}; 0 \leq i < n \wedge a[i] != b[i];
```

in behavior `some_not_equal`. Therefore, these two behaviors complement each other. Also note that the variable `i` is not of type `int`, but of type `integer`. While the former type comprises finitely many (often just 4294967296) distinct numbers available on the target platform hardware, the latter type contains numbers of arbitrary size, and is allowed only in ACSL specifications. Using type `integer` becomes a real issue e.g. in Sect. 7.1.

The complete behaviors-clause in Listing 3.1 expresses the fact that for all ranges `a` and `b` that satisfy the preconditions of the contract *at least one* of the specified named behaviors, in this case `all_equal` and `some_not_equal`, applies.

The disjoint behaviors-clause in Listing 3.1 formalizes the fact that for all ranges `a` and `b` that satisfy the preconditions of the contract *at most one* of the specified named behaviors, in this case `all_equal` and `some_not_equal`, applies.

3.1.2. The `EqualRanges` predicate

The fact that two arrays `a[0]..a[n-1]` and `b[0]..b[n-1]` are equal when compared element by element, is a property we might need again in other specifications, as it describes a very basic behavior.

The motto *don't repeat yourself* is not just good programming practice.¹⁶ It is also true for concise and easy to understand specifications. We will therefore introduce specification elements that we can apply to the `equal` algorithm as well as to other specifications and implementations with the described behavior.

In Listing 3.2 we introduce the predicate `EqualRanges`.

```
/*@
  predicate
    EqualRanges(A,B){value_type* a, integer n, value_type* b) =
      \forall i: \text{integer}; 0 \leq i < n \implies \text{at}(a[i], A) == \text{at}(b[i], B);
*/
```

Listing 3.2: The predicate `EqualRanges`

This predicate formalizes the fact that the arrays `a[0]..a[n-1]` and `b[0]..b[n-1]` are equal when compared element by element. The letters `A` and `B` are *labels*¹⁷ that must be supplied when using the predicate `EqualRanges`. We use labels in the definition of `EqualRanges` to extend its applicability. The expression `\at(a[i], A)` means that `a[i]` is evaluated at the label `A`. Frama-C defines several standard labels, e.g. `Old` and `Post`, a programmer can use to refer to the pre-state or post-state, respectively, of a function. For more details on labels we refer to the ACSL specification [9, p. 39].

Using this predicate we can reformulate the specification of `equal` in Listing 3.1 as shown in Listing 3.3. Here we use the predefined label `Here`. When used in an `ensures` clause the label `Here` refers to the pre-state of a function.

Note that the equivalence is needed in the `ensures` clause. Putting an equality instead is not legal in ACSL, because `EqualRanges` is a predicate.

¹⁶ Compare http://en.wikipedia.org/wiki/Don't_repeat_yourself

¹⁷ Labels are used in C to name the target of the `goto` jump statement.

```

/*@
  requires \valid_read(a + (0..n-1));
  requires \valid_read(b + (0..n-1));

  assigns \nothing;

  ensures \result <==> EqualRanges{Here,Here}(a, n, b);
*/
bool equal(const value_type* a, size_type n, const value_type* b);

```

Listing 3.3: Formal specification of `equal` using the `EqualRanges` predicate

3.1.3. Implementation of `equal`

Listing 3.4 shows one way to implement the function `equal`. In our description, we concentrate on the *loop annotations*.

```

bool equal(const value_type* a, size_type n, const value_type* b)
{
  /*@
    loop invariant 0 <= i <= n;
    loop invariant \forall integer k; 0 <= k < i ==> a[k] == b[k];
    loop assigns i;
    loop variant n-i;
  */
  for (size_type i = 0; i < n; i++) {
    if (a[i] != b[i]) {
      return false;
    }
  }

  return true;
}

```

Listing 3.4: Implementation of `equal`

The first loop *invariant* is needed to prove that all accesses to `a` and `b` occur with valid indices. However, we may *not* require simply

```
loop invariant 0 <= i < n;
```

since the very last loop iteration would violate this formula. Therefore, we have to weaken the formula to that shown in the implementation of Listing 3.4, which is preserved by *all* iterations of the loop. Note that $0 \leq i < n$ is still valid immediately before the array accesses in, since we may assume there in addition that the loop condition $i < n$ holds. However, $0 \leq i < n$ is invalid after completion of the loop, while the loop invariant is guaranteed to hold there, too, cf. the loop rule in Figure 2.12 on Page 23.

Most important is the last loop *invariant*. It complies with the postcondition of the specification in Listing 3.3 and is needed to prove that for each iteration all elements of `a` and `b` up to that iteration step are equal. The loop assigns-clause in Listing 3.4 expresses that only the loop index is modified in any iteration. This is in accordance with the fact that `equal` is a *non-mutating* algorithm. The loop *variant* is needed to generate correct verification conditions for the termination of the `for`-loop. In order to prove the termination of the loop, Frama-C needs to know an expression whose value is decreased by each and every loop cycle and is always positive¹⁸[9, Subsections 2.4.2, 2.5.1]. For a `for` loop as simple as that the expression `n-i` is

¹⁸ Except for possibly the very last iteration.

sufficient for that purpose. Again, we can use the predicate `EqualRanges` in order to simplify the second loop invariant, which complies our postcondition. Listing 3.5 shows the modified implementation using the predicate `EqualRanges`.

```
bool equal(const value_type* a, size_type n, const value_type* b)
{
    /*@
        loop invariant 0 <= i <= n;
        loop invariant EqualRanges{Here,Here}(a, i, b);
        loop assigns    i;
        loop variant    n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        if (a[i] != b[i]) {
            return false;
        }
    }

    return true;
}
```

Listing 3.5: Implementation of `equal` using the `EqualRanges` predicate

3.2. The mismatch algorithm

The `mismatch` algorithm is closely related to the negation of `equal` from Section 3.1. Its signature reads

```
size_type mismatch(const value_type* a, int n,
                  const value_type* b);
```

The function `mismatch` returns the smallest index where the two ranges `a` and `b` differ. If no such index exists, that is, if both ranges are equal then `mismatch` returns the length `n` of the two ranges.¹⁹

3.2.1. Formal specification of `mismatch`

We use the `EqualRanges` predicate defined in Listing 3.2 also for the formal specification of `mismatch` that is shown in Listing 3.6.

Note in particular the use of `EqualRanges` in the specification shown in Listing 3.6 in order to express that `mismatch` returns the *smallest* index where the two arrays differ. Note also that the completeness and disjointedness of the behaviors `all_equal` and `some_not_equal` has now become immediately obvious, since their `assumes` clauses are just literal negations of each other.

```
/*@
requires \valid_read(a + (0..n-1));
requires \valid_read(b + (0..n-1));

assigns \nothing;

behavior all_equal:
  assumes EqualRanges{Here,Here}(a, n, b);
  ensures \result == n;

behavior some_not_equal:
  assumes !EqualRanges{Here,Here}(a, n, b);
  ensures 0 <= \result < n;
  ensures a[\result] != b[\result];
  ensures EqualRanges{Here,Here}(a, \result, b);

complete behaviors;
disjoint behaviors;
*/
size_type mismatch(const value_type* a, size_type n,
                  const value_type* b);
```

Listing 3.6: Formal specification of `mismatch`

¹⁹ See also <http://www.sgi.com/tech/stl/mismatch.html>

3.2.2. Implementation of mismatch

Listing 3.7 shows an implementation of `mismatch` that we have enriched with some loop annotations to support the deductive verification.

```
size_type mismatch(const value_type* a, size_type n,
                  const value_type* b)
{
    /*@
      loop invariant 0 <= i <= n;
      loop invariant EqualRanges{Here,Here}(a, i, b);
      loop assigns i;
      loop variant n-i;
    */
    for (size_type i = 0; i < n; i++) {
        if (a[i] != b[i]) {
            return i;
        }
    }

    return n;
}
```

Listing 3.7: Implementation of `mismatch`

We use the predicate `EqualRanges` as shown in Listing 3.7 in order to express that all indices k that are less than the current index i satisfy the condition $a[k] == b[k]$. This is necessary to prove that `mismatch` indeed returns the smallest index where the two ranges differ.

3.2.3. Implementation of `equal` by calling `mismatch`

Listing 3.8 shows an implementation of the `equal` algorithm by a simple call of `mismatch`.²⁰

```
bool equal(const value_type* a, size_type n, const value_type* b)
{
    return mismatch(a, n, b) == n;
}
```

Listing 3.8: Implementation of `equal` with `mismatch`

²⁰ See also the note on the relationship of `equal` and `mismatch` on <http://www.sgi.com/tech/stl/equal.html>

3.3. The find algorithm

The `find` algorithm in the C++ standard library implements *sequential search* for general sequences.²¹ We have modified the generic implementation, which relies heavily on C++ templates, to that of a range of type `value_type`. The signature now reads:

```
size_type find(const value_type* a, size_type n, value_type val);
```

The function `find` returns the least *valid* index `i` of `a` where the condition `a[i] == val` holds. If no such index exists then `find` returns the length `n` of the array.

3.3.1. Formal specification of find

The formal specification of `find` in ACSL is shown in Listing 3.9.

```
/*@
  requires  \valid_read(a + (0..n-1));

  assigns  \nothing;

  behavior some:
    assumes \exists integer i; 0 <= i < n && a[i] == val;
    ensures 0 <= \result < n;
    ensures a[\result] == val;
    ensures \forall integer i; 0 <= i < \result ==> a[i] != val;

  behavior none:
    assumes \forall integer i; 0 <= i < n ==> a[i] != val;
    ensures \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type find(const value_type* a, size_type n, value_type val);
```

Listing 3.9: Formal specification of `find`

The `requires`-clause indicates that `n` is non-negative and that the pointer `a` points to `n` contiguously allocated objects of type `value_type` (see Section 1.2).

The `assigns`-clause indicates that `find` (as a non-mutating algorithm), does not modify any memory location outside its scope (see Page 27).

We have subdivided the specification of `find` into two behaviors (named `some` and `none`). The behavior `some` applies if the sought-after value is contained in the array. We express this condition by using the `assumes`-clause. The next line expresses that if the assumptions of the behavior are satisfied then `find` will return a valid index. The algorithm also ensures that the returned (valid) index `i`, `a[i] == val` holds. Therefore we define this property in the second postcondition of behavior `some`. Finally, it is important to express that `find` return the smallest index `i` for which `a[i] == val` holds (see last postcondition of behavior `some`).

The behavior `none` covers the case that the sought-after value is *not* contained in the array (see `assumes`-clause of behavior `none` in Listing 3.9). In this case, `find` must return the length `n` of the range `a`.

²¹ See <http://www.sgi.com/tech/stl/find.html>

Note that the formula in the `assumes`-clause of the behavior `some` is the negation of the `assumes`-clause of the behavior `none`. Therefore, we can express that these two behaviors are *complete* and *disjoint*.

3.3.2. Implementation of `find`

Listing 3.10 shows a straightforward implementation of `find`. The only noteworthy elements of this implementation are the *loop annotations*.

```
size_type find(const value_type* a, size_type n, value_type val)
{
    /*@
      loop invariant 0 <= i <= n;
      loop invariant \forall integer k; 0 <= k < i ==> a[k] != val;
      loop assigns i;
      loop variant n-i;
    */
    for (size_type i = 0; i < n; i++) {
        if (a[i] == val) {
            return i;
        }
    }

    return n;
}
```

Listing 3.10: Implementation of `find`

The first loop *invariant* is needed to prove that accesses to `a` only occur with valid indices. With the second loop *invariant* is needed for the proof of the postconditions of the behavior `some` (see Listing 3.9). It expresses that for each iteration the sought-after value is not yet found up to that iteration step.

Finally, the loop *variant* `n-i` is needed to generate correct verification conditions for the termination of the loop.

3.4. The `find` algorithm reconsidered

In this section we specify the `find` algorithm in a slightly different way when compared to Section 3.3. Our approach is motivated by a considerable number of closely related formulas. We have in Listings 3.9 and 3.10 the following formulas

<code>\exists</code>	<code>integer i; 0 <= i < n</code>	<code>&& a[i] == val;</code>
<code>\forall</code>	<code>integer i; 0 <= i < \result</code>	<code>=> a[i] != val;</code>
<code>\forall</code>	<code>integer i; 0 <= i < n</code>	<code>=> a[i] != val;</code>
<code>\forall</code>	<code>integer k; 0 <= k < i</code>	<code>=> a[k] != val;</code>

Note that the first formula is the negation of the third one.

In order to be more explicit about the commonalities of these formulas we define a predicate, called `HasValue` (see Listing 3.11), which describes the situation that there is a valid index `i` such that

$$a[i] == val$$

holds.

```
/*@
  predicate
    HasValue{A}(value_type* a, integer n, value_type v) =
      \exists integer i; 0 <= i < n && a[i] == v;
*/
```

Listing 3.11: The predicate `HasValue`

Note that we needed to provide a label, viz. `A`, to the predicate, since its evaluation depends on a memory state, viz. then contents of `a[]`. ACSL allows to abbreviate `\at(a[i], A)` by `a[i]` if, as in our predicate body, `A` is the only available label.

With this predicate we can encapsulate all uses of the \forall and \exists quantifiers in both the specification of the function contract of `find` and in the loop annotations. The result is shown in Listings 3.12 and 3.13.

3.4.1. Formal specification of `find`

This approach leads to a specification of `find` that is more readable than the one described in Section 3.3.

In particular, it can be seen immediately that the conditions in the assumes clauses of the two behaviors `some` and `none` are mutually exclusive since one is the literal negation of the other. Moreover, the requirement that `find` returns the smallest index can also be expressed using the `HasValue` predicate, as depicted with the last postcondition of behavior `some` as shown in Listing 3.12.

```

/*@
requires   \valid_read(a + (0..n-1));

assigns   \nothing;

behavior some:
  assumes   HasValue(a, n, val);
  ensures   0 <= \result < n;
  ensures   a[\result] == val;
  ensures   !HasValue(a, \result, val);

behavior none:
  assumes   !HasValue(a, n, val);
  ensures   \result == n;

complete behaviors;
disjoint behaviors;
*/
size_type find(const value_type* a, size_type n, value_type val);

```

Listing 3.12: Formal specification of find using the HasValue predicate

3.4.2. Implementation of find

The predicate HasValue is also used in the loop annotation inside the implementation of find.

```

size_type find(const value_type* a, size_type n, value_type val)
{
  /*@
    loop invariant 0 <= i <= n;
    loop invariant !HasValue(a, i, val);
    loop assigns i;
    loop variant n-i;
  */
  for (size_type i = 0; i < n; i++) {
    if (a[i] == val) {
      return i;
    }
  }

  return n;
}

```

Listing 3.13: Implementation of find with loop annotations based on HasValue

3.5. The `find_first_of` algorithm

The `find_first_of` algorithm²² is closely related to `find` (see Sections 3.3 and 3.4).

```
size_type find_first_of(const value_type* a, size_type m,
                       const value_type* b, size_type n);
```

As in `find` it performs a sequential search. However, whereas `find` searches for a particular value, `find_first_of` returns the least index i such that $a[i]$ is equal to one of the values $b[0..n-1]$.

3.5.1. Formal specification of `find_first_of`

Similar to our approach in Section 3.4, we define a predicate `HasValueOf` that formalizes the fact that there are valid indices i for a and j of b such that $a[i] == b[j]$ hold. We have chosen to reuse the predicate `HasValue` (Listing 3.11) to define `HasValueOf` (Listing 3.14).

```
/*@
predicate
  HasValueOf{A}(value_type* a, integer m, value_type* b, integer n) =
    \exists integer i; 0 <= i < m && HasValue{A}(b, n, a[i]);
*/
```

Listing 3.14: The predicate `HasValueOf`

Both the predicates `HasValueOf` and `HasValue` occur in the formal specification of `find_first_of` (see Listing 3.15). Note how similar the specification of `find_first_of` becomes to that of `find` (Listing 3.12) when using these predicates.

```
/*@
requires \valid_read(a + (0..m-1));
requires \valid_read(b + (0..n-1));

assigns \nothing;

behavior found:
  assumes HasValueOf(a, m, b, n);
  ensures bound: 0 <= \result < m;
  ensures result: HasValue(b, n, a[\result]);
  ensures first: !HasValueOf(a, \result, b, n);

behavior not_found:
  assumes !HasValueOf(a, m, b, n);
  ensures result: \result == m;

complete behaviors;
disjoint behaviors;
*/
size_type find_first_of(const value_type* a, size_type m,
                       const value_type* b, size_type n);
```

Listing 3.15: Formal specification of `find_first_of`

²² See http://www.sgi.com/tech/stl/find_first_of.html

3.5.2. Implementation of `find_first_of`

Our implementation of `find_first_of` is shown in Listing 3.16.

Note the call of the `find` function shown in the Listing above. In the original implementation²³, `find_first_of` does not call `find` but rather inlines it. The reason for this were probably efficiency considerations. We opted for an implementation of `find_first_of` that emphasizes reuse. Besides, leading to a more concise implementation, we also have to write less loop annotations.

```
size_type find_first_of (const value_type* a, size_type m,
                        const value_type* b, size_type n)
{
    /*@
    loop invariant bound: 0 <= i <= m;
    loop invariant not_found: !HasValueOf(a, i, b, n);
    loop assigns i;
    loop variant m-i;
    */
    for (size_type i = 0; i < m; i++) {
        if (find(b, n, a[i]) < n) {
            return i;
        }
    }

    return m;
}
```

Listing 3.16: Implementation of `find_first_of`

²³ See http://www.sgi.com/tech/stl/stl_algo.h

3.6. The adjacent_find algorithm

The adjacent_find algorithm²⁴

```
size_type adjacent_find(const value_type* a, size_type n);
```

returns the smallest valid index i , such that $i+1$ is also a valid index and such that

$$a[i] == a[i+1]$$

holds. The adjacent_find algorithm returns n if no such index exists.

3.6.1. Formal specification of adjacent_find

As in the case of other search algorithms, we first define a predicate HasEqualNeighbors (see Listing 3.17) that captures the essence of finding two adjacent indices at which the array holds equal values.

```
/*@
  predicate
    HasEqualNeighbors(A)(value_type* a, integer n) =
      \exists integer i; 0 <= i < n-1 && a[i] == a[i+1];
*/
```

Listing 3.17: The predicate HasEqualNeighbors

```
/*@
  requires \valid_read(a + (0..n-1));

  assigns \nothing;

  behavior some:
    assumes HasEqualNeighbors(a, n);
    ensures 0 <= \result < n-1;
    ensures a[\result] == a[\result+1];
    ensures !HasEqualNeighbors(a, \result);

  behavior none:
    assumes !HasEqualNeighbors(a, n);
    ensures \result == n;

  complete behaviors;
  disjoint behaviors;
*/
size_type adjacent_find(const value_type* a, size_type n);
```

Listing 3.18: Formal specification of adjacent_find

We use the predicate HasEqualNeighbors to define the formal specification of adjacent_find (see Listing 3.18).

²⁴ See http://www.sgi.com/tech/stl/adjacent_find.html

3.6.2. Implementation of adjacent_find

The implementation of `adjacent_find`, including loop (in)variants is shown in Listing 3.19. Please note the use of the predicate `HasEqualNeighbors` in the loop invariant to match the similar postcondition of behavior `some`.

```
size_type
adjacent_find(const value_type* a, size_type n)
{
    if (n == 0) {
        return n;
    }

    /*@
        loop invariant 0 <= i < n;
        loop invariant !HasEqualNeighbors(a, i+1);
        loop assigns i;
        loop variant n-i;
    */
    for (size_type i = 0; i < n - 1; i++) {
        if (a[i] == a[i + 1]) {
            return i;
        }
    }

    return n;
}
```

Listing 3.19: Implementation of `adjacent_find`

3.7. The search algorithm

The `search` algorithm in the C++ standard library finds a subsequence that is identical to a given sequence when compared element-by-element. For our purposes we have modified the generic implementation to that of an array of type `value_type`.²⁵ The signature now reads:

```
size_type search(const value_type* a, size_type m,
                 const value_type* b, size_type n)
```

The function `search` returns the first index `i` of the array `a` where the condition `a[i+k]==b[k]` for each $0 \leq k < n$ holds (see Figure 3.20). If no such index exists then `search` returns the length `m` of the array `a`.




Figure 3.20.: Matching `b[0..n-1]` in `a[0..m-1]`

3.7.1. Formal specification of `search`

Our specification of `search` starts with introducing the predicate `HasSubRange` in Listing 3.21. This predicate formalizes the fact that the sequence `a` contains a subsequence which is identical to the sequence `b`.

```
/*@
  predicate
    HasSubRange{A} (value_type* a, integer m,
                    value_type* b, integer n) =
      \exists size_type k;
        (0 <= k <= m-n) && EqualRanges{A,A}(a+k, n, b);
*/
```

Listing 3.21: The predicate `HasSubRange`

²⁵ See <http://www.sgi.com/tech/stl/search.html>

The ACSL specification of `search` is shown in Listing 3.22. The behavior `has_match` applies if the sequence `a` contains a subsequence, which is identical to the sequence `b`. We express this condition with `assumes` by using the predicate `HasSubRange`.

```
/*@
  requires \valid_read(a + (0..m-1));
  requires \valid_read(b + (0..n-1));

  assigns \nothing;

  ensures (n == 0 || m == 0) ==> \result == 0;

  behavior has_match:
    assumes HasSubRange(a, m, b, n);
    ensures 0 <= \result <= m-n;
    ensures EqualRanges{Here,Here}(a+\result, n, b);
    ensures !HasSubRange(a, \result+n-1, b, n);

  behavior no_match:
    assumes !HasSubRange(a, m, b, n);
    ensures \result == m;

  complete behaviors;
  disjoint behaviors;
*/
size_type search(const value_type* a, size_type m,
                 const value_type* b, size_type n);
```

Listing 3.22: Formal specification of `search`

The first `ensures` clause of behavior `has_match` indicates that the return value must be in the range $[0 \dots m-n]$. The second one expresses that `search` returns the smallest index where `b` can be found in `a`. Finally, in the last line under behavior `has_match` we indicate that the sequence `a` contains a subsequence (from the position `\result`), which is identical to the sequence `b`.

The behavior `no_match` covers the case that there is no such subsequence in sequence `a`, which equals to the sequence `b`. In this case, `search` must return the length `m` of the range `a`. In any case, if the ranges `a` or `b` are empty, then the return value will be 0. We express this fact with the following line:

```
ensures (n == 0 || m == 0) ==> \result == 0;
```

The formula in the `assumes` clause of the behavior `has_match` is the negation of the `assumes` clause of the behavior `no_match`. Therefore, we can express that these two behaviors are *complete* and *disjoint*.

3.7.2. Implementation of search

Our implementation of `search` is shown in Listing 3.23. It follows the C++ standard library implementation in being easy to understand, but needing an order of magnitude of $m \cdot n$ operations. In contrast, the sophisticated algorithm from [13] needs only $m+n$ operations.²⁶

```
size_type search(const value_type* a, size_type m,
                 const value_type* b, size_type n)
{
    if ((n == 0) || (m == 0)) {
        return 0;
    }

    if (n > m) {
        return m;
    }

    /*@
    loop invariant 0 <= i <= m-n+1;
    loop invariant !HasSubRange(a, n+i-1, b, n);
    loop assigns i;
    loop variant m-i;
    */
    for (size_type i = 0; i <= m - n; i++) {
        if (equal(a + i, n, b)) { // Is there a match?
            return i;
        }
    }

    return m;
}
```

Listing 3.23: Implementation of `search`

The second loop *invariant* is needed for the proof of the postconditions of the behavior `has_match` (see Listing 3.22). It expresses that for each iteration the subsequence, which equals to the sequence `b`, is not yet found up to that iteration step.

²⁶ This question has been also discussed by the C++ standardization committee, see <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n3905.html>

3.8. The count algorithm

The `count` algorithm in the C++ standard library counts the frequency of occurrences for a particular element in a sequence. For our purposes we have modified the generic implementation²⁷ to that of arrays of type `value_type`. The signature now reads:

```
size_type count(const value_type* a, size_type n, value_type val);
```

Informally, the function returns the number of occurrences of `val` in the array `a`.

3.8.1. An axiomatic definition of counting

The specification of `count` will be fairly short because it employs the *logic function* `Count` whose (considerably) longer definition is given in Listing 3.24.²⁸ We will reuse this axiomatic definition of counting for the specification of other algorithms, e.g., `remove_copy` (Section 6.11).

```
/*@
axiomatic CountAxiomatic
{
  logic integer Count{L}(value_type* a, integer n, value_type v)
    reads a[0..n-1];

  axiom CountEmpty:
    \forall value_type *a, v, integer n;
      n <= 0 ==> Count(a, n, v) == 0;

  axiom CountOneHit:
    \forall value_type *a, v, integer n;
      a[n] == v ==> Count(a, n + 1, v) == Count(a, n, v) + 1;

  axiom CountOneMiss:
    \forall value_type *a, v, integer n;
      a[n] != v ==> Count(a, n + 1, v) == Count(a, n, v);

  axiom CountRead{L1,L2}:
    \forall value_type *a, v, integer n;
      Unchanged{L1,L2}(a, n) ==>
        Count{L1}(a, n, v) == Count{L2}(a, n, v);
}
*/
```

Listing 3.24: The logic function `Count`

²⁷ See <http://www.sgi.com/tech/stl/count.html>

²⁸ This definition of `Count` is a generalization of the *logic function* `nb_occ` of the ACSL specification [9].

The logic function `Count` in Listing 3.24 determines the number of occurrences of a value v in the index range $[0..n-1]$ of an array of type `value_type`.

- The ACSL keyword `axiomatic` is used to gather the logic function `Count` and its defining axioms. Note that the interval bound n and the return value for `Count` are of type `integer`.
- Axiom `CountEmpty` covers the case of an empty range.
- Axioms `CountOneHit` and `CountOneMiss` reduce counting of a range of length $n + 1$ to a range of length n .
- The `reads` clause in the axiomatic definition of `Count` specifies the set of memory locations on which `Count` depends.

Axiom `CountRead` makes this claim explicit by ensuring that `Count` produces the same result if the values $a[0..n-1]$ do not change between two program states indicated by the labels `L1` and `L2`. We use predicate `Unchanged` (Listing 6.1 in Section 6.1) to express the premise of Axiom `CountRead`.

Axiom `CountRead` is helpful if one has to verify *mutating* algorithms that are specified with `Count`, e.g., `remove_copy` in Sections 6.11 and 6.12.

The following properties of `Count` can be verified with the help of the axioms given in Listing 3.24.

```
/*@
lemma CountOne:
  \forall value_type *a, v, integer n;
    Count(a, n + 1, v) == Count(a, n, v) + Count(a + n, 1, v);

lemma CountUnion:
  \forall value_type *a, v, integer n, k;
    0 <= k ==> Count(a, n + k, v) == Count(a, n, v) + Count(a + n, k, v);

lemma CountBounds:
  \forall value_type *a, v, integer n;
    0 <= n ==> 0 <= Count(a, n, v) <= n;

lemma CountMonotonic:
  \forall value_type *a, v, integer m, n;
    0 <= m <= n ==> Count(a, m, v) <= Count(a, n, v);
*/
```

Listing 3.25: Some lemmas for `Count`

3.8.2. Formal specification of count

Listing 3.26 shows how we use the logic function `count` to specify `count` in ACSL. Note that our specification also states that the result of `count` is non-negative and less than or equal the size of the array.

```
/*@
  requires \valid_read(a + (0..n-1));

  assigns \nothing;

  ensures \result == Count(a, n, val);
  ensures 0 <= \result <= n;
*/
size_type count(const value_type* a, size_type n, value_type val);
```

Listing 3.26: Formal specification of `count`

3.8.3. Implementation of count

Listing 3.27 shows a possible implementation of `count`. Note that we refer to the logic function `Count` in one of the loop invariants.

```
size_type
count(const value_type* a, size_type n, value_type val)
{
  size_type counted = 0;

  /*@
    loop invariant 0 <= i <= n;
    loop invariant 0 <= counted <= i;
    loop invariant counted == Count(a, i, val);
    loop assigns i, counted;
    loop variant n-i;
  */
  for (size_type i = 0; i < n; ++i) {
    if (a[i] == val) {
      counted++;
    }
  }

  return counted;
}
```

Listing 3.27: Implementation of `count`

4. Maximum and minimum algorithms

In this chapter we discuss the formal specification of algorithms that compute the maximum or minimum values of their arguments. As the algorithms in Chapter 3, they also do not modify any memory locations outside their scope. The most important new feature of the algorithms in this chapter is that they compare values using binary operators such as $<$.

We consider in this chapter the following algorithms.

- `max_element` (Section 4.2, on Page 52) returns an index to a maximum element in range. Similar to `find` it also returns the smallest of all possible indices. An alternative specification which relies on user-defined predicates will be introduced in Section 4.3, on Page 54).
- `max_seq` (Section 4.4, on Page 56) is very similar to `max_element` and will serve as an example of *modular verification*. It returns the maximum value itself rather than an index to it.
- `min_element` which can be used to find the smallest element in an array (Section 4.5).

First, however, we discuss in Section 4.1 general properties that must be satisfied by the relational operators.

4.1. A note on relational operators

Note that in order to compare values, the algorithms in the C++ standard library usually rely solely on the *less than* operator `<` or special function objects.²⁹ To be precise, the operator `<` must be a *partial order*,³⁰ which means that the following rules hold.

$$\begin{array}{lll} \text{irreflexivity} & \forall x & : \neg(x < x) \\ \text{asymmetry} & \forall x, y & : x < y \implies \neg(y < x) \\ \text{transitivity} & \forall x, y, z & : x < y \wedge y < z \implies x < z \end{array}$$

If you wish check that the operator `<` of our `value_type`³¹ satisfies this properties you can formulate lemmas in ACSL and verify them with Frama-C. (see Listing 4.1).

```
/*@
lemma LessIrreflexivity:
  \forallall value_type a; !(a < a);

lemma LessAntisymmetry:
  \forallall value_type a, b; (a < b) ==> !(b < a);

lemma LessTransitivity:
  \forallall value_type a, b, c; (a < b) && (b < c) ==> (a < c);
*/
```

Listing 4.1: Requirements for a partial order on `value_type`

It is of course possible to specify and implement the algorithms of this chapter by only using operator `<`. For example, `a <= b` can be written as `a < b | a == b`, or, for our particular ordering on `value_type`, as `!(b < a)`. However, for the purpose of this introductory document we have opted for a more user friendly representation.

Listing 4.2 formulates condition on the semantics of the derived operator `>`, `<=` and `>=`.

```
/*@
lemma Greater:
  \forallall value_type a, b; (a > b) <==> (b < a);

lemma LessOrEqual:
  \forallall value_type a, b; (a <= b) <==> !(b < a);

lemma GreaterOrEqual:
  \forallall value_type a, b; (a >= b) <==> !(a < b);
*/
```

Listing 4.2: Semantics of derived comparison operators

²⁹ See <http://www.sgi.com/tech/stl/LessThanComparable.html>.

³⁰ See http://en.wikipedia.org/wiki/Partially_ordered_set

³¹ See Section 1.2

We also provide a group of predicates that concisely express the comparison of the elements in an array segment with a given value (see Listing 4.3). We will use these predicates both in this chapter and in Chapter binary-search.

```
/*@
predicate ConstantRange(value_type* a, integer first,
                        integer last, value_type val) =
    \forall integer i; first <= i < last ==> a[i] == val;

predicate StrictLowerBound(value_type* a, integer first,
                           integer last, value_type val) =
    \forall integer i; first <= i < last ==> val < a[i];

predicate LowerBound(value_type* a, integer first,
                     integer last, value_type val) =
    \forall integer i; first <= i < last ==> !(a[i] < val);

predicate StrictUpperBound(value_type* a, integer first,
                            integer last, value_type val) =
    \forall integer i; first <= i < last ==> a[i] < val;

predicate UpperBound(value_type* a, integer first,
                     integer last, value_type val) =
    \forall integer i; first <= i < last ==> !(val < a[i]);
*/
```

Listing 4.3: Predicates for comparing array elements with a given value

4.2. The max_element algorithm

The `max_element` algorithm in the C++ Standard Template Library³² searches the maximum of a general sequence. The signature of our version of `max_element` reads:

```
size_type max_element(const value_type* a, size_type n);
```

The function finds the largest element in the range $a[0, n)$. More precisely, it returns the unique valid index i such that

1. for each index k with $0 \leq k < n$ the condition $a[k] \leq a[i]$ holds and
2. for each index k with $0 \leq k < i$ the condition $a[k] < a[i]$ holds.

The return value of `max_element` is n if and only if there is no maximum, which can only occur if $n == 0$.

4.2.1. Formal specification of max_element

A formal specification of `max_element` in ACSL is shown in Listing 4.4.

```
/*@
  requires  \valid_read(a + (..n-1));

  assigns  \nothing;

  behavior empty:
    assumes n == 0;
    ensures \result == 0;

  behavior not_empty:
    assumes 0 < n;

    ensures 0 <= \result < n;
    ensures \forall integer i; 0 <= i < n ==> a[i] <= a[\result];
    ensures \forall integer i; 0 <= i < \result ==> a[i] < a[\result];

  complete behaviors;
  disjoint behaviors;
*/
size_type max_element(const value_type* a, size_type n);
```

Listing 4.4: Formal specification of `max_element`

We have subdivided the specification of `max_element` into two behaviors (`empty` and `not_empty`). The behavior `empty` contains the specification for the case that the range contains no elements. The behavior `not_empty` applies if the range has a positive length.

The second `ensures` clause of behavior `not_empty` indicates that the returned valid index k refers to a maximum value of the array. The third one expresses that k is indeed the *first* occurrence of a maximum value in the array.

³² See http://www.sgi.com/tech/stl/max_element.html

4.2.2. Implementation of `max_element`

Listing 4.5 shows an implementation of `max_element`. In our description, we concentrate on the *loop annotations*.

```
size_type max_element(const value_type* a, size_type n)
{
    if (n == 0) {
        return 0;
    }

    size_type max = 0;

    /*@
    loop invariant 0 <= i <= n;
    loop invariant 0 <= max < n;
    loop invariant \forall integer k; 0 <= k < i ==> a[k] <= a[max];
    loop invariant \forall integer k; 0 <= k < max ==> a[k] < a[max];
    loop assigns max, i;
    loop variant n-i;
    */
    for (size_type i = 1; i < n; i++) {
        if (a[max] < a[i]) {
            max = i;
        }
    }

    return max;
}
```

Listing 4.5: Implementation of `max_element`

The second loop invariant is needed to prove the first postcondition of behavior `not_empty` in Listing 4.4. Using the next loop invariant we prove the second postcondition of behavior `not_empty` in Listing 4.4. Finally, the last postcondition of this behavior can be proved with the endmost loop *invariant*.

4.3. The `max_element` algorithm with predicates

In this section we present another specification of the `max_element` algorithm. The main difference is that we employ two user defined predicates. First we define the predicate `MaxElement` by using the previously introduced predicate `UpperBound` (Listing 4.3) by stating that it is an upper bound that belongs to the sequence $a[0..n - 1]$.

```
/*@  
  predicate MaxElement{L}(value_type* a, integer n, integer max) =  
    0 <= max < n && UpperBound(a, 0, n, a[max]);  
*/
```

Listing 4.6: Definition of the `MaxElement` predicate

4.3.1. Formal specification of `max_element`

The new formal specification of `max_element` in ACSL is shown in Listing 4.7. Note that we also use the predicate `StrictUpperBound` (Listing 4.3) in order to express that `max_element` returns the *first* maximum position in $[0..n - 1]$.

```
/*@  
  requires  \valid_read(a + (0..n-1));  
  
  assigns  \nothing;  
  
  behavior empty:  
    assumes n == 0;  
  
    ensures result:  \result == 0;  
  
  behavior not_empty:  
    assumes 0 < n;  
  
    ensures result:  0 <= \result < n;  
    ensures maximum: MaxElement(a, n, \result);  
    ensures first:   StrictUpperBound(a, 0, \result, a[\result]);  
  
  complete behaviors;  
  disjoint behaviors;  
*/  
size_type max_element(const value_type* a, size_type n);
```

Listing 4.7: Formal specification of `max_element`

4.3.2. Implementation of `max_element`

Listing 4.8 shows implementation of `max_element` with rewritten loop invariants. In the loop invariants we also employ the predicates `UpperBound` and `StrictUpperBound` that we have used in the specification.

```
size_type max_element(const value_type* a, size_type n)
{
    if (n == 0) {
        return 0;
    }

    size_type max = 0;

    /*@
    loop invariant bound:  0 <= i <= n;
    loop invariant min:    0 <= max < n;
    loop invariant lower:  UpperBound(a, 0, i, a[max]);
    loop invariant first:  StrictUpperBound(a, 0, max, a[max]);

    loop assigns max, i;
    loop variant n-i;
    */
    for (size_type i = 0; i < n; i++) {
        if (a[max] < a[i]) {
            max = i;
        }
    }

    return max;
}
```

Listing 4.8: Implementation of `max_element`

4.4. The `max_seq` algorithm

In this section we consider the function `max_seq` (see Chapter 3, [8]) that is very similar to the `max_element` function of Section 4.2. The main difference between `max_seq` and `max_element` is that `max_seq` returns the maximum value (not just the index of it). Therefore, it requires a *non-empty* range as an argument.

Of course, `max_seq` can easily be implemented using `max_element` (see Listing 4.10). Moreover, using only the formal specification of `max_element` in Listing 4.7 we are also able to deductively verify the correctness of this implementation. Thus, we have a simple example of *modular verification* in the following sense:

Any implementation of `max_element` that is separately proven to implement the contract in Listing 4.7 makes `max_seq` behave correctly. Once the contracts have been defined, the function `max_element` could be implemented in parallel, or just after `max_seq`, without affecting the verification of `max_seq`.

4.4.1. Formal specification of `max_seq`

A formal specification of `max_seq` in ACSL is shown in Listing 4.9.

```
/*@
  requires n > 0;
  requires \valid_read(p + (0..n-1));

  assigns \nothing;

  ensures \forall integer i; 0 <= i <= n-1 ==> \result >= p[i];
  ensures \exists integer e; 0 <= e <= n-1 && \result == p[e];
*/
value_type max_seq(const value_type* p, size_type n);
```

Listing 4.9: Formal specification of `max_seq`

Using the first `requires`-clause we express that `max_seq` needs a *non-empty* range as input. By using the `ensures`-clause we express our postconditions. They formalize that `max_seq` indeed returns the maximum value of the range.

4.4.2. Implementation of `max_seq`

Listing 4.10 shows the trivial implementation of `max_seq` using `max_element`. Since `max_seq` requires a non-empty range the call of `max_element` returns an index to a maximum value in the range. The fact that `max_element` returns the smallest index is of no importance in this context.

```
value_type max_seq(const value_type* p, size_type n)
{
    return p[max_element(p, n)];
}
```

Listing 4.10: Implementation of `max_seq`

4.5. The min_element algorithm

The `min_element` algorithm in the C++ standard library³³ searches the minimum in a general sequence. The signature of our version of `min_element` reads:

```
size_type min_element(const value_type* a, size_type n);
```

The function `min_element` finds the smallest element in the range $a[0..n-1]$. More precisely, it returns the unique valid index i such that The return value of `min_element` is n if and only if $n == 0$. First we define the predicate `MinElement` by using the previously introduced predicate `LowerBound` (Listing 4.3) by stating that it is an lower bound that belongs to the sequence $a[0..n-1]$.

```
/*@
  predicate MinElement{L}(value_type* a, integer n, integer min) =
    0 <= min < n && LowerBound(a, 0, n, a[min]);
*/
```

Listing 4.11: Definition of the `MinElement` predicate

4.5.1. Formal specification of min_element

```
/*@
  requires \valid_read(a + (0..n-1));

  assigns \nothing;

  behavior empty:
    assumes n == 0;

    ensures result: \result == 0;

  behavior not_empty:
    assumes 0 < n;

    ensures result: 0 <= \result < n;
    ensures minimum: MinElement(a, n, \result);
    ensures first: StrictLowerBound(a, 0, \result, a[\result]);

  complete behaviors;
  disjoint behaviors;
*/
size_type min_element(const value_type* a, size_type n);
```

Listing 4.12: Formal specification of `min_element`

The ACSL specification of `min_element` is shown in Listing 4.12. Note that we also use the predicate `StrictLowerBound` (Listing 4.3) in order to express that `min_element` returns the *first* minimum position in $[0..n-1]$.

³³ See http://www.sgi.com/tech/stl/min_element.html

4.5.2. Implementation of min_element

Listing 4.13 shows implementation of min_element with rewritten loop invariants. In the loop invariants we also employ the predicates LowerBound and StrictLowerBound that we have used in the specification.

```
size_type min_element(const value_type* a, size_type n)
{
    if (0 == n) {
        return n;
    }

    size_type min = 0;

    /*@
    loop invariant bound:    0 <= i    <= n;
    loop invariant min:      0 <= min <  n;
    loop invariant lower:    LowerBound(a, 0, i, a[min]);
    loop invariant first:    StrictLowerBound(a, 0, min, a[min]);

    loop assigns min, i;
    loop variant n-i;
    */
    for (size_type i = 0; i < n; i++) {
        if (a[i] < a[min]) {
            min = i;
        }
    }

    return min;
}
```

Listing 4.13: Implementation of min_element

5. Binary search algorithms

In this chapter, we consider the four *binary search* algorithms of the C++ standard library, namely

- `lower_bound` in Section 5.1
- `upper_bound` in Section 5.2
- `equal_range` in Section 5.3
- `binary_search` in Section 5.4

All binary search algorithms require that their input array is sorted in ascending order. The predicate `Sorted` in Listing 5.1 formalizes these requirements.

```
/*@  
  predicate  
  Sorted{L}(value_type* a, integer n) =  
    \forall i, j; 0 <= i < j < n ==> a[i] <= a[j];  
*/
```

Listing 5.1: The predicate `Sorted`

As in the case of the of maximum/minimum algorithms from Chapter 4 the binary search algorithms primarily use the less-than operator `<` (and the derived operators `<=`, `>` and `>=`) to determine whether a particular value is contained in a sorted range. Thus, different to the `find` algorithm in Section 3.3, the equality operator `==` will play only a supporting part in the specification of binary search.

In order to make the specifications of the binary search algorithms more compact and (arguably) more readable we use the predicates from Listing 4.3.

5.1. The `lower_bound` algorithm

The `lower_bound` algorithm is one of the four binary search algorithms of the C++ standard library. For our purposes we have modified the generic implementation³⁴ to that of an array of type `value_type`. The signature now reads:

```
size_type lower_bound(const value_type* a, size_type n,
                     value_type val);
```

As with the other binary search algorithms `lower_bound` requires that its input array is sorted in ascending order. Specifically, `lower_bound` will return the *largest* index i with $0 \leq i \leq n$ such that for each index k with $0 \leq k < i$ the condition $a[k] < \text{val}$ holds. This specification makes `lower_bound` a bit tricky to use as a search algorithm:

- If `lower_bound` returns n then for each index i with $0 \leq i < n$ holds $a[i] < \text{val}$. Thus, val is not contained in a .
- If, however, `lower_bound` returns an index r with $0 \leq r < n$ then we can only be sure that $a[i] < \text{val}$ holds for $0 \leq i < r$ and that $\text{val} \leq a[i]$ holds for $r \leq i < n$.

5.1.1. Formal specification of `lower_bound`

The ACSL specification of `lower_bound` is shown in Listing 5.2.

```
/*@
requires \valid_read(a + (0..n-1));
requires Sorted(a, n);

assigns \nothing;

ensures result: 0 <= \result <= n;
ensures left:   StrictUpperBound(a, 0, \result, val);
ensures right:  LowerBound(a, \result, n, val);
*/
size_type
lower_bound(const value_type* a, size_type n, value_type val);
```

Listing 5.2: Formal specification of `lower_bound`

- The preconditions express, by using the predicate `Sorted`, that the values in the (valid) array need to be sorted in ascending order.
- The postconditions formalize the central properties, mentioned above, of the return value of `lower_bound`.

³⁴ See http://www.sgi.com/tech/stl/lower_bound.html

5.1.2. Implementation of `lower_bound`

Our implementation of `lower_bound` is shown in Listing 5.3. Each iteration step narrows down the range that contains the sought-after result. The loop invariants express that in each iteration step all indices less than the temporary left bound `left` contain values smaller than `val` and all indices not less than the temporary right bound `right` contain values not smaller than `val`.

```
size_type
lower_bound(const value_type* a, size_type n, value_type val)
{
    size_type left = 0;
    size_type right = n;
    size_type middle = 0;

    /*@
    loop invariant bound:  0 <= left <= right <= n;
    loop invariant left:   StrictUpperBound(a, 0, left, val);
    loop invariant right:  LowerBound(a, right, n, val);

    loop assigns middle, left, right;
    loop variant right - left;
    */
    while (left < right) {
        middle = left + (right - left) / 2;

        if (a[middle] < val) {
            left = middle + 1;
        } else {
            right = middle;
        }
    }

    return left;
}
```

Listing 5.3: Implementation of `lower_bound`

5.2. The upper_bound algorithm

The `upper_bound`³⁵ algorithm is a version of the `binary_search` algorithm closely related to `lower_bound` of Section 5.1.

The signature reads:

```
size_type upper_bound(const value_type* a, size_type n,
                     value_type val)
```

In contrast to the `lower_bound` algorithm the `upper_bound` algorithm locates the *largest* index i with $0 \leq i \leq n$ such that for each index k with $0 \leq k < i$ the condition $a[k] \leq \text{val}$ holds. This means:

- If `upper_bound` returns n then we can only be sure that for each index $0 \leq i < n$ the relationship $a[i] \leq \text{val}$.
- If `upper_bound` returns an index r with $0 \leq r < n$ then we can be sure that $\text{val} < a[i]$ holds for i where $r \leq i < n$. Thus, if `upper_bound` returns 0 then we know that val is not contained in a .

5.2.1. Formal specification of upper_bound

The ACSL specification of `upper_bound` is shown in Listing 5.4.

```
/*@
  requires \valid_read(a + (0..n-1));
  requires Sorted(a, n);

  assigns \nothing;

  ensures result: 0 <= \result <= n;
  ensures left:   UpperBound(a, 0, \result, val);
  ensures right:  StrictLowerBound(a, \result, n, val);
*/
size_type
upper_bound(const value_type* a, size_type n, value_type val);
```

Listing 5.4: Formal specification of `upper_bound`

The specification is quite similar to the specification of `lower_bound` (see Listing 5.2). The difference can be seen in the postconditions. As we are searching for the upper bound this time, `upper_bound` has to ensure that

- all indices less than the returned one belong to elements are less than or equal to val
- all indices greater than or equal to the returned one belong to elements that are greater than val .

³⁵ See http://www.sgi.com/tech/stl/upper_bound.html

5.2.2. Implementation of upper_bound

Our implementation of `upper_bound` is shown in Listing 5.5.

The loop invariants express that for each iteration step all indices less than the temporary left bound `left` contain values not greater than `val` and all indices not less than the temporary right bound `right` contain values greater than `val`.

```
size_type
upper_bound(const value_type* a, size_type n, value_type val)
{
    size_type left = 0;
    size_type right = n;
    size_type middle = 0;

    /*@
    loop invariant bound:  0 <= left <= right <= n;
    loop invariant left:   UpperBound(a, 0, left, val);
    loop invariant right:  StrictLowerBound(a, right, n, val);

    loop assigns middle, left, right;
    loop variant right - left;
    */
    while (left < right) {
        middle = left + (right - left) / 2;

        if (a[middle] <= val) {
            left = middle + 1;
        } else {
            right = middle;
        }
    }

    return right;
}
```

Listing 5.5: Implementation of `upper_bound`

5.3. The `equal_range` algorithm

The `equal_range` algorithm is one of the four binary search algorithms of the C++ standard library. For our purposes we have modified the generic implementation³⁶ to that of an array of type `value_type`. Moreover, instead of a pair of iterators, our version of `equal_range` returns a pair of indices. To be more precisely, the return type of `equal_range` is the struct `size_type_pair` from Listing 5.6. Thus, the signature of `equal_range` now reads:

```
size_type_pair equal_range(const value_type* a, size_type n,
                           value_type val);
```

As with the other binary search algorithms `equal_range` requires that its input array is sorted in ascending order. The specification of `equal_range` states that it *combines* the results of the algorithms `lower_bound` (Section 5.1) and `upper_bound` (Section 5.2).

5.3.1. Formal specification of `equal_range`

The ACSL specification of `equal_range` is shown in Listing 5.6.

```
struct spair {
    size_type first;
    size_type second;
};

typedef struct spair size_type_pair;

/*@
    requires \valid_read(a + (0..n-1));
    requires Sorted(a, n);

    assigns \nothing;

    ensures result: 0 <= \result.first <= \result.second <= n;
    ensures left:   StrictUpperBound(a, 0, \result.first, val);
    ensures middle: ConstantRange(a, \result.first,
                                   \result.second, val);
    ensures right:  StrictLowerBound(a, \result.second, n, val);
*/
size_type_pair
equal_range(const value_type* a, size_type n, value_type val);
```

Listing 5.6: Formal specification of `equal_range`

The preconditions express that the values in the (valid) array need to be sorted in ascending order.

The postconditions express that the pair of indices (f, s) returned by `equal_range` satisfy the following properties:

- $0 \leq f \leq s \leq n$
- the set of indices $[f, s) = \{i \mid f \leq i < s\}$ is the *largest* set for which $a[i] = \text{val}$ holds

³⁶ See http://www.sgi.com/tech/stl/equal_range.html

5.3.2. Implementation of `equal_range`

Our implementation of `equal_range` is shown in Listing 5.7. We call the two functions `lower_bound` and `upper_bound` and return their respective results as a pair. However, instead of doing this straightforward, we use the auxiliary function `make_pair`³⁷ and formulate an assertion for its arguments `first` \leq `second`. Using this assertion simplifies the task of *automatically* proving the postcondition in Listing 5.6.

```
/*@
  assigns \nothing;

  ensures \result.first == first;
  ensures \result.second == second;
*/
size_type_pair make_pair(size_type first, size_type second)
{
  size_type_pair pair;
  pair.first = first;
  pair.second = second;

  return pair;
}

size_type_pair
equal_range(const value_type* a, size_type n, value_type val)
{
  size_type first = lower_bound(a, n, val);
  size_type second = upper_bound(a, n, val);
  //@ assert aux: second < n ==> val < a[second];

  return make_pair(first, second);
}
```

Listing 5.7: Implementation of `equal_range`

In an earlier version of this document we had proven the similar assertion `first` \leq `second` with the interactive theorem prover Coq. After reviewing this proof we formulated the new assertion `aux` that uses a fact from the postcondition of `upper_bound` (Listing 5.4). The benefit of this reformulation is that both the assertion `aux` and the postcondition `first` \leq `second` can now be verified automatically.

³⁷ This function is modelled after the C++ template function `std::make_pair`.

5.4. The `binary_search` algorithm

The `binary_search` algorithm is one of the four binary search algorithms of the C++ standard library. For our purposes we have modified the generic implementation³⁸ to that of an array of type `value_type`. The signature now reads:

```
bool binary_search(const value_type* a, size_type n,
                  value_type val);
```

Again, `binary_search` requires that its input array is sorted in ascending order. It will return `true` if there exists an index `i` in `a` such that `a[i] == val` holds.³⁹

5.4.1. Formal specification of `binary_search`

The ACSL specification of `binary_search` is shown in Listing 5.8.

```
/*@
  requires \valid_read(a + (0..n-1));
  requires Sorted(a, n);

  assigns \nothing;

  ensures result: \result <==>
                  \exists integer i; 0 <= i < n && a[i] == val;
*/
bool binary_search(const value_type* a, size_type n, value_type val);
```

Listing 5.8: Formal specification of `binary_search`

Note that we can use our previously introduced predicate `HasValue` (see Page 36) in Listing 5.9. It is interesting to compare this specification with that of `find` in Listing 3.12. Both `find` and `binary_search` allow to determine whether a value is contained in an array. The fact that the C++ standard library requires that `find` has *linear* complexity whereas `binary_search` must have a *logarithmic* complexity can currently not be expressed with ACSL.

```
/*@
  requires \valid_read(a + (0..n-1));
  requires Sorted(a, n);

  assigns \nothing;

  ensures result: \result <==> HasValue(a, n, val);
*/
bool binary_search(const value_type* a, size_type n, value_type val);
```

Listing 5.9: Formal specification of `binary_search` using the `HasValue` predicate

³⁸ See http://www.sgi.com/tech/stl/binary_search.html

³⁹ To be more precise: The C++ standard library requires that $(a[i] \leq val) \wedge (val \leq a[i])$ holds. For our definition of `value_type` (see Section 1.2) this means that `val` equals `a[i]`.

5.4.2. Implementation of `binary_search`

Our implementation of `binary_search` is shown in Listing 5.10.

```
bool binary_search(const value_type* a, size_type n, value_type val)
{
    size_type i = lower_bound(a, n, val);
    return i < n && a[i] <= val;
}
```

Listing 5.10: Implementation of `binary_search`

The function `binary_search` first calls `lower_bound` from Section 5.1. Remember that if `lower_bound` returns an index $0 \leq i < n$ then we can be sure that $\text{val} \leq a[i]$ holds.

6. Mutating algorithms

Let us now turn our attention to another class of algorithms, viz. *mutating* algorithms, i.e., algorithms that change one or more ranges. In Frama-C, you can explicitly specify that, e.g., entries in an array `a` may be modified by a function `f`, by including the following *assigns clause* into the `f`'s specification:

```
assigns a[0..length-1];
```

The expression `length-1` refers to the value of `length` when `f` is entered, see [9, Section 2.3.2]. Below are the algorithms we will discuss in this chapter. First, however, we introduce in Section 6.1 the auxiliary predicate `Unchanged`.

- `swap` in Section 6.2 exchanges two values.
- `fill` in Section 6.3 initializes each element of an array by a given fixed value.
- `swap_ranges` in Section 6.4 exchanges the contents of the arrays of equal length, element by element. We use this example to present “modular verification”, as `swap_ranges` reuses the verified properties of `swap`.
- `copy` in Section 6.5 copies a source array to a destination array.
- `reverse_copy` and `reverse` in Sections 6.6 and 6.7, respectively, reverse an array. Whereas `reverse_copy` copies the result to a separate destination array, the `reverse` algorithm works in place.
- `rotate_copy` in Section 6.8 rotates a source array by `m` positions and copies the results to a destination array.
- `replace_copy` and `replace` in Sections 6.9 and 6.10, respectively, substitute each occurrence of a value by a given new value. Whereas `replace_copy` copies the result to a separate array, the `replace` algorithm works in place.
- `remove_copy` copies a source array to a destination array, but omits each occurrence of a given value. We provide two specifications for `remove_copy`:
 - In Section 6.11 we provide a relatively simple contract that omits, however, an important aspect of the informal specification
 - In Section 6.12 we show how the missing part of the specification can be expressed.

6.1. The predicate Unchanged

Many of the algorithms in this section iterate sequentially over one or several sequences. For the verification of such algorithms it is often important to express that a section of an array, or the complete array, have remained *unchanged*. In Listing 6.1 we therefore introduce the overloaded predicate `Unchanged` together with some simple lemmas. The expression `Unchanged{K,L}(a, f, l)` is true if the memory area of `a[f..l-1]` agrees in states `K` and `L`.

```
/*@
  predicate
    Unchanged{K,L}(value_type* a, integer first, integer last) =
      \forall integer i; first <= i < last ==>
        \at(a[i],K) == \at(a[i],L);

  predicate
    Unchanged{K,L}(value_type* a, integer n) =
      Unchanged{K,L}(a, 0, n);

  lemma
    UnchangedStep{K,L}:
      \forall value_type *a, integer n;
        0 <= n ==>
          Unchanged{K,L}(a, n) ==>
            \at(a[n],K) == \at(a[n],L) ==>
              Unchanged{K,L}(a, n+1);

  lemma
    UnchangedSection{K,L}:
      \forall value_type *a, integer m, n;
        0 <= m <= n ==>
          Unchanged{K,L}(a, n) ==>
            Unchanged{K,L}(a, m);
*/
```

Listing 6.1: Predicate Unchanged

6.2. The swap algorithm

The swap algorithm⁴⁰ in the C++ standard library exchanges the contents of two variables. Similarly, the `iter_swap` algorithm⁴¹ exchanges the contents referenced by two pointers. Since C and hence ACSL, does not support an `&` type constructor (“declarator”), we will present an algorithm that processes pointers and refer to it as `swap`.

⁴⁰ See <http://www.sgi.com/tech/stl/swap.html>

⁴¹ See http://www.sgi.com/tech/stl/iter_swap.html

6.2.1. Formal specification of swap

The ACSL specification for the `swap` function is shown in Listing 6.2. The preconditions are formalized by the `requires`-clauses which state that both pointer arguments of the `swap` function must be dereferenceable.

```
/*@
  requires \valid(p);
  requires \valid(q);

  assigns *p;
  assigns *q;

  ensures *p == \old(*q);
  ensures *q == \old(*p);
*/
void swap(value_type* p, value_type* q);
```

Listing 6.2: Formal specification of `swap`

The `assigns`-clauses formalize that the `swap` algorithm modifies only the entries referenced by the pointers `p` and `q`. Nothing else may be altered. In general, when more than one *assigns clause* appears in a function's specification, it is permitted to modify any of the referenced locations. However, if no *assigns clause* appears at all, the function is free to modify any memory location, see [9, Section 2.3.2]. To forbid a function to do any modifications outside its scope, a clause

```
assigns \nothing;
```

must be used, as we practised in the example specifications in Chapter 3.

Upon termination of `swap` the entries must be mutually exchanged. We can express those postconditions by using the `ensures`-clause. The expression `\old(*p)` refers to the pre-state of the function contract, whereas by default, a postcondition refers to the values after the function has been terminated.

6.2.2. Implementation of swap

Listing 6.3 shows the usual straight-forward implementation of `swap`. No interspersed ACSL is needed to get it verified by Frama-C.

```
void swap(value_type* p, value_type* q)
{
  value_type save = *p;
  *p = *q;
  *q = save;
}
```

Listing 6.3: Implementation of `swap`

6.3. The `fill` algorithm

The `fill` algorithm in the C++ Standard Library initializes general sequences with a particular value. For our purposes we have modified the generic implementation⁴² to that of an array of type `value_type`. The signature now reads:

```
void fill(value_type* a, size_type n, value_type val);
```

6.3.1. Formal specification of `fill`

Listing 6.4 shows the formal specification of `fill` in ACSL. We can express the postcondition of `fill` simply by using the predicate `ConstantRange` from Listing 4.3.

```
/*@
  requires valid: \valid(a + (0..n-1));

  assigns a[0..n-1];

  ensures constant: ConstantRange(a, 0, n, val);
*/
void fill(value_type* a, size_type n, value_type val);
```

Listing 6.4: Formal specification of `fill`

⁴² See <http://www.sgi.com/tech/stl/fill.html>

6.3.2. Implementation of `fill`

Listing 6.5 shows an implementation of `fill`.

```
void fill(value_type* a, size_type n, value_type val)
{
    /*@
      loop invariant bound:    0 <= i <= n;
      loop invariant constant: ConstantRange(a, 0, i, val);

      loop assigns i, a[0..n-1];
      loop variant n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        a[i] = val;
    }
}
```

Listing 6.5: Implementation of `fill`

The loop invariant `bound` is necessary to prove that each access to the range `a` occurs with valid indices. The loop invariant `constant` expresses that for each iteration the array is filled with the value of `val` up to the index `i` of the iteration. Note that we use here again the predicate `ConstantRange` from Listing 4.3.

6.4. The `swap_ranges` algorithm

The `swap_ranges` algorithm⁴³ in the C++ standard library exchanges the contents of two expressed ranges element-wise. After translating C++ reference types and iterators to C, our version of the original signature reads:

```
void swap_ranges(value_type* a, size_type n, value_type* b);
```

We do not return a value since it would equal `n`, anyway.

This function refers to the previously discussed algorithm `swap`. Thus, `swap_ranges` serves as another example for “modular verification”. The specification of `swap` will be automatically integrated into the proof of `swap_ranges`.

6.4.1. Formal specification of `swap_ranges`

Listing 6.6 shows an ACSL specification for the `swap_ranges` algorithm.

```
/*@
requires valid_a: \valid(a + (0..n-1));
requires valid_b: \valid(b + (0..n-1));
requires sep:      \separated(a+(0..n-1), b+(0..n-1));

assigns a[0..n-1];
assigns b[0..n-1];

ensures equal_a: EqualRanges{Here,Old}(a, n, b);
ensures equal_b: EqualRanges{Old,Here}(a, n, b);
*/
void swap_ranges(value_type* a, size_type n, value_type* b);
```

Listing 6.6: Formal specification of `swap_ranges`

The `swap_ranges` algorithm works correctly only if `a` and `b` do not overlap. Because of that fact we use the `separated`-clause to tell Frama-C that `a` and `b` must not overlap.

With the `assigns`-clause we postulate that the `swap_ranges` algorithm alters the elements contained in two distinct ranges, modifying the corresponding elements and nothing else.

The postconditions of `swap_ranges` specify that the content of each element in its post-state must equal the pre-state of its counterpart. We can use the predicate `EqualRanges` (see Listing 3.2) together with the label `Old` and `Here` to express the postcondition of `swap_ranges`. In our specification in Listing 6.6, for example, we specify that the array `a` in the memory state that corresponds to the label `Here` is equal to the array `b` at the label `Old`. Since we are specifying a postcondition `Here` refers to the post-state of `swap_ranges` whereas `Old` refers to the pre-state.

⁴³ See http://www.sgi.com/tech/stl/swap_ranges.html

6.4.2. Implementation of `swap_ranges`

Listing 6.7 shows an implementation of `swap_ranges` together with the necessary loop annotations.

```
void swap_ranges(value_type* a, size_type n, value_type* b)
{
    /*@
    loop invariant bound: 0 <= i <= n;
    loop invariant equal_a: EqualRanges{Here,Pre}(a, i, b);
    loop invariant equal_b: EqualRanges{Here,Pre}(b, i, a);

    loop invariant unchanged_a: Unchanged{Here,Pre}(a, i, n);
    loop invariant unchanged_b: Unchanged{Here,Pre}(b, i, n);

    loop assigns i, a[0..n-1], b[0..n-1];
    loop variant n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        swap(a + i, b + i);
    }
}
```

Listing 6.7: Implementation of `swap_ranges`

For the postcondition of the specification in Listing 6.6 to hold, our loop invariants must ensure that at each iteration all of the corresponding elements that have already been visited are swapped.

Note that there are two additional loop invariants which claim that all the elements that have not visited yet equal their original values. This is a workaround that allows us to prove the postconditions of `swap_ranges` despite the fact that the loop assigns is coarser than it should be. The predicate `Unchanged` from Listing 6.1 is used to express this property.

6.5. The copy algorithm

The `copy` algorithm in the C++ Standard Library implements a duplication algorithm for general sequences. For our purposes we have modified the generic implementation⁴⁴ to that of a range of type `value_type`. The signature now reads:

```
void copy(const value_type* a, size_type n, value_type* b);
```

Informally, the function copies every element from the source range `a` to the destination range `b`.

6.5.1. Formal specification of `copy`

The ACSL specification of `copy` is shown in Listing 6.8. The `copy` algorithm expects that the ranges `a` and `b` are valid for reading and writing, respectively. Also important is that the ranges do not overlap, this property is expressed with the `separated`-clause in our specification.

```
/*@
requires valid: \valid_read(a + (0..n-1));
requires valid: \valid(b + (0..n-1));
requires sep:   \separated(a + (0..n-1), b + (0..n-1));

assigns b[0..n-1];

ensures equal:   EqualRanges{Here,Here}(a, n, b);
ensures unchanged: Unchanged{Here,Here}(a, n);
*/
void copy(const value_type* a, const size_type n, value_type* b);
```

Listing 6.8: Formal specification of `copy`

Furthermore the function `copy` assigns the elements from the source range `a` to the destination range `b`, modifying the memory of the elements pointed to by `b`. Again, we can use the `EqualRanges` predicate from Section 3.1 to express that the array `a` equals `b` after `copy` has been called. Nothing else must be altered. To state this we use the `assigns`-clause.

Note the similarities in the specifications of `copy` and `swap_ranges` (Section 6.4).

⁴⁴ See <http://www.sgi.com/tech/stl/copy.html>

6.5.2. Implementation of copy

Listing 6.9 shows an implementation of the `copy` function.

```
void copy(const value_type* a, size_type n, value_type* b)
{
    /*@
        loop invariant bound: 0 <= i <= n;
        loop invariant equal: EqualRanges{Here,Here}(a, i, b);
        loop assigns    i, b[0..n-1];
        loop variant    n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        b[i] = a[i];
    }
}
```

Listing 6.9: Implementation of `copy`

Here are some remarks on its loop invariants.

For the postcondition to be true, we must ensure that for every element `i`, the comparison `a[i] == b[i]` is **true**. This can be expressed by using the `EqualRanges` predicate.

The `assigns` clause ensures that nothing but the range `b[0..i-1]` and the loop variable `i` is modified. In order to prove the termination of the loop for every possible `n` we use the loop variant `n-i` and cover it with an assertion.

6.6. The reverse_copy algorithm

The `reverse_copy`⁴⁵ algorithm of the C++ Standard Library inverts the order of elements in a sequence. `reverse_copy` does not change the input sequence, and copies its result to the output sequence. For our purposes we have modified the generic implementation to that of a range of type `value_type`. The signature now reads:

```
void reverse_copy(const value_type* a, size_type n, value_type* b);
```

6.6.1. The predicate Reverse

Informally, `reverse_copy` copies the elements from the array `a` into array `b` such that the copy is a reverse of the original array. Thus, after calling `reverse_copy` the following conditions shall be satisfied.

$$\begin{array}{rcl} b[0] & == & a[n-1] \\ b[1] & == & a[n-2] \\ \vdots & & \vdots \\ b[n-1] & == & a[0] \end{array}$$

In order to concisely formalize these conditions we define the overloaded predicates `Reverse` that are shown in Listing 6.10.

```
/*@
predicate
Reverse{A,B}(value_type* a, integer n, value_type* b,
              integer first, integer last) =
  \forall integer k; first <= k < last ==>
    \at(a[k], A) == \at(b[n-1-k], B);

predicate
Reverse{A,B}(value_type* a, integer n, value_type* b) =
  Reverse{A,B}(a, n, b, 0, n);

predicate
Reverse{A,B}(value_type* a, integer n) = Reverse{A,B}(a, n, a);
*/
```

Listing 6.10: The predicate `Reverse`

⁴⁵ See http://www.sgi.com/tech/stl/reverse_copy.html

Figure 6.11 graphically represents the first version of predicate `Reverse` in Listing 6.10. The second version of `Reverse` in Listing 6.10 captures the reverse relation between two complete arrays while the third version is useful when a single array shall be reversed (see Section 6.7).




Figure 6.11.: Sketch of predicate `Reverse`

6.6.2. Formal specification of `reverse_copy`

The ACSL specification of `reverse_copy` is shown in Listing 6.12. We use the second version of predicate `Reverse` from Listing 6.10 in order to formulate the postcondition of `reverse_copy`.

```

/*@
  requires valid_a: \valid_read(a + (0..n-1));
  requires valid_b:  \valid(b + (0..n-1));
  requires sep:      \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..(n-1)];

  ensures reverse:   Reverse{Here,Here}(a, n, b);
  ensures unchanged: Unchanged{Pre,Here}(a, n);
*/
void reverse_copy(const value_type* a, size_type n, value_type* b);

```

Listing 6.12: Formal specification of `reverse_copy`

6.6.3. Implementation of `reverse_copy`

Listing 6.13 shows an implementation of the `reverse_copy` function. For the postcondition to be true, we must ensure that for every element i , the comparison $b[i] == a[n-1-i]$ holds. This is formalized by the loop invariant `reverse` where we employ the first version of `Reverse` from Listing 6.10.

```
void reverse_copy(const value_type* a, size_type n, value_type* b)
{
    /*@
        loop invariant bound:    0 <= i <= n;
        loop invariant reverse: Reverse{Here,Here}(b, n, a, 0, i);
        loop assigns i, b[0..n-1];
        loop variant n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        b[i] = a[n - 1 - i];
    }
}
```

Listing 6.13: Implementation of `reverse_copy`

6.7. The reverse algorithm

The `reverse`⁴⁶ algorithm of the C++ Standard Library inverts the order of elements in a sequence. The `reverse` algorithm works in place, meaning that it modifies its input sequence. For our purposes we have modified the generic implementation to that of a range of type `value_type`. The signature now reads:

```
void reverse(value_type* a, size_type n);
```

6.7.1. Formal specification of `reverse`

The ACSL specification for the `reverse` function is shown in listing 6.14. In the postcondition we use the third version of predicate `Reverse` from Listing 6.10.

```
/*@
    requires valid: \valid(a + (0..n-1));

    assigns a[0..(n-1)];

    ensures reverse: Reverse{Here,Old}(a, n);
*/
void reverse(value_type* a, size_type n);
```

Listing 6.14: Formal specification of `reverse`

⁴⁶ See <http://www.sgi.com/tech/stl/reverse.html>

6.7.2. Implementation of reverse

Listing 6.15 shows an implementation of the `reverse` function. Since the `reverse` algorithm operates *in place* we use the `swap` function from Section 6.2 in order to exchange the elements of the first half of the array with the corresponding elements of the second half.

```
void reverse(value_type* a, size_type n)
{
    const size_type half = n / 2;

    /*@
        loop invariant bound: 0 <= i <= half;

        loop invariant left:  Reverse{Here,Pre}(a, n, a, 0, i);
        loop invariant middle: Unchanged{Here,Pre}(a, i, n-i);
        loop invariant right: Reverse{Here,Pre}(a, n, a, n-i, n);

        loop assigns i, a[0..n-1];
        loop variant half - i;
    */
    for (size_type i = 0; i < half; ++i) {
        swap(&a[i], &a[n - 1 - i]);
    }
}
```

Listing 6.15: Implementation of `reverse`

We reuse the predicates `Reverse` (Listing 6.10) and `Unchanged` (Listing 6.1) in order to write concise loop invariants.

6.8. The `rotate_copy` algorithm

The `rotate_copy` algorithm in the C++ Standard Library rotates a sequence downwards by m positions and copies the results to another same-sized sequence. For our purposes we have modified the generic implementation⁴⁷ to that of a range of type `value_type`. The signature now reads:

```
void rotate_copy(const value_type* a, size_type m,
                 size_type n, value_type* b);
```

Informally, the last $n-m$ elements of the array $a[0..n-1]$ are moved m places downwards and stored as the first $n-m$ elements of the array $b[0..n-1]$, whereas the first m elements of the array $a[0..n-1]$ are wrapped around and stored as the last m elements of the array $b[0..n-1]$.




Figure 6.16.: Effects of `rotate_copy`

Figure 6.16 illustrates the effects of `rotate_copy` by highlighting how the initial and final segments of the array $a[0..n-1]$ are mapped to corresponding segments of the array $b[0..n-1]$.

⁴⁷ See http://www.sgi.com/tech/stl/rotate_copy.html

6.8.1. Formal specification of rotate_copy

The ACSL specification of rotate_copy is shown in Listing 6.17.

```
/*@
  requires bound: 0 <= m <= n;
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep: \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..(n-1)];

  ensures equal_first: EqualRanges{Here,Here}(a, m, b+(n-m));
  ensures equal_last: EqualRanges{Here,Here}(a+m, n-m, b);
  ensures unchanged: Unchanged{Old,Here}(a, n);
*/
void rotate_copy(const value_type* a, size_type m, size_type n,
                 value_type* b);
```

Listing 6.17: Formal specification of rotate_copy

6.8.2. Implementation of rotate_copy

Listing 6.18 shows an implementation of the rotate_copy function. The implementation simply calls the function copy twice.

```
void rotate_copy(const value_type* a, size_type m, size_type n,
                 value_type* b)
{
  copy(a, m, b + (n - m));
  copy(a + m, n - m, b);
}
```

Listing 6.18: Implementation of rotate_copy

6.9. The `replace_copy` algorithm

The `replace_copy` algorithm of the C++ Standard Library substitutes specific elements from general sequences. Here, the general implementation⁴⁸ has been altered to process `value_type` ranges. The new signature reads:

```
size_type replace_copy(const value_type* a, size_type n,
                      value_type* b,
                      value_type v, value_type w);
```

The `replace_copy` algorithm copies the elements from the range `a[0..n]` to range `b[0..n]`, substituting every occurrence of `v` by `w`. The return value is the length of the range. As the length of the range is already a parameter of the function this return value does not contain new information.

6.9.1. The predicate `Replace`

We start with defining in Listing 6.19 the predicate `Replace` that describes the intended relationship between the input array `a[0..n-1]` and the output array `b[0..n-1]`. Note the introduction of *local bindings* `\let ai = ...` and `\let bi = ...` in the definition of `Replace` (see [9, §2.2]).

```
/*@
  predicate
    Replace{K,L}(value_type* a, integer n, value_type* b, value_type v, value_type w)
    =
    \forall integer i; 0 <= i < n ==>
      \let ai = \at(a[i],K); \let bi = \at(b[i],L);
      (ai == v ==> bi == w) && (ai != v ==> bi == ai) ;

  predicate
    Replace{K,L}(value_type* a, integer n, value_type v, value_type w) =
    Replace{K,L}(a, n, a, v, w);
*/
```

Listing 6.19: The predicate `Replace`

Listing 6.19 also contains a second, overloaded version of `Replace` which we will use for the specification of the related in-place algorithm `replace` in Section 6.10.

⁴⁸ See http://www.sgi.com/tech/stl/replace_copy.html

6.9.2. Formal specification of `replace_copy`

Using predicate `Replace` the ACSL specification of `replace_copy` is as simple as in Listing 6.20. Note that we require that the arrays `a` and `b` are non-overlapping.

```
/*@
  requires valid_a: \valid_read(a + (0..n-1));
  requires valid_b:  \valid(b + (0..n-1));
  requires sep:      \separated(a + (0..n-1), b + (0..n-1));

  assigns b[0..n-1];

  ensures replace:   Replace{Old,Here}(a, n, b, oldv, newv);
  ensures unchanged: Unchanged{Old,Here}(a, n);
  ensures result:    \result == n;
*/
size_type replace_copy(const value_type* a, size_type n,
                      value_type* b,
                      value_type oldv, value_type newv);
```

Listing 6.20: Formal specification of the `replace_copy`

6.9.3. Implementation of `replace_copy`

An implementation (including loop annotations) of `replace_copy` is shown in Listing 6.21. Note how the structure of the loop annotations resembles the specification of Listing 6.20.

```
size_type replace_copy(const value_type* a, size_type n,
                      value_type* b,
                      value_type oldv, value_type newv)
{
  /*@
    loop invariant bounds:  0 <= i <= n;
    loop invariant replace: Replace{Pre,Here}(a, i, b, oldv, newv);
    loop assigns i, b[0..n-1];
    loop variant n-i;
  */
  for (size_type i = 0; i < n; ++i) {
    b[i] = (a[i] == oldv ? newv : a[i]);
  }

  return n;
}
```

Listing 6.21: Implementation of the `replace_copy` algorithm

6.10. The `replace` algorithm

The `replace` algorithm of the C++ Standard Library substitutes specific values in a general sequence. Here, the general implementation⁴⁹ has been altered to process `value_type` ranges. The new signature reads

```
void replace(value_type* a, size_type n, value_type oldv, value_type newv);
```

The `replace` algorithm substitutes all elements from the range `a[0..n-1]` that equal `oldv` by `newv`.

6.10.1. Formal specification of `replace`

Using the second predicate `Replace` from Listing 6.19 the ACSL specification of `replace` can be expressed as in Listing 6.22.

```
/*@
  requires valid: \valid(a + (0..n-1));

  assigns a[0..n-1];

  ensures replace: Replace{Old,Here}(a, n, oldv, newv);
*/
void replace(value_type* a, size_type n, value_type oldv, value_type newv);
```

Listing 6.22: Formal specification of the `replace`

⁴⁹ See <http://www.sgi.com/tech/stl/replace.html>

6.10.2. Implementation of replace

An implementation of `replace` is shown in Listing 6.23. The loop invariant `unchanged` expresses that when entering iteration `i` the elements `a[i..n-1]` have not yet changed.

```
void replace(value_type* a, size_type n, value_type oldv, value_type newv)
{
    /*@
        loop invariant bounds:    0 <= i <= n;
        loop invariant replace:   Replace{Pre,Here}(a, i, oldv, newv);
        loop invariant unchanged: Unchanged{Pre,Here}(a, i, n);
        loop assigns i, a[0..n-1];
        loop variant n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        if (a[i] == oldv) {
            a[i] = newv;
        }
    }
}
```

Listing 6.23: Implementation of the `replace` algorithm

6.11. The `remove_copy` algorithm

The `remove_copy` algorithm of the C++ Standard Library copies all elements of a sequence other than a given value. Here, the general implementation has been altered to process `value_type` ranges.⁵⁰ The new signature reads:

```
size_type remove_copy(const value_type* a, size_type n,
                     value_type* b, value_type v);
```

The most important facts of this algorithms are:

1. The return value is the length of the resulting range.
2. The `remove_copy` algorithm copies elements that are not equal to `v` from range `a[0..n-1]` to the range `b[0..\result-1]`.
3. The algorithm is stable, that is, the relative order of the elements in `b` is the same as in `a`.




Figure 6.24.: Effects of `remove_copy`

Figure 6.24 shows how `remove_copy` is supposed to copy elements that differ from `v` from range `a` to `b`.

6.11.1. The predicate `RetainAllButOne`

In order to achieve a concise specification we introduce the auxiliary predicate `RetainAllButOne` (see Listing 6.25). The expression `RetainAllButOne(a, m, b, n, v)` is true if range `a[0..n-1]` contains the same elements as `b[0..m-1]`, except possibly for occurrences of `v`; the elements' order may differ in `a` and `b`.

```
/*@
predicate
  RetainAllButOne(value_type* a, integer m,
                 value_type* b, integer n, value_type v) =
    \forall value_type x;
      x != v ==> Count(a, m, x) == Count(b, n, x);
*/
```

Listing 6.25: The predicate `RetainAllButOne`

⁵⁰ See http://www.sgi.com/tech/stl/remove_copy.html

6.11.2. Formal specification of `remove_copy`

Listing 6.26 now shows our first attempt to specify `remove_copy`.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires sep: \separated(a + (0..n-1), b+(0..n-1));

  assigns b[0..(n-1)];

  ensures bound: 0 <= \result <= n;
  ensures result: \result == n - Count(a, n, v);
  ensures retain: RetainAllButOne(a, n, b, \result, v);
  ensures discard: !HasValue(b, \result, v);
  ensures unchanged: Unchanged{Here,Old}(b, \result, n);
  ensures unchanged: Unchanged{Here,Old}(a, n);
*/
size_type remove_copy(const value_type* a, size_type n,
                      value_type* b, value_type v);
```

Listing 6.26: Formal specification of `remove_copy`

- We use the predicate `RetainAllButOne` in order to express that the number of elements different from `v` is the same in the source and target range.

Note that this property does not guarantee the stability of `remove_copy` because given e.g. a range `{1,0,5,2,0,5}` and the value `v=0` the expected result of `remove_copy` is the range `{1,5,2,5}`. However, since `Count` is invariant under permutations the specification in Listing 6.26 would also allow e.g. the result `{5,5,1,2}`. In Section 6.12 we will discuss how the stability of `remove_copy` can be captured in an ACSL specification.

- The predicate `Unchanged` from Listing 6.1 is used to express that `remove_copy` does neither change `b[\result..n-1]` nor `a[0..n-1]`.
- Note the re-use of predicate `HasValue` (Listing 3.11) to express that the target range does not contain the value `v`.

6.11.3. Implementation of `remove_copy`

An implementation of `remove_copy` is shown in Listing 6.27.

```
size_type remove_copy(const value_type* a, size_type n,
                      value_type* b, value_type v)
{
    size_type j = 0;

    /*@
    loop invariant bound:      0 <= j <= i <= n;
    loop invariant result:     j == i - Count(a, i, v);
    loop invariant retain:     RetainAllButOne(a, i, b, j, v);
    loop invariant discard:     !HasValue(b, j, v);
    loop invariant unchanged:   Unchanged{Here,Pre}(b, j, n);

    loop assigns i, j, b[0..n-1];
    loop variant n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        /*@
        requires discard:      !HasValue(b, j, v);
        requires bound:        0 <= j <= i < n;
        assigns                 j, b[j];
        ensures bound:         0 <= j <= i+1 <= n;
        ensures result:        j == i+1 - Count(a, i+1, v);
        ensures discard:       !HasValue(b, \old(j), v);
        ensures retain:        RetainAllButOne(a, i+1, b, j, v);
        ensures unchanged:     Unchanged{Here,Pre}(b, j, n);

        behavior not_equal:
            assumes a[i] != v;
            assigns j, b[j];
            ensures discard: j == \old(j) + 1;
            ensures discard: b[\old(j)] == a[i];
            ensures discard: b[\old(j)] != v;
            ensures discard: !HasValue(b, \old(j), v);
            ensures retain: RetainAllButOne(a, i+1, b, j, v);

        behavior equal:
            assumes a[i] == v;
            assigns \nothing;
            ensures discard: !HasValue(b, j, v);
            ensures retain: RetainAllButOne(a, i+1, b, j, v);

        complete behaviors;
        disjoint behaviors;
        */
        if (a[i] != v) {
            b[j++] = a[i];
        }
        /*@ assert retain: a[i] != v ==> RetainAllButOne(a, i+1, b, j, v);
        /*@ assert discard: a[i] != v ==> b[j-1] != v;
        /*@ assert discard: a[i] != v ==> !HasValue(b, j, v);
    }

    return j;
}
```

Listing 6.27: Implementation of `remove_copy`

Not surprisingly, the logical function `Count` and the predicates `RetainAllButOne`, `HasValue`, and `Unchanged` also appear in the loop invariants of `remove_copy`. In order to automatically discharge the various loop invariants we introduce a (fairly large) *statement contract* (cf. [9, § 2.4.4]) for the if-statement in `remove_copy` together with some assertions.

The verification also relies on the Lemma `RetainAllButOneMiss`, shown in Listing 6.28. It gives a sufficient condition to infer `RetainAllButOne{L}(a, m+1, b, n+1, v)` from `RetainAllButOne{K}(a, m, b, n, v)`.

```
/*@
lemma RetainAllButOneMiss{K,L}:
  \forall value_type *a, *b, v, integer m, n;
    RetainAllButOne{K}(a, m, b, n, v) ==>
      \at(a[m], K) != v ==>
      \at(a[m], L) == \at(b[n], L) ==>
      Unchanged{K,L}(a, m+1) ==>
      Unchanged{K,L}(b, n) ==>
      RetainAllButOne{L}(a, m+1, b, n+1, v);
*/
```

Listing 6.28: A lemma for `RetainAllButOne`

6.12. Capturing the stability of `remove_copy`

In this section, we have a closer look at the *stability* of `remove_copy` and its expression in the ACSL language.

Figure 6.29 shows, with respect to array indices, how the elements different from `v` “slide” to smaller or equal positions. The main observation here is that *an element slides as many positions down as there are occurrences of `v` below it*.




Figure 6.29.: Stability of `remove_copy` with respect to indices

As it turns out, it is relatively easy to express this property using the previously introduced logic function `Count` (see Listing 3.24 in Section 3.8). We simply define in Listing 6.30 a logic function `RemoveCount` which subtracts from every position `i` the number of occurrences of `v` below `i`. In addition, we provide four lemmas that easily follow from corresponding properties of `Count`, but are needed for the verification.

```
/*@
logic
integer RemoveCount{L}(value_type* a, integer i, value_type v) =
    i - Count{L}(a, i, v);

lemma RemoveCountEmpty:
    \forall value_type *a, v, integer i;
        i <= 0 ==> RemoveCount(a, i, v) == i;

lemma RemoveCountHit:
    \forall value_type *a, v, integer i; a[i] == v ==>
        RemoveCount(a, i+1, v) == RemoveCount(a, i, v);

lemma RemoveCountMiss:
    \forall value_type *a, v, integer i; a[i] != v ==>
        RemoveCount(a, i+1, v) == RemoveCount(a, i, v) + 1;

lemma RemoveCountRead{L1,L2}:
    \forall value_type *a, v, integer i; Unchanged{L1,L2}(a, i) ==>
        RemoveCount{L1}(a, i, v) == RemoveCount{L2}(a, i, v);
*/
```

Listing 6.30: The logic function `RemoveCount`

The value `RemoveCount(a, v, i)` equals the number of elements of `a[0..i-1]` that are copied to

the destination range `b[0..n-1]` by `remove_copy`.

The stability of `remove_copy` depends, strictly speaking, on monotonicity properties of `RemoveCount`. Listing 6.31 formulates two lemmas that describe such properties. These lemmas, whose verification depends on the lemmas for `Count` from Listing 3.25, will play an important role in the formal verification of this version of `remove_copy`.

```
/*@
  lemma RemoveCountMonotonic:
    \forallall value_type *a, v, integer m, n; 0 <= m <= n ==>
      RemoveCount(a, m, v) <= RemoveCount(a, n, v);

  lemma RemoveCountStrictlyMonotonic:
    \forallall value_type *a, v, integer m, n; 0 <= m < n ==>
      a[m] != v ==> RemoveCount(a, m, v) < RemoveCount(a, n, v);
*/
```

Listing 6.31: Additional lemmas for `RemoveCount`

Also, note that `RemoveCount` is defined for all integers, including those indices `i` where `a[i]` equals `v` (see the dashed lines in Figure 6.29). In the specification of `remove_copy` we will, however, only use `RemoveCount` for indices where `a[i]` is different from `v`. This can be seen in the definition of predicate `RemoveMapping` (Listing 6.32) that uses `RemoveCount` to formally capture the stability with respect to corresponding elements of the source and target ranges.

```
/*@
  predicate
    RemoveMapping{L}(value_type* a, integer n, value_type* b, value_type v) =
    \forallall integer i; 0 <= i < n ==>
      a[i] != v ==> b[RemoveCount(a, i, v)] == a[i];
*/
```

Listing 6.32: The predicate `RemoveMapping`

6.12.1. Formal specification of `remove_copy`

Listing 6.33 shows an improved specification of `remove_copy` that also captures the required stability.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires valid:   \valid(b + (0..n-1));
  requires sep:     \separated(a + (0..n-1), b+(0..n-1));

  assigns  b[0..(n-1)];

  ensures  bound:      0 <= \result <= n;
  ensures  result:     \result == RemoveCount(a, n, v);
  ensures  retain:     RetainAllButOne(a, n, b, \result, v);
  ensures  discard:    !HasValue(b, \result, v);
  ensures  stable:     RemoveMapping(a, n, b, v);
  ensures  unchanged:  Unchanged{Here,Old}(b, \result, n);
  ensures  unchanged:  Unchanged{Here,Old}(a, n);
*/
size_type remove_copy(const value_type* a, size_type n,
                      value_type* b, value_type v);
```

Listing 6.33: Improved formal specification of `remove_copy`

This specification of `RemoveCount` differs from the one in Listing 6.26 in the following points.

1. We now use `RemoveCount` in order to specify the expected return value in postcondition `result`.
2. We use `RemoveMapping` in the new postcondition `stable`. Here we exactly specify to which element in the output range `b[0..n-1]` an element of the input range `a[0..n-1]`, that is different from `v`, is copied.

Properties `bound`, `retain`, and `discard` are redundant.

6.12.2. Implementation of `remove_copy`

Listing 6.34 shows the additional loop annotations, assertions and statement contract that were necessary to verify the stronger specification of `remove_copy`.

In order to prove the additional loop invariant `stable` we rely also on the monotonicity properties of `RemoveCount` in Listing 6.31


```

size_type remove_copy(const value_type* a, size_type n,
                     value_type* b, value_type v)
{
    size_type j = 0;

    /*@
    loop invariant bound:      0 <= j <= i <= n;
    loop invariant result:     j == RemoveCount(a, i, v);
    loop invariant retain:     RetainAllButOne(a, i, b, j, v);
    loop invariant discard:    !HasValue(b, j, v);
    loop invariant stable:     RemoveMapping(a, i, b, v);
    loop invariant unchanged:  Unchanged{Here,Pre}(b, j, n);
    loop invariant unchanged:  Unchanged{Here,Pre}(a, n);

    loop assigns i, j, b[0..n-1];
    loop variant n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        /*@
        requires discard:      !HasValue(b, j, v);
        requires bound:        0 <= j <= i < n;
        assigns                j, b[j];
        ensures bound:          0 <= j <= i+1 <= n;
        ensures result:         j == RemoveCount(a, i+1, v);
        ensures discard:        !HasValue(b, \old(j), v);
        ensures retain:         RetainAllButOne(a, i+1, b, j, v);
        ensures stable:         RemoveMapping(a, i+1, b, v);
        ensures unchanged:     Unchanged{Here,Pre}(b, j, n);

        behavior not_equal:
            assumes a[i] != v;
            assigns j, b[j];
            ensures discard: j == \old(j) + 1;
            ensures discard: b[\old(j)] == a[i];
            ensures discard: b[\old(j)] != v;
            ensures discard: !HasValue(b, \old(j), v);
            ensures retain: RetainAllButOne(a, i+1, b, j, v);
            ensures stable: RemoveMapping(a, i+1, b, v);

        behavior equal:
            assumes a[i] == v;
            assigns \nothing;
            ensures discard: !HasValue(b, j, v);
            ensures retain: RetainAllButOne(a, i+1, b, j, v);
            ensures stable: RemoveMapping(a, i+1, b, v);

        complete behaviors;
        disjoint behaviors;
        */
        if (a[i] != v) {
            b[j++] = a[i];
        }
        /*@ assert retain: a[i] != v ==> RetainAllButOne(a, i+1, b, j, v);
        /*@ assert discard: a[i] != v ==> b[j-1] != v;
        /*@ assert discard: a[i] != v ==> !HasValue(b, j, v);
        /*@ assert stable: a[i] != v ==> j == RemoveCount(a, i, v) + 1;

    }

    return j;
}

```

Listing 6.34: Implementation of `remove_copy` with additional loop invariants

7. Numeric algorithms

The algorithms that we considered so far only *compared*, *read* or *copied* values in sequences. In this chapter, we consider so-called *numeric* algorithms of the C++ standard library that use arithmetic operations on `value_type` to combine the elements of sequences.

In order to refer to potential arithmetic overflows we introduce the two constants

```
#define VALUE_TYPE_MAX INT_MAX
#define VALUE_TYPE_MIN INT_MIN
```

Listing 7.1: Limits of `value_type`

which refer to the numeric limits of `value_type` (see also Section 1.2.1).

We consider the following algorithms.

- `iota` writes sequentially increasing values into a range (Section 7.1 on Page 100)
- `accumulate` computes the sum of the elements in a range (Section 7.2 on Page 102)
- `inner_product` computes the inner product of two ranges (Section 7.3 on Page 106)
- `partial_sum` computes the sequence of partial sums of a range (Section 7.4 on Page 109)
- `adjacent_difference` computes the differences of adjacent elements in a range (Section 7.5 on Page 113)

The formal specifications of these algorithms raise new questions. In particular, we now have to deal with arithmetic overflows in `value_type`.

7.1. The `iota` algorithm

The `iota` algorithm in the C++ standard library assigns sequentially increasing values to a range, where the initial value is user-defined. Our version of the original signature⁵¹ reads:

```
void iota(value_type* a, size_type n, value_type val);
```

Starting at `val`, the function assigns consecutive integers to the elements of the range `a`. When specifying `iota` we must be careful to deal with possible overflows of the argument `val`.

7.1.1. Formal specification of `iota`

The specification of `iota` relies on the logic function `Iota` that is defined in Listing 7.2.

```
/*@
  predicate
  Iota(value_type* a, integer n, value_type v) =
    \forall integer i; 0 <= i < n ==> a[i] == v + i;
*/
```

Listing 7.2: Logic function `Iota`

The ACSL specification of `iota` is shown in Listing 7.3. It uses the logic function `Iota` in order to express the postcondition increment.

```
/*@
  requires valid: \valid(a + (0..n-1));
  requires limit: val + n <= VALUE_TYPE_MAX;

  assigns a[0..n-1];

  ensures increment: Iota(a, n, val);
*/
void iota(value_type* a, size_type n, value_type val);
```

Listing 7.3: Formal specification of `iota`

The specification of `iota` refers to `VALUE_TYPE_MAX` which is the maximum value of the underlying integer type (see Listing 7.1). In order to avoid integer overflows the sum `val+n` must not be greater than the constant `VALUE_TYPE_MAX`.

⁵¹ See <http://www.sgi.com/tech/stl/iota.html>

7.1.2. Implementation of `iota`

Listing 7.4 shows an implementation of the `iota` function.

```
void iota(value_type* a, size_type n, value_type val)
{
    /*@
    loop invariant bound:      0 <= i <= n;
    loop invariant limit:      val == \at(val, Pre) + i;
    loop invariant increment: Iota(a, i, \at(val, Pre));

    loop assigns i, val, a[0..n-1];
    loop variant n-i;
    */
    for (size_type i = 0; i < n; ++i) {
        a[i] = val++;
    }
}
```

Listing 7.4: Implementation of `iota`

The loop invariant `increment` describes that in each iteration of the loop the current value `val` is equal to the sum of the value `val` in state of function entry and the loop index `i`. We have to refer here to `\at(val, Pre)` which is the value on entering `iota`.

7.2. The accumulate algorithm

The `accumulate` algorithm in the C++ standard library computes the sum of an given initial value and the elements in a range. Our version of the original signature⁵² reads:

```
value_type
accumulate(const value_type* a, size_type n, value_type init);
```

The result of `accumulate` shall equal the value

$$\text{init} + \sum_{i=0}^{n-1} a[i]$$

This implies that `accumulate` will return `init` for an empty range.

7.2.1. Axiomatic definition of accumulating over an array

As in the case of `count` (see Section 3.8) we specify `accumulate` by first defining a *logic function* `Accumulate` that formally defines the summation of elements in an array.

```
/*@
  axiomatic AccumulateAxiomatic
  {
    logic value_type Accumulate{L}(value_type* a, integer n,
                                   value_type init) reads a[0..n-1];

    axiom AccumulateEmpty:
      \forall value_type *a, init, integer n;
        n <= 0 ==> Accumulate(a, n, init) == init;

    axiom AccumulateNext:
      \forall value_type *a, init, integer n;
        n >= 0 ==>
          Accumulate(a, n+1, init) == Accumulate(a, n, init) + a[n];

    axiom AccumulateRead{L1,L2}:
      \forall value_type *a, init, integer n;
        Unchanged{L1,L2}(a, n) ==>
          Accumulate{L1}(a, n, init) == Accumulate{L2}(a, n, init);
  }
*/
```

Listing 7.5: The logic function `Accumulate`

With this definition the following equation holds

$$\text{Accumulate}(a, n + 1, \text{init}) = \text{init} + \sum_{i=0}^n a[i] \quad (7.1)$$

Both the `reads` clause and the axiom `AccumulateRead` in Listing 7.5 express that the result of `Accumulate` only depends on the values of `a[0..n-1]`.

⁵² See <http://www.sgi.com/tech/stl/accumulate.html>

Listing 7.6 shows an overloaded version of `Accumulate` that uses 0 as default value of `init`. Included in this listing is also a property corresponding to axiom `AccumulateNext` from Listing 7.5, here given as a lemma. We will use this version for the specification of the algorithm `partial_sum` (see Section 7.4).

Thus, for the overloaded version of `Accumulate` we have

$$\text{Accumulate}(a, n + 1) = \sum_{i=0}^n a[i] \quad (7.2)$$

```
/*@
logic value_type Accumulate{L}(value_type* a, integer n) =
  Accumulate{L}(a, n, (value_type) 0);

lemma AccumulateDefault0{L}:
  \forall value_type* a;
  Accumulate(a, 0) == 0;

lemma AccumulateDefault1{L}:
  \forall value_type* a;
  Accumulate(a, 1) == a[0];

lemma AccumulateDefaultNext{L}:
  \forall value_type* a, integer n;
  n >= 0 ==> Accumulate(a, n+1) == Accumulate(a, n) + a[n];

lemma AccumulateDefaultRead{L1,L2}:
  \forall value_type* a, integer n;
  Unchanged{L1,L2}(a, n) ==>
  Accumulate{L1}(a, n) == Accumulate{L2}(a, n);

*/
```

Listing 7.6: An overloaded version of `Accumulate`

7.2.2. Preventing numeric overflows for `accumulate`

Before we present our formal specification of `accumulate` we introduce in Listing 7.7 a predicate `AccumulateBounds` that we will subsequently use in order to compactly express requirements that exclude numeric overflows while accumulating value.

```
/*@
  predicate
    AccumulateBounds{L}(value_type* a, integer n, value_type init) =
      \forall integer i; 0 <= i <= n ==>
        VALUE_TYPE_MIN <= Accumulate(a, i, init) <= VALUE_TYPE_MAX;

  predicate
    AccumulateBounds{L}(value_type* a, integer n) =
      AccumulateBounds{L}(a, n, (value_type) 0);
*/
```

Listing 7.7: The overloaded predicate `AccumulateBounds`

Predicate `AccumulateBounds` expresses that for $0 \leq i < n$ the *partial sums*

$$\text{init} + \sum_{k=0}^i a[k] \tag{7.3}$$

do not overflow. If one of them did, one couldn't guarantee that the result of `accumulate` equals the mathematical description of `Accumulate`.

Note that we also provide a second (overloaded) version of `AccumulateBounds` which uses a default value 0 for `init`.

7.2.3. Formal specification of accumulate

Using the logic function `Accumulate` and the predicate `AccumulateBounds`, the ACSL specification of `accumulate` is then as simple as shown in Listing 7.8.

```
/*@
  requires valid: \valid_read(a + (0..n-1));
  requires bounds: AccumulateBounds(a, n, init);

  assigns \nothing;

  ensures result: \result == Accumulate(a, n, init);
*/
value_type
accumulate(const value_type* a, size_type n, value_type init);
```

Listing 7.8: Formal specification of `accumulate`

7.2.4. Implementation of accumulate

Listing 7.9 shows an implementation of the `accumulate` function with corresponding loop annotations.

```
value_type
accumulate(const value_type* a, size_type n, value_type init)
{
  /*@
    loop invariant index:    0 <= i <= n;
    loop invariant partial:  init == Accumulate(a, i, \at(init,Pre));
    loop assigns i, init;
    loop variant n-i;
  */
  for (size_type i = 0; i < n; ++i) {
    init = init + a[i];
  }

  return init;
}
```

Listing 7.9: Implementation of `accumulate`

Note that loop invariant `partial` claims that in the i -th iteration step `result` equals the accumulated value of Equation (7.3). This depends on the property `bounds` in Listing 7.8 which expresses that there is no numeric overflow when updating the variable `init`.

7.3. The inner_product algorithm

The `inner_product` algorithm in the C++ standard library computes the *inner product*⁵³ of two ranges. Our version of the original signature⁵⁴ reads:

```
value_type
inner_product(const value_type* a, const value_type* b,
              size_type n, value_type init);
```

The result of `inner_product` equals the value

$$\text{init} + \sum_{i=0}^{n-1} a[i] \cdot b[i]$$

thus, `inner_product` will return `init` for empty ranges.

7.3.1. The logic function InnerProduct

As in the case of `accumulate` (see Section 7.2) we specify `inner_product` by first defining a *logic function* `InnerProduct` that formally defines the summation of the element-wise product of two arrays.

```
/*@
  axiomatic InnerProductAxiomatic
  {
    logic integer
      InnerProduct{L}(value_type* a, value_type* b, integer n,
                     value_type init) reads a[0..n-1], b[0..n-1];

    axiom InnerProductEmpty:
      \forall value_type *a, *b, init, integer n;
        n <= 0 ==> InnerProduct(a, b, n, init) == init;

    axiom InnerProductNext:
      \forall value_type *a, *b, init, integer n;
        n >= 0 ==> InnerProduct(a, b, n + 1, init) ==
          InnerProduct(a, b, n, init) + (a[n] * b[n]);

    axiom InnerProductRead{L1,L2}:
      \forall value_type *a, *b, init, integer n;
        Unchanged{L1,L2}(a, n) && Unchanged{L1,L2}(b, n) ==>
          InnerProduct{L1}(a, b, n, init) ==
          InnerProduct{L2}(a, b, n, init);
  }
*/
```

Listing 7.10: The logic function `InnerProduct`

Both Axiom `InnerProductRead` and the `reads` clause serve the same purpose in that they express that the result of the `InnerProduct` only depends on the values of `a[0..n-1]` and `b[0..n-1]`.

⁵³ Also referred to as *dot product*, see http://en.wikipedia.org/wiki/Dot_product

⁵⁴ See http://www.sgi.com/tech/stl/inner_product.html

7.3.2. Preventing numeric overflows for `inner_product`

Before we present our formal specification of `inner_product` we introduce in Listing 7.11 two predicates that we will use subsequently in order to compactly express requirements that exclude numeric overflows while computing the inner product.

```
/*@
  predicate
    ProductBounds(value_type* a, value_type* b, integer n) =
      \forall integer i; 0 <= i < n ==>
        VALUE_TYPE_MIN <= a[i] * b[i] <= VALUE_TYPE_MAX;

  predicate
    InnerProductBounds(value_type* a, value_type* b, integer n,
                       value_type init) =
      \forall integer i; 0 <= i <= n ==>
        VALUE_TYPE_MIN <= InnerProduct(a, b, i, init) <= VALUE_TYPE_MAX;
*/
```

Listing 7.11: The predicates `ProductBounds` and `InnerProductBounds`

Predicate `ProductBounds` expresses that for $0 \leq i < n$ the products

$$a[i] \cdot b[i] \tag{7.4}$$

do not overflow. Predicate `InnerProductBounds`, on the other hand, states that for $0 \leq i < n$ the *partial sums*

$$\text{init} + \sum_{k=0}^i a[k] \cdot b[k] \tag{7.5}$$

do not overflow.

Otherwise, one cannot guarantee that the result of `inner_product` equals the mathematical description of `InnerProduct`.

7.3.3. Formal specification of `inner_product`

Using the logic function `InnerProduct`, we specify `inner_product` as shown in Listing 7.12. Note that we needn't require that `a` and `b` are separated.

```
/*@
  requires valid_a: \valid_read(a + (0..n-1));
  requires valid_b: \valid_read(b + (0..n-1));
  requires bounds: ProductBounds(a, b, n);
  requires bounds: InnerProductBounds(a, b, n, init);

  assigns \nothing;

  ensures result: \result == InnerProduct(a, b, n, init);
  ensures unchanged: Unchanged{Here,Pre}(a, n);
  ensures unchanged: Unchanged{Here,Pre}(b, n);
*/
value_type
inner_product(const value_type* a, const value_type* b, size_type n,
              value_type init);
```

Listing 7.12: Formal specification of `inner_product`

7.3.4. Implementation of `inner_product`

Listing 7.13 shows an implementation of `inner_product` with corresponding loop annotations.

```
value_type
inner_product(const value_type* a, const value_type* b, size_type n,
              value_type init)
{
  /*@
    loop invariant index: 0 <= i <= n;
    loop invariant inner: init == InnerProduct(a, b, i, \at(init,Pre));
    loop assigns i, init;
    loop variant n-i;
  */
  for (size_type i = 0; i < n; ++i) {
    init = init + a[i] * b[i];
  }

  return init;
}
```

Listing 7.13: Implementation of `inner_product`

Note that the loop invariant `inner` claims that in the i -th iteration step the current value of `init` equals the accumulated value of Equation (7.5). This depends of course on the properties `bounds` in Listing 7.12, which express that there is no arithmetic overflow when computing the updates of the variable `init`.

7.4. The `partial_sum` algorithm

The `partial_sum` algorithm in the C++ standard library computes the sum of a given initial value and the elements in a range. Our version of the original signature⁵⁵ reads:

```
size_type  
partial_sum(const value_type* a, size_type n, value_type* b);
```

After executing the function `partial_sum` the array `b[0..n-1]` holds the following values

$$\begin{aligned}b[0] &= a[0] \\ b[1] &= a[0] + a[1] \\ &\vdots \\ b[n-1] &= a[0] + a[1] + \dots + a[n-1]\end{aligned}$$

More concisely, for $0 \leq i < n$ holds

$$b[i] = \sum_{k=0}^i a[k] \quad (7.6)$$

7.4.1. The predicate `PartialSum`

Equations (7.6) and (7.2) suggest that we define the ACSL predicate `PartialSum` in Listing 7.14 by using the logic function `Accumulate` from Listing 7.6. Listing 7.14.

```
/*@  
  predicate  
    PartialSum{L}(value_type* a, integer n, value_type* b) =  
      \forall i; 0 <= i < n ==> Accumulate(a, i+1) == b[i];  
*/
```

Listing 7.14: The predicate `PartialSum`

⁵⁵ See http://www.sgi.com/tech/stl/partial_sum.html

7.4.2. Formal specification of `partial_sum`

Using the predicates `PartialSum` and `AccumulateBounds`, we specify `partial_sum` as shown in Listing 7.15.

```
/*@
requires valid:      \valid_read(a + (0..n-1));
requires valid:      \valid(b + (0..n-1));
requires separated: \separated(a + (0..n-1), b + (0..n-1));
requires bounds:     AccumulateBounds(a, n+1);

assigns b[0..n-1];

ensures result:      \result == n;
ensures partialsum: PartialSum(a, n, b);
ensures unchanged:   Unchanged{Here,Pre}(a, n);
*/
size_type
partial_sum(const value_type* a, size_type n, value_type* b);
```

Listing 7.15: Formal specification of `partial_sum`

Our specification requires that the arrays `a[0..n-1]` and `b[0..n-1]` are separated, that is, they do not overlap. Note that is a stricter requirement than in the case of the original C++ version of `partial_sum`, which allows that `a` equals `b`, thus allowing the computation of partial sums *in place*.⁵⁶

⁵⁶ See Note [1] at http://www.sgi.com/tech/stl/partial_sum.html

7.4.3. Implementation of `partial_sum`

Listing 7.16 shows an implementation of `partial_sum` with corresponding loop annotations.

```
size_type
partial_sum(const value_type* a, size_type n, value_type* b)
{
    if (n > 0) {
        b[0] = a[0];

        /*@
        loop invariant bound:      1 <= i <= n;
        loop invariant unchanged:  Unchanged{Here,Pre}(a, n);
        loop invariant accumulate: b[i-1] == Accumulate(a, i);
        loop invariant partialsum: PartialSum(a, i, b);
        loop assigns i, b[1..n-1];
        loop variant n - i;
        */
        for (size_type i = 1u; i < n; ++i) {
            //@ ghost Enter:
            b[i] = b[i - 1u] + a[i];
            //@ assert unchanged: a[i] == \at(a[i],Enter);
            //@ assert unchanged: Unchanged{Enter,Here}(a, i);
            //@ assert unchanged: Unchanged{Enter,Here}(b, i);
        }
    }

    return n;
}
```

Listing 7.16: Implementation of `partial_sum`

In order to facilitate the automatic verification of `partial_sum`, we had to add the assertions `unchanged` and provide the lemmas of Listing 7.17.

7.4.4. Additional lemmas

The lemmas shown in Listing 7.17 are needed for the verification of `partial_sum` and the algorithms in Sections 7.6 and 7.7.

```
/*@
lemma PartialSumSection{K}:
  \forallall value_type *a, *b, integer m, n;
    0 <= m <= n ==>
    PartialSum{K}(a, n, b) ==>
    PartialSum{K}(a, m, b);

lemma PartialSumUnchanged{K,L}:
  \forallall value_type *a, *b, integer n;
    0 <= n ==>
    PartialSum{K}(a, n, b) ==>
    Unchanged{K, L}(a, n) ==>
    Unchanged{K, L}(b, n) ==>
    PartialSum{L}(a, n, b);

lemma PartialSumStep{L}:
  \forallall value_type *a, *b, integer n;
    1 <= n ==>
    PartialSum(a, n, b) ==>
    b[n] == Accumulate(a, n+1) ==>
    PartialSum(a, n+1, b);

lemma PartialSumStep2{K,L}:
  \forallall value_type *a, *b, integer n;
    1 <= n ==>
    PartialSum{K}(a, n, b) ==>
    Unchanged{K,L}(a, n+1) ==>
    Unchanged{K,L}(b, n) ==>
    \at (b[n] == Accumulate(a, n+1), L) ==>
    PartialSum{L}(a, n+1, b);
*/
```

Listing 7.17: The lemma `PartialSumStep`

7.5. The `adjacent_difference` algorithm

The `adjacent_difference` algorithm in the C++ standard library computes the differences of adjacent elements in a range. Our version of the original signature⁵⁷ reads:

```
size_type  
adjacent_difference(const value_type* a, size_type n, value_type* b);
```

After executing the function `adjacent_difference` the array `b[0..n-1]` holds the following values

$$\begin{aligned}b[0] &= a[0] \\ b[1] &= a[1] - a[0] \\ &\vdots \\ b[n-1] &= a[n-1] - a[n-2]\end{aligned}$$

If we form the partial sums of the sequence `b` we find that

$$\begin{aligned}a[0] &= b[0] \\ a[1] &= b[0] + b[1] \\ &\vdots \\ a[n-1] &= b[0] + b[1] + \dots + b[n-1]\end{aligned}$$

Thus, we have for $0 \leq i < n$

$$a[i] = \sum_{k=0}^i b[k] \tag{7.7}$$

which means that applying `partial_sum` on the output of `adjacent_difference` produces the original input of `adjacent_difference`.

Conversely, if `a[0..n-1]` and `b[0..n-1]` are the input and output of `partial_sum`, then we have

$$\begin{aligned}b[0] &= a[0] \\ b[1] &= a[0] + a[1] \\ &\vdots \\ b[n-1] &= a[0] + b[1] + \dots + b[n-1]\end{aligned}$$

from which we can conclude

$$\begin{aligned}a[0] &= b[0] \\ a[1] &= b[1] - b[0] \\ &\vdots \\ a[n-1] &= b[n-1] - b[n-2]\end{aligned} \tag{7.8}$$

We will verify these claims in Sections 7.6 and 7.7.

⁵⁷ See http://www.sgi.com/tech/stl/adjacent_difference.html

7.5.1. The predicate `AdjacentDifference`

We define the predicate `AdjacentDifference` in Listing 7.19 by first introducing the logic function `Difference` (Listing 7.18).

```
/*@
  axiomatic DifferenceAxiomatic
  {
    logic value_type
      Difference{L}(value_type* a, integer n) reads a[0..n];

    axiom DifferenceEmptyOrSingle:
      \forall value_type *a, integer n;
        n <= 0 ==> Difference(a, n) == a[0];

    axiom DifferenceNext:
      \forall value_type *a, integer n;
        n >= 1 ==> Difference(a, n) == a[n] - a[n-1];

    axiom DifferenceRead{K,L}:
      \forall value_type *a, integer n;
        Unchanged{K,L}(a, 1+n) ==>
          Difference{K}(a, n) == Difference{L}(a, n);
  }
*/
```

Listing 7.18: The logic function `Difference`

```
/*@
  predicate
    AdjacentDifference{L}(value_type* a, integer n, value_type* b) =
      \forall integer i; 0 <= i < n ==> b[i] == Difference(a, i);
*/
```

Listing 7.19: The predicate `AdjacentDifference`

7.5.2. Formal specification of `adjacent_difference`

We introduce here the predicate `AdjacentDifferenceBounds` (Listing 7.20) that captures conditions that prevent numeric overflows while computing difference of the form $a[i] - a[i-1]$.

```
/*@
  predicate
    AdjacentDifferenceBounds(value_type* a, integer n) =
      \forall integer i; 1 <= i < n ==>
        VALUE_TYPE_MIN <= Difference(a, i) <= VALUE_TYPE_MAX;
*/
```

Listing 7.20: The predicate `AdjacentDifferenceBounds`

Using the predicates `AdjacentDifference` and `AdjacentDifferenceBounds` we can provide a concise formal specification of `adjacent_difference` (Listing 7.21). As in the case of the specification of `partial_sum` we require that the arrays $a[0..n-1]$ and $b[0..n-1]$ are separated.

```

/*@
requires valid:      \valid_read(a + (0..n-1));
requires valid:      \valid(b + (0..n-1));
requires separated:  \separated(a + (0..n-1), b + (0..n-1));
requires bounds:    AdjacentDifferenceBounds(a, n);

assigns b[0..n-1];

ensures result:      \result == n;
ensures difference:  AdjacentDifference(a, n, b);
ensures unchanged:   Unchanged{Here,Pre}(a, n);
*/
size_type
adjacent_difference(const value_type* a, size_type n, value_type* b);

```

Listing 7.21: Formal specification of adjacent_difference

7.5.3. Implementation of adjacent_difference

Listing 7.22 shows an implementation of adjacent_difference with corresponding loop annotations.

In order to achieve the verification of the loop invariant difference we added

- the assertions bound and difference
- the lemmas AdjacentDifferenceStep and AdjacentDifferenceSection from Listing 7.23
- a statement contract with the two postconditions labeled as step

```

size_type
adjacent_difference(const value_type* a, size_type n, value_type* b)
{
    if (n > 0) {
        b[0] = a[0];

        /*@
            loop invariant index:      1 <= i <= n;
            loop invariant unchanged:  Unchanged{Here,Pre}(a, n);
            loop invariant difference: AdjacentDifference(a, i, b);
            loop assigns i, b[1..n-1];
            loop variant n - i;
        */
        for (size_type i = 1u; i < n; ++i) {
            //@ assert bound: VALUE_TYPE_MIN <= Difference(a, i) <= VALUE_TYPE_MAX;
            /*@
                assigns b[i];
                ensures step: Unchanged{Old,Here}(b, i);
                ensures step: b[i] == Difference(a, i);
            */
            b[i] = a[i] - a[i - 1u];
            //@ assert difference: AdjacentDifference{Here}(a, i+1, b);
        }
    }

    return n;
}

```

Listing 7.22: Implementation of adjacent_difference

7.5.4. Additional Lemmas

The lemmas shown in Listing 7.23 are also needed for the verification of the algorithm in Section 7.7.

```
/*@  
  lemma AdjacentDifferenceStep{K,L}:  
    \forallall value_type *a, *b, integer n;  
      AdjacentDifference{K}(a, n, b) ==>  
      Unchanged{K,L}(b, n) ==>  
      Unchanged{K,L}(a, n+1) ==>  
      \at(b[n],L) == Difference{L}(a, n) ==>  
      AdjacentDifference{L}(a, 1+n, b);  
  
  lemma AdjacentDifferenceSection{K}:  
    \forallall value_type *a, *b, integer m, n;  
      0 <= m <= n ==>  
      AdjacentDifference{K}(a, n, b) ==>  
      AdjacentDifference{K}(a, m, b);  
*/
```

Listing 7.23: The lemma AdjacentDifferenceStep

7.6. Inverting `partial_sum` with `adjacent_difference`

In Section 7.5 we had informally argued that `partial_sum` and `adjacent_difference` are inverse to each other (see Equations (7.7) and (7.8)). In the current section, we are going to verify the second of these claims with the help of Frama-C, viz. that applying `adjacent_difference` to the output of `partial_sum` produces the original array. In Section 7.7, we will verify the converse first claim.

Listing 7.24 expresses the property from Equation (7.8) as lemma, on the ACSL logical level. This lemma is verified by Frama-C with the help of automatic theorem provers.

```
/*@
  lemma PartialSumInverse{K,L}:
    \forallall value_type *a, *b, integer n;
      0 <= n ==>
        PartialSum{K}(a, n, b) ==>
        Unchanged{K,L}(b, n) ==>
        AdjacentDifference{L}(b, n, a) ==>
        Unchanged{K,L}(a, n);
*/
```

Listing 7.24: The lemma `PartialSumInverse`

Since the lemma does not deal with arithmetic overflows or potential aliasing of data, we give a corresponding auxiliary C function which takes these issues into account.

Function `partial_sum_inverse`, shown in Listing 7.25, calls first `partial_sum` and then `adjacent_difference`. The contract of this function formulates preconditions that shall guarantee that during the computation neither arithmetic overflows (property `bound`) nor unintended aliasing of arrays (property `separated`) occur. Under these precondition, Frama-C automatically verifies that the final `adjacent_difference` call just restores the original contents of `a` used for the initial `partial_sum` call.

```
/*@
  requires valid: \valid(a + (0..n-1));
  requires valid: \valid(b + (0..n-1));
  requires separated: \separated(a + (0..n-1), b + (0..n-1));
  requires bounds: AccumulateBounds(a, n+1);

  assigns a[0..n-1], b[0..n-1];

  ensures unchanged: Unchanged{Here,Pre}(a, n);
*/
void partial_sum_inverse(value_type* a, size_type n, value_type* b)
{
  partial_sum(a, n, b);
  adjacent_difference(b, n, a);
}
```

Listing 7.25: `partial_sum` and then `adjacent_difference`

7.7. Inverting adjacent_difference with partial_sum

In this section, we prove the converse property, viz. that applying `adjacent_difference`, and thereafter `partial_sum`, restores the original data array. Listing 7.26 expresses this property as a lemma on the level of ACSL predicates. It had to be proven interactively with Coq, by induction on `n`.

```
/*@
  lemma AdjacentDifferenceInverse{K,L}:
    \forallall value_type *a, *b, integer n;
      0 <= n ==>
        AdjacentDifference{K}(a, n, b) ==>
          Unchanged{K,L}(b, n) ==>
            PartialSum{L}(b, n, a) ==>
              Unchanged{K,L}(a, n);
*/
```

Listing 7.26: The lemma `AdjacentDifferenceInverse`

As in the case discussed in Section 7.6, we give a corresponding C function in order to account for possible arithmetic overflows and potential aliasing of data. Function `adjacent_difference_inverse`, shown in Listing 7.27, calls first `adjacent_difference` and then `partial_sum`. The contract of this function formulates preconditions that shall guarantee that during the computation neither arithmetic overflows (property `bound`) nor unintended aliasing of arrays (property `separated`) occur.

```
/*@
  requires valid:      \valid(a + (0..n-1));
  requires valid:      \valid(b + (0..n-1));
  requires separated:  \separated(a + (0..n-1), b + (0..n-1));
  requires bounds:     AdjacentDifferenceBounds(a, n+1);

  assigns a[0..n-1], b[0..n-1];

  ensures unchanged:   Unchanged{Here,Pre}(a, n);
*/
void adjacent_difference_inverse(value_type* a, size_type n, value_type* b)
{
  adjacent_difference(a, n, b);
  partial_sum(b, n, a);
}
```

Listing 7.27: `adjacent_difference` and then `partial_sum`

In order to improve the automatic verification rate of the function `adjacent_difference_inverse` we also supplied the following lemma (Listing 7.28) which captures the transitivity of the predicate `Unchanged`. Both the lemma itself and (with its additional help) the contract of `adjacent_difference_inverse` are proven by Frama-C without further manual intervention. This finishes the formal proof of our inversivity claims from Section 7.5.

```
/*@  
  lemma UnchangedTransitive{K,L,M}:  
    \forall a, integer n;  
      0 <= n ==>  
        Unchanged{K,L}(a, n) ==>  
        Unchanged{L,M}(a, n) ==>  
        Unchanged{K,M}(a, n);  
*/
```

Listing 7.28: The lemma `UnchangedTransitive`

8. The Stack data type

Originally, ACSL is tailored to the task of specifying and verifying one single C function at a time. However, in practice we are also faced with the task to implement a family of functions, usually around some sophisticated data structure, which have to obey certain rules of interdependence. In this kind of task, we are not interested in the properties of a single function (usually called “*implementation details*”), but in properties describing how several function play together (usually called “*abstract interface description*”, or “*abstract data type properties*”).

This chapter introduces a methodology to formally denote and verify the latter property sets using ACSL. For a more detailed discussion of our approach to the formal verification of `Stack` we refer to this thesis [14].

A *stack* is a data type that can hold objects and has the property that, if an object *a* is *pushed* on a stack *before* object *b*, then *a* can only be removed (*popped*) after *b*. A stack is, in other words, a *first-in, last-out* data type (see Figure 8.1). The *top* function of a stack returns the last element that has been pushed on a stack.




Figure 8.1.: Push and pop on a stack

We consider only stacks that have a finite *capacity*, that is, that can only hold a maximum number *c* of elements that is constant throughout their lifetime. This restriction allows us to define a stack without relying on dynamic memory allocation. When a stack is *created* or *initialized*, it contains no elements, i.e., its *size* is 0. The function *push* and *pop* increases and decreases the size of a stack by at most one, respectively.

8.1. Methodology overview

Figure 8.2 gives an overview of our methodology to specify and verify abstract data types (verification of one axiom shown only).

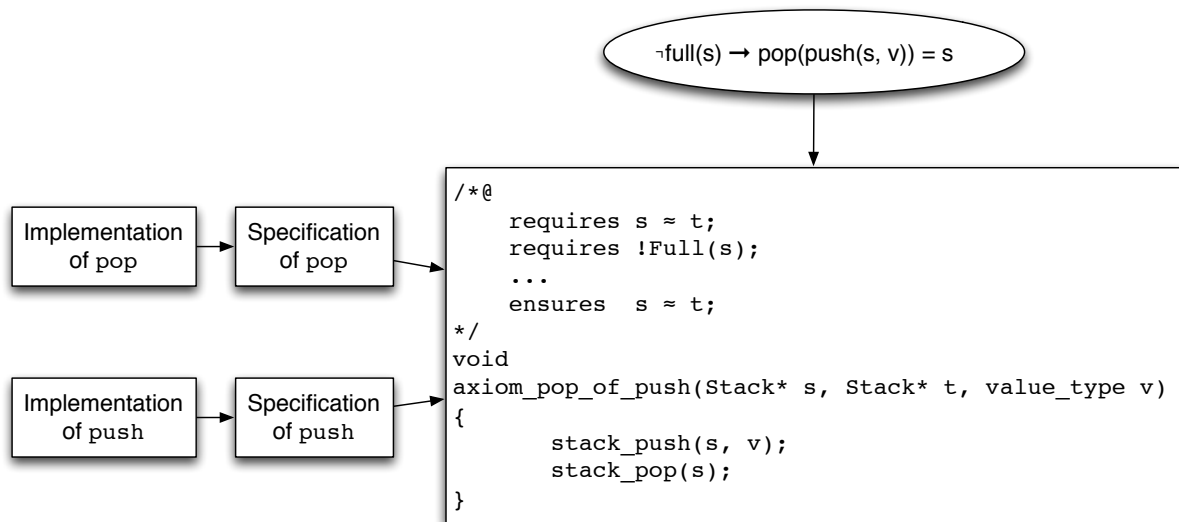


Figure 8.2.: Methodology Overview

What we will basically do is:

1. specify axioms about how the stack functions should interact with each other (Section 8.2),
2. define a basic implementation of C data structures (only one in our example, viz. `struct Stack`; see Section 8.3) and some invariants the instances of them have to obey (Section 8.4),
3. provide for each stack function an ACSL contract and a C implementation (Section 8.7),
4. verify each function against its contract (Section 8.7),
5. transform the axioms into ACSL-annotated C code (Section 8.8), and
6. verify that code, using access function contracts and data-type invariants as necessary (Section 8.8).

Section 8.5 provides an ACSL-predicate deciding whether two instances of a `struct Stack` are considered to be equal (indicated by “ \approx ” in Figure 8.2), while Section 8.6 gives a corresponding C implementation. The issue of an appropriate definition of equality of data instances is familiar to any C programmer who had to replace a faulty comparison `if (s1 == s2)` by the correct `if (strcmp(s1, s2) == 0)` to compare two strings `char *s1, *s2` for equality.

8.2. Stack axioms

To specify the interplay of the stack access functions, we use a set of axioms⁵⁸, all but one of them having the form of a conditional equation.

Let V denote an arbitrary type. We denote by S_c the type of stacks with capacity $c > 0$ of elements of type V . The aforementioned functions then have the following signatures.

$$\begin{aligned} \text{init} &: S_c \rightarrow S_c, \\ \text{push} &: S_c \times V \rightarrow S_c, \\ \text{pop} &: S_c \rightarrow S_c, \\ \text{top} &: S_c \rightarrow V, \\ \text{size} &: S_c \rightarrow \mathbb{N}. \end{aligned}$$

With \mathbb{B} denoting the *boolean* type we will also define two auxiliary functions

$$\begin{aligned} \text{empty} &: S_c \rightarrow \mathbb{B}, \\ \text{full} &: S_c \rightarrow \mathbb{B}. \end{aligned}$$

To qualify as a stack these functions must satisfy the following rules which are also referred to as *stack axioms*.

8.2.1. Stack initialization

After a stack has been initialized its size is 0.

$$\text{size}(\text{init}(s)) = 0. \quad (8.1)$$

The auxiliary functions *empty* and *full* are defined as follows

$$\text{empty}(s), \quad \text{iff} \quad \text{size}(s) = 0, \quad (8.2)$$

$$\text{full}(s), \quad \text{iff} \quad \text{size}(s) = c. \quad (8.3)$$

We expect that for every stack s the following condition holds

$$0 \leq \text{size}(s) \leq c. \quad (8.4)$$

8.2.2. Adding an element to a stack

To push an element v on a stack the stack must not be full. If an element has been pushed on an eligible stack, its size increases by 1

$$\text{size}(\text{push}(s, v)) = \text{size}(s) + 1, \quad \text{if} \quad \neg \text{full}(s). \quad (8.5)$$

Moreover, the element pushed on a stack is the top element of the resulting stack

$$\text{top}(\text{push}(s, v)) = v, \quad \text{if} \quad \neg \text{full}(s). \quad (8.6)$$

⁵⁸ There is an analogy in geometry: Euclid (e.g. [15]) invented the use of axioms there, but still kept definitions of *point*, *line*, *plane*, etc. Hilbert [16] recognized that the latter are not only unformalizable, but also unnecessary, and dropped them, keeping only the formal descriptions of relations between them.

8.2.3. Removing an element from a stack

An element can only be removed from a non-empty stack. If an element has been removed from an eligible stack the stack size decreases by 1

$$\text{size}(\text{pop}(s)) = \text{size}(s) - 1, \quad \text{if } \neg \text{empty}(s). \quad (8.7)$$

If an element is pushed on a stack and immediately afterwards an element is removed from the resulting stack then the final stack is equal to the original stack

$$\text{pop}(\text{push}(s, v)) = s, \quad \text{if } \neg \text{full}(s). \quad (8.8)$$

Conversely, if an element is removed from a non-empty stack and if afterwards the top element of the original stack is pushed on the new stack then the resulting stack is equal to the original stack.

$$\text{push}(\text{pop}(s), \text{top}(s)) = s, \quad \text{if } \neg \text{empty}(s). \quad (8.9)$$

8.2.4. A note on exception handling

We don't impose a requirement on $\text{push}(s, v)$ if s is a full stack, nor on $\text{pop}(s)$ or $\text{top}(s)$ if s is an empty stack. Specifying the behavior in such *exceptional* situations is a problem by its own; a variety of approaches is discussed in the literature. We won't elaborate further on this issue, but only give an example to warn about "innocent-looking" exception specifications that may lead to undesired results.

If we'd introduce an additional error value `err` in the element type V and require $\text{top}(s) = \text{err}$ if s is empty, we'd be faced with the problem of specifying the behavior of $\text{push}(s, \text{err})$. At first glance, it would seem a good idea to have `err` just been ignored by `push`, i.e. to require

$$\text{push}(s, \text{err}) = s. \quad (8.10)$$

However, we then could derive for any non-full and non-empty stack s , that

$$\begin{aligned} \text{size}(s) &= \text{size}(\text{pop}(\text{push}(s, \text{err}))) && \text{by 8.8} \\ &= \text{size}(\text{pop}(s)) && \text{as assumed in 8.10} \\ &= \text{size}(s) - 1 && \text{by 8.7} \end{aligned}$$

i.e. no such stacks could exist, or all `int` values would be equal.

8.3. The structure `Stack` and its associated functions

We now introduce one possible C implementation of the above axioms. It is centred around the C structure `Stack` shown in Listing 8.3.

```
struct Stack
{
    value_type* obj;

    size_type   capacity;

    size_type   size;
};

typedef struct Stack Stack;
```

Listing 8.3: Definition of type `Stack`

This struct holds an array `obj` of positive length called `capacity`. The capacity of a stack is the maximum number of elements this stack can hold. The field `size` indicates the number elements that are currently in the stack. See also Figure 8.4 which attempts to interpret this definition according to Figure 8.1.




Figure 8.4.: Interpreting the data structure `Stack`

Based on the stack functions from Section 8.2, we declare in Listing 8.5 the following functions as part of our Stack data type.

```

void      stack_init(Stack* s, value_type* a, size_type n);
bool      stack_equal(const Stack* s, const Stack* t);
size_type stack_size(const Stack* s);
bool      stack_empty(const Stack* s);
bool      stack_full(const Stack* s);
value_type stack_top(const Stack* s);
void      stack_push(Stack* s, value_type v);
void      stack_pop(Stack* s);

```

Listing 8.5: Declaration of functions of type Stack

Most of these functions directly correspond to methods of the C++ `std::stack` template class.⁵⁹ The function `stack_equal` corresponds to the comparison operator `==`, whereas one use of `stack_init` is to bring a stack into a well-defined initial state. The function `stack_full` has no counterpart in `std::stack`. This reflects the fact that we avoid dynamic memory allocation, while `std::stack` does not.

8.4. Stack invariants

Not every possible instance of type `Stack` is considered a valid one, e.g., with our definition of `Stack` in Listing 8.3, `Stack s = {{0, 0, 0, 0}, 4, 5}` is not. Below, we will define an ACSL-predicate `Valid` that discriminates valid and invalid instances.

Before, we introduce in Listing 8.6 the auxiliary logical function `Capacity` and `Size` which we can use in specifications to refer to the fields `capacity` and `size` of `Stack`, respectively. This listing also contains the logical function `Top` which defines the array element with index `size-1` as the top place of a stack. The reader can consider this as an attempt to hide implementation details from the specification.

```

/*@ logic size_type Capacity{L}(Stack* s) = s->capacity;

/*@ logic size_type Size{L}(Stack* s) = s->size;

/*@ logic value_type* Storage{L}(Stack* s) = s->obj;

/*@ logic value_type Top{L}(Stack* s) = s->obj[s->size-1];

```

Listing 8.6: The logical functions `Capacity`, `Size` and `Top`

We also introduce in Listing 8.7 two predicates that express the concepts of empty and full stacks by referring to a stack's size and capacity (see Equations (8.2) and (8.3)).

⁵⁹ See <http://www.sgi.com/tech/stl/stack.html>

```
//@ predicate Empty{L}(Stack* s) = Size(s) == 0;

//@ predicate Full{L}(Stack* s) = Size(s) == Capacity(s);
```

Listing 8.7: Predicates for empty and full stacks

There are some obvious invariants that must be fulfilled by every valid object of type `Stack`:

- The stack capacity shall be strictly greater than zero (an empty stack is ok but a stack that cannot hold anything is not useful).
- The pointer `obj` shall refer to an array of length `capacity`.
- The number of elements `size` of a stack must be non-negative and not greater than its capacity.

These invariants are formalized in the predicate `Valid` of Listing 8.8.

```
/*@
predicate Valid{L}(Stack* s) =
    \valid(s) &&
    0 < Capacity(s) &&
    0 <= Size(s) <= Capacity(s) &&
    \valid(Storage(s) + (0..Capacity(s)-1)) &&
    \separated(s, Storage(s) + (0..Capacity(s)-1));
*/
```

Listing 8.8: The predicate `Valid`

Note how the use of the previously defined logical functions and predicates allows us to define the stack invariant without directly referring to the fields of `Stack`. As we usually have to deal with a pointer `s` of type `Stack` we add the necessary `\valid(s)` to the predicate `Valid`.

8.5. Equality of stacks

Defining equality of instances of non-trivial data types, in particular in object-oriented languages, is not an easy task. The book *Programming in Scala*[17, Chapter 28] devotes to this topic a whole chapter of more than twenty pages. In the following two sections we give a few hints how ACSL and Frama-C can help to correctly define equality for a simple data type.

We consider two stacks as equal if they have the same size and if they contain the same objects. To be more precise, let s and t two pointers of type `Stack`, then we define the predicate `Equal` as in Listing 8.9.

```
/*@
  predicate Equal{S,T}(Stack* s, Stack* t) =
    Size{S}(s) == Size{T}(t) &&
    EqualRanges{S,T}(Storage{S}(s), Size{S}(s), Storage{T}(t));
*/
```

Listing 8.9: Equality of stacks

Our use of labels in Listing 8.9 makes the specification somewhat hard to read (in particular in the last line where we reuse the predicate `EqualRanges` from Page 29). However, this definition of `Equal` will allow us later to compare the same stack object at different points of a program. The logical expression `Equal{A,B}(s,t)` reads informally as: The stack object $*s$ at program point A equals the stack object $*t$ at program point B .

The reader might wonder why we exclude the capacity of a stack into the definition of stack equality. This approach can be motivated with the behavior of the method `capacity` of the class `std::vector<T>`. There, equal instances of type `std::vector<T>` may very well have different capacities.⁶⁰

If equal stacks can have different capacities then, according to our definition of the predicate `Full` in Listing 8.7, we can have to equal stacks where one is full and the other one is not.

A finer, but very important point in our specification of equality of stacks is that the elements of the arrays $s \rightarrow \text{obj}$ and $t \rightarrow \text{obj}$ are compared only up to $s \rightarrow \text{size}$ and *not* up to $s \rightarrow \text{capacity}$. Thus the two stacks s and t in Figure 8.10 are considered equal although there are obvious differences in their internal arrays.




Figure 8.10.: Example of two equal stacks

⁶⁰ See <http://www.cplusplus.com/reference/vector/vector/capacity>

If we define an equality relation ($=$) of objects for a data type such as `Stack`, we have to make sure that the following rules hold.

$$\text{reflexivity} \quad \forall s \in S : s = s, \quad (8.11a)$$

$$\text{symmetry} \quad \forall s, t \in S : s = t \implies t = s, \quad (8.11b)$$

$$\text{transitivity} \quad \forall s, t, u \in S : s = t \wedge t = u \implies s = u. \quad (8.11c)$$

Any relation that satisfies the conditions (8.11) is referred to as an *equivalence relation*. The mathematical set of all instances that are considered equal to some given instance s is called the equivalence class of s with respect to that relation.

Listing 8.11 shows a formalization of these three rules for the relation `Equal`; it can be automatically verified that they are a consequence of the definition of `Equal` in Listing 8.9.

```
/*@
lemma StackEqualReflexive{S} :
  \forallall Stack* s; Equal{S,S}(s, s);

lemma StackEqualSymmetric{S,T} :
  \forallall Stack *s, *t;
    Equal{S,T}(s, t) ==> Equal{T,S}(t, s);

lemma StackEqualTransitive{S,T,U}:
  \forallall Stack *s, *t, *u;
    Equal{S,T}(s, t) && Equal{T,U}(t, u) ==> Equal{S,U}(s, u);
*/
```

Listing 8.11: Equality of stacks is an equivalence relation

The two stacks in Figure 8.10 show that an equivalence class of `Equal` can contain more than one element.⁶¹ The stacks s and t in Figure 8.10 are also referred to as two *representatives* of the same equivalence class. In such a situation, the question arises whether a function that is defined on a set with an equivalence relation can be defined in such a way that its definition is *independent of the chosen representatives*.⁶² We ask, in other words, whether the function is *well-defined* on the set of all equivalence classes of the relation `Equal`.⁶³ The question of well-definition will play an important role when verifying the functions of the `Stack` (see Section 8.7).

⁶¹ This is a common situation in mathematics. For example, the equivalence class of the rational number $\frac{1}{2}$ contains infinitely many elements, viz. $\frac{1}{2}, \frac{2}{4}, \frac{7}{14}, \dots$

⁶² This is why mathematicians have to *prove* that $\frac{1}{2} + \frac{3}{5}$ equals $\frac{7}{14} + \frac{3}{5}$.

⁶³ See <http://en.wikipedia.org/wiki/Well-definition>.

8.6. Runtime equality of stacks

The function `stack_equal` is the C equivalent for the `Equal` predicate. The specification of `stack_equal` is shown in Listing 8.12. Note that this specifications explicitly refers to valid stacks.

```
/*@  
  requires Valid(s);  
  requires Valid(t);  
  
  assigns \nothing;  
  
  ensures \result == 1 <==> Equal{Here,Here}(s, t);  
  ensures \result == 0 <==> !Equal{Here,Here}(s, t);  
*/  
bool stack_equal(const Stack* s, const Stack* t);
```

Listing 8.12: Specification of `stack_equal`

The implementation of `stack_equal` in Listing 8.13 compares two stacks according to the same rules of predicate `Equal`.

```
bool stack_equal(const Stack* s, const Stack* t)  
{  
  return (s->size == t->size) && equal(s->obj, s->size, t->obj);  
}
```

Listing 8.13: Implementation of `stack_equal`

8.7. Verification of stack functions

In this section we verify the functions `stack_init` (Section 8.7.1), `stack_size` (Section 8.7.2), `stack_empty` (Section 8.7.3), `stack_full` (Section 8.7.4), `stack_top` (Section 8.7.5), and `stack_push` (Section 8.7.6) `stack_pop` (Section 8.7.7), of the data type `Stack`. To be more precise, we provide for each of function `stack_foo`:

- an ACSL specification of `stack_foo`
- a C implementation of `stack_foo`
- a C function `stack_foo_wd`⁶⁴ accompanied by an ACSL contract that expresses that the implementation of `stack_foo` is well-defined. Figure 8.14 shows our methodology for the verification of well-definition in the `pop` example, (\approx) again indicating the user-defined `Stack` equality.

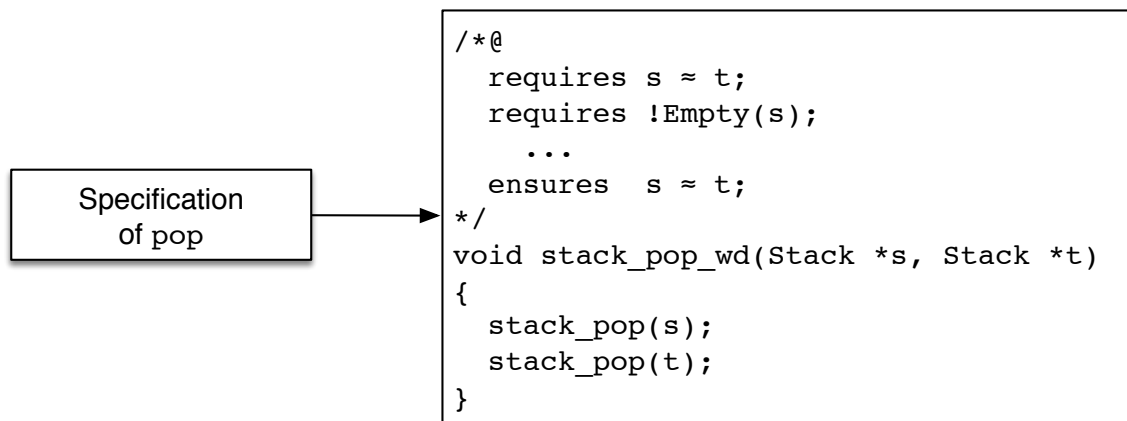


Figure 8.14.: Methodology for the verification of well-definition

Note that the specifications of the various functions will explicitly refer to the *internal state* of `Stack`. In Section 8.8 we will show that the *interplay* of these functions satisfy the stack axioms from Section 8.2.

⁶⁴ The suffix `_wd` stands for *well definition*

8.7.1. The function `stack_init`

Listing 8.15 shows the ACSL specification of `stack_init`. Note that our specification of the post-conditions contains a redundancy because a stack is empty if and only if its size is zero.

```
/*@
  requires \valid(s);
  requires 0 < capacity;
  requires \valid(storage + (0..capacity-1));
  requires \separated(s, storage + (0..capacity-1));

  assigns s->obj;
  assigns s->capacity;
  assigns s->size;

  ensures Valid(s);
  ensures Capacity(s) == capacity;
  ensures Size(s) == 0;
  ensures Empty(s);
  ensures Storage(s) == storage;
*/
void stack_init(Stack* s, value_type* storage, size_type capacity);
```

Listing 8.15: Specification of `stack_init`

Listing 8.15 shows the implementation of `stack_init`. It simply initializes `obj` and `capacity` with the respective value of the array and sets the field `size` to zero.

```
void stack_init(Stack* s, value_type* storage, size_type capacity)
{
    s->obj      = storage;
    s->capacity = capacity;
    s->size     = 0;
}
```

Listing 8.16: Implementation of `stack_init`

8.7.2. The function `stack_size`

The function `stack_size` is the runtime version of the logical function `Size` from Listing 8.6 on Page 126. The specification of `stack_size` in Listing 8.17 simply states that `stack_size` produces the same result as `Size`.

```
/*@
  requires Valid(s);

  assigns \nothing;

  ensures \result == Size(s);
*/
size_type stack_size(const Stack* s);
```

Listing 8.17: Specification of `stack_size`

As in the definition of the logical function `Size` the implementation of `stack_size` in Figure 8.18 simply returns the field `size`.

```
size_type stack_size(const Stack* s)
{
  return s->size;
}
```

Listing 8.18: Implementation of `stack_size`

Listing 8.19 shows our check whether `stack_size` is well-defined. Since `stack_size` neither modifies the state of its `Stack` argument nor that of any global variable we only check whether it produces the same result for equal stacks. Note that we simply may use operator `==` to compare integers since we didn't introduce a nontrivial equivalence relation on that data type.

```
/*@
  requires Valid(s) && Valid(t);
  requires Equal{Here,Here}(s, t);

  assigns \nothing;

  ensures \result;
*/
bool stack_size_wd(const Stack* s, const Stack* t)
{
  return stack_size(s) == stack_size(t);
}
```

Listing 8.19: Well-definition of `stack_size`

8.7.3. The function `stack_empty`

The function `stack_empty` is the runtime version of the predicate `Empty` from Listing 8.7 on Page 127.

```
/*@
  requires Valid(s);

  assigns \nothing;

  ensures \result == 1 <==> Empty(s);
  ensures \result == 0 <==> !Empty(s);
*/
bool stack_empty(const Stack* s);
```

Listing 8.20: Specification of `stack_empty`

As in the definition of the predicate `Empty` the implementation of `stack_empty` in Figure 8.21 simply checks whether the size of the stack is zero.

```
bool stack_empty(const Stack* s)
{
  return stack_size(s) == 0;
}
```

Listing 8.21: Implementation of `stack_empty`

Listing 8.22 shows our check whether `stack_empty` is well-defined.

```
/*@
  requires Valid(s);
  requires Valid(t);
  requires Equal{Here,Here}(s, t);

  assigns \nothing;

  ensures \result;
*/
bool stack_empty_wd(const Stack* s, const Stack* t)
{
  return stack_empty(s) == stack_empty(t);
}
```

Listing 8.22: Well-definition of `stack_empty`

8.7.4. The function `stack_full`

The function `stack_full` is the runtime version of the predicate `Full` from Listing 8.7 on Page 127.

```
/*@
  requires Valid(s);

  assigns \nothing;

  ensures \result == 1 <==> Full(s);
  ensures \result == 0 <==> !Full(s);
*/
bool stack_full(const Stack* s);
```

Listing 8.23: Specification of `stack_full`

As in the definition of the predicate `Full` the implementation of `stack_full` in Figure 8.24 simply checks whether the size of the stack equals its capacity.

```
bool stack_full(const Stack* s)
{
  return stack_size(s) == s->capacity;
}
```

Listing 8.24: Implementation of `stack_full`

Note that with our definition of stack equality (Section 8.5) there can be equal stack with different capacities. Accordingly, there can exist equal stacks where one is full while the other is not.

8.7.5. The function `stack_top`

The function `stack_top` is the runtime version of the logical function `Top` from Listing 8.6 on Page 126. The specification of `stack_top` in Listing 8.25 simply states that for non-empty stacks `stack_top` produces the same result as `Top` which in turn just returns the element `obj[size-1]` of `Stack`.

```
/*@
  requires Valid(s);

  assigns \nothing;

  ensures !Empty(s) ==> \result == Top(s);
*/
value_type stack_top(const Stack* s);
```

Listing 8.25: Specification of `stack_top`

For a non-empty stack the implementation of `stack_top` in Figure 8.26 simply returns the element `obj[size-1]`. Note that our implementation of `stack_top` does not crash when it is applied to an empty stack. In this case we return the first element of the internal, non-empty array `obj`. This is consistent with our specification of `stack_top` which only refers to non-empty stacks.

```
value_type stack_top(const Stack* s)
{
  if (!stack_empty(s)) {
    return s->obj[s->size - 1];
  } else {
    return s->obj[0];
  }
}
```

Listing 8.26: Implementation of `stack_top`

Listing 8.27 shows our check whether `stack_top` well-defined for non-empty stacks.

```
/*@
  requires Valid(s) && !Empty(s);
  requires Valid(t) && !Empty(t);
  requires Equal{Here,Here}(s, t);

  assigns \nothing;

  ensures \result;
*/
bool stack_top_wd(const Stack* s, const Stack* t)
{
  return stack_top(s) == stack_top(t);
}
```

Listing 8.27: Well-definition of `stack_top`

Since our axioms in Section 8.2 did not impose any behavior on the behavior of `stack_top` for empty stacks, we prove the well-definition of `stack_top` only for nonempty stacks.

8.7.6. The function `stack_push`

Listing 8.28 shows the ACSL specification of the function `stack_push`. In accordance with Axiom (8.5), `stack_push` is supposed to increase the number of elements of a non-full stack by one. The specification also demands that the value that is pushed on a non-full stack becomes the top element of the resulting stack (see Axiom (8.6)).

```
/*@
  requires Valid(s);

  assigns s->size;
  assigns s->obj[s->size];

  behavior not_full:
    assumes !Full(s);

    assigns s->size;
    assigns s->obj[s->size];

    ensures Valid(s);
    ensures Size(s) == Size{Old}(s) + 1;
    ensures Top(s) == v;
    ensures !Empty(s);
    ensures Unchanged{Pre,Here}(Storage(s), Size{Pre}(s));
    ensures Storage(s) == Storage{Old}(s);
    ensures Capacity(s) == Capacity{Old}(s);

  behavior full:
    assumes Full(s);

    assigns \nothing;

    ensures Valid(s);
    ensures Full(s);
    ensures Unchanged{Pre,Here}(Storage(s), Size(s));
    ensures Size(s) == Size{Old}(s);
    ensures Storage(s) == Storage{Old}(s);
    ensures Capacity(s) == Capacity{Old}(s);

  complete behaviors;
  disjoint behaviors;
*/
void stack_push(Stack* s, value_type v);
```

Listing 8.28: Specification of `stack_push`

The implementation of `stack_push` is shown in Listing 8.29. It checks whether its argument is a non-full stack in which case it increases the field `size` by one but only after it has assigned the function argument to the element `obj[size]`.

```
void stack_push(Stack* s, value_type v)
{
  if (!stack_full(s)) {
    s->obj[s->size++] = v;
  }
}
```

Listing 8.29: Implementation of `stack_push`

The function `stack_push` does not return a value but rather modifies its argument. For the well-definition of `stack_push` we therefore check whether it turns equal stacks into equal stacks. However, equality of the stack arguments is not sufficient for a proof that `stack_push` is well-defined. We must also ensure that there is no *aliasing* between the two stacks. Otherwise modifying one stack could modify the other stack in unexpected ways. In order to express that there is no aliasing between two stacks, we define in Listing 8.30 the predicate `Separated`.

```
/*@
  predicate Separated(Stack* s, Stack* t) =
    \separated(s, s->obj + (0..s->capacity-1),
              t, t->obj + (0..t->capacity-1));
*/
```

Listing 8.30: The predicate `Separated`

Listing 8.31 shows our formalization of the well-definition for `stack_push`.

```
/*@
  requires valid:      Valid(s) && Valid(t);
  requires equal:      Equal{Here,Here}(s, t);
  requires not_full:   !Full(s) && !Full(t);
  requires separated:  Separated(s, t);

  assigns s->size, s->obj[s->size];
  assigns t->size, t->obj[t->size];

  ensures valid:      Valid(s) && Valid(t);
  ensures equal:      Equal{Here,Here}(s, t);
*/
void stack_push_wd(Stack* s, Stack* t, value_type v)
{
  stack_push(s, v);
  stack_push(t, v);
  //@ assert top:      Top(s) == v;
  //@ assert top:      Top(t) == v;
  //@ assert equal:    EqualRanges{Here,Here}(Storage(s), Size{Pre}(s), Storage(t));
}
```

Listing 8.31: Well-definition of `stack_push`

In order to achieve an automatic verification of the well-definition of `stack_push` we added in Listing 8.31 the assertions `top` and `equal` and introduced the lemma `StackPushEqual` from Listing 8.32.

```
/*@
  lemma StackPushEqual{K,L}:
    \forallall Stack *s, *t;
      Equal{K,K}(s,t) ==>
      Size{L}(s) == Size{K}(s) + 1 ==>
      Size{L}(s) == Size{L}(t) ==>
      Top{L}(s) == Top{L}(t) ==>
      EqualRanges{L,L}(Storage{L}(s), Size{K}(s), Storage{L}(t)) ==>
      Equal{L,L}(s,t);
*/
```

Listing 8.32: The lemma `StackPushEqual`

8.7.7. The function `stack_pop`

Listing 8.33 shows the ACSL specification of the function `stack_pop`. In accordance with Axiom (8.7) `stack_pop` is supposed to reduce the number of elements in a non-empty stack by one. In addition to the requirements imposed by the axioms, our specification demands that `stack_pop` changes no memory location if it is applied to an empty stack.

```
/*@
  requires Valid(s);

  assigns s->size;

  ensures Valid(s);

  behavior not_empty:
    assumes !Empty(s);

    assigns s->size;

    ensures Size(s) == Size{Old}(s) - 1;
    ensures !Full(s);
    ensures Unchanged{Pre,Here}(Storage(s), Size(s));
    ensures Storage(s) == Storage{Old}(s);
    ensures Capacity(s) == Capacity{Old}(s);

  behavior empty:
    assumes Empty(s);

    assigns \nothing;

    ensures Empty(s);
    ensures Unchanged{Pre,Here}(Storage(s), Size(s));
    ensures Size(s) == Size{Old}(s);
    ensures Storage(s) == Storage{Old}(s);
    ensures Capacity(s) == Capacity{Old}(s);

  complete behaviors;
  disjoint behaviors;
*/
void stack_pop(Stack* s);
```

Listing 8.33: Specification of `stack_pop`

The implementation of `stack_pop` is shown in Listing 8.34. It checks whether its argument is a non-empty stack in which case it decreases the field `size` by one.

```
void stack_pop(Stack* s)
{
  if (!stack_empty(s)) {
    --s->size;
  }
}
```

Listing 8.34: Implementation of `stack_pop`

Listing 8.35 shows our check whether `stack_pop` is well-defined. As in the case of `stack_push` we use the predicate `Separated` (Listing 8.30) in order to express that there is no aliasing between the two stack arguments.

```
/*@
  requires Valid(s);
  requires Valid(t);
  requires Equal{Here,Here}(s, t);
  requires Separated(s, t);

  assigns s->size;
  assigns t->size;

  ensures Valid(s);
  ensures Valid(t);
  ensures Equal{Here,Here}(s, t);
*/
void stack_pop_wd(Stack* s, Stack* t)
{
  stack_pop(s);
  stack_pop(t);
}
```

Listing 8.35: Well-definition of `stack_pop`

8.8. Verification of stack axioms

In this section we show that the stack functions defined in Section 8.7 satisfy the stack Axioms of Section 8.2.

The annotated code has been obtained from the axioms in a fully systematical way. In order to transform a condition equation $p \rightarrow s = t$:

- Generate a clause `requires p`.
- Generate a clause `requires x1 == ... == xn` for each variable x with n occurrences in s and t .
- Change the i -th occurrence of x to x_i in s and t .
- Translate both terms s and t to reversed polish notation.
- Generate a clause `ensures y1 == y2`, where y_1 and y_2 denote the value corresponding to the translated s and t , respectively.

This makes it easy to implement a tool that does the translation automatically, but yields a slightly longer contract in our example.

8.8.1. Resetting a stack

Our formulation in ACSL/C of the Axiom in Equation (8.1) on Page 123 is shown in Listing 8.36.

```
/*@
  requires \valid(s);
  requires 0 < n;
  requires \valid(a + (0..n-1));
  requires \separated(s, a + (0..n-1));

  assigns s->obj, s->capacity, s->size;

  ensures Valid(s);
  ensures \result == 0;
*/
size_type axiom_size_of_init(Stack* s, value_type* a, size_type n)
{
  stack_init(s, a, n);
  return stack_size(s);
}
```

Listing 8.36: Specification of Axiom (8.1)

8.8.2. Adding an element to a stack

Axioms (8.5) and (8.6) describe the behavior of a stack when an element is added.

```
/*@
  requires Valid(s);
  requires !Full(s);

  assigns s->size;
  assigns s->obj[s->size];

  ensures Valid(s);
  ensures \result == Size{Old}(s) + 1;
*/
size_type axiom_size_of_push(Stack* s, value_type v)
{
  stack_push(s, v);
  return stack_size(s);
}
```

Listing 8.37: Specification of Axiom (8.5)

Except for the `assigns` clauses, the ACSL-specification refers only to encapsulating logical functions and predicates defined in Section 8.4. If ACSL would provide a means to define encapsulating logical functions returning also sets of memory locations, the expressions in `assigns` clauses would not need to refer to the details of our `Stack` implementation.⁶⁵ As an alternative, `assigns` clauses could be omitted, as long as the proofs are only used to convince a human reader.

```
/*@
  requires Valid(s);
  requires !Full(s);

  assigns s->size;
  assigns s->obj[s->size];

  ensures \result == v;
*/
value_type axiom_top_of_push(Stack* s, value_type v)
{
  stack_push(s, v);
  return stack_top(s);
}
```

Listing 8.38: Specification of Axiom (8.6)

⁶⁵ In [9, § 2.3.4], a powerful sublanguage to build memory location set expressions is defined. We will explore its capabilities in a later version.

8.8.3. Removing an element from a stack

This section shows the Listings for Axioms 8.7, 8.8 and 8.9 which describe the behavior of a stack when an element is removed.

```
/*@
  requires Valid(s) && !Empty(s);
  assigns s->size;
  ensures \result == Size{Old}(s) - 1;
*/
size_type axiom_size_of_pop(Stack* s)
{
    stack_pop(s);
    return stack_size(s);
}
```

Listing 8.39: Specification of Axiom (8.7)

```
/*@
  requires Valid(s) && !Full(s);
  assigns s->size, s->obj[s->size];
  ensures Equal{Pre,Here}(s, s);
*/
void axiom_pop_of_push(Stack* s, value_type v)
{
    stack_push(s, v);
    stack_pop(s);
}
```

Listing 8.40: Specification of Axiom (8.8)

```
/*@
  requires Valid(s) && !Empty(s);
  assigns s->size, s->obj[s->size-1];
  ensures Equal{Here,Old}(s, s);
*/
void axiom_push_of_pop_top(Stack* s)
{
    const value_type val = stack_top(s);
    stack_pop(s);
    stack_push(s, val);
}
```

Listing 8.41: Specification of Axiom (8.9)

9. Results of formal verification

In this chapter we introduce the formal verification tools used in this tutorial. We will afterwards present to what extent the examples from Chapters 3–8 could be deductively verified.

Within Frama-C, the WP plug-in [2] enables deductive verification of C programs that have been annotated with the ANSI/ISO-C Specification Language (ACSL)[1]. The WP plug-in uses weakest precondition computations to generate proof obligations. To formally prove the ACSL properties, these proof obligations can be submitted to external automatic theorem provers or interactive proof assistants. The precise settings for WP and the associated provers that we used in this release we refer to Page 5.

For each algorithm we list in the following tables the number of generated verification conditions (VC), as well as the percentage of proven verification conditions. The tables show that all verification conditions could be verified. Please note that the number of proven verification conditions do *not* reflect on the quality/strength of the individual provers. The reason for that is that we “pipe” each verification condition sequentially through a list of provers. If one prover succeeds, then the remaining provers are not called.

The tables show that the majority of verification conditions could be verified by automatic provers. In Table 9.1 we explicitly list the ACSL lemmas that required induction proofs performed with Coq.

ACSL Lemma	Listing
CountBounds	3.25
CountMonotonic	3.25
CountUnion	3.25
RemoveCountMonotonic	6.31
AdjacentDifferenceInverse	7.26

Table 9.1.: ACSL lemmas that were proved with Coq

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
equal	3.1	18	18	100	8	0	10	0	0
equal (IsEqual)	3.1	15	15	100	7	0	8	0	0
equal (mismatch)	3.2	9	9	100	6	0	3	0	0
mismatch	3.2	20	20	100	9	0	11	0	0
find	3.3	19	19	100	8	0	11	0	0
find (2)	3.4	19	19	100	9	0	10	0	0
find_first_of	3.5	28	28	100	16	0	12	0	0
adjacent_find	3.6	23	23	100	11	0	11	1	0
search	3.7	33	33	100	19	0	14	0	0
count	3.8	19	19	100	8	0	11	0	0

Table 9.2.: Results for non-mutating algorithms

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
properties of operator <	4.1	6	6	100	4	0	2	0	0
max_element	4.2	25	25	100	13	0	12	0	0
max_element (2)	4.3	25	25	100	12	0	13	0	0
max_seq	4.4	8	8	100	5	0	3	0	0
min_element	4.5	25	25	100	12	0	13	0	0

Table 9.3.: Results for maximum and minimum algorithms

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
lower_bound	5.1	22	22	100	10	0	12	0	0
upper_bound	5.2	22	22	100	8	0	14	0	0
equal_range	5.3	20	20	100	15	0	5	0	0
binary_search	5.4	10	10	100	8	0	2	0	0
binary_search (2)	5.4	11	11	100	8	0	3	0	0

Table 9.4.: Results for binary search algorithms

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
fill	6.3	14	14	100	4	0	10	0	0
swap	6.2	8	8	100	8	0	0	0	0
swap_ranges	6.4	27	27	100	5	0	22	0	0
copy	6.5	18	18	100	4	0	14	0	0
reverse_copy	6.6	21	21	100	4	0	17	0	0
reverse	6.7	29	29	100	7	0	22	0	0
rotate_copy	6.8	19	19	100	4	0	15	0	0
replace_copy	6.9	22	22	100	6	0	16	0	0
replace	6.10	19	19	100	4	0	15	0	0
remove_copy	6.11	86	86	100	52	0	33	1	0
remove_copy (2)	6.12	108	108	100	56	1	47	0	4

Table 9.5.: Results for mutating algorithms

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
iota	7.1	18	18	100	7	0	11	0	0
accumulate	7.2	16	16	100	6	0	10	0	0
inner_product	7.3	21	21	100	6	0	12	3	0
partial_sum	7.4	41	41	100	8	2	30	1	0
adjacent_difference	7.5	35	35	100	10	0	25	0	0
partial_sum_inverse	7.6	20	20	100	4	0	16	0	0
adjacent_difference_inverse	7.7	27	27	100	3	1	21	1	1

Table 9.6.: Results for numeric algorithms

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
stack_equal	8.6	20	20	100	7	0	13	0	0
stack_init	8.7.1	16	16	100	3	0	13	0	0
stack_size	8.7.2	8	8	100	1	0	7	0	0
stack_empty	8.7.3	12	12	100	5	0	7	0	0
stack_full	8.7.4	13	13	100	5	0	8	0	0
stack_top	8.7.5	18	18	100	6	0	12	0	0
stack_push	8.7.6	45	45	100	28	0	17	0	0
stack_pop	8.7.7	34	34	100	20	0	14	0	0

Table 9.7.: Results for Stack functions

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
stack_size_wd	8.7.2	14	14	100	8	0	6	0	0
stack_empty_wd	8.7.3	14	14	100	8	0	6	0	0
stack_top_wd	8.7.5	14	14	100	8	0	6	0	0
stack_push_wd	8.7.6	17	17	100	3	0	13	1	0
stack_pop_wd	8.7.7	14	14	100	6	0	8	0	0

Table 9.8.: Results for the well-definition of the Stack functions

Algorithm	Section	VCs			Individual Provers				
		All	Proven	(%)	Qed	Alt-Ergo	CVC4	Z3	Coq
axiom_size_of_init	8.8.1	17	17	100	12	0	5	0	0
axiom_size_of_push	8.8.2	14	14	100	9	0	5	0	0
axiom_top_of_push	8.8.2	13	13	100	8	0	5	0	0
axiom_pop_of_push	8.8.3	12	12	100	6	0	6	0	0
axiom_size_of_pop	8.8.3	13	13	100	8	0	5	0	0
axiom_push_of_pop_top	8.8.3	17	17	100	9	0	7	1	0

Table 9.9.: Results for Stack axioms

A. Changes in previous releases

This chapter describes the changes in previous versions of this document. For the most recent changes see Section .

The version numbers of this document are related to the versioning of Frama-C [3]. The versions of Frama-C are named consecutively after the elements of the periodic table. Therefore, our version numbering (X.Y.Z) are constructed as follows:

X the major number of our tutorial is the atomic number⁶⁶ of the chemical element after which Frama-C is named.

Y the Frama-C subrelease number

Z the subrelease number of this tutorial

A.1. New in Version 11.1.1 (June 2015)

- add Chapter 7 on numeric algorithms
 - move `iota` algorithm to numeric algorithms (Section 7.1)
 - add `accumulate` algorithm (Section 7.2)
 - add `inner_product` algorithm (Section 7.3)
 - add `partial_sum` algorithm (Section 7.4)
 - add `adjacent_difference` algorithm (Section 7.5)

A.2. New in Version 11.1.0 (March 2015)

- Use built-in predicates `\valid` and `\valid_read` instead of `IsValidRange`.
- Simplify loop invariants of `find_first_of`.
- Replace two loop invariants of `remove_copy` by ACSL lemmas.
- Rename several predicates
 - `IsEqual` \mapsto `EqualRanges`.
 - `IsMaximum` \mapsto `MaxElement`.
 - `IsMinimum` \mapsto `MinElement`.
 - `Reverse` \mapsto `Reversed`.
 - `IsSorted` \mapsto `Sorted`.

⁶⁶See http://en.wikipedia.org/wiki/Atomic_number

- Several changes for `Stack`:
 - Rename `Stack` functions from `foo_stack` to `stack_foo`.
 - Equality of stacks now ignores the `capacity` field. This is similar to how equality for objects of type `std::vector<T>` is defined. As a consequence `stack_full` is not well-defined any more. Other stack functions are not effected.
 - Remove all assertions from stack functions (including in axioms).
 - Describe predicate `Separated` in text.

A.3. New in Version 10.1.1 (January 2015)

- use option `-wp-split` to create simpler (but more) proof obligations
- simplify definition of predicate `Count`
- add new predicates for lower and upper bounds of ranges and use it in
 - `max_element`
 - `min_element`
 - `lower_bound`
 - `upper_bound`
 - `equal_range`
 - `fill`
- use a new auxiliary assertion in `equal_range` to enable the complete *automatic* verification of this algorithm
- add predicate `Unchanged` and use it to simplify the specification of several algorithms
 - `swap_ranges`
 - `reverse`
 - `remove_copy`
 - `stack_push` and `stack_push_wd`
 - `stack_pop` and `stack_pop_wd`
- add predicate `Reverse` and use it for more concise specifications of
 - `reverse_copy`
 - `reverse`
- several changes in the two versions of `remove_copy`
 - use predicate `HasValue` instead of logic function `Count`
 - add predicate `PreserveCount`
 - reformulate logic function `RemoveCount`
 - add predicate `StableRemove`

- add predicate `RemoveCountMonotonic`
- add predicate `RemoveCountJump`
- use overloading in ACSL to create shorter logic names for `Stack`
- remove unnecessary labels in several `Stack` functions

A.4. New in Version 10.1.0 (September 2014)

- remove additional labels in the `assumes` clauses of some stack function that were necessary due to an error in Oxygen
- provide a second version of `remove_copy` in order to explain the specification of the *stability* of the algorithms
- coarsen loop assigns of mutating algorithms
- temporarily remove the `unique_copy` algorithm

A.5. New in Version 9.3.1 (not published)

- specify bounds of the return value of `count` and fix reads clause of `Count` predicate
- use an auxiliary function `make_pair` in the implementation of `equal_range`
- provide more precise loop assigns clauses for the mutating algorithms
 - simplify implementation of `fill`
 - removed the `ensures \valid(p)` clause in specification of `swap`
 - simplify implementation of `swap_ranges`
 - simplify implementation of `copy`
 - fix implementation of `reverse_copy` after discovering an undefined behavior
 - new implementation of `reverse` that uses a simple `for`-loop
 - simplify implementation of `replace_copy`
 - refactor specification and simplify implementation of `remove_copy`
- remove work-around with `Pre`-label in `assumes` clauses of `stack_push` and `stack_pop`

A.6. New in Version 9.3.0 (December 2013)

- adjustments for *Fluorine* release of Frama-C
- `swap` now ensures that its pointer arguments are valid after the function has been called
- change definition of `size_type` to `unsigned int`
- change implementation of the `iota` algorithm . The content of the field `a` is calculated by increasing the value `val` instead of `sum val+i`.

- change implementation of `fill`.
- The specification/implementation of `Stack` has been revised by Kim Völlinger [14] and now has a much better verification rate.

A.7. New in Version 8.1.0 (not published)

- simplified specification and loop annotations of `replace_copy`
- add binary search variant `equal_range`
- greatly simplified specification of `remove_copy` by using the logic function `Count`
- remove chapter on heap operations

A.8. New in Version 7.1.1 (August 2012)

- improvements with respect to several suggestions and comments of Yannick Moy, e.g., specification refinements of `remove_copy`, `reverse_copy` and `iota`
- restricted verification of algorithms to WP with Alt-Ergo
- replaced deprecated `\valid_range` by `\valid` in definition of `IsValidRange`
- fixed inconsistencies in the description of the `Stack` data type
- binary search algorithms can now be proven without additional axioms for integer division
- changed axioms into lemmas to document that provability is expected, even if not currently granted
- adopted new Fraunhofer logo and contact email

A.9. New in Version 7.1.0 (December 2011)

- changed to Frama-C Nitrogen
- changed to Why 2.30
- discussed both plug-ins WP and Jessie
- removed `swap_values` algorithm

A.10. New in Version 6.1.0 (not published)

- changed definition of `Stack`
- renamed `reset_stack` to `init_stack`

A.11. New in Version 5.1.1 (February 2011)

- prepared algorithms for checking by the new WP plug-in of Frama-C

- changed to Alt-Ergo Version 0.92, Z3 Version 2.11 and Why 2.27
- added List of user-defined predicates and logic functions
- added remarks on the relation of logical values in C and ACSL
- rewrote section on `equal` and `mismatch`
- used a simpler logical function to count elements in an array
- added `search` algorithm
- added chapter to unite the maximum/minimum algorithms
- added chapter for the new `lower_bound`, `upper_bound` and `binary_search` algorithms
- added `swap_values` algorithm
- used `IsEqual` predicate for `swap_ranges` and `copy`
- added `reverse_copy` and `reverse` algorithms
- added `rotate_copy` algorithm
- added `unique_copy` algorithm
- added chapter on specification of the data type `Stack`

A.12. New in Version 5.1.0 (May 2010)

- adaption to Frama-C Boron and Why 2.26 releases
- changed from the `-jessie-no-regions` command-line option to using the pragma `SeparationPolicy(value)`

A.13. New in Version 4.2.2 (May 2010)

- changed to latest version of CVC3 2.2
- added additional remarks to our implementation of `find_first_of`
- changed `size_type` (`int`) to `integer` in all specifications
- removed casts in `fill` and `iota`
- renamed `is_valid_range` as `IsValidRange`
- renamed `has_value` as `HasValue`
- renamed predicate `all_equal` as `IsEqual`
- extended timeout to 30 sec.

A.14. New in Version 4.2.1 (April 2010)

- added alternative specification of `remove_copy` algorithm that uses ghost variables

- added Chapter on heap operations
- added `mismatch` algorithm
- moved algorithms `adjacent_find` and `min_element` from the appendix to chapter on non-mutating algorithms
- added typedefs `size_type` and `value_type` and used them in all algorithms
- renamed `is_valid_int_range` as `is_valid_range`

A.15. New in Version 4.2.0 (January 2010)

- complete rewrite of previous release
- adaption to Frama-C Beryllium 2 release

Bibliography

- [1] ANSI/ISO C Specification Language. <http://frama-c.com/acsl.html>.
- [2] WP Plug-in. <http://frama-c.com/wp.html>.
- [3] Frama-C Software Analyzers. <http://frama-c.com>.
- [4] CEA LIST, Laboratory of Applied Research on Software-Intensive Technologies. http://www-list.cea.fr/gb/index_gb.htm.
- [5] INRIA-Saclay, French National Institute for Research in Computer Science and Control . <http://www.inria.fr/saclay/>.
- [6] LRI, Laboratory for Computer Science at Université Paris-Sud. <http://www.lri.fr/>.
- [7] Fraunhofer-Institut für Offene Kommunikationssysteme (FOKUS). <http://www.fokus.fraunhofer.de>.
- [8] Virgile Prevosto. ACSL Mini-Tutorial. <http://frama-c.com/download/acsl-tutorial.pdf>.
- [9] Patrick Baudin, Pascal Cuoq, Jean-Christophe Filliâtre, Claude Marché, Benjamin Monate, Yannick Moy, and Virgile Prevosto. ACSL 1.10 Implementation in Magnesium-20151002. <http://frama-c.com/download/acsl-implementation-Magnesium-20151002.pdf>, October 2015.
- [10] Programming languages – C, Committee Draft. <http://www.open-std.org/JTC1/SC22/WG14/www/docs/n1362.pdf>, 2009.
- [11] C.A.R. Hoare. An axiomatic basis for computer programming. *Communications of the ACM*, 12:576–583, 1969.
- [12] Robert W. Floyd. Assigning meanings to programs. In J. T. Schwartz, editor, *Proc. Symposium on Applied Mathematics*, volume 19 of *Mathematical Aspects of Computer Science*, pages 19–32, Providence, RI, 1967. American Mathematical Society.
- [13] Donald E. Knuth, James H. Morris, and Vaughan R. Pratt. Fast pattern matching in strings. *SIAM J Comput*, 6(2):323–350, Jun 1977.
- [14] Kim Völlinger. *Einsatz des Beweisassistenten Coq zur deduktiven Programmverifikation*. Diplomarbeit, Humboldt Universität zu Berlin, Germany, August 2013.
- [15] Richard Fitzpatrick J.L. Heiberg. *Euclid's Elements of Geometry*. <http://farside.ph.utexas.edu/euclid.html>, Austin/TX, 2008.
- [16] David Hilbert. *Grundlagen der Geometrie*. B.G.Teubner, Stuttgart, 1968.
- [17] Martin Odersky, Lex Spoon, and Bill Venners. *Programming in Scala*. Artima, 2008.